

Netflix in the Cloud

Nov 6, 2010

Adrian Cockcroft

@adrianco #netflixcloud

<http://www.linkedin.com/in/adriancockcroft>



Combined Slides

For both Developer and Operations Audiences (2-3 hours of content)

Teaser intro slides also at slideshare.net/adrianco

Oct 14th Beta #devops subset - Cloud Computing Meetup

Nov 3rd GA – QConSF Developer Oriented subset

Nov 6th, 2010 – Combined slides at slideshare.net/adrianco



With more than 16 million subscribers in the United States and Canada, Netflix, Inc. is the world's leading Internet subscription service for enjoying movies and TV shows.

Source: <http://ir.netflix.com>



Why Give This Talk?



Netflix is Path-finding

The Cloud ecosystem is evolving very fast
Share with and learn from the cloud community





adrianco adrian cockcroft

This week's choice: buy/sell cloud at [#CloudExpo](#) or build/use cloud
at [#QConSF](#), I'm a user, you?

15 hours ago

We want to use clouds, not build them

Cloud technology should be a commodity
Public cloud and open source for agility and scale



We are looking for talent

Netflix wants to connect with the very best
engineers



Why Use AWS?



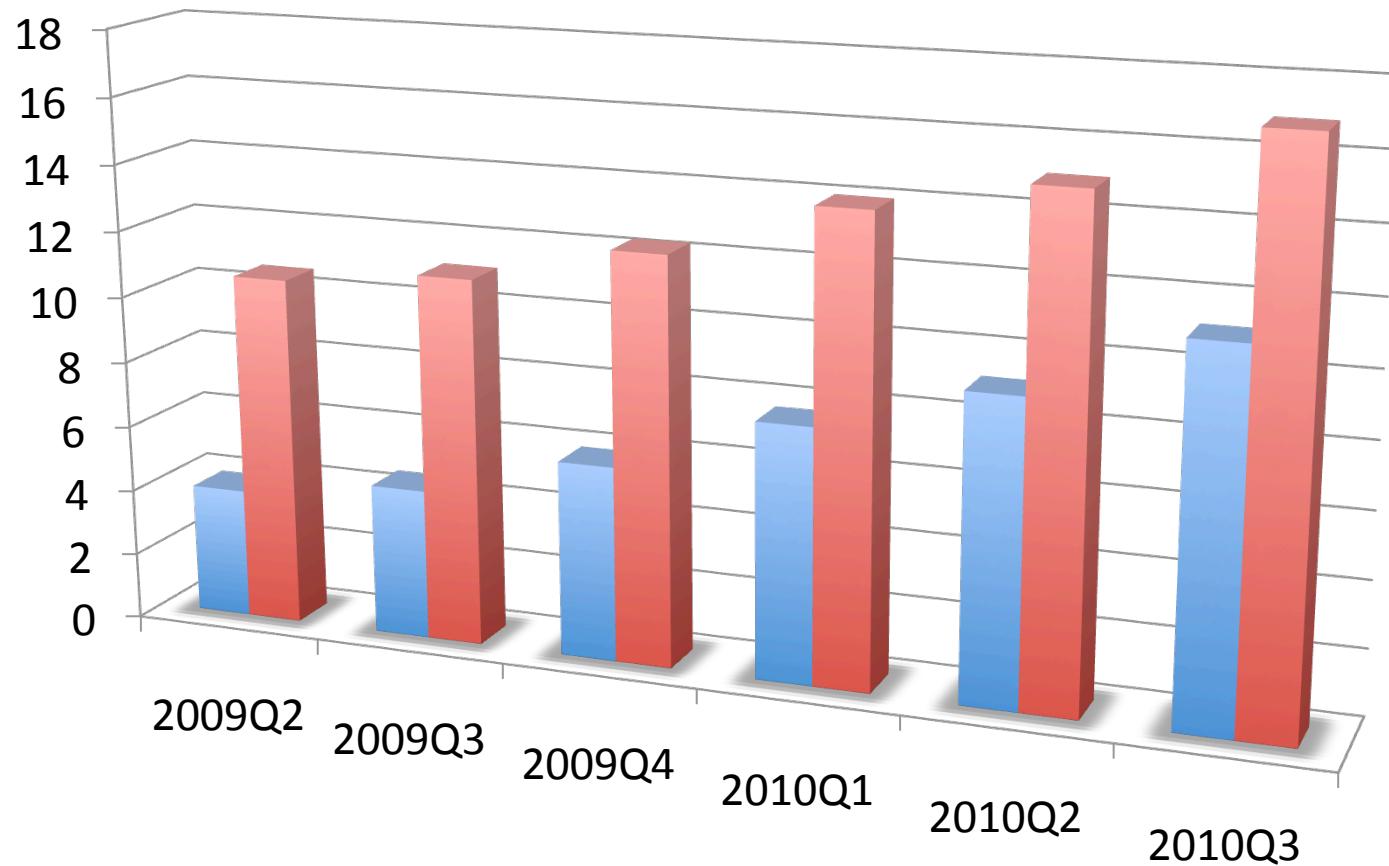
We stopped building our own datacenters

Capacity growth rate is accelerating, unpredictable
Product launch spikes - iPhone, Wii, PS3, XBox
Datacenter is large inflexible capital commitment



Customers

Q3 year/year +52% Total and +145% Streaming



Source: <http://ir.netflix.com>



Leverage AWS Scale “the biggest public cloud”

AWS investment in tooling and automation

AWS zones for high availability, scalability

AWS skills are common on resumes...



Leverage AWS Feature Set “two years ahead of the others”

EC2, S3, SDB, SQS, EBS, EMR, ELB, ASG, IAM, RDB...



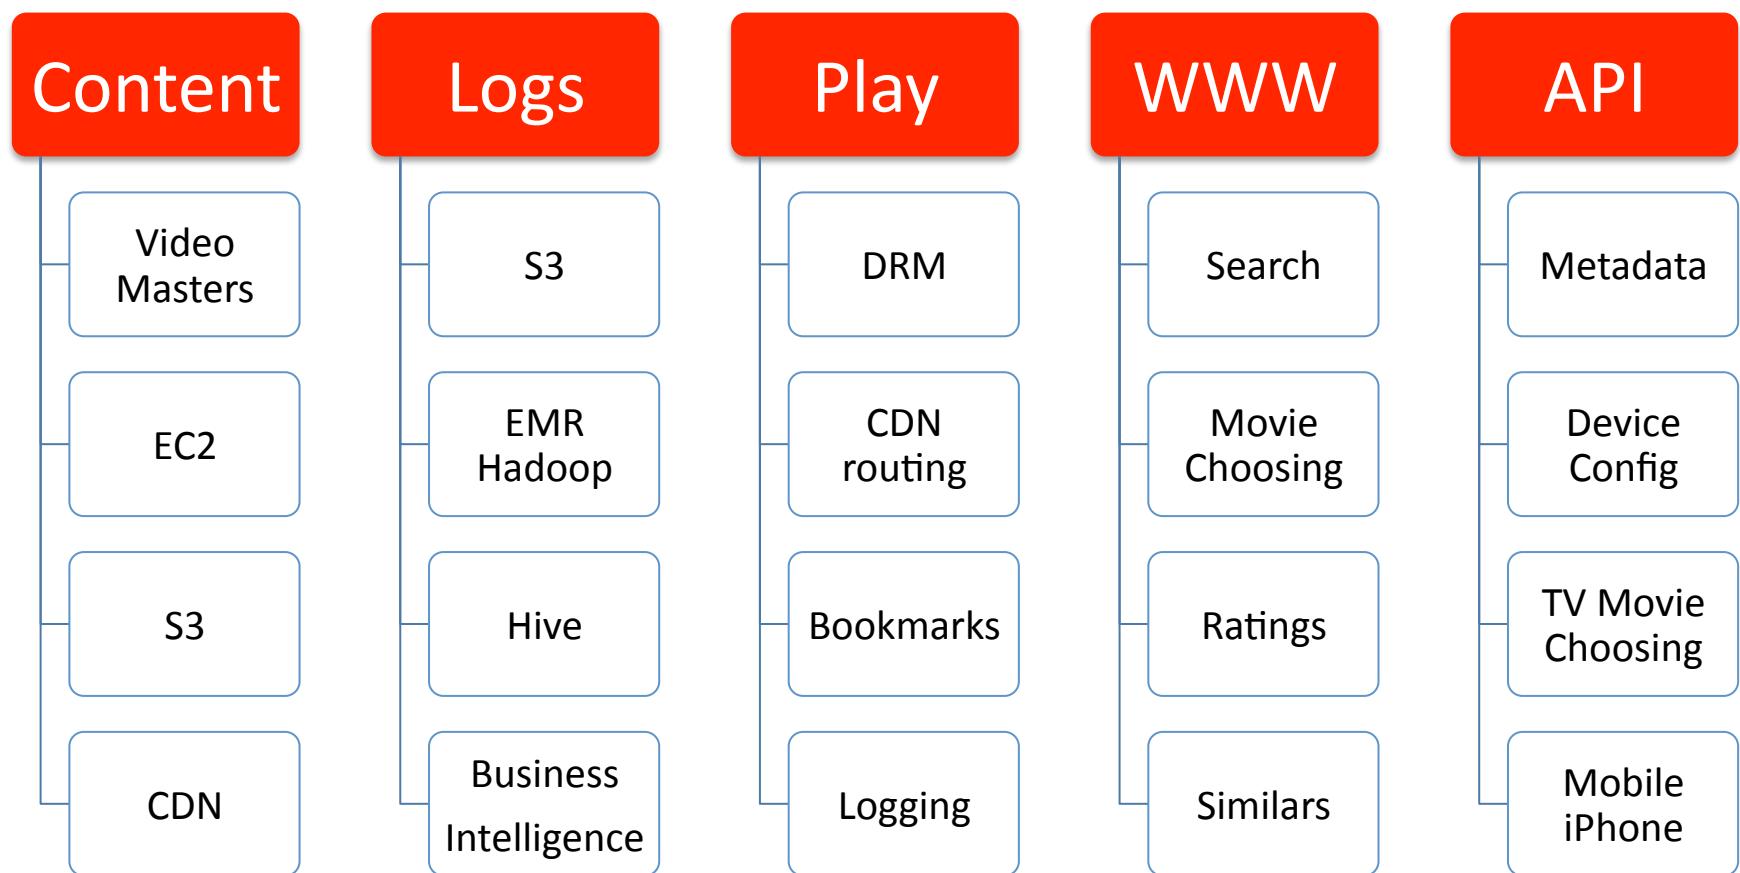
*“The cloud lets its users focus on delivering differentiating business value instead of wasting valuable resources on the **undifferentiated heavy lifting** that makes up most of IT infrastructure.”*



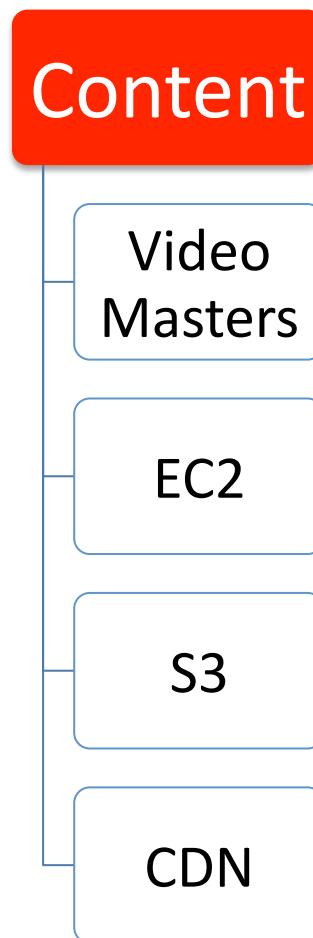
Werner Vogels
Amazon CTO



Netflix Deployed on AWS



Movie Encoding farm (2009)

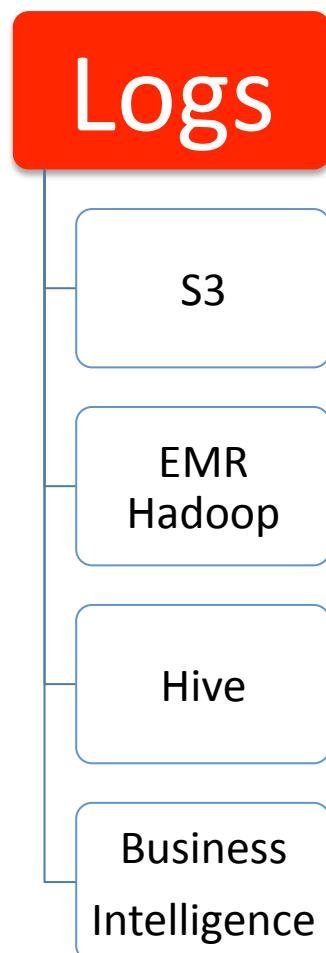


- Tens of thousands of videos
- Thousands of EC2 instances
- Encoding apps on MS Windows
- ~100 speed/format permutations
- Petabytes of S3
- Content Delivery Networks

“Netflix is one of the largest customers of the biggest CDNs Akamai and Limelight”



Hadoop - Elastic Map-Reduce (2009)



- Web Access Logs
- Streaming Service Logs
- Terabyte per day scale
- Easy Hadoop via Amazon EMR
- Hive SQL “Data Mart”
- Gateway to Datacenter BI

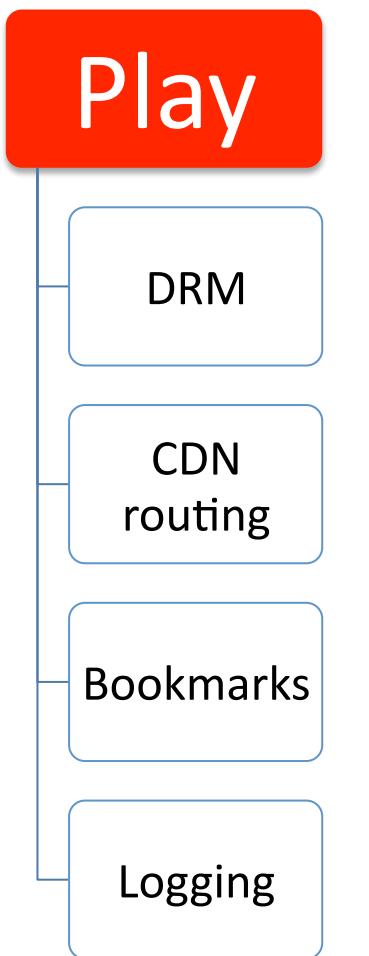
Slideshare.net talks

evamtse “Netflix: Hive User Group” <http://slidesha.re/aqJLAC>

adrianco “Crunch Your Data In The Cloud” <http://slidesha.re/dx4oCK>



Streaming Service Back-end (early 2010)



- PC/Mac Silverlight Player Support
- Highly available “play button”
- DRM Key Management
- Generate route to stream on CDN
- Lookup bookmark for user/movie
- Update bookmark for user/movie
- Log quality of service



Web site, a page at a time (through 2010)

WWW

Search

Movie
Choosing

Ratings

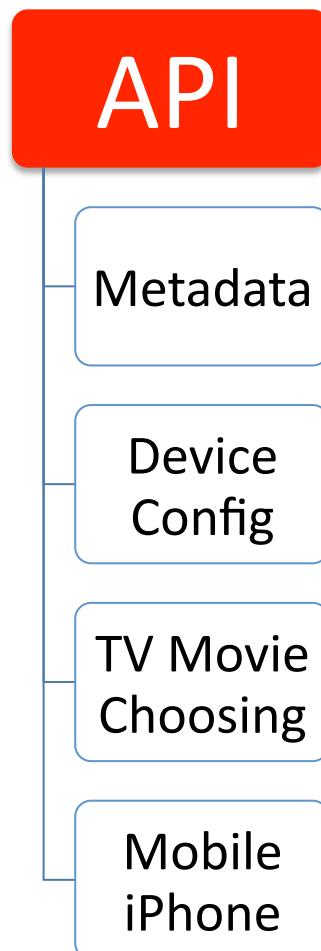
Similar

- Clean presentation layer rewrite
- Search auto-complete
- Search backend and landing page
- Movie and genre choosing
- Star ratings and recommendations
- Similar movies
- Page by page to 80% of views

(leave account signup in Datacenter for now)



API for TV devices and iPhone etc. (2010)



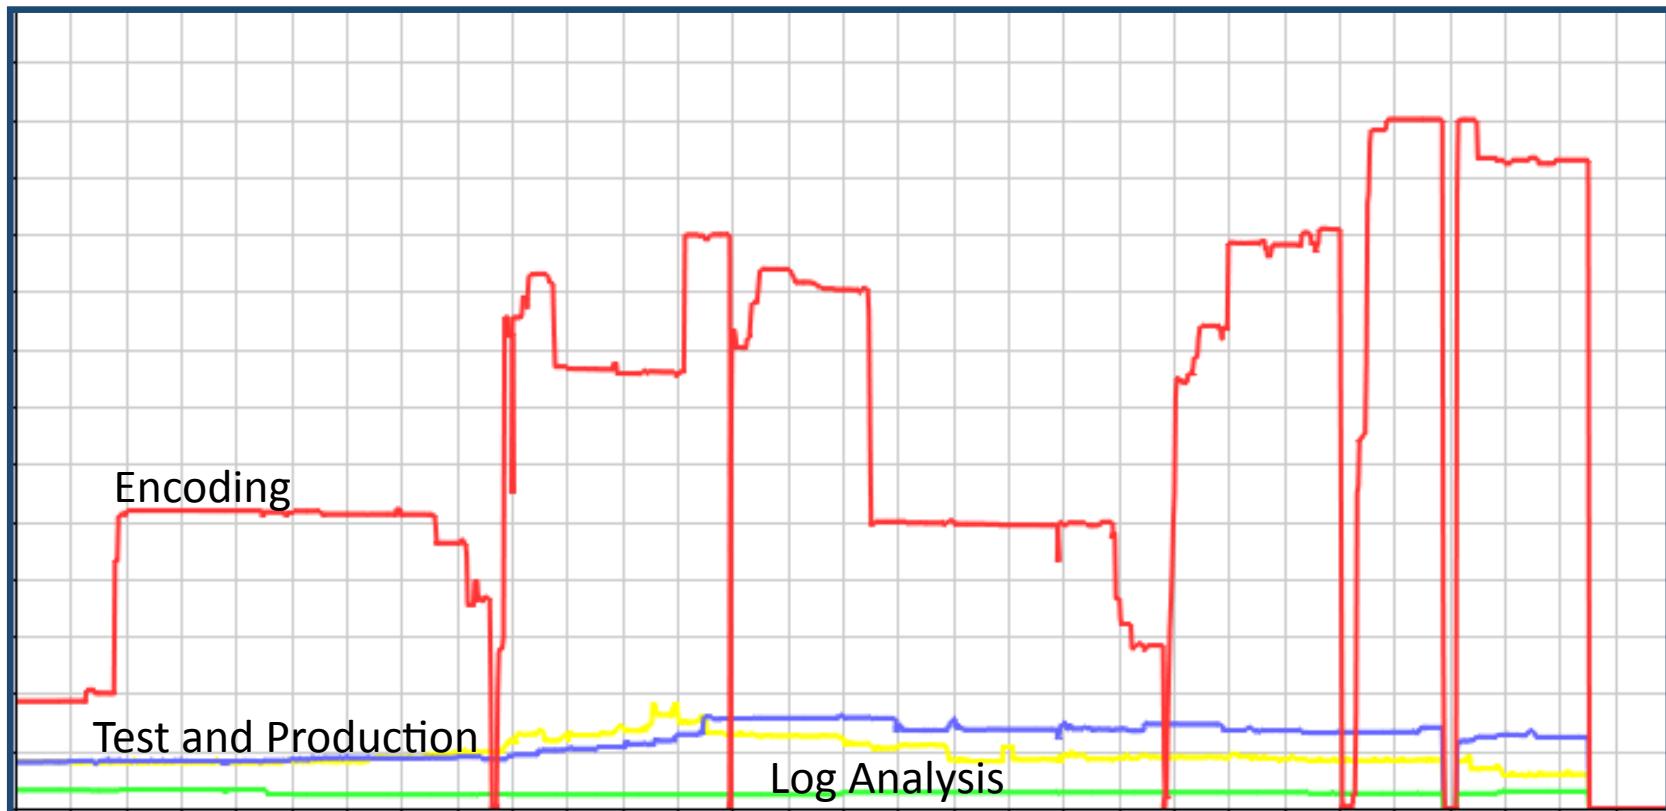
- REST API: developer.netflix.com
- Interfaces to everything else
- TV Device Configuration
- Personalized movie choosing
- iPhone Launch in the cloud only

“Netflix is an API for streaming to TVs

(we also do DVD's and a web site)”



Netflix EC2 Instances per Account



Learnings...



Datacenter oriented tools don't work

Ephemeral instances

High rate of change



Cloud Tools Don't Scale for Enterprise

Too many are “Startup” oriented
Built our own tools
Drove vendors hard



“fork-lifted” apps don’t work well

Fragile

Too many datacenter oriented
assumptions



Faster to re-code from scratch

- Opportunity to pay down technical debt
- Re-architected and re-wrote most of the code
- Fine grain web services
- Leveraged many open source Java projects
- Systematically instrumented
- “NoSQL” SimpleDB backend



“In the datacenter, robust code is best practice. In the cloud, it’s essential.”



Takeaway

Netflix is path-finding the use of public AWS cloud to replace in-house IT for non-trivial applications with hundreds of developers and thousands of systems.

(Pause for questions before we dive into details)



What, Why and How?

The details...



Synopsis

- The Goals
 - Faster, Scalable, Available and Productive
- Anti-patterns and Cloud Architecture
 - The things we wanted to change and why
- Cloud Bring-up Strategy
 - Developer Transitions and Tools
- Roadmap and Next Steps



Goals

- Faster
 - Lower latency than the equivalent datacenter web pages and API calls
 - Measured as mean and 99th percentile
 - For both first hit (e.g. home page) and in-session hits for the same user
- Scalable
 - Avoid needing any more datacenter capacity as subscriber count increases
 - No central vertically scaled databases
 - Leverage AWS elastic capacity effectively
- Available
 - Substantially higher robustness and availability than datacenter services
 - Leverage multiple AWS availability zones
 - No scheduled down time, no central database schema to change
- Productive
 - Optimize agility of a large development team with automation and tools
 - Leave behind complex tangled datacenter code base (~8 year old architecture)
 - Enforce clean layered interfaces and re-usable components



Cloud Architecture Patterns

Where do we start?



Datacenter Anti-Patterns

What do we currently do in the datacenter that prevents us from meeting our goals?



Architecture

- Software Architecture
 - The abstractions and interfaces that developers build against
- Systems Architecture
 - The service instances that define availability, scalability
- Compose-ability
 - software architecture that is independent of the systems architecture
 - decoupled flexible building block components



Rewrite from Scratch

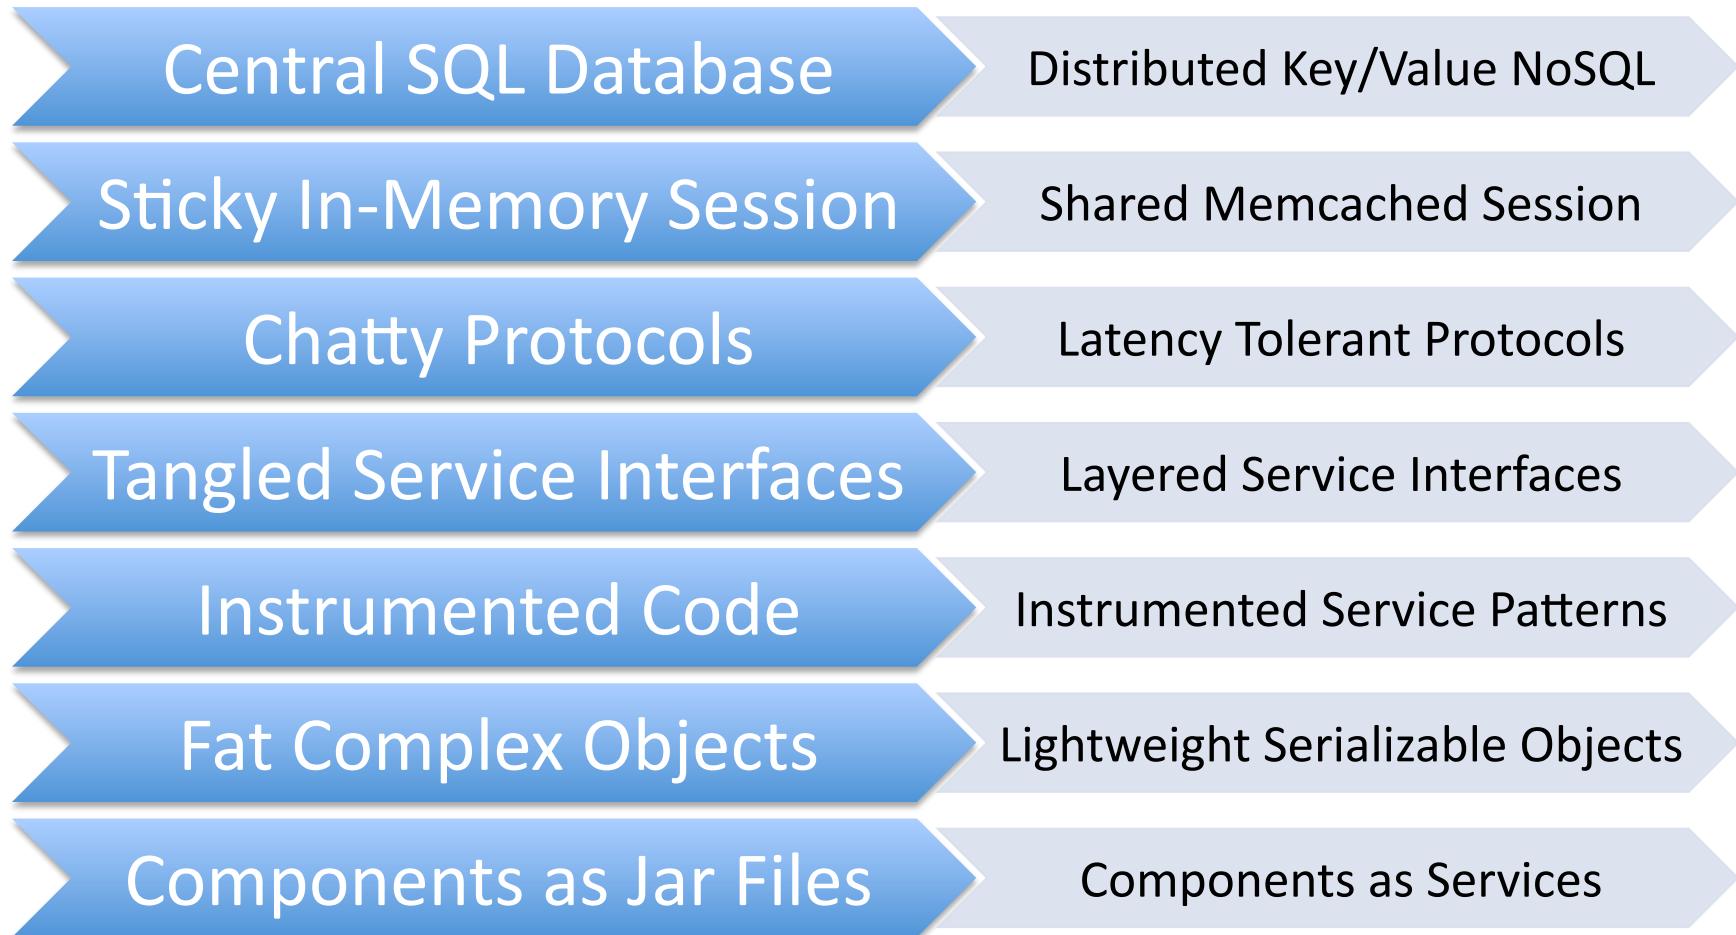
Not everything is cloud specific

Pay down technical debt

Robust patterns



Old Datacenter vs. New Cloud Arch



The Central SQL Database

- Datacenter has a central database
 - Everything in one place is convenient until it fails
 - Customers, movies, history, configuration
- Schema changes require downtime

Anti-pattern impacts scalability, availability



The Distributed Key-Value Store

- Cloud has many key-value data stores
 - More complex to keep track of, do backups etc.
 - Each store is much simpler to administer
 - Joins take place in java code
- No schema to change, no scheduled downtime
- Latency for Memcached vs. Oracle vs. SimpleDB
 - Memcached is dominated by network latency <1ms
 - Oracle for simple queries is a few milliseconds
 - SimpleDB has replication and REST overheads >10ms



The Sticky Session

- Datacenter Sticky Load Balancing
 - Efficient caching for low latency
 - Tricky session handling code
 - Middle tier load balancer has issues in practice
- Encourages concentrated functionality
 - one service that does everything

Anti-pattern impacts productivity, availability



The Shared Session

- Cloud Uses Round-Robin Load Balancing
 - Simple request-based code
 - External shared caching with memcached
- More flexible fine grain services
 - Works better with auto-scaled instance counts



Chatty Opaque and Brittle Protocols

- Datacenter service protocols
 - Assumed low latency for many simple requests
- Based on serializing existing java objects
 - Inefficient formats
 - Incompatible when definitions change

Anti-pattern causes productivity, latency and availability issues



Robust and Flexible Protocols

- Cloud service protocols
 - JSR311/Jersey is used for REST/HTTP service calls
 - Custom client code includes service discovery
 - Support complex data types in a single request
- Apache Avro
 - Evolved from Protocol Buffers and Thrift
 - Includes JSON header defining key/value protocol
 - Avro serialization is **half the size** and several times faster than Java serialization, more work to code



Persisted Protocols

- Persist Avro in Memcached
 - Save space/latency (zigzag encoding, half the size)
 - Less brittle across versions
 - New keys are ignored
 - Missing keys are handled cleanly
- Avro protocol definitions
 - Can be written in JSON or generated from POJOs
 - It's hard, needs better tooling



Tangled Service Interfaces

- Datacenter implementation is exposed
 - Oracle SQL queries mixed into business logic
- Tangled code
 - Deep dependencies, false sharing
- Data providers with sideways dependencies
 - Everything depends on everything else

Anti-pattern affects productivity, availability



Untangled Service Interfaces

- New Cloud Code With Strict Layering
 - Compile against interface jar
 - Can use spring runtime binding to enforce
- Service interface **is** the service
 - Implementation is completely hidden
 - Can be implemented locally or remotely
 - Implementation can evolve independently



Untangled Service Interfaces

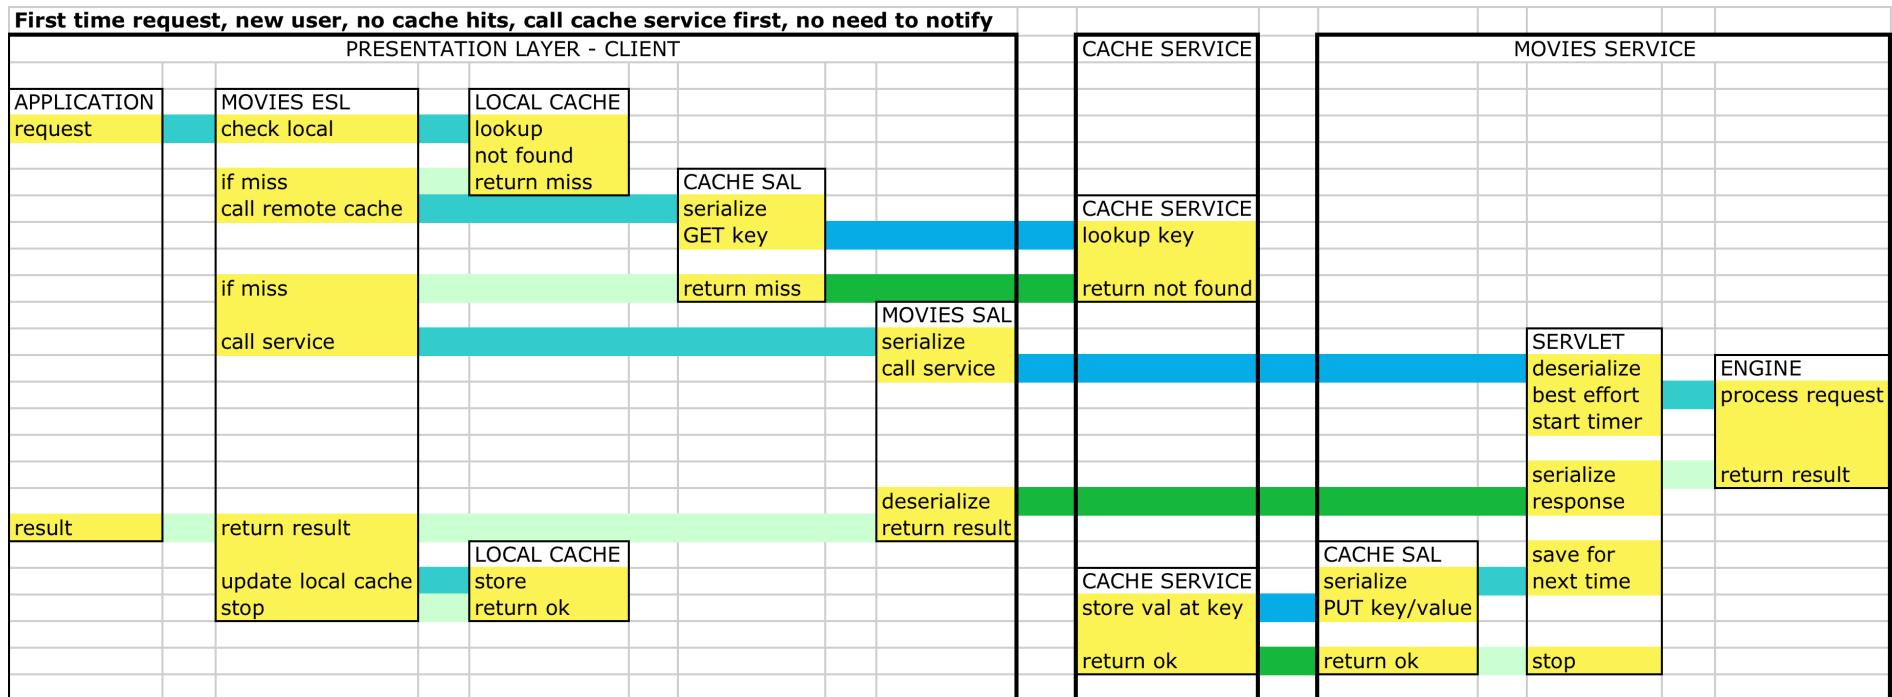
Two layers:

- SAL - Service Access Library
 - Basic serialization and error handling
 - REST or POJO's defined by data provider
- ESL - Extended Service Library
 - Caching, conveniences
 - Can combine several SALs
 - Exposes faceted type system (described later)
 - Interface defined by data consumer in many cases



Service Interaction Pattern

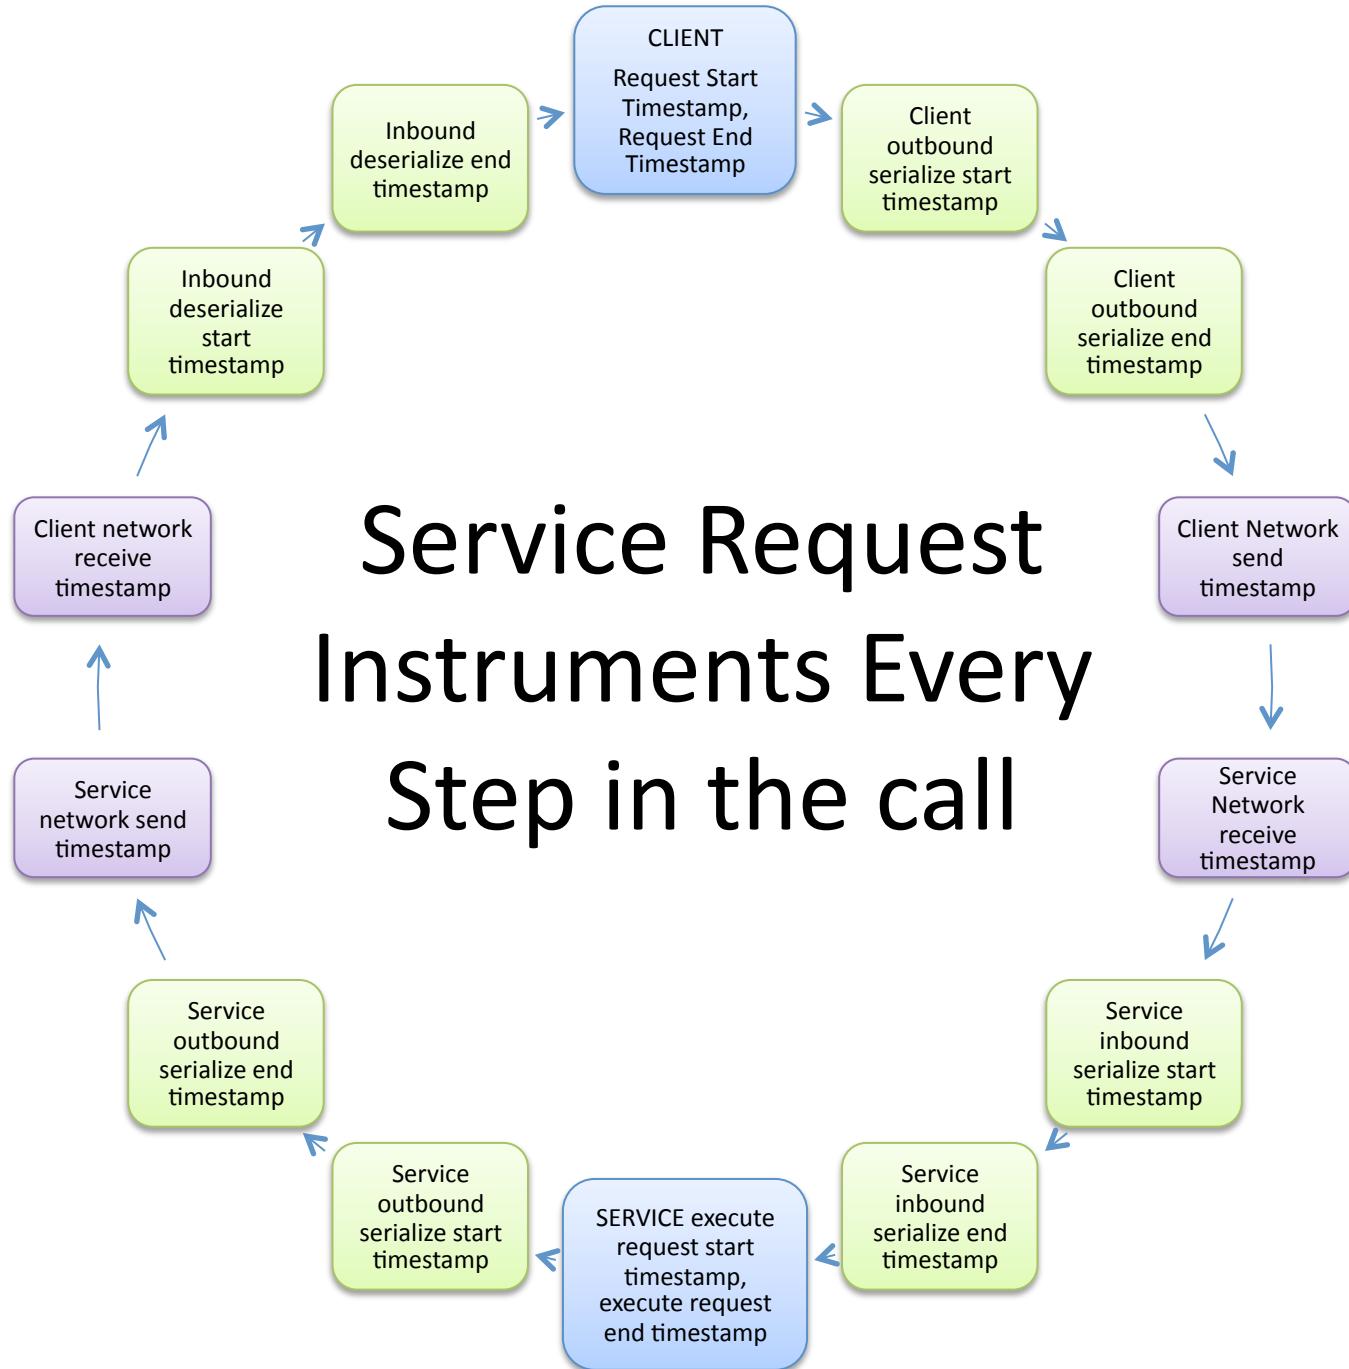
Sample Swimlane Diagram



Service Architecture Patterns

- Internal Interfaces Between Services
 - Common patterns as templates
 - Highly instrumented, observable, analytics
 - Service Level Agreements – SLAs
- Library templates for generic features
 - Instrumented Netflix Base Servlet template
 - Instrumented generic client interface template
 - Instrumented S3, SimpleDB, Memcached clients





Boundary Interfaces

- Isolate teams from external dependencies
 - Fake SAL built by cloud team
 - Real SAL provided by data provider team later
 - ESL built by cloud team using faceted objects
- Fake data sources allow development to start
 - e.g. Fake Identity SAL for a test set of customers
 - Development solidifies dependencies early
 - Helps external team provide the right interface



One Object That Does Everything

- Datacenter uses a few big complex objects
 - Movie and Customer objects are the foundation
 - Good choice for a small team and one instance
 - Problematic for large teams and many instances
- False sharing causes tangled dependencies
 - Unproductive re-integration work

Anti-pattern impacting productivity and availability



An Interface For Each Component

- Cloud uses faceted Video and Visitor
 - Basic types hold only the identifier
 - Facets scope the interface you actually need
 - Each component can define its own facets
- No false-sharing and dependency chains
 - Type manager converts between facets as needed
 - `video.asA(PresentationVideo)` for www
 - `video.asA(MerchableVideo)` for middle tier



Software Architecture Patterns

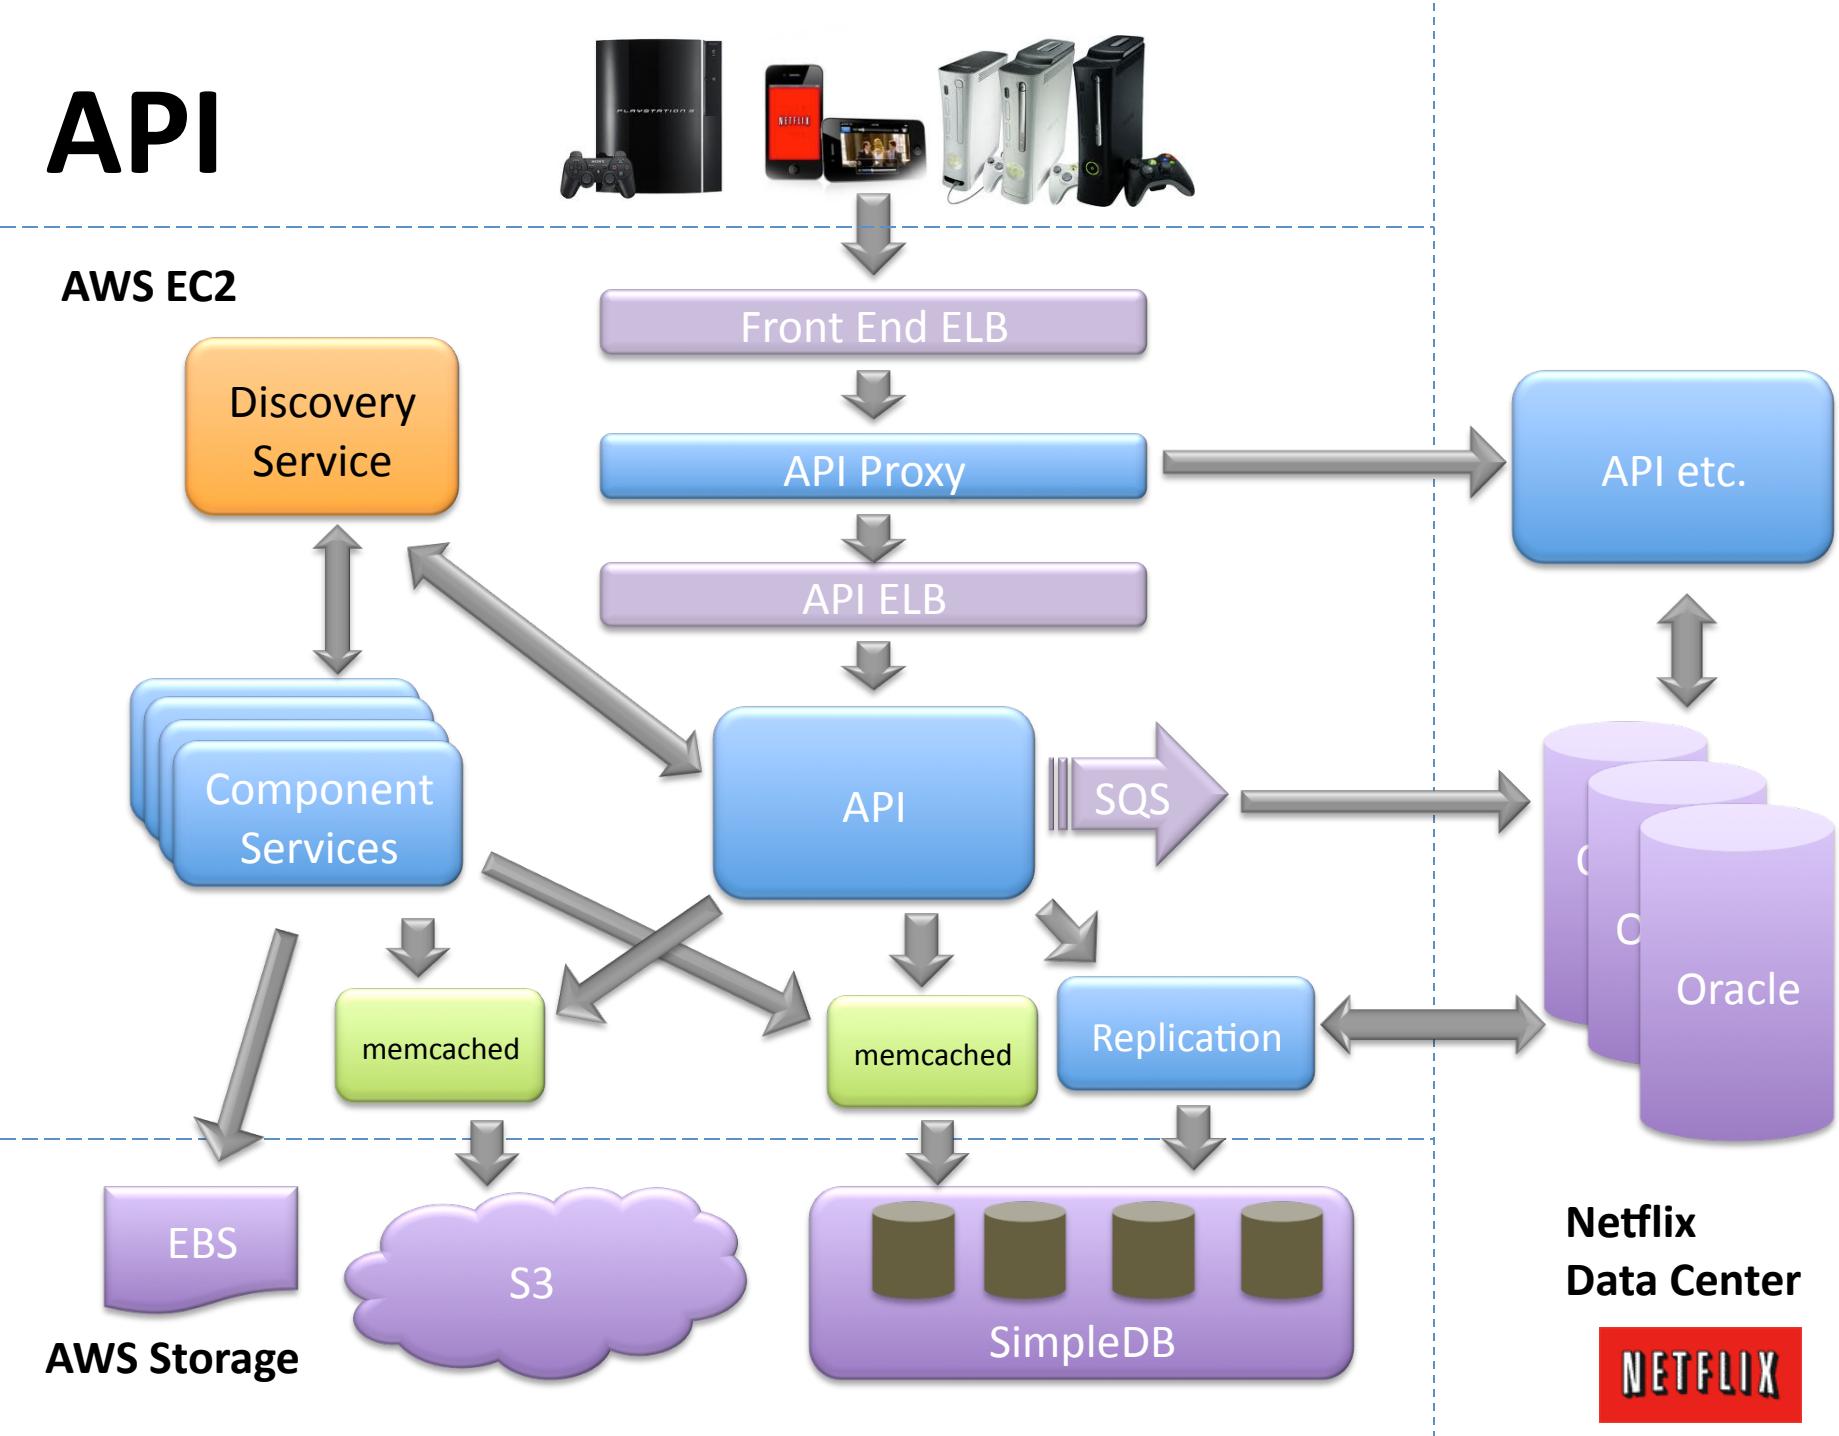
- Object Models
 - Basic and derived types, facets, serializable
 - Pass by reference within a service
 - Pass by value between services
- Computation and I/O Models
 - Service Execution using Best Effort
 - Common thread pool management



Netflix Systems Architecture



API



Netflix
Data Center



Netflix Undifferentiated Lifting

- Middle Tier Load Balancing
- Discovery (local DNS)
- Encryption Services
- Caching
- Distributed App Management



We want cloud vendors to do all this for us as well!



Load Balancing in AWS

- Middle tier currently not supported in AWS
 - ELB are public-facing only
 - Cannot apply security group settings
- ELB vertical scalability for concentrated clients
 - Too few proxy IP addresses leads to hot spots
- ELB needs support for balancing heuristics
 - Proportional balance across Availability Zones
 - Weighted Least connections, Weighted Round Robin
- Zone aware routing
 - Default to instances in the same Availability Zone
 - Falls back to cross-zone on failure



Discovery

- Discovery Service (Redundant instances per zone)
 - Simple REST interface
 - Cloud apps register with Discovery
- Apps send heartbeats every 30 sec to renew lease
 - App evicted after 3 missed heartbeats
 - Can re-register if the problem was transient
- Apps can store custom metadata
 - Version number, AMI id, Availability Zone, etc.
- Software Round-robin Load Balancer
 - Query Discovery for instances of specific application
 - Baked into Netflix REST client (JSR311/Jersey based)

AWS Middle-tier ELB would eliminate most use cases



Database Migration

- Why SimpleDB?
 - No DBA's in the cloud, Amazon hosted service
 - Work started two years ago, fewer viable options
 - Worked with Amazon to speed up and scale SimpleDB
- Alternatives?
 - Investigating adding Cassandra and Membase to the mix
 - Need several options to match use cases well
- Detailed SimpleDB Advice
 - Sid Anand - QConSF Nov 5th – Netflix' Transition to High Availability Storage Systems
 - Blog - <http://practicalcloudcomputing.com/>
 - Download Paper PDF - <http://bit.ly/bhOTLu>



Oracle to SimpleDB

(See Sid's paper for details)

- SimpleDB Domains
 - De-normalize multiple tables into a single domain
 - Work around size limits (10GB per domain, 1KB per key)
 - Shard data across domains to scale
 - Key – Use distributed sequence generator, GUID or natural unique key such as customer-id
 - Implement a schema validator to catch bad attributes
- Application layer support
 - Do GROUP BY and JOIN operations in the application
 - Compose relations in the application layer
 - Check constraints on read, and repair data as a side effect
- Do without triggers, PL/SQL, clock operations



Tools and Automation

- Developer and Build Tools
 - Jira, Eclipse, Hudson, Ivy, Artifactory
 - Builds, creates .war file, .rpm, bakes AMI and launches
- Custom Netflix Application Console
 - AWS Features at Enterprise Scale (hide the keys!)
 - Auto Scaler Group is unit of deployment to production
- Open Source + Support
 - Apache, Tomcat, OpenJDK, CentOS
- Monitoring Tools
 - Keynote – service monitoring and alerting
 - AppDynamics – Developer focus for cloud
 - EpicNMS – flexible data collection and plots <http://epicnms.com>
 - Nimsoft NMS – ITOps focus for Datacenter + Cloud alerting



Netflix App Console

The screenshot shows the Netflix Application Console interface. At the top, there is a red header bar with the Netflix logo and the text "Application Console (test)". To the right of the header is a search bar labeled "CMC:" with a clear button. Below the header is a navigation menu with icons and labels: Home, Apps, Images, Auto Scaling, Load Balancers, Instances, EBS, RDS, and Tasks. The "Auto Scaling" tab is currently selected.

The main content area has three tabs: "Netflix Abstractions", "AWS Objects", and "NAC Tasks". The "AWS Objects" tab is active and displays a list of management links:

- Manage Images
- Manage Auto Scaling Groups
- Manage Load Balancers
- Manage Launch Configurations
- Manage Security Groups
- Manage Running Instances
- View Reservations

The "NAC Tasks" tab contains a single link: "Monitor NAC Background Tasks".

Below the tabs, there is a section titled "All NAC Links" containing two buttons: "NAC for test account" and "NAC for prod account". To the right of these buttons is a search bar with the placeholder "Jump to an instance:" and a "Go" button.



Auto Scaling Groups

 **NETFLIX Application Console (test)** CMC: [Clear](#)

[!\[\]\(3717f9b70da7b77f6af849cc519cd6de_img.jpg\) Home](#) [!\[\]\(c8551bf3f4f352e1d5cbd21e2c1666e8_img.jpg\) Apps](#) [!\[\]\(17ca22000f00aedb731eebddd06ab5df_img.jpg\) Images](#) [!\[\]\(ec0b8ade592dc59b37022b0e9b8e50d8_img.jpg\) Auto Scaling](#) [!\[\]\(25367e20c9876bcb85e9f5df83bd6056_img.jpg\) Load Balancers](#) [!\[\]\(552011001188f0fd63491de20f5ed09c_img.jpg\) Instances](#) [!\[\]\(fcb765d385c4c166c89c348f5324d55a_img.jpg\) EBS](#) [!\[\]\(cfc0828d1ba0daa260f149fe055bb997_img.jpg\) RDS](#) [!\[\]\(d749717472f683b92025dd17644bb5ba_img.jpg\) Tasks](#)

Auto Scaling Groups

| Create New Auto Scaling Group | | | | | | | | | | Filter: <input type="text" value="web"/> |
|-------------------------------|--------------------------------|-----|-----|-----|------|--------------------------|-----------------------|--|-------------------------|--|
| Name | Launch Configuration | Min | Max | Des | Cool | Avail Zones | Load Balancers | Instances | Created Time | |
| gps-merchweb | gps-merchweb-201010052146 | 0 | 0 | 0 | 0 | us-east-1c us-east-1d | | | 2010-08-23 17:44:48 PDT | |
| merchweb | merchweb-201010131658 | 2 | 2 | 2 | 0 | us-east-1a us-east-1c | merchweb-frontend | i-62d6280f us-east-1a InService i-10ce307d us-east-1c InService | 2010-03-25 18:56:40 PDT | |
| merchweb-abhayani | merchweb-abhayani-201010011350 | 1 | 1 | 1 | 0 | us-east-1c us-east-1d | | i-615f410b us-east-1c InService | 2010-09-13 18:23:55 PDT | |
| merchweb-dev | merchweb-dev-201010081534 | 1 | 1 | 1 | 0 | us-east-1a us-east-1c | merchweb-frontend-dev | i-d0dc31bd us-east-1a InService | 2010-10-08 15:18:15 PDT | |
| merchweb-gpstest | merchweb-gpstest-201009291050 | 1 | 1 | 1 | 10 | us-east-1c us-east-1d | | i-43262d29 us-east-1d InService | 2010-09-29 10:49:20 PDT | |



ASG Configuration

 **NETFLIX Application Console (test)** CMC:

Home Apps Images Auto Scaling Load Balancers Instances EBS RDS Tasks

Auto Scaling Group Details

| | |
|-----------------------|--|
| Name: | merchweb-dev |
| Launch Configuration: | merchweb-dev-201010081534 |
| Min Instances: | 1 |
| Max Instances: | 1 |
| Desired Capacity: | 1 |
| Cool Down: | 0 |
| Availability Zones: | [us-east-1a, us-east-1c] |
| Created Time: | 2010-10-08 15:18:15 PDT |
| Load Balancers: | merchweb-frontend-dev |
| Instances: | Instance Count: 1 i-d0dc31bd us-east-1a InService |
| Activities: | At 2010-10-08T22:38:10Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 1. : Launching a new EC2 instance: i-d0dc31bd At 2010-10-08T22:37:38Z an instance was terminated in response to a system health-check. : Terminating EC2 instance i-b2f31edf At 2010-10-08T22:18:15Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 1. At 2010-10-08T22:20:04Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 1. : Launching a new EC2 instance: i-b2f31edf |

Pattern Matches

| | |
|--------------|----------|
| Application: | merchweb |
|--------------|----------|

 [Edit Auto Scaling Group](#)  [Delete Auto Scaling Group](#)  [Create a new Launch Config for this Auto Scaling Group](#)  [Prepare to Push Image into Auto Scaling Group](#)



Monitoring Tools



Monitoring Vision

- Problem
 - Too many tools, each with a good reason to exist
 - Hard to get an integrated view of a problem
 - Too much manual work building dashboards
 - Tools are not discoverable, views are not filtered
- Solution
 - Get vendors to add deep linking and embedding
 - Integration “portal” ties everything together
 - Dynamic portal generation, relevant data, all tools



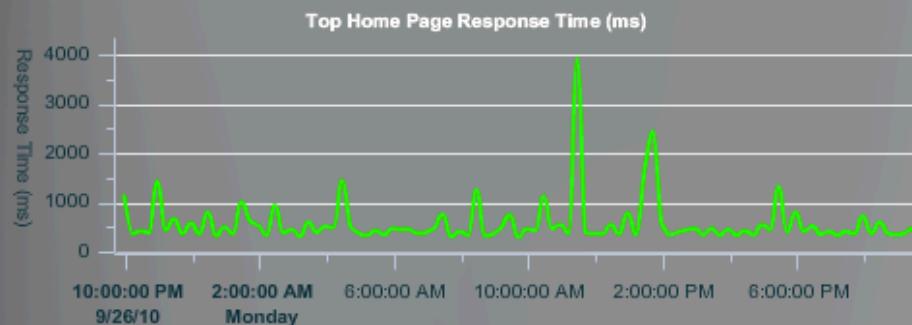
Cloud Monitoring Mechanisms

- Keynote
 - External URL monitoring
- Amazon CloudWatch
 - Metrics for ELB and Instances
- AppDynamics
 - End to end transaction view showing resources used
 - Powerful real time debug tools for latency, CPU and Memory
- Nimsoft NMS
 - Scalable and reliable monitoring and alerting, integration portal
- Epic
 - Flexible and easy to use to extend and embed plots
- Logs
 - High capacity logging and analysis framework
 - Hadoop (log4j -> chukwa -> EMR)



HEALTH STATUS DASHBOARD FOR SELECT REMOTE COMPUTING AND APPLICATION SERVICES

| Website | Home Page Response (ms) | Home Page Response (ms) |
|---------------------|-------------------------|-------------------------|
| Intacct | | 1110.00 |
| Salesforce.com | | 483.00 |
| Trinet HR Services | | 466.00 |
| Rackspace Cloud | | 465.00 |
| Google Apps Engine | | 272.00 |
| Google Apps | | 214.00 |
| Amazon Web Services | | 168.00 |



Home Page Response - Sf.com (ms)



Overall Status

Response Time

Double Click for Details >>>



Overall Status

Response Time

Double Click for Details >>>



Overall Status

Response Time

Double Click for Details >>>

NIMSOFT DATA CENTER - NETWORK, VOIP, VIRTUAL SERVERS



Overall Status
 Response Time
 Double Click for Details >>>



Overall Status
 Response Time
 Double Click for Details >>>

SAAS SERVICES



Overall Status
 Response Time
 Double Click for Details >>>



Overall Status
 Response Time
 Double Click for Details >>>

CLOUD/IaaS SERVICES



Nimsoft's Hosted MS Exchange
 Overall Status
 Response Time
 Double Click for Details >>>



Overall Status
 Response Time
 Double Click for Details >>>

HOSTED SERVICES



Overall Status
 Response Time
 Double Click for Details >>>

PAAS SERVICES

Snapshots for a Business Transaction

Sort Call Graphs to Top, pick a slow one

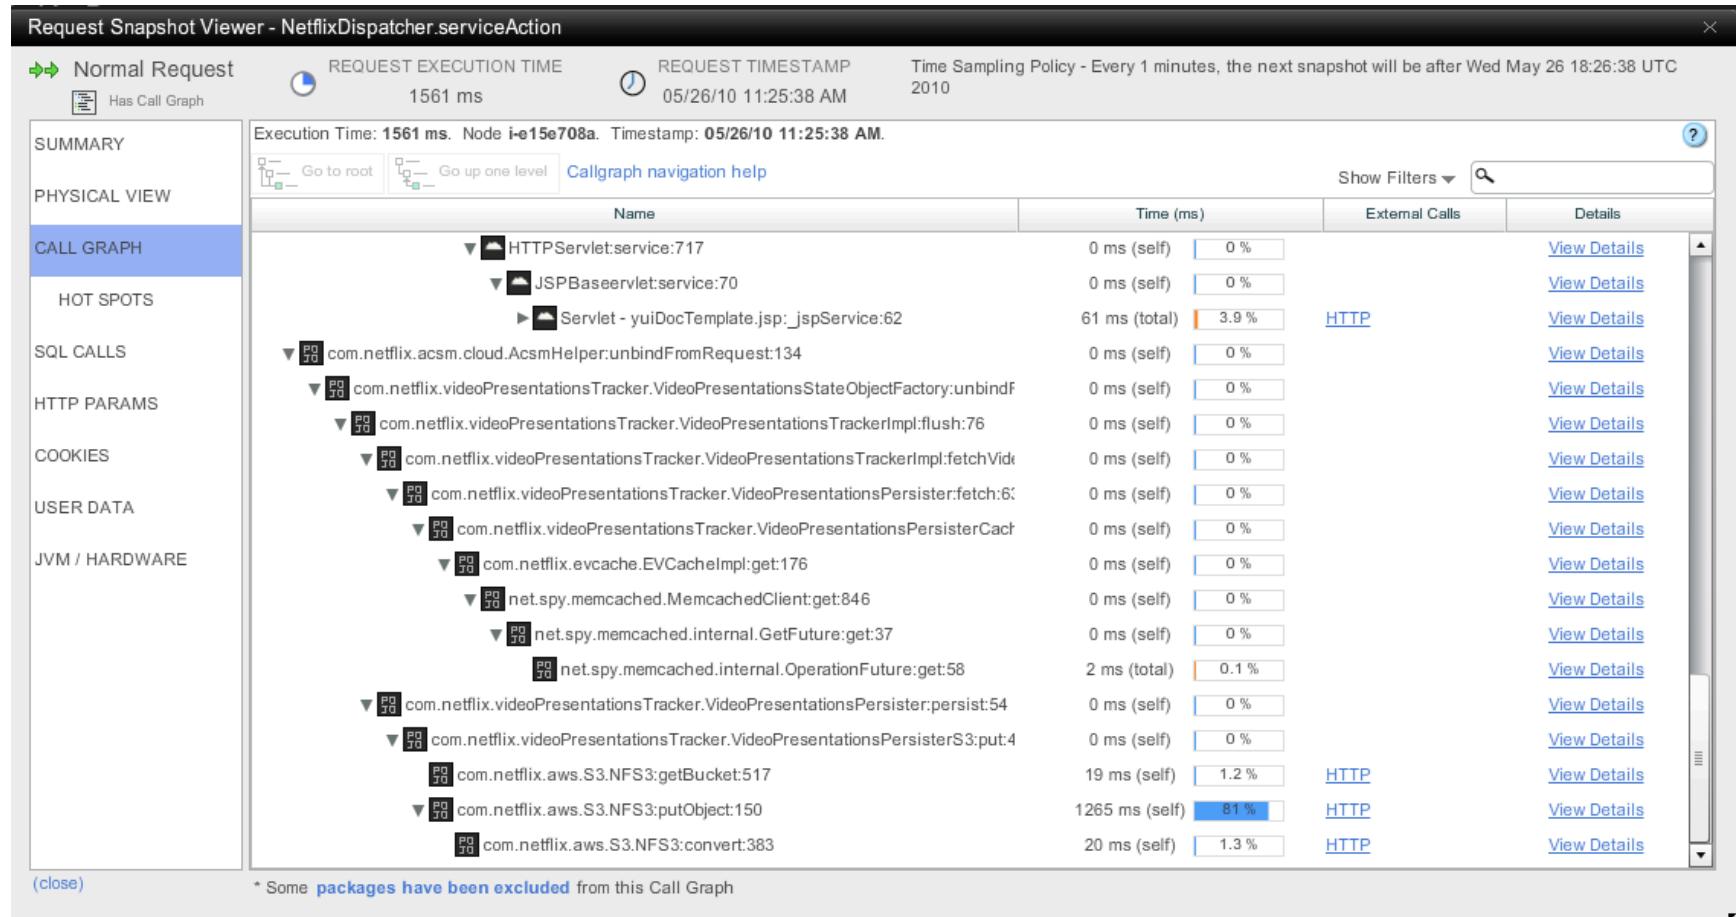
The screenshot shows the AppDynamics interface for monitoring business transactions. The top navigation bar includes 'Applications', 'Infrastructure', 'Settings', 'Help', and 'Logout (bob)'. The main content area displays '4664 Total Requests' with a breakdown by request status: Normal (97.8%), Slow (1.0%), Very Slow (1.1%), Errors (0.1%), and Stalls (0.0%). Below this is a section titled 'Request Snapshots' with a table listing 201 snapshots. The table columns include 'User Experience', 'Time', 'Execution Time(ms)', and 'Has Call Gra'. A red box highlights the 'Has Call Gra' column. The table rows show various requests with their execution times and associated sampling policies.

| User Experience | Time | Execution Time(ms) | Has Call Gra | Summary |
|-----------------|----------------------|--------------------|--------------|--|
| Normal | 05/26/10 11:25:21 AM | 379 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-6d0f2d06 |
| Normal | 05/26/10 11:25:57 AM | 275 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-a30e5bc8 |
| Normal | 05/26/10 11:29:00 AM | 259 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-a30e5bc8 |
| Normal | 05/26/10 11:24:20 AM | 539 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-8b0a24e0 |
| Normal | 05/26/10 11:29:24 AM | 384 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-8b0a24e0 |
| Normal | 05/26/10 11:32:27 AM | 419 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-6d0f2d06 |
| Normal | 05/26/10 11:28:24 AM | 1309 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-6d0f2d06 |
| Normal | 05/26/10 11:34:14 AM | 455 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-adf8dac6 |
| Normal | 05/26/10 11:25:38 AM | 1561 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-e15e708a |
| Normal | 05/26/10 11:31:27 AM | 489 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-a3f8dac8 |
| Normal | 05/26/10 11:32:27 AM | 389 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-a3f8dac8 |
| Normal | 05/26/10 11:26:26 AM | 648 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-538ba438 |
| Normal | 05/26/10 11:26:05 AM | 423 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-5d8ba436 |
| Normal | 05/26/10 11:24:03 AM | 346 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-adf8dac6 |
| Normal | 05/26/10 11:32:38 AM | 553 ms | [Icon] | Time Sampling Policy - Every 1 minutes, the next snapshot will be after Wed May i-538ba438 |



Drill in to Slow Call

Slow Asynchronous S3 Write – no big deal...



Cloud Bring-Up Strategy

Simplest and Soonest

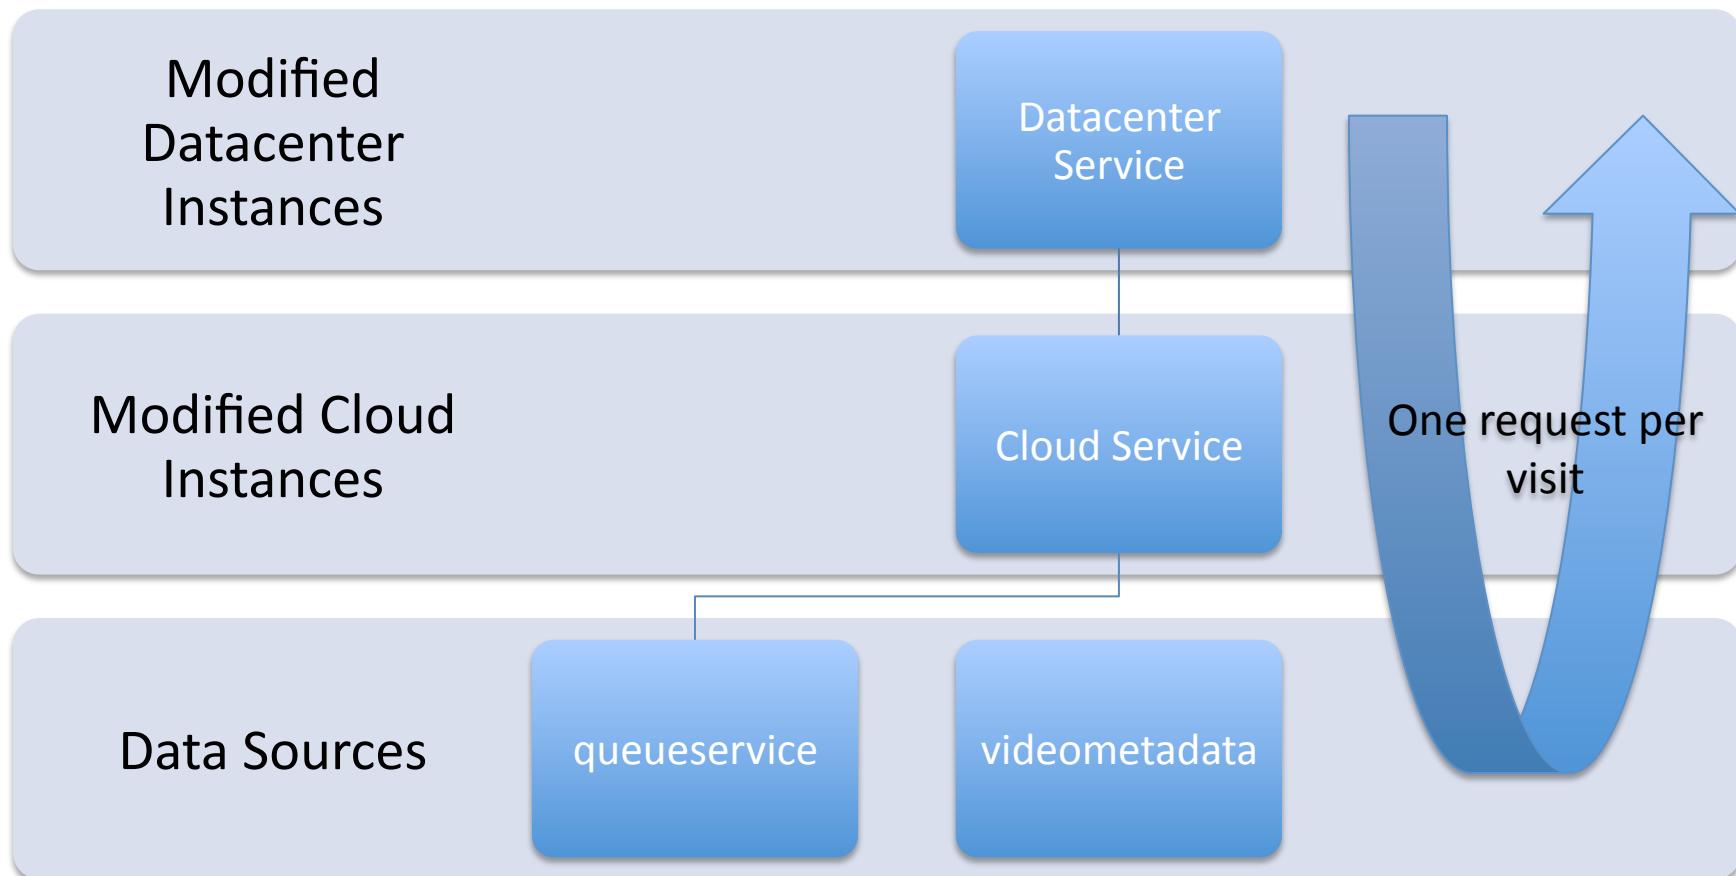


Shadow Traffic Redirection

- Early attempt to send traffic to cloud
 - Real traffic stream to validate cloud back end
 - Uncovered lots of process and tools issues
 - Uncovered Service latency issues
- TV Device calls Datacenter API
 - Returns Genre/movie list for a customer
 - Asynchronously duplicates request to cloud
 - Start with send-and-forget mode, ignore response



Shadow Redirect Instances



First Web Pages in the Cloud



Starz Page

NETFLIX

Adrian Cockcroft | Your Account & Help

Watch Instantly | Browse DVDs | Your Queue | Movies You'll ❤️

Movies, TV shows, actors, directors, genres | Search

Home | Genres | New Arrivals | Starz Play | Instantly to your TV | Help

Watch Movies & TV Episodes Instantly

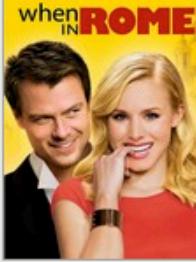
starz PLAY

- Approximately 1,000 titles available from Starz Play
- Recent Hollywood hits, contemporary favorites, and classics
- Original series, specials, and more

Show rated/seen titles

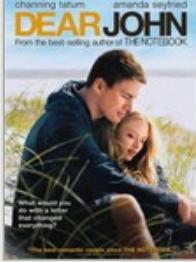
Gallery View | Sortable List | [← prev](#) | 1 | 2 | 3 | 4 | ... | 46 | [next →](#)

When in Rome | The Princess and the Frog | Astro Boy | Dear John


Play
★★★☆☆
 Not Interested


Play
★★★☆☆
 Not Interested


Play
★★★☆☆
 Not Interested


Play
★★★☆☆
 Not Interested

starz PLAY

Live Starz TV Channel | About Starz Play

NETFLIX

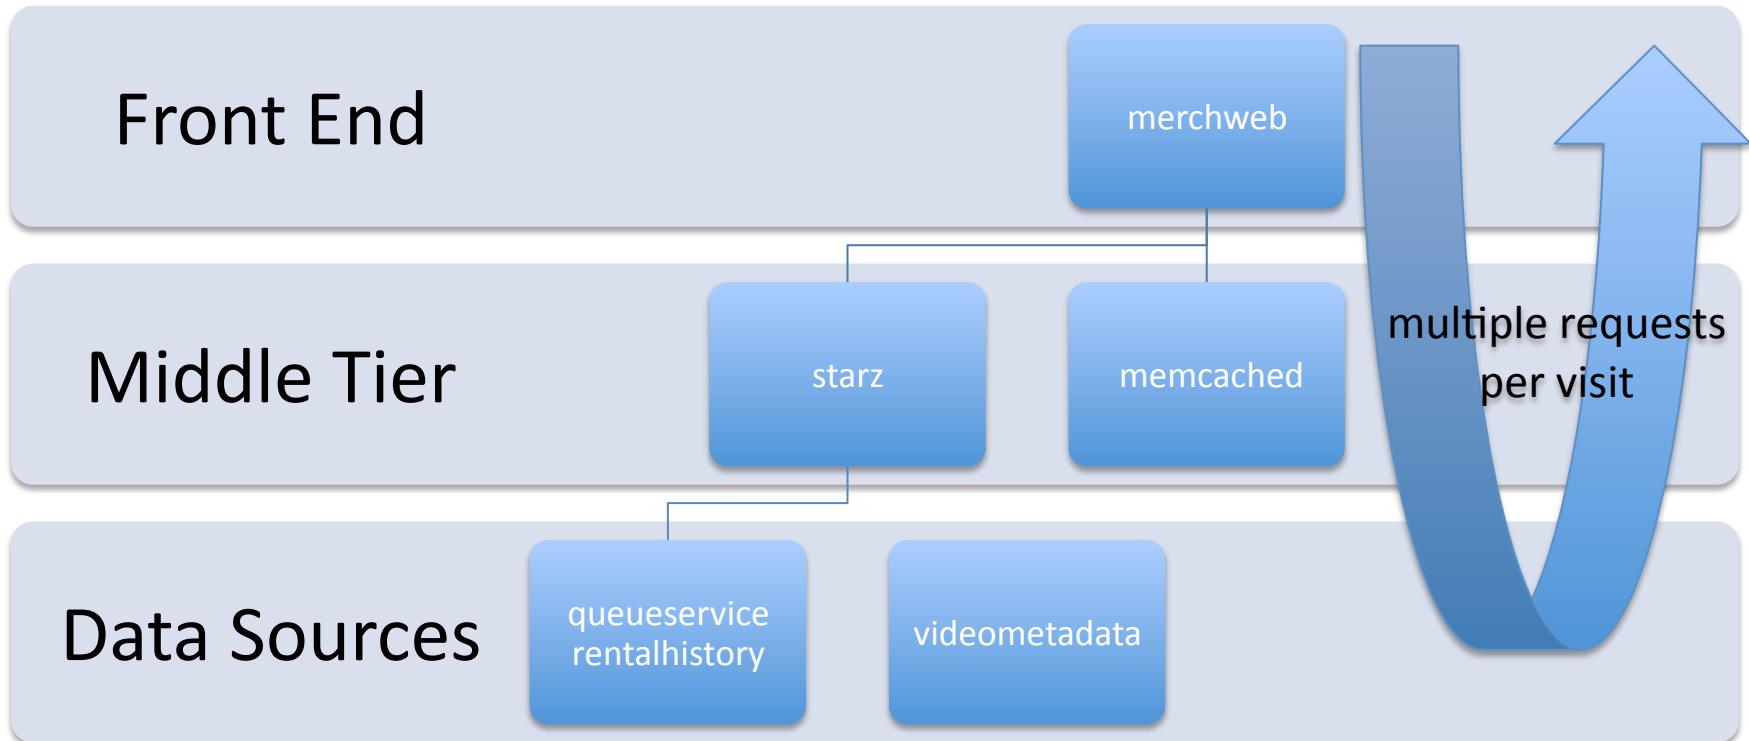


First Page

- First full page – Starz Channel Genre
 - Simplest page, no sub-genres, minimal personalization
 - Lots of investment in new Struts based page design
 - Uses identity cookie to lookup in member info svc
- New “merchweb” front end instance
 - movies.netflix.com points to merchweb instance
- Uncovered lots of latency issues
 - Used memcached to hide S3 and SimpleDB latency
 - Improved from slower to faster than Datacenter



Starz Page Cloud Instances



Controlled Cloud Transition

- WWW calling code chooses who goes to cloud
 - Filter out corner cases, send percentage of users
 - The URL that customers see is
<http://movies.netflix.com/WiContentPage?csid=1>
 - If problem, redirect to old Datacenter page
<http://www.netflix.com/WiContentPage?csid=1>
- Play Button and Star Rating Action redirect
 - Point URLs for actions that create/modify data back to datacenter to start with



Developers



Cloud Developer Setup

- Cloud Boot Camp
 - Room full of engineers sharing the pain for 1-2 days
 - Built a very rough prototype working web site
 - Get everyone hands-on in the cloud with a new code base
 - Debug lots of tooling and conceptual issues very fast
 - Member info in SimpleDB with developer's accounts only
- Cloud Specific Key Setup
 - It's a pain, need to configure your IDE's JVM
 - Needed to integrate with AWS security model
- Startup Guide Wiki Pages
 - What object facets already exist, how to make your own
 - What components already exist or are work in progress



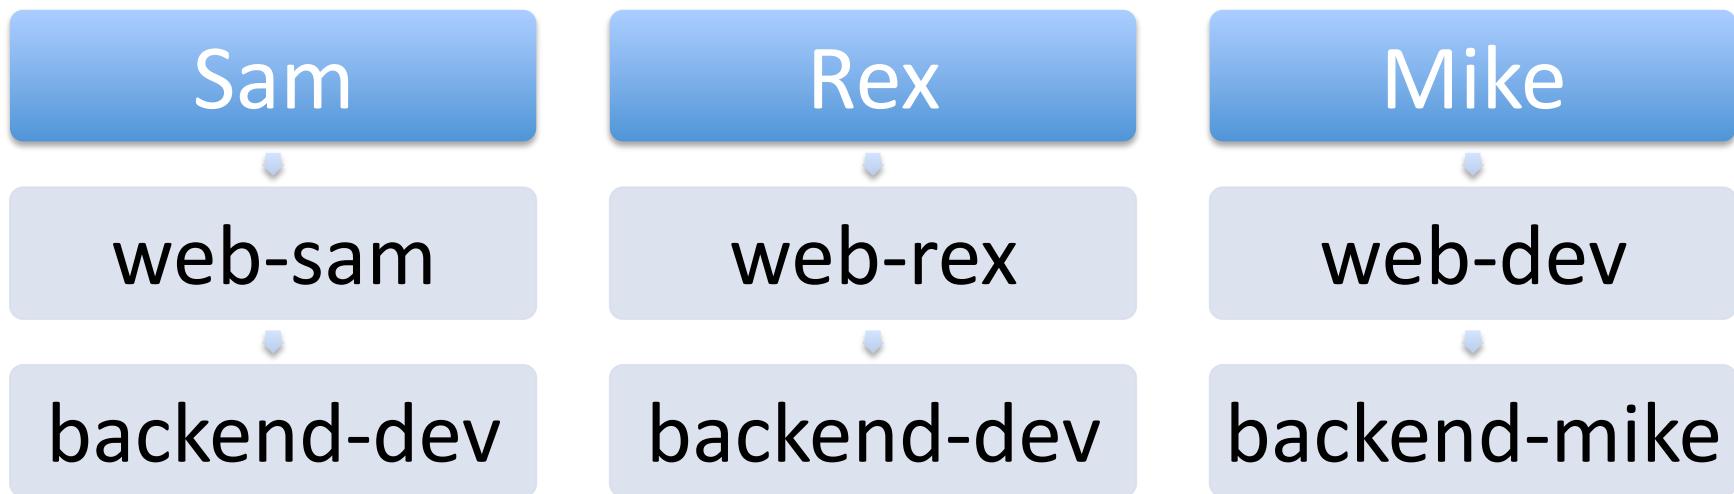
Developer Instances Collision

Sam and Rex both want to deploy web front end for development



Per-Service Namespace Routing

Developers choose what to share



Developer Service Namespaces

- Developer specific service instances
 - Configured via Java properties at runtime
 - Routing implemented by REST client library
- Server Configuration
 - Configure discovery service version string
 - Registers as <appname>-<namespace>
- Client Configuration
 - Route traffic on per-service basis including namespace



Current Status



WWW Page by Page during Q2/Q3/Q4

- Simplest possible page first
 - Minimal dependencies
- Add pages as dependent services come online
- Home page – most complex and highest traffic
- Leave low traffic pages for later cleanup

gradual migration from Datacenter pages



Big-Bang Transition

- iPhone Launch (August/Sept)
 - No capacity in the datacenter, cloud only
 - App Store gates release, not gradual, can't back out
 - Market is huge (existing and new customers)
 - Has to work at large scale on day one
- Datacenter Shadow Redirect Technique
 - Used to stress back-end and data sources
- SOASTA Cloud Based Load Generation
 - Used to stress test API and end-to-end functionality



Current Work for Cloud Platform

- Drive latency and availability goals
 - More Aggressive caching
 - Improving Fault and latency robustness
- Logging and monitoring portal/dashboards
 - Working to integrate tools and data sources
 - Need better observability and automation
- Evaluating a range of NoSQL choices
 - Broad set of use cases, no single winner
 - Good topic for another talk...

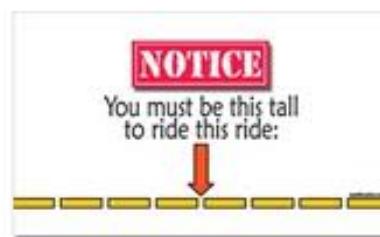


Wrap Up



Next Few Years...

- “System of Record” moves to Cloud
 - Master copies of data live only in the cloud, with backups etc.
 - Cut the datacenter to cloud replication link
- International Expansion – Global Clouds
 - Rapid deployments to new markets
- GPU Clouds optimized for video encoding
- Cloud Standardization
 - Cloud features and APIs should be a commodity not a differentiator
 - Differentiate on scale and quality of service
 - Competition also drives cost down
 - Higher resilience
 - Higher scalability



We would prefer to be an insignificant customer in a giant cloud



Remember the Goals

Faster

Scalable

Available

Productive

Track progress against these goals



Takeaway

Netflix is path-finding the use of public AWS cloud to replace in-house IT for non-trivial applications with hundreds of developers and thousands of systems.

<http://www.linkedin.com/in/adriancockcroft>

@adrianco #netflixcloud

