

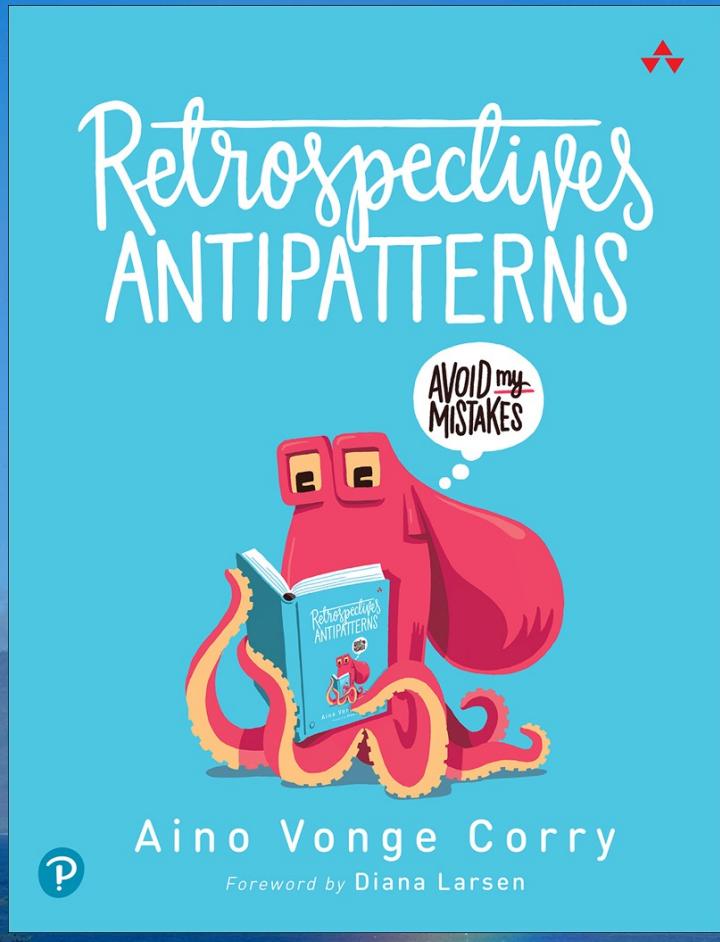
# **Microservices Retrospective - What we learned (and didn't learn) from Netflix**

A photograph of a coastal landscape. In the foreground, there's a dark, low-lying landmass or a small island. Beyond it is a wide expanse of deep blue ocean. In the background, a large, rugged mountain range rises, its peaks obscured by thick, white clouds. A vibrant rainbow arches across the sky, its colors reflected in the water below. The overall scene is bright and colorful, with the deep blues of the sea and sky contrasting with the white of the clouds and the greenish-blue of the rainbow.

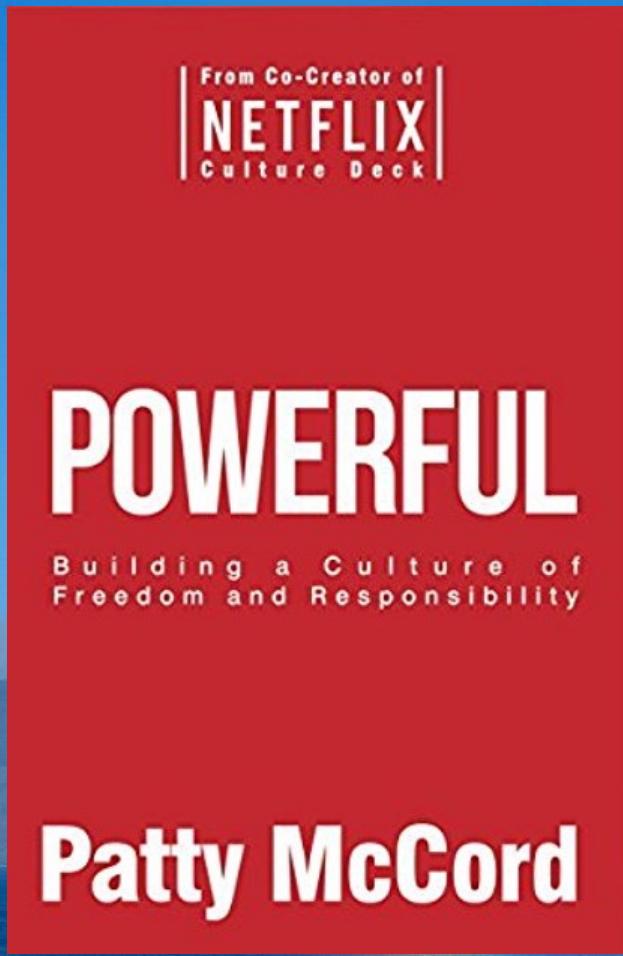
**Adrian Cockcroft | Qcon London | March 2023**

**@adrianco@mastodon.social**

All photos taken by Adrian, mostly in Maui



Retrospective?  
Time for me to  
learn some  
Agile rituals...



Netflix was  
extremely agile,  
without the rituals  
of Extreme or Agile

## **Story Points**

Retrospectives and Netflix Culture

Netflix in the Cloud – QconSF 2010, Cassandra Summit 2011

Cloud Native Applications 2013

Microservices – Dockercon and MicroXchg Berlin 2016

Why don't microservices work for some people?

## **Backlog**

Systems for Innovation 2015



## Culture

Seven aspects  
of Netflix culture

1. Values are what we value
2. High performance
3. Freedom and responsibility
4. Context, not control
5. Highly aligned, loosely coupled
6. Pay top of market
7. Promotions and development

# NETFLIX

## Systems Thinking

Optimized for agility, minimized processes  
Ability to evolve rapidly  
Pioneers, happy to be first to do something  
Comfortable with ambiguity  
Creativity and self discipline combined  
Flexibility is more important than efficiency  
Steer any pain to where it can be helpful

# Netflix in the Cloud

Nov 6, 2010

Adrian Cockcroft

@adrianco #netflixcloud

<http://www.linkedin.com/in/adriancockcroft>



*“The cloud lets its users focus on delivering differentiating business value instead of wasting valuable resources on the **undifferentiated heavy lifting** that makes up most of IT infrastructure.”*



Werner Vogels  
Amazon CTO



# What, Why and How?

The details...

NETFLIX



# Goals

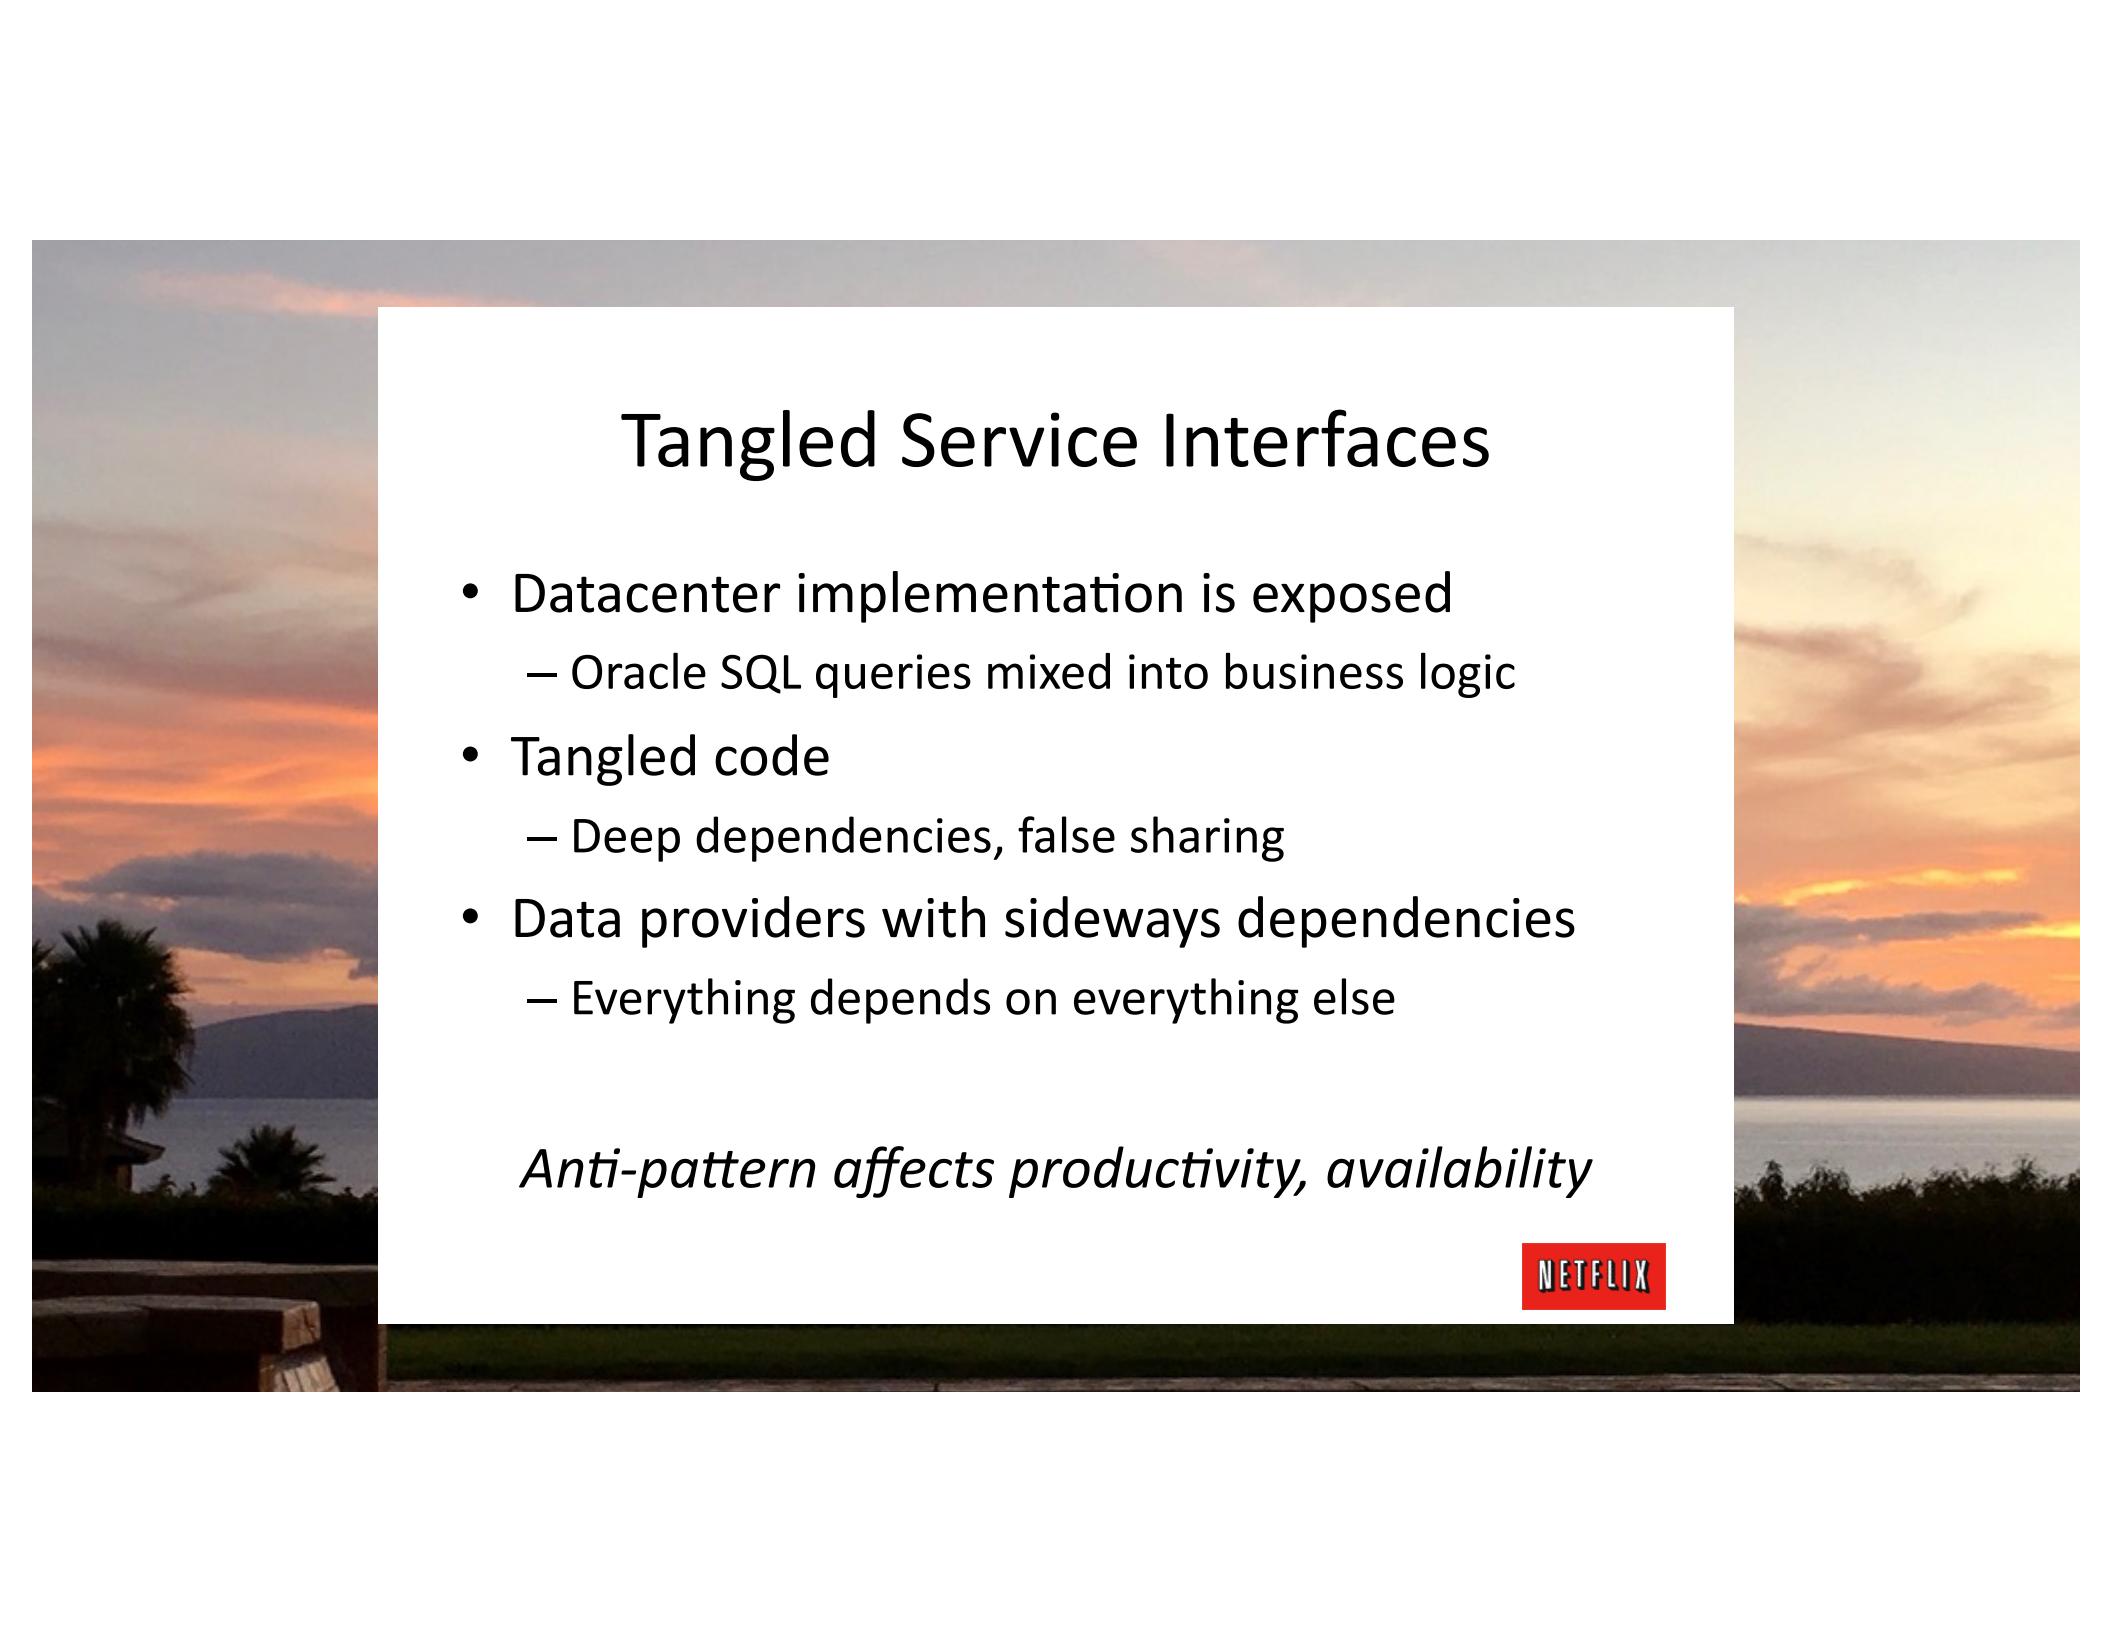
- Faster
  - Lower latency than the equivalent datacenter web pages and API calls
  - Measured as mean and 99<sup>th</sup> percentile
  - For both first hit (e.g. home page) and in-session hits for the same user
- Scalable
  - Avoid needing any more datacenter capacity as subscriber count increases
  - No central vertically scaled databases
  - Leverage AWS elastic capacity effectively
- Available
  - Substantially higher robustness and availability than datacenter services
  - Leverage multiple AWS availability zones
  - No scheduled down time, no central database schema to change
- Productive
  - Optimize agility of a large development team with automation and tools
  - Leave behind complex tangled datacenter code base (~8 year old architecture)
  - Enforce clean layered interfaces and re-usable components



# Old Datacenter vs. New Cloud Arch

- |                            |                                  |
|----------------------------|----------------------------------|
| Central SQL Database       | Distributed Key/Value NoSQL      |
| Sticky In-Memory Session   | Shared Memcached Session         |
| Chatty Protocols           | Latency Tolerant Protocols       |
| Tangled Service Interfaces | Layered Service Interfaces       |
| Instrumented Code          | Instrumented Service Patterns    |
| Fat Complex Objects        | Lightweight Serializable Objects |
| Components as Jar Files    | Components as Services           |





# Tangled Service Interfaces

- Datacenter implementation is exposed
  - Oracle SQL queries mixed into business logic
- Tangled code
  - Deep dependencies, false sharing
- Data providers with sideways dependencies
  - Everything depends on everything else

*Anti-pattern affects productivity, availability*

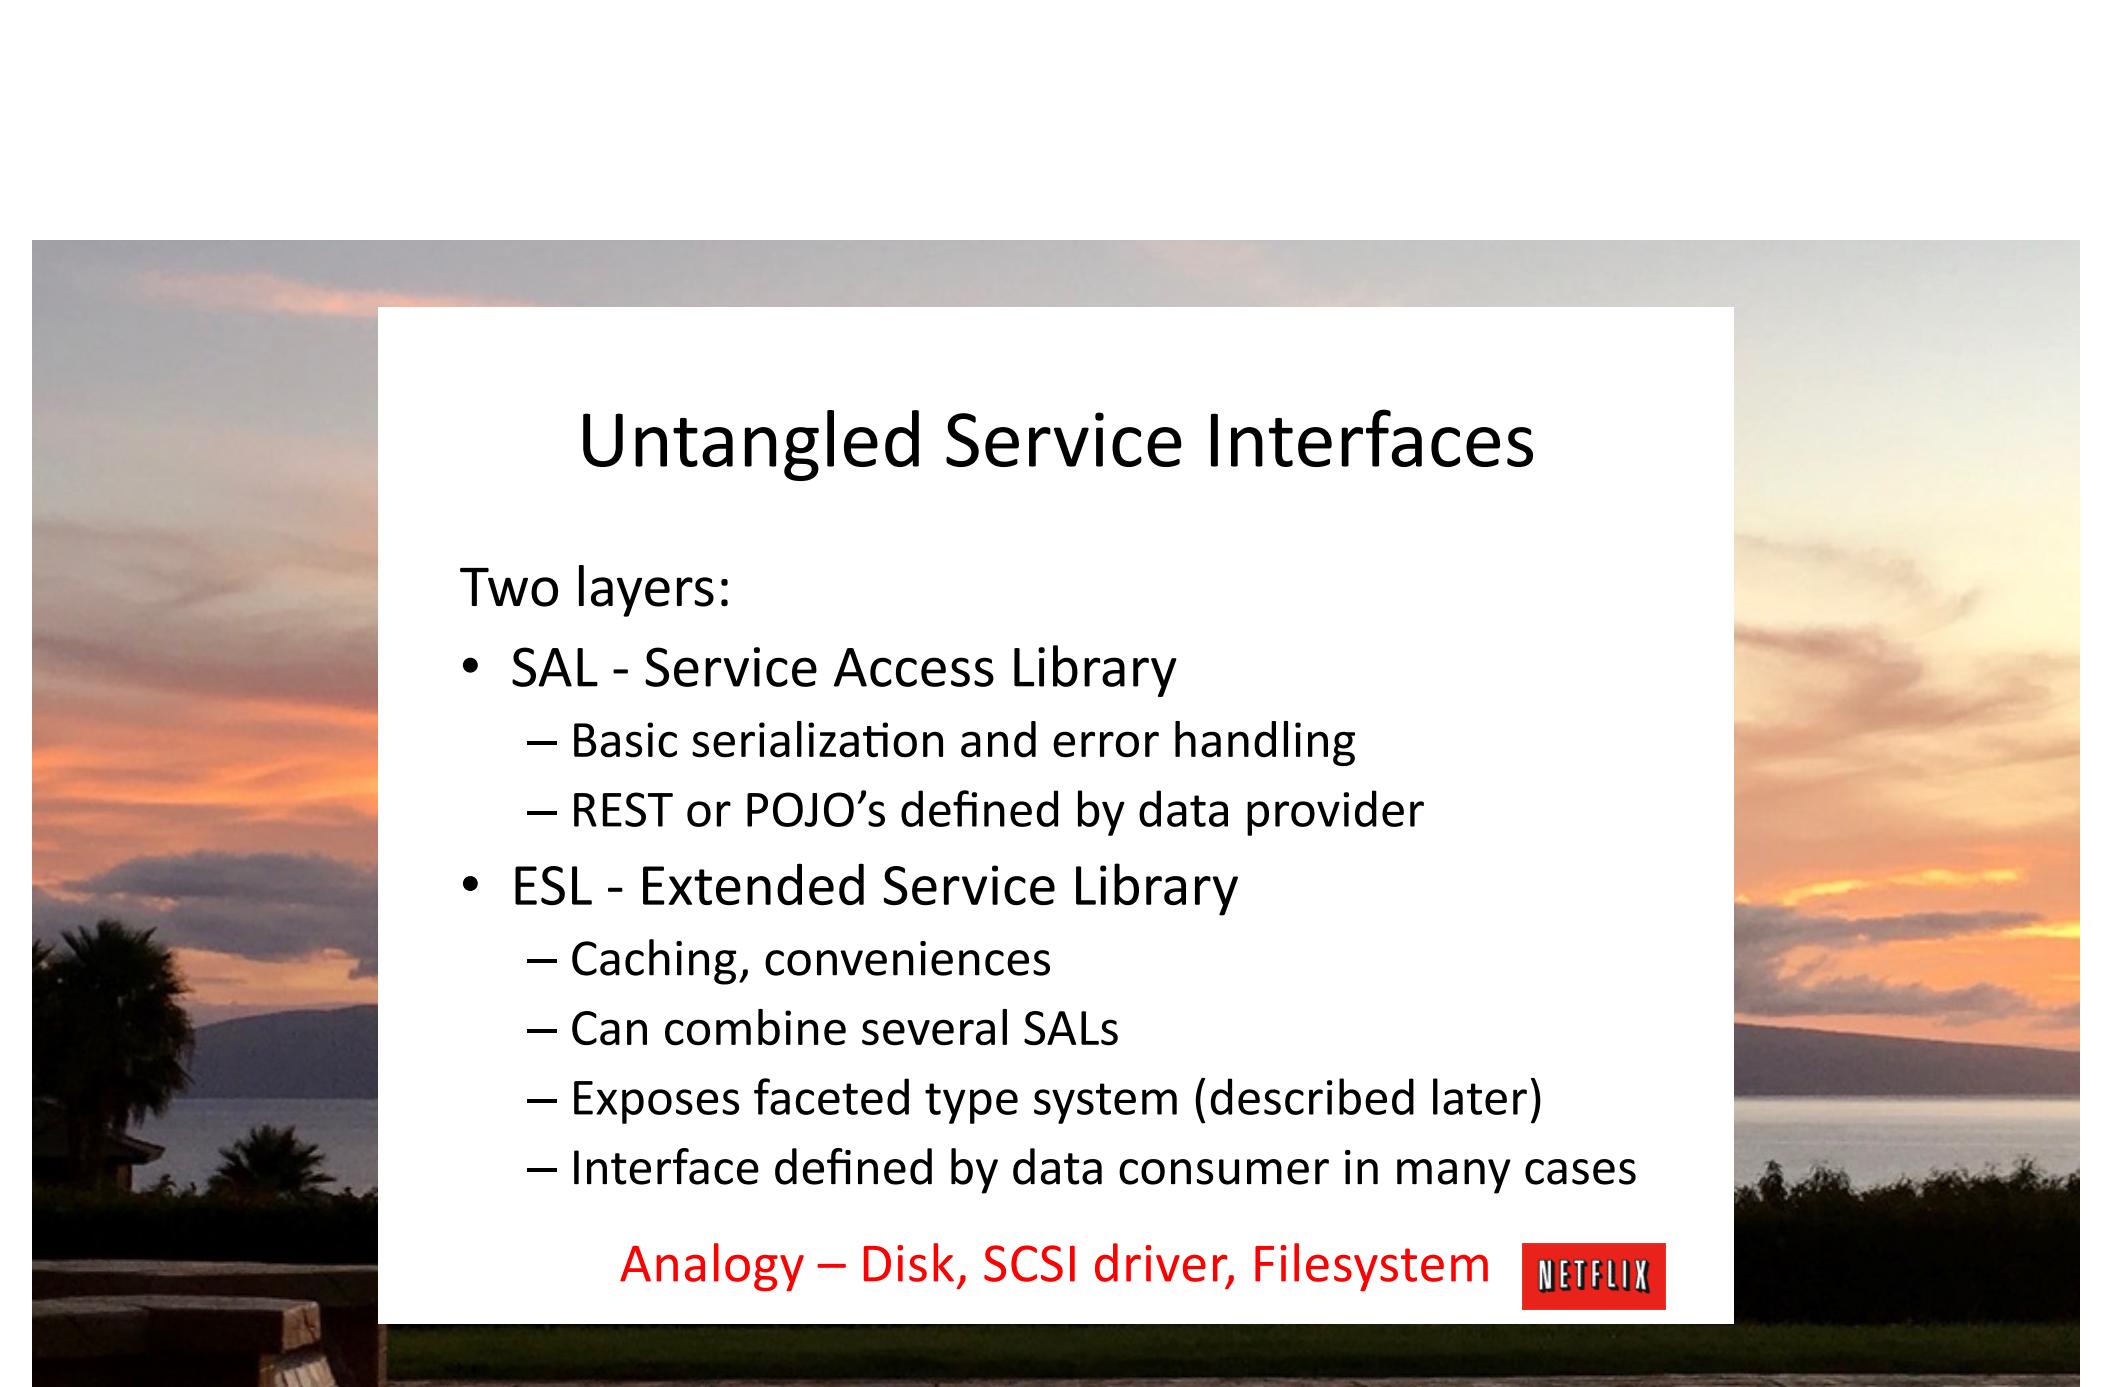


# Untangled Service Interfaces

- New Cloud Code With Strict Layering
  - Compile against interface jar
  - Can use spring runtime binding to enforce
- Service interface **is** the service
  - Implementation is completely hidden
  - Can be implemented locally or remotely
  - Implementation can evolve independently

This lesson wasn't learned – the wire protocol shouldn't be the interface between microservices





# Untangled Service Interfaces

Two layers:

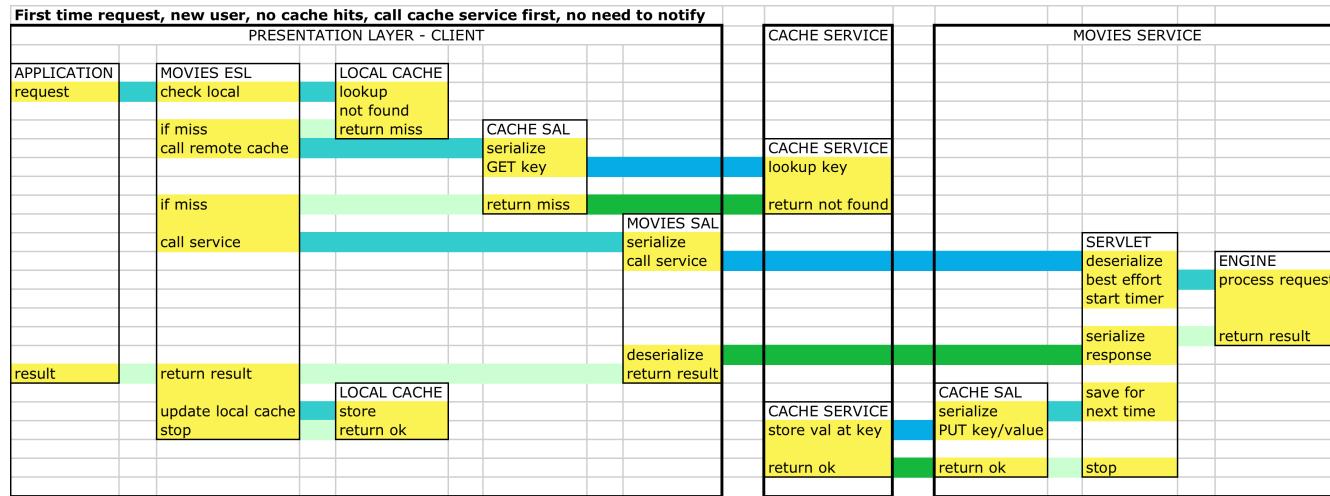
- SAL - Service Access Library
  - Basic serialization and error handling
  - REST or POJO's defined by data provider
- ESL - Extended Service Library
  - Caching, conveniences
  - Can combine several SALs
  - Exposes faceted type system (described later)
  - Interface defined by data consumer in many cases

Analogy – Disk, SCSI driver, Filesystem

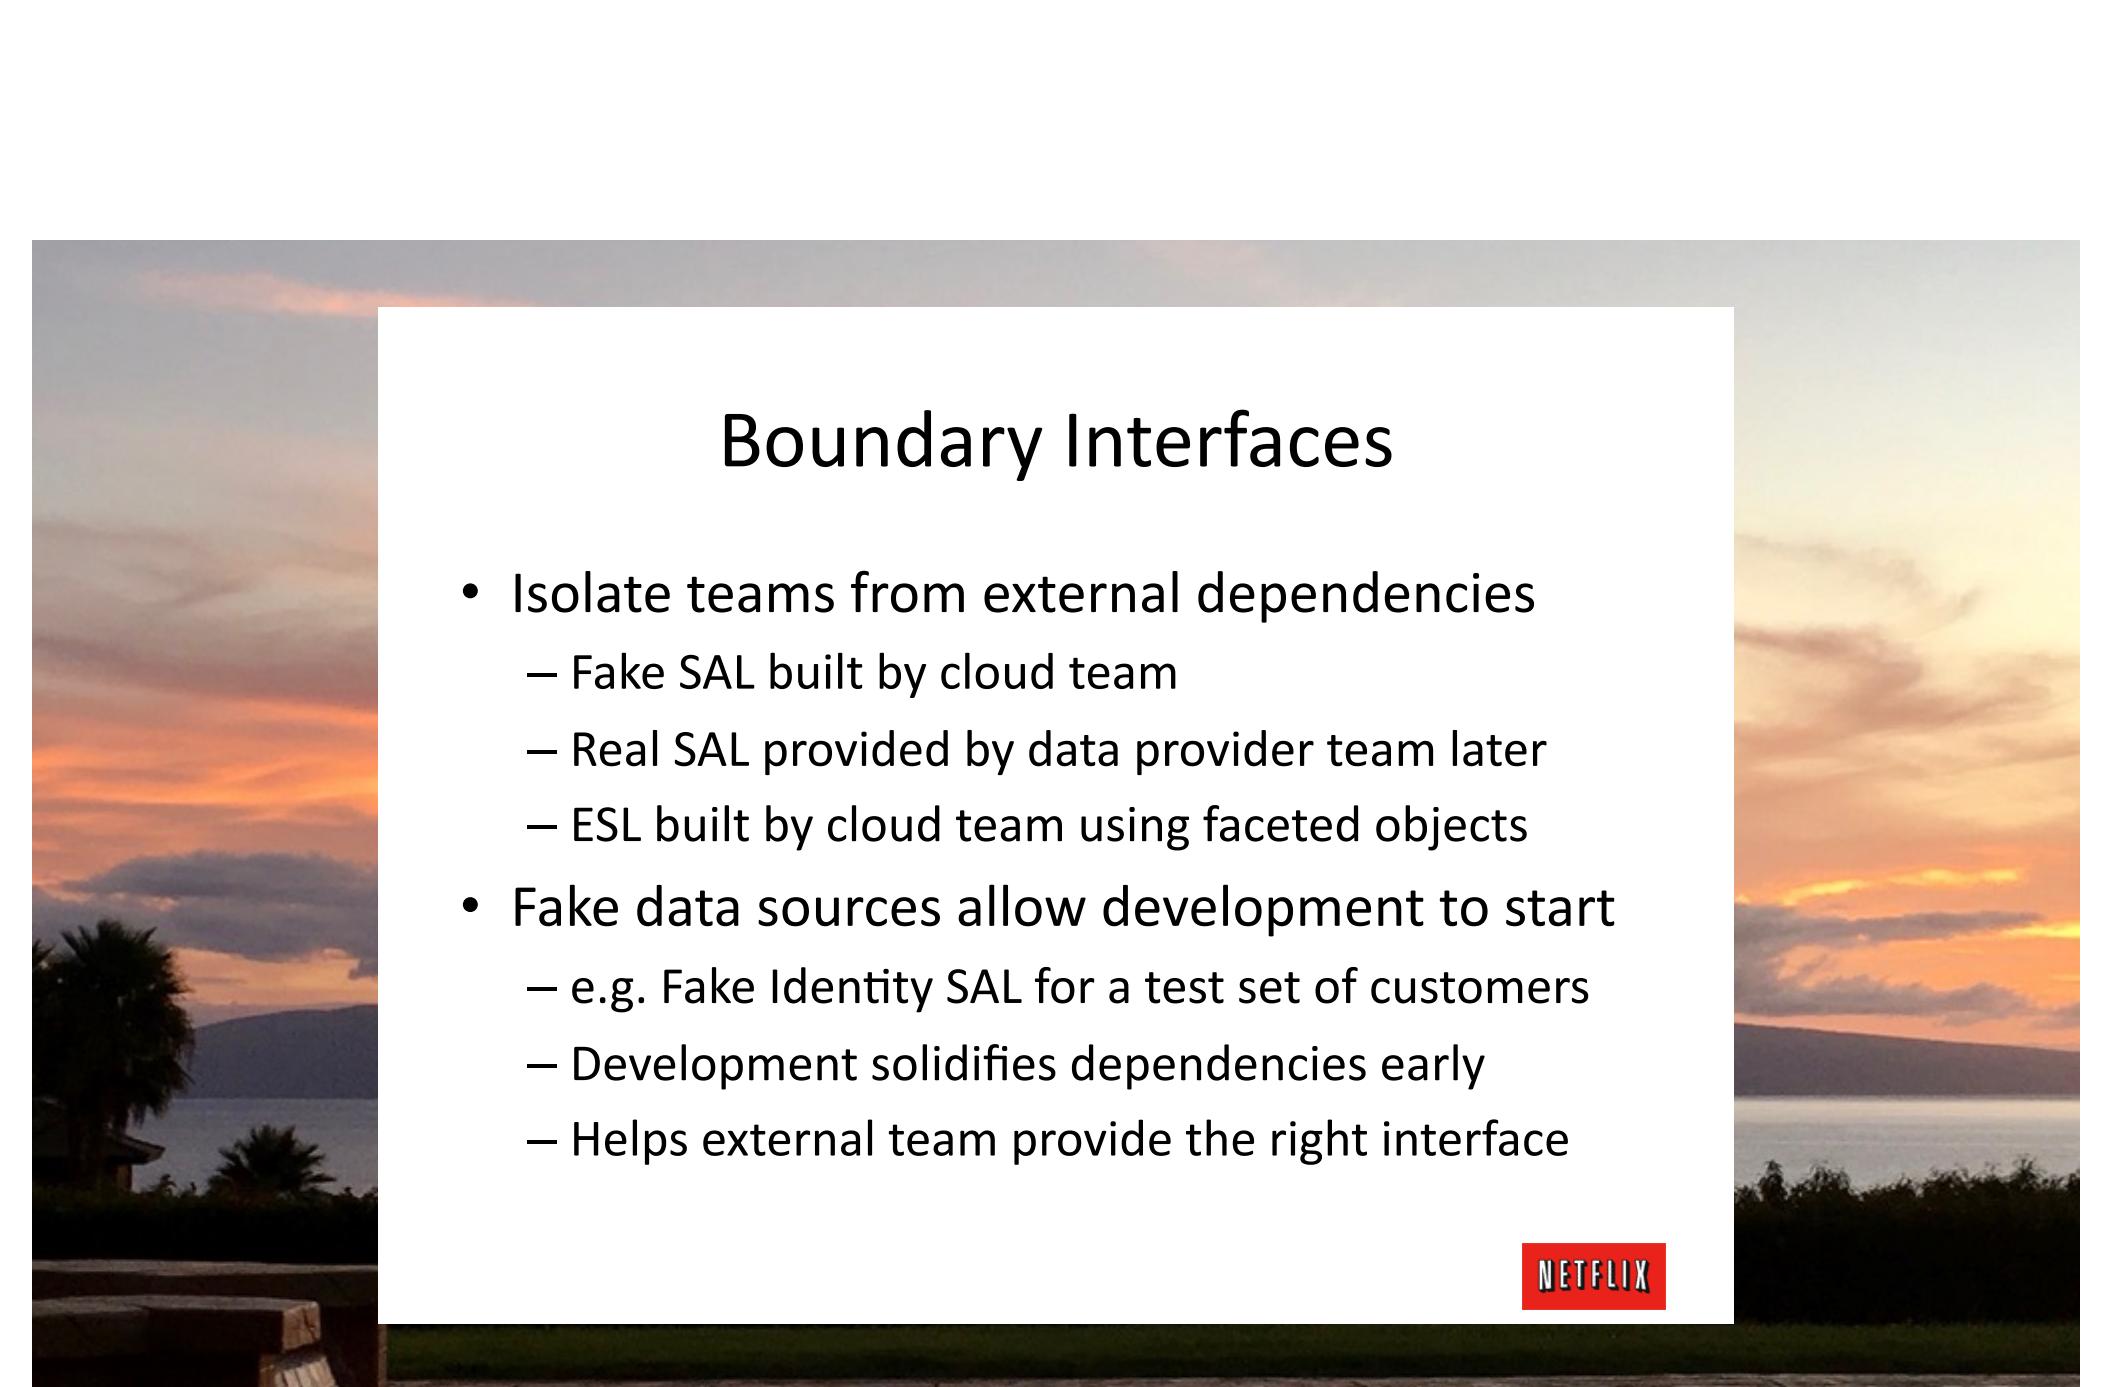


# Service Interaction Pattern

## Sample Swimlane Diagram



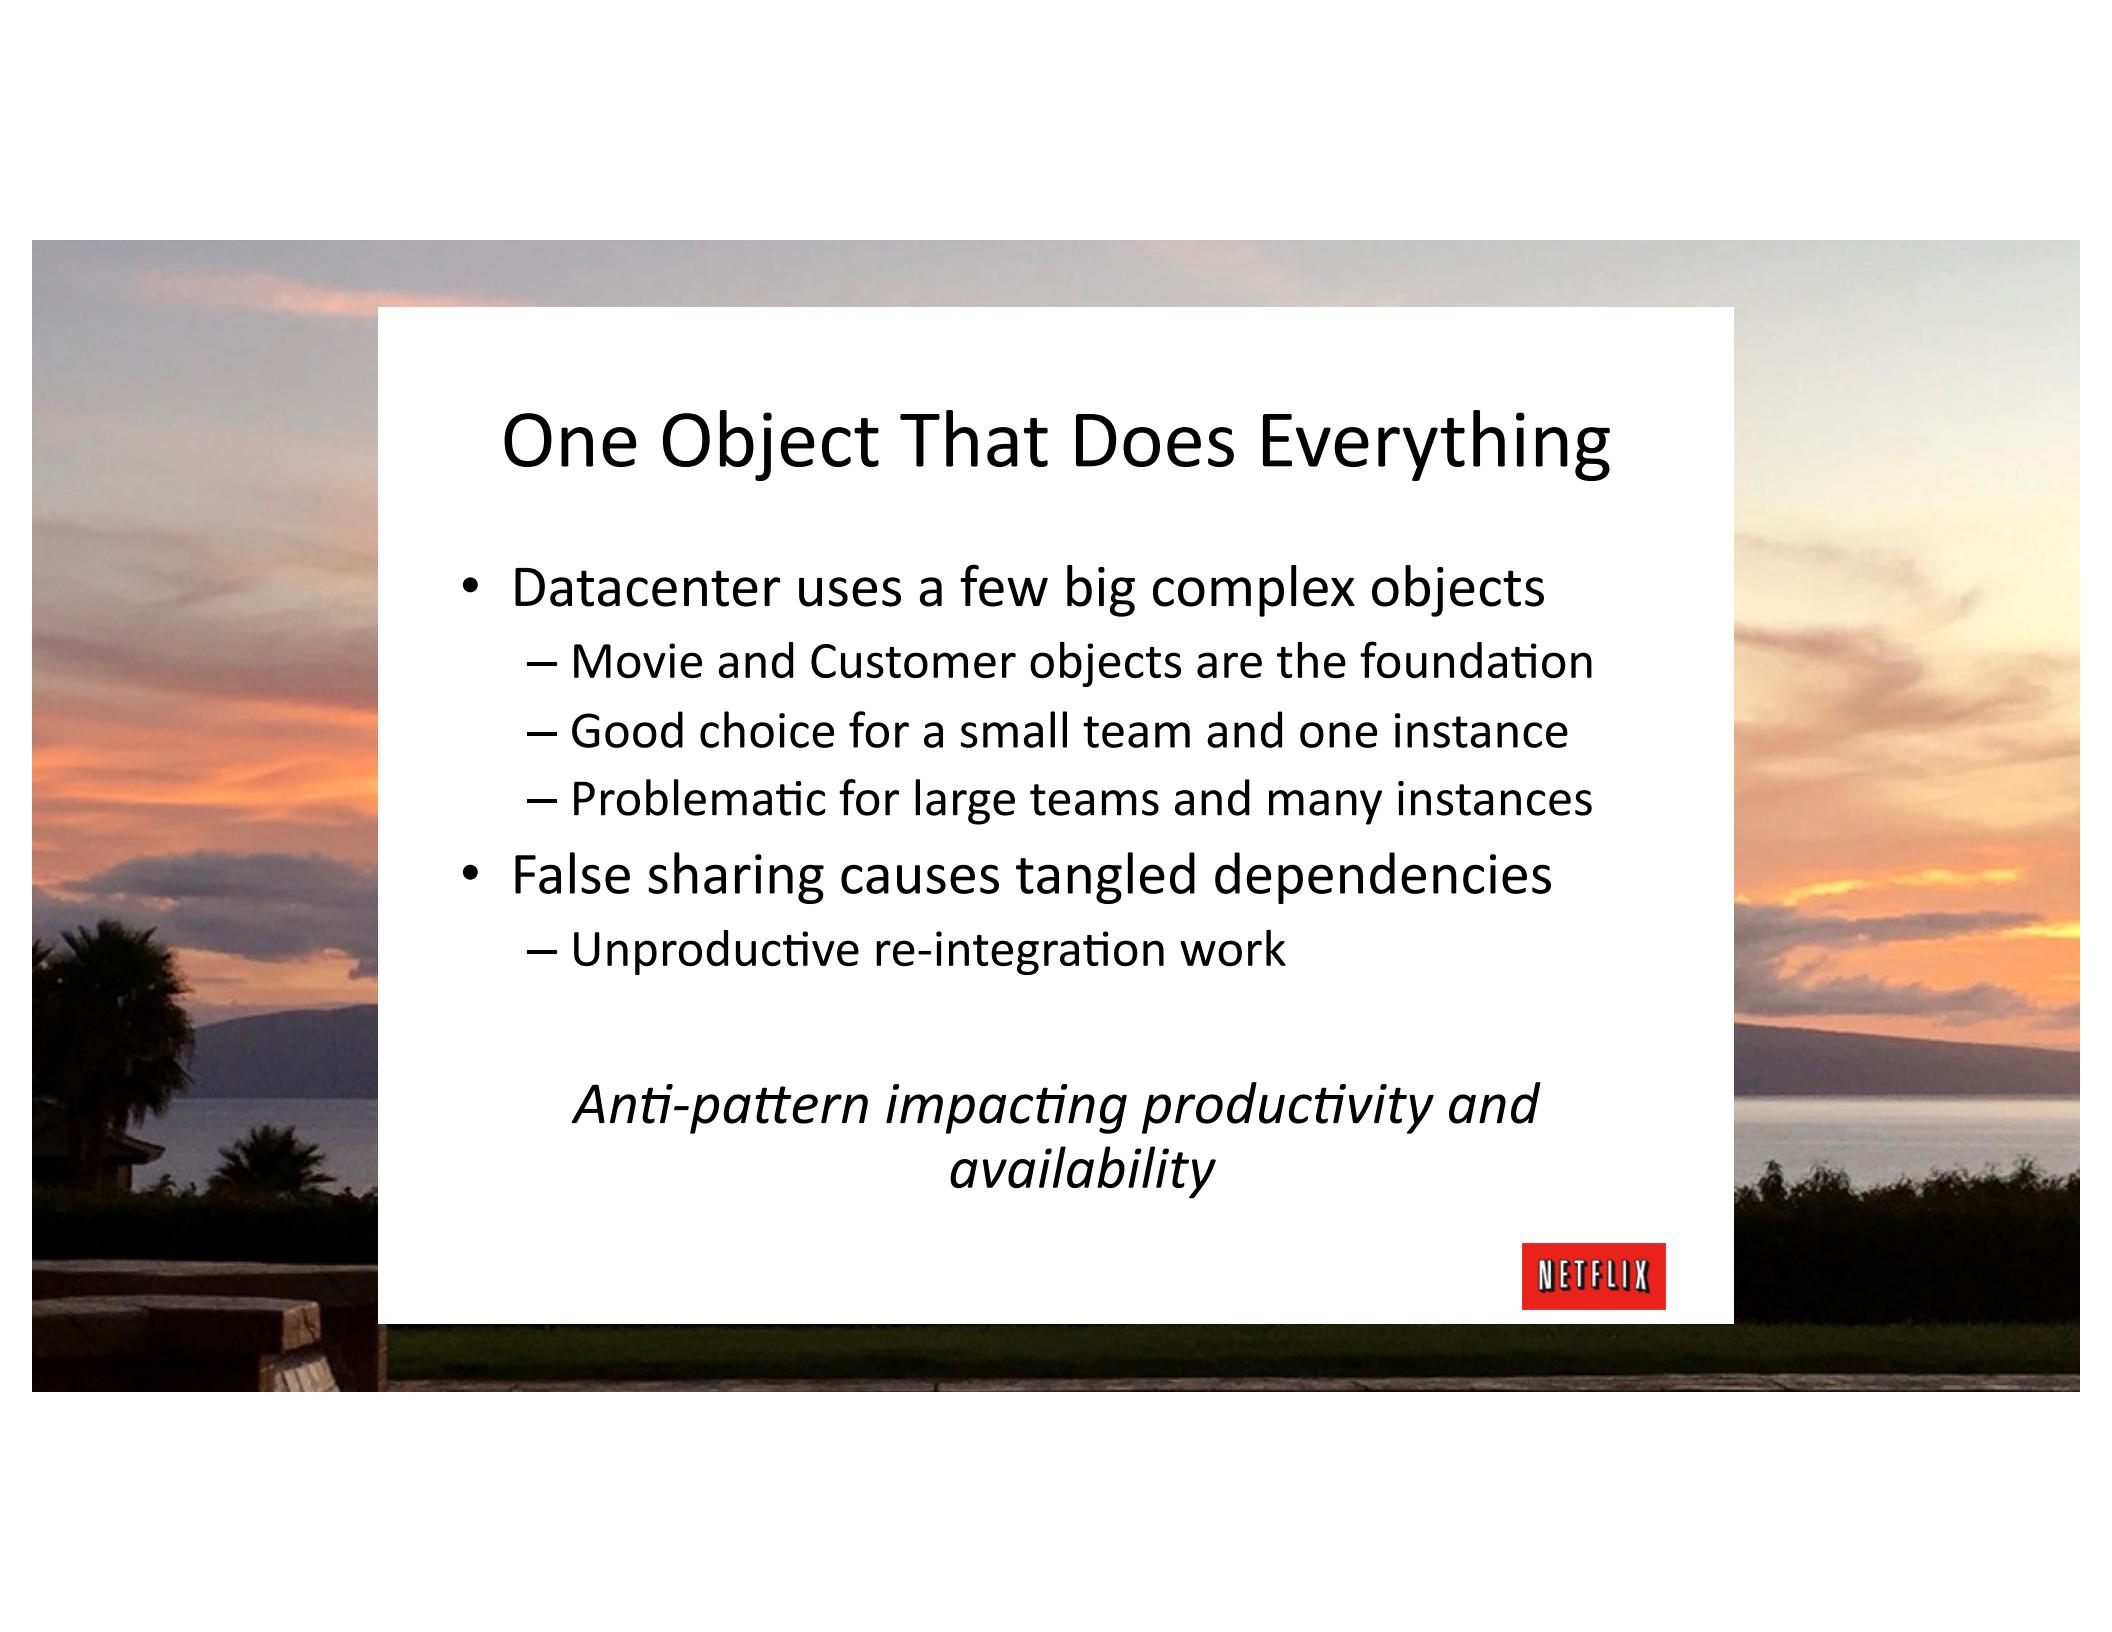
NETFLIX



# Boundary Interfaces

- Isolate teams from external dependencies
  - Fake SAL built by cloud team
  - Real SAL provided by data provider team later
  - ESL built by cloud team using faceted objects
- Fake data sources allow development to start
  - e.g. Fake Identity SAL for a test set of customers
  - Development solidifies dependencies early
  - Helps external team provide the right interface





# One Object That Does Everything

- Datacenter uses a few big complex objects
  - Movie and Customer objects are the foundation
  - Good choice for a small team and one instance
  - Problematic for large teams and many instances
- False sharing causes tangled dependencies
  - Unproductive re-integration work

*Anti-pattern impacting productivity and availability*



# An Interface For Each Component

- Cloud uses faceted Video and Visitor
  - Basic types hold only the identifier
  - Facets scope the interface you actually need
  - Each component can define its own facets
- No false-sharing and dependency chains
  - Type manager converts between facets as needed
  - `video.asA(PresentationVideo)` for www
  - `video.asA(MerchableVideo)` for middle tier



The background image shows a stunning sunset or sunrise over a landscape. The sky is filled with wispy clouds colored in shades of orange, yellow, and red. In the foreground, the dark silhouette of a hillside is visible, with several small figures of people standing on top, appearing as dark shapes against the bright sky. A thin crescent moon is visible in the upper left corner.

**Response to 2010 talk was a mixture of  
incomprehension and confusion. Most  
people thought we were crazy and would be  
back in our datacenters when it failed...**

A wide-angle photograph of a sunset or sunrise. The sky is filled with dynamic, wispy clouds that transition from deep orange and yellow near the horizon to darker blues and purples higher up. In the dark foreground, the silhouettes of hills and a few small structures are visible. A single crescent moon is positioned in the upper left corner of the frame.

Less than a year later...

# **Replacing Datacenter Oracle with Global Apache Cassandra on AWS**

July 11, 2011

Adrian Cockcroft

@adrianco #netflixcloud

<http://www.linkedin.com/in/adriancockcroft>





Get stuck with wrong config

Wait Wait File tickets

Ask permission Wait Wait

Wait Things We Don't Do Wait

Run out of space/power

Plan capacity in advance

Have meetings with IT Wait





Netflix could not  
build new  
datacenters fast  
enough

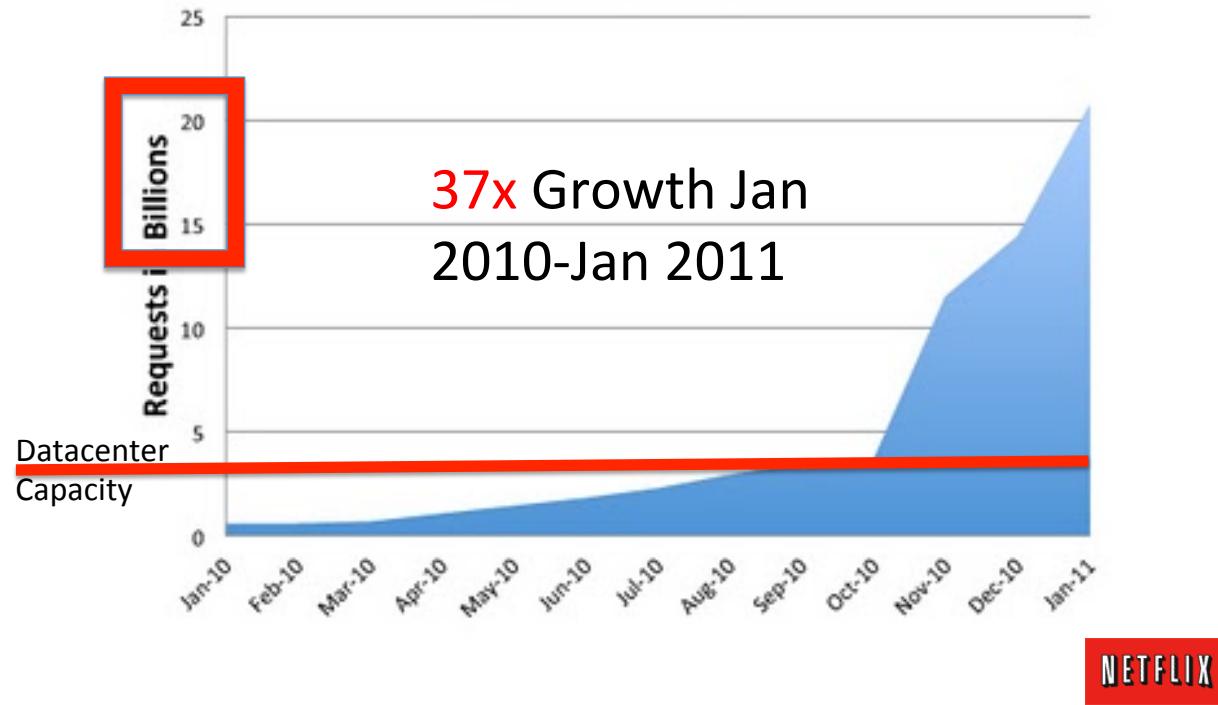
Capacity growth is accelerating, unpredictable  
Product launch spikes - iPhone, Wii, PS3, XBox



# Out-Growing Data Center

<http://techblog.netflix.com/2011/02/redesigning-netflix-api.html>

**Netflix API : Growth in Requests**



# High Availability

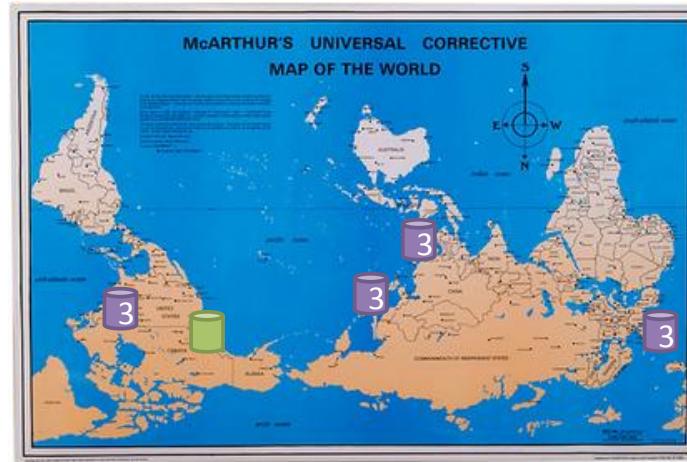
- Cassandra stores 3 local copies, 1 per zone
  - Synchronous access, durable, highly available
  - Read/Write One fastest, least consistent - ~1ms
  - Read/Write Quorum 2 of 3, consistent - ~3ms
- AWS Availability Zones
  - Separate buildings
  - Separate power etc.
  - Close together

Actually 10-100km apart



# Remote Copies

- Cassandra duplicates across AWS regions
  - Asynchronous write, replicates at destination
  - Doesn't directly affect local read/write latency
- Global Coverage
  - Business agility
  - Follow AWS...
- Local Access
  - Better latency
  - Fault Isolation



NETFLIX

# Chaos Monkey



- Make sure systems are resilient
  - Allow any instance to fail without customer impact
- Chaos Monkey hours
  - Monday-Thursday 9am-3pm random instance kill
- Application configuration option
  - Apps now have to opt-out from Chaos Monkey
- Computers (Datacenter or AWS) randomly die
  - Fact of life, but too infrequent to test resiliency

Architecture design control

Be sure you can auto-scale down!

NETFLIX

A wide-angle photograph of a tropical sunset or sunrise. The sky is filled with dramatic, wispy clouds colored in shades of orange, yellow, and pink. In the foreground, silhouettes of palm trees are visible against the bright horizon. The overall atmosphere is serene and majestic.

**Response to 2011 progress was that  
Netflix was a Unicorn, and while it  
might work for us, it wasn't relevant to  
others**

A wide-angle photograph of a tropical sunset or sunrise. The sky is filled with dramatic, colorful clouds in shades of orange, yellow, pink, and purple. In the foreground, silhouettes of palm trees are visible against the bright horizon. The overall atmosphere is serene and beautiful.

**18 Months later in 2013**  
**Early use of Cloud Native and**  
**Microservices terminology**

# Cloud Native Applications: What Changed?

February 2013

Adrian Cockcroft

@adrianco #netflixcloud @NetflixOSS

<http://www.linkedin.com/in/adriancockcroft>



---

## Cloud Native?

---

## Agile Infrastructure

---

## NetflixOSS – Cloud Native On-Ramp

August 25<sup>th</sup>, 2006

# AWS EC2 Launched

Infrastructure as a Service

# Forklift Old Code to the Cloud

Faster than datacenter deployments but fragile

# New Anti-Fragile Patterns

Micro-services

Chaos engines

HA systems composed from ephemeral components

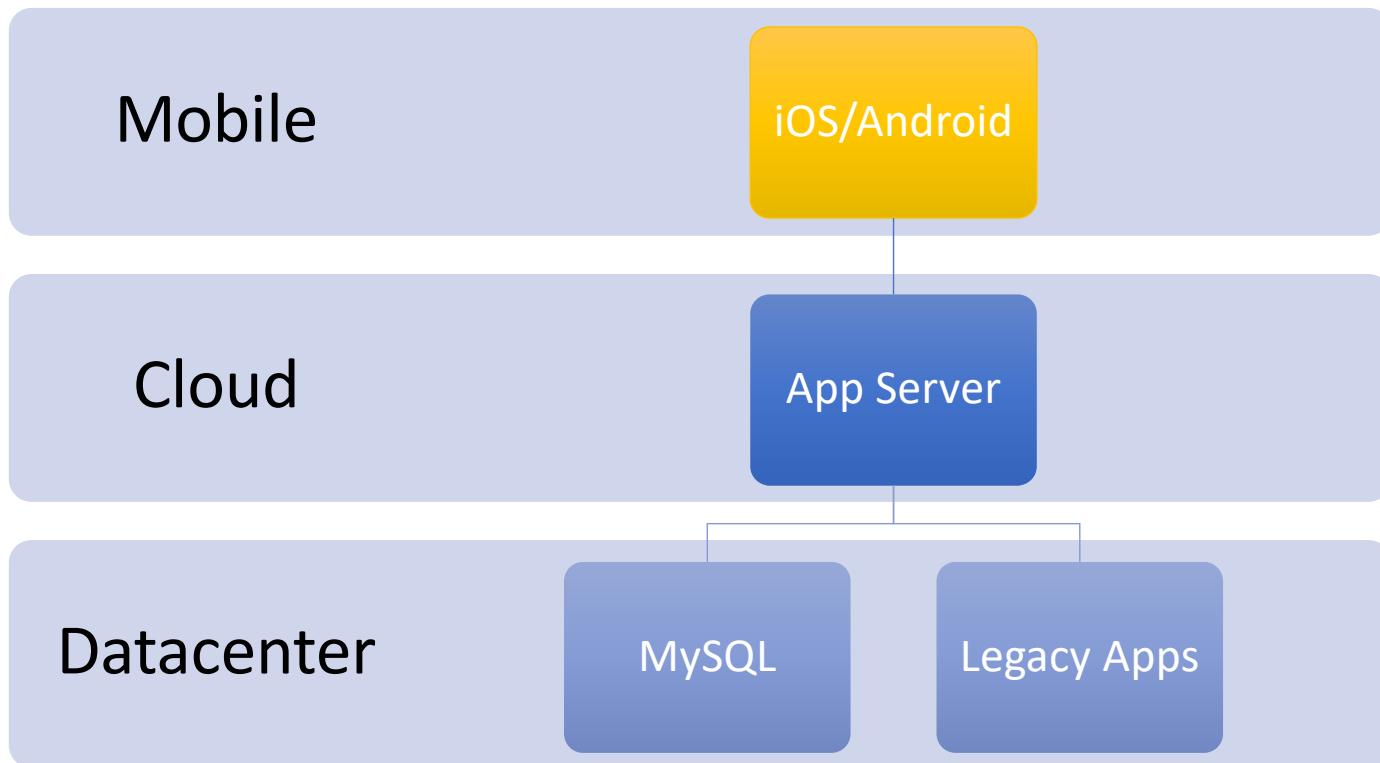
# Open Source Platforms

Techniques spread rapidly

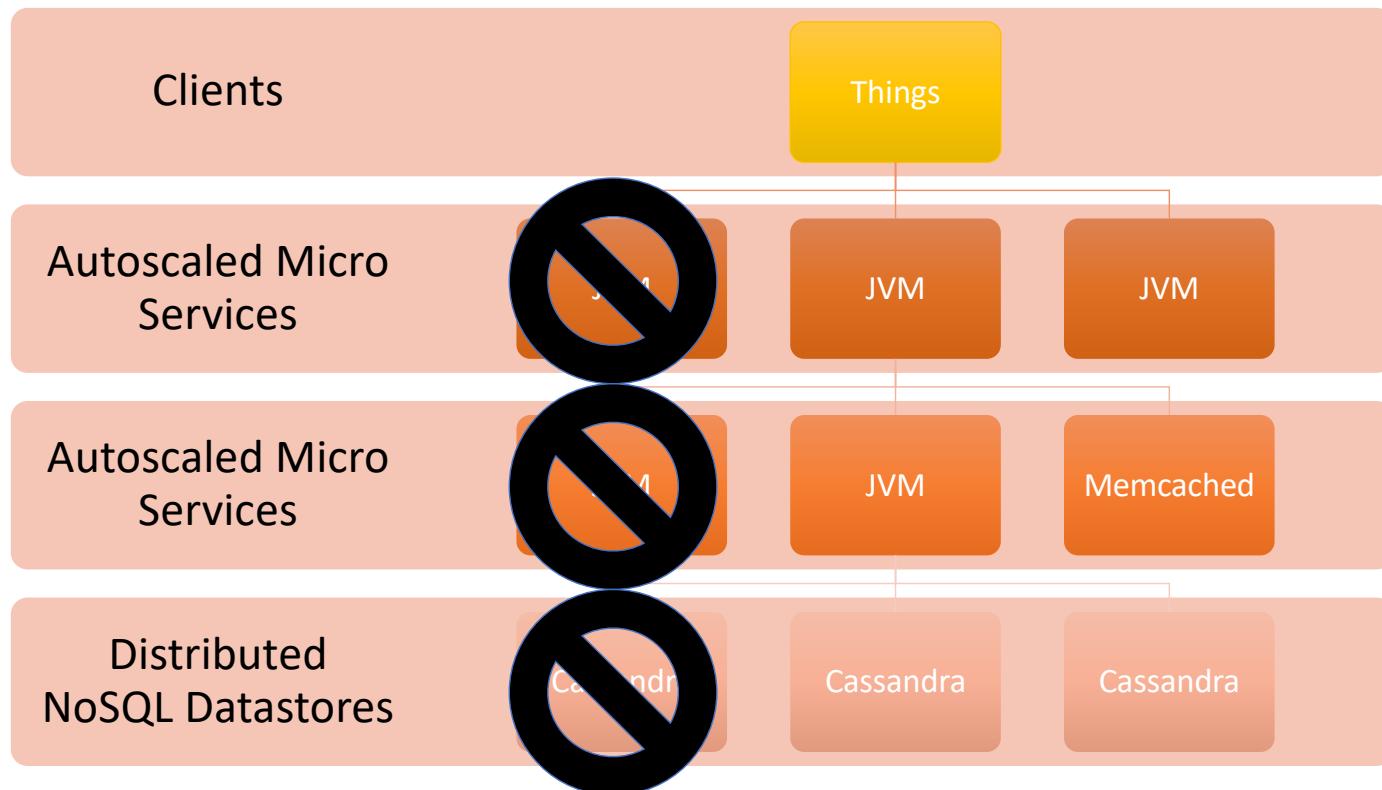
# Totally new development patterns and tools

(Slow enterprise adoption)

# Hybrid Cloud Architecture



# Cloud Native Architecture



# Cloud Native

Master copies of data are cloud resident

Everything is dynamically provisioned

All services are ephemeral

Cloud native meant that the application code was optimized to only run in the cloud

The CNCF later co-opted the term to mean code for running clouds...

# Availability

Is it running yet?

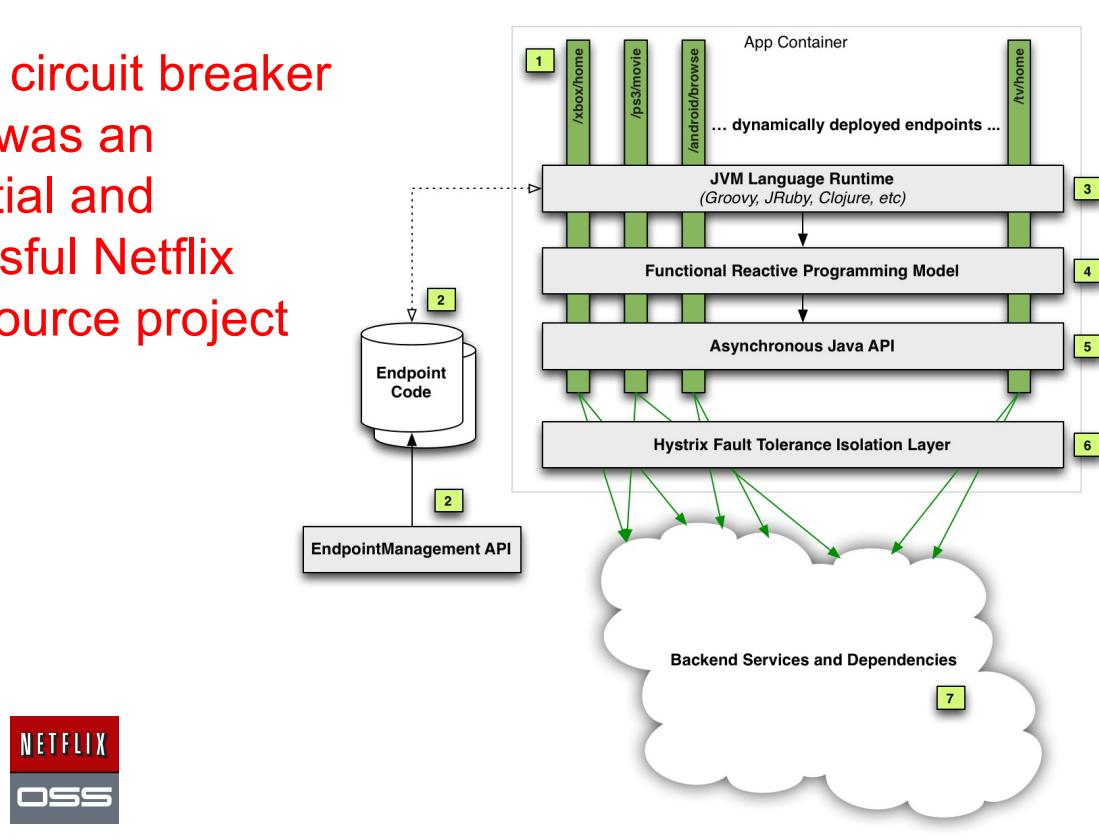
How many places is it running in?

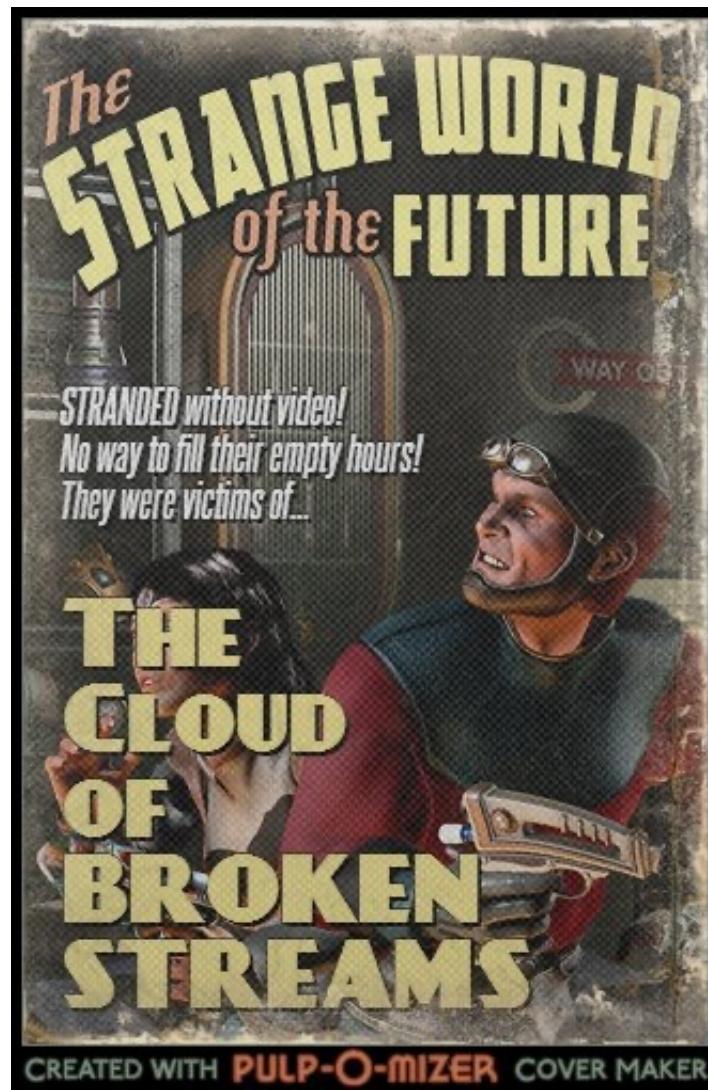
How far apart are those places?

# Antifragile Development Patterns

Functional Reactive with Circuit Breakers and Bulkheads

Hystrix circuit breaker library was an influential and successful Netflix open source project



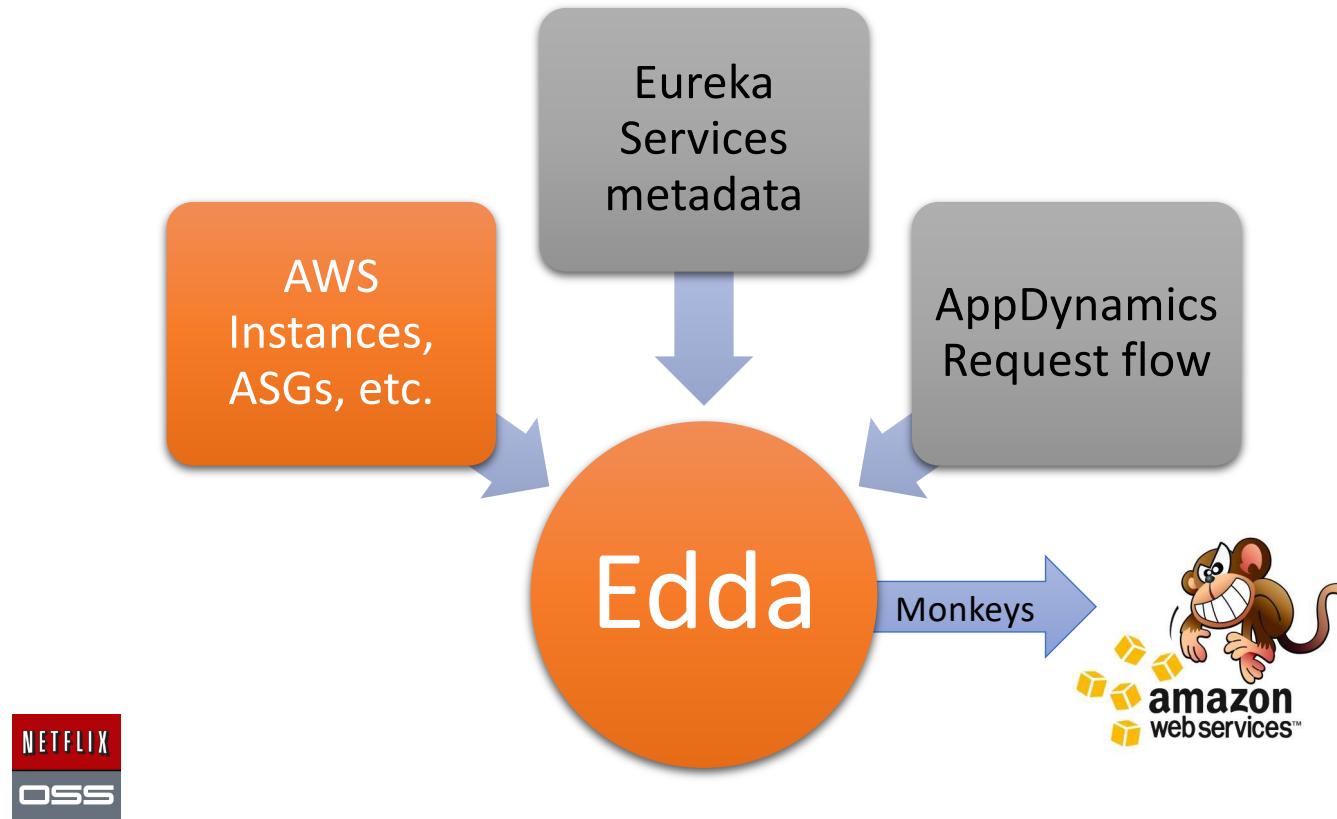


# Configuration Management

Datacenter CMDB's woeful  
Cloud is a model driven architecture  
Dependably complete

# Edda – Configuration History

<http://techblog.netflix.com/2012/11/edda-learn-stories-of-your-cloud.html>



# Edda Query Examples

Find any instances that have ever had a specific public IP address

```
$ curl "http://edda/api/v2/view/instances;publicIpAddress=1.2.3.4;_since=0"  
["i-0123456789","i-012345678a","i-012345678b"]
```

Show the most recent change to a security group

```
$ curl "http://edda/api/v2/aws/securityGroups/sg-0123456789;_diff;_all;_limit=2"  
--- /api/v2/aws.securityGroups/sg-0123456789;_pp;_at=1351040779810  
+++ /api/v2/aws.securityGroups/sg-0123456789;_pp;_at=1351044093504  
@@ -1,33 +1,33 @@  
{  
...  
    "ipRanges" : [  
        "10.10.1.1/32",  
        "10.10.1.2/32",  
+       "10.10.1.3/32",  
-       "10.10.1.4/32"  
...  
}
```

Not sure if this lesson was learned... but the ability to query the entire history of your infrastructure is very powerful

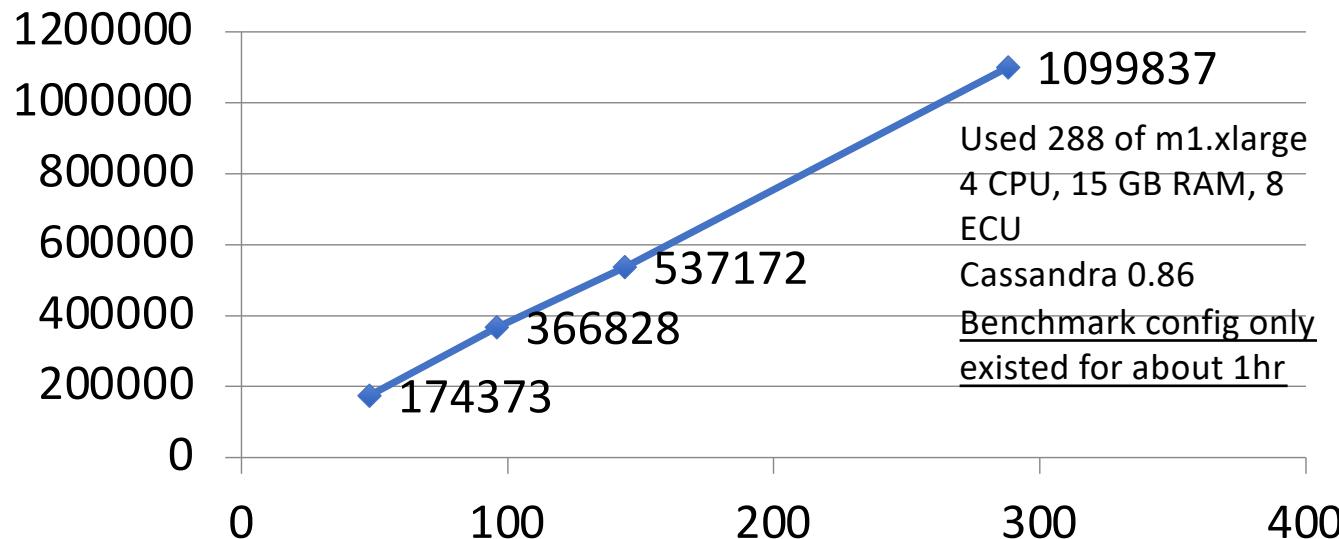


Frictionless infrastructure experiments!

## Scalability from 48 to 288 nodes on AWS

<http://techblog.netflix.com/2011/11/benchmarking-cassandra-scalability-on.html>

### Client Writes/s by node count – Replication Factor = 3



This benchmark was used by Datastax in a “Million Writes Per Second” ad campaign



A Cloud Native Open Source Platform

Open source has become a common way to build  
a technology brand and influence the industry

# Three Questions

Why is Netflix doing this?

How does it all fit together?

What is coming next?

# Platform Evolution

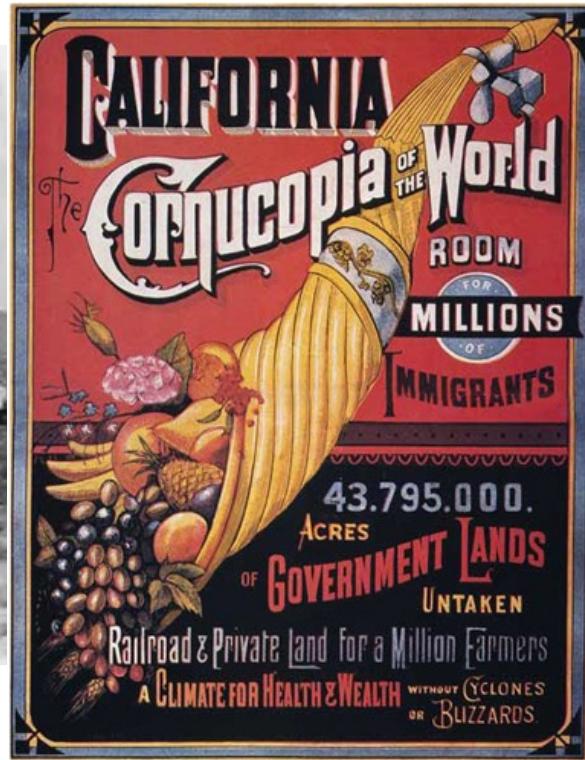


Netflix ended up several years ahead of the industry, but it's not a sustainable position

# Making it easy to follow

Exploring the wild west each time

vs. laying down a shared route



Establish our  
solutions as Best  
Practices / Standards

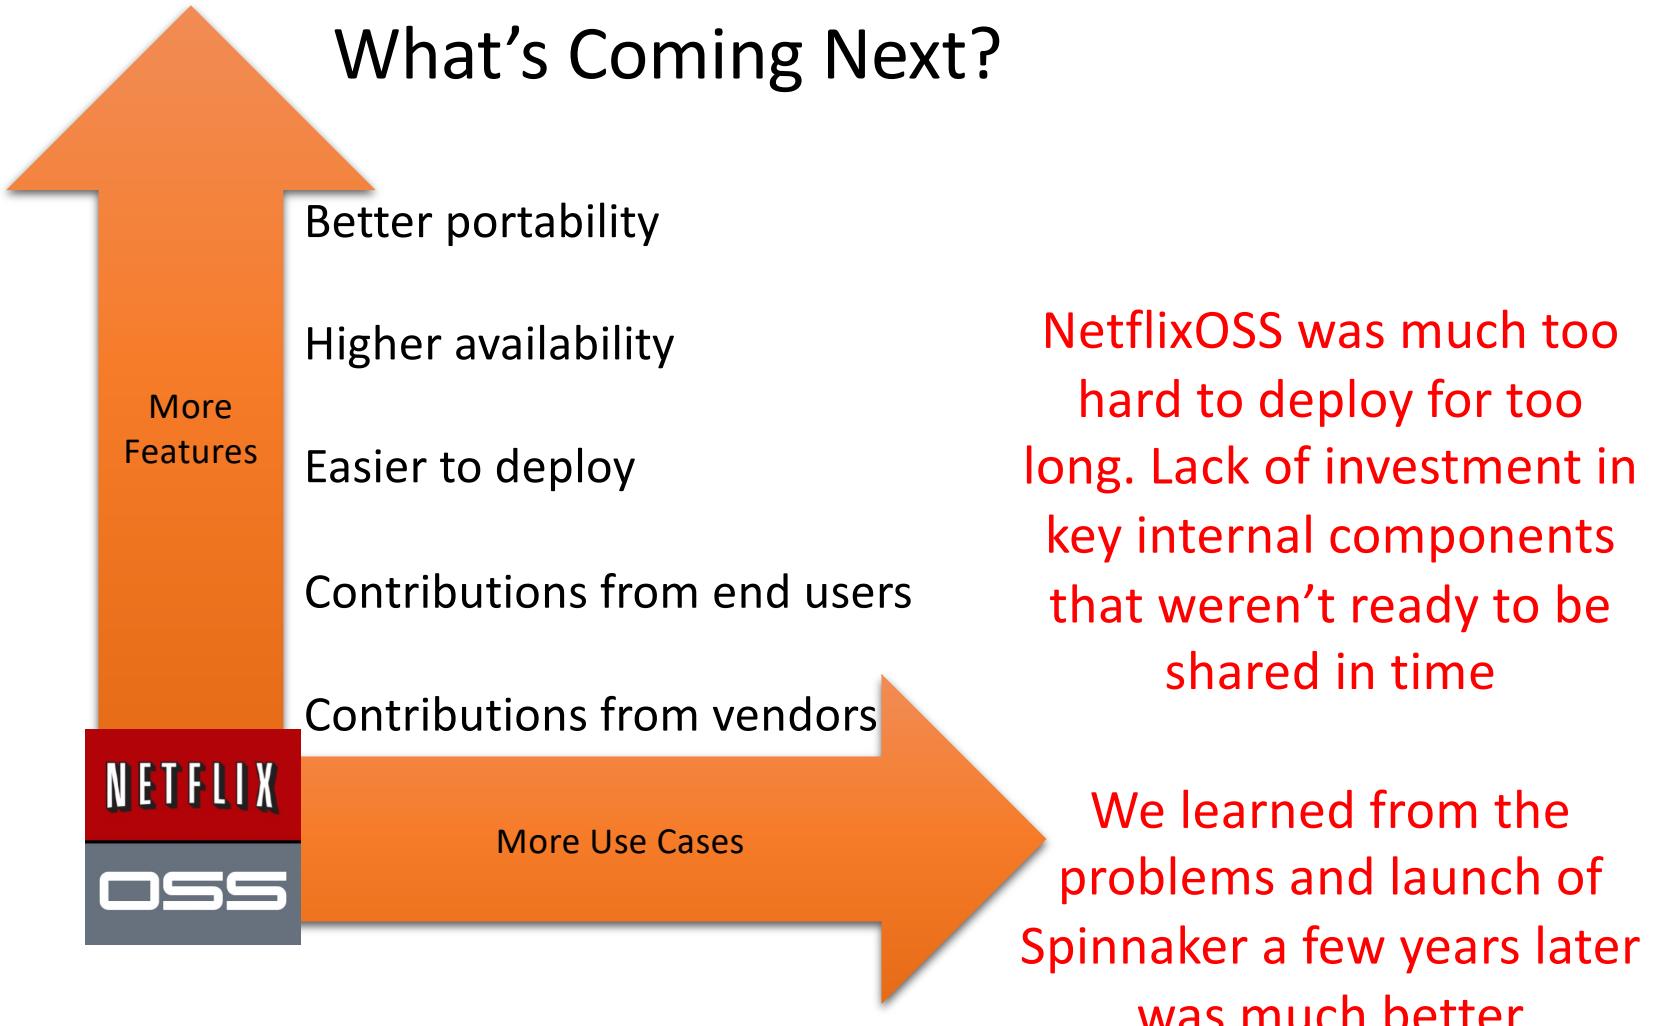
Hire, Retain and  
Engage Top  
Engineers

## Goals

Build up Netflix  
Technology Brand

Benefit from a  
shared ecosystem

# What's Coming Next?





Functionality and scale now, portability coming

Moving from parts to a platform in 2013

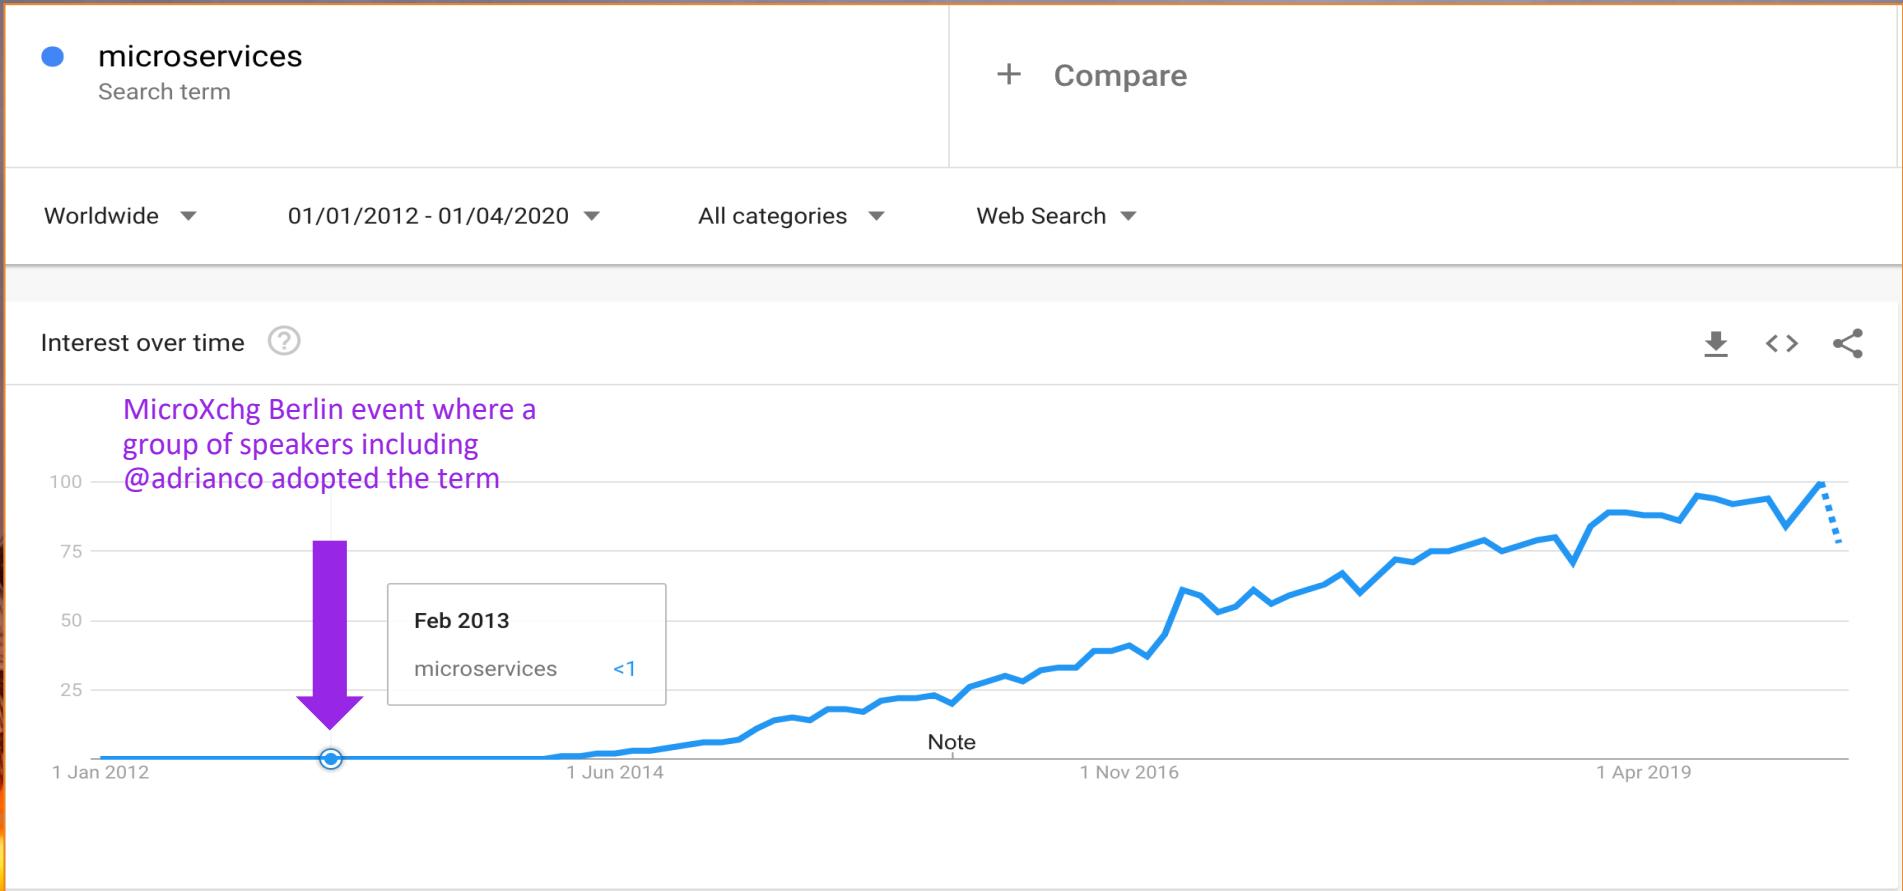
Netflix is fostering an ecosystem

Rapid Evolution - Low MTBIAMSH  
(Mean Time Between Idea And Making Stuff Happen)

A photograph of a sunset over a mountain range. The sky is filled with warm orange and red hues, transitioning into cooler blues and purples at the top. In the foreground, the silhouettes of palm trees are visible against the bright horizon. A prominent dark mountain peak rises in the center. The overall atmosphere is peaceful and scenic.

**Adrian left Netflix in January 2014  
and joined Battery Ventures just as  
Docker containers began to take off as a  
microservices pattern...**

# Trends: Microservices





## Typical reactions to my Netflix talks...



“You guys are crazy! Can’t believe it”  
– 2009

“What Netflix is doing won’t work”  
– 2010

It only works for ‘Unicorns’ like Netflix  
– 2011

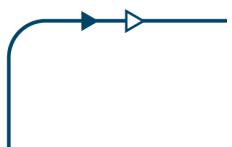
“We’d like to do that but can’t”  
– 2012

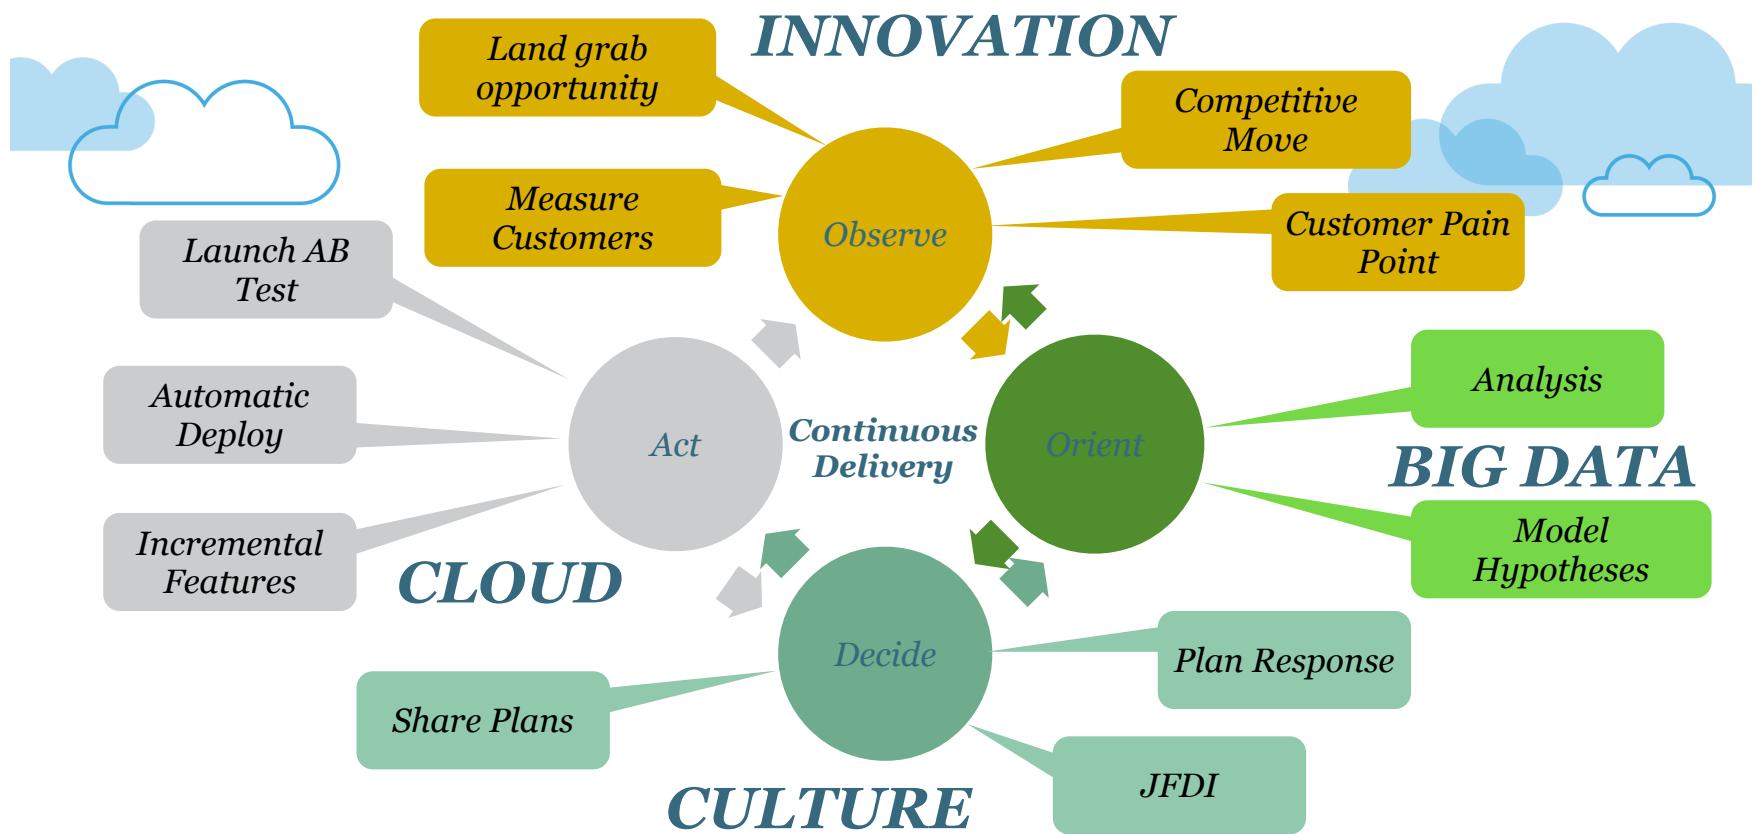
“We’re on our way using Netflix OSS code”  
– 2013



## **What I learned from my time at Netflix**



- *Speed wins in the marketplace*
  - *Remove friction from product development*
  - *High trust, low process, no hand-offs between teams*
  - *Freedom and responsibility culture*
  - *Don't do your own undifferentiated heavy lifting*
  - *Use simple patterns automated by tooling*
  - *Self service cloud makes impossible things instant*
- 



Netflix was accused of “Ruining the cloud” because we didn’t use Chef/Puppet to update our code, we just started a new instance

## Non-Destructive Production Updates

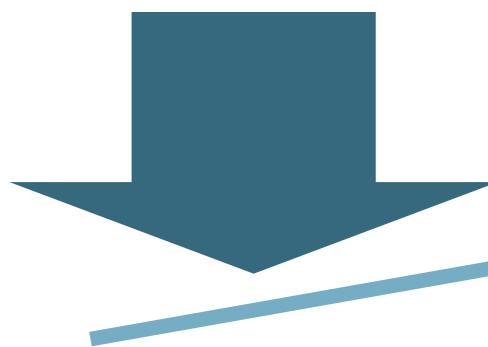


- “*Immutable Code*” Service Pattern
  - Existing services are unchanged, old code remains in service
  - New code deploys as a new service group
  - No impact to production until traffic routing changes
- A|B Tests, Feature Flags and Version Routing control traffic
  - First users in the test cell are the developer and test engineers
  - A cohort of users is added looking for measurable improvement

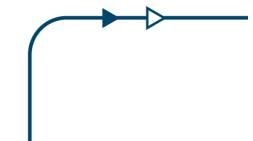
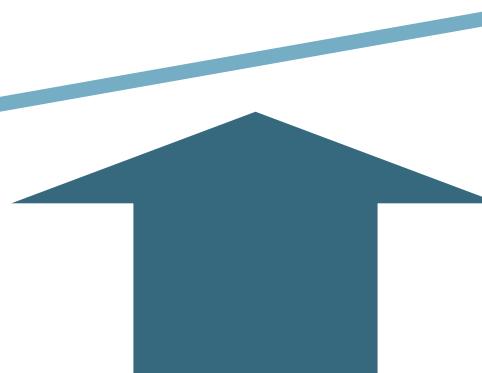
The Netflix model used version-aware routing between services, another idea that seemed to get lost along the way...

# What Happened?

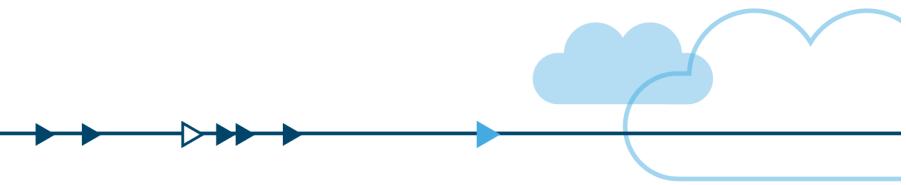
*Rate of change  
increased*



*Cost and size and  
risk of change  
reduced*

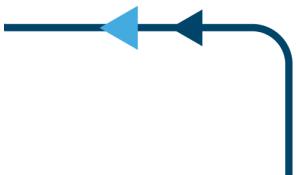


*If every service has to be updated at the same time it's not loosely coupled*



## *A Microservice Definition*

*Loosely coupled service oriented architecture with bounded contexts*



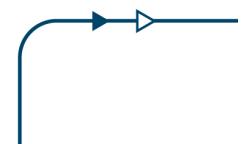
*If you have to know too much about surrounding services you don't have a bounded context. See the Domain Driven Design book by Eric Evans.*

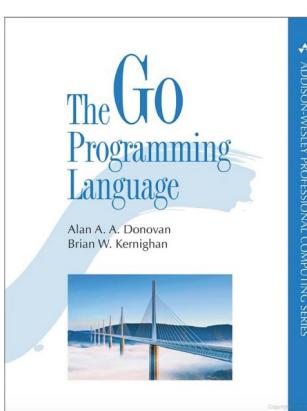
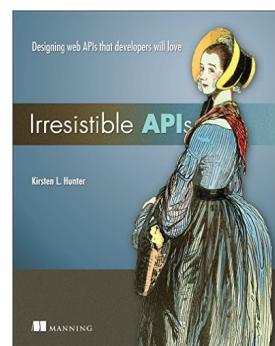
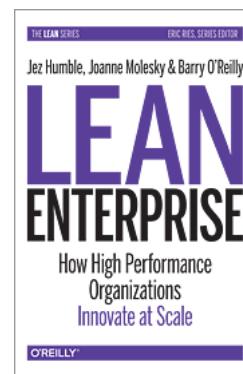
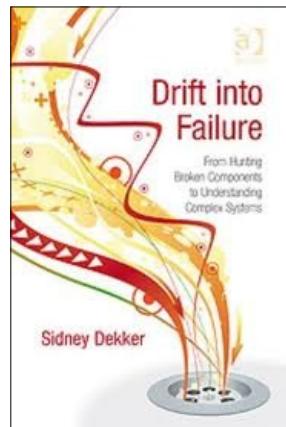
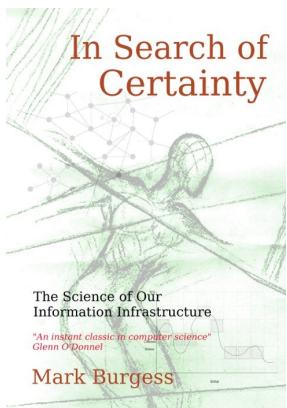
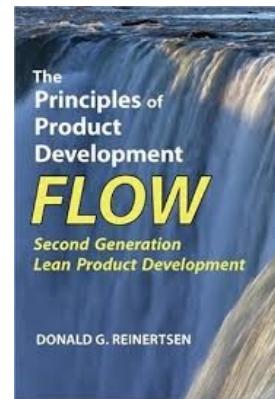
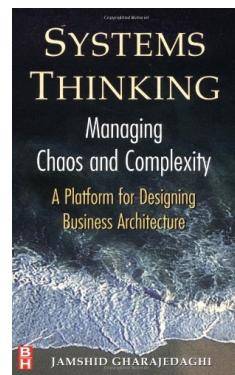
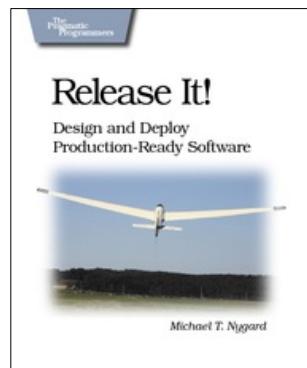
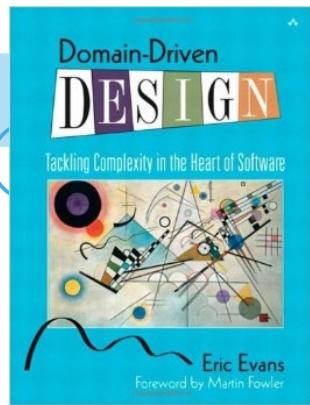
## Separate Concerns with Microservices

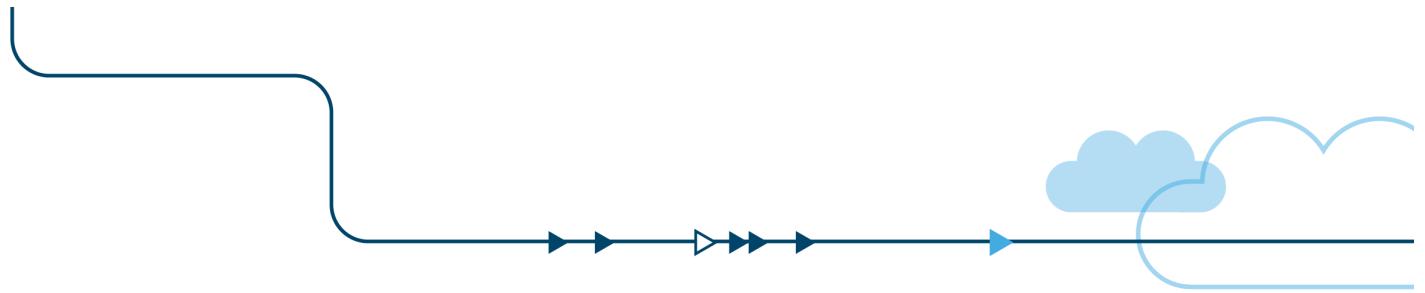


- *Invert Conway's Law – teams own service groups and backend stores*
- *One “verb” per single function micro-service, size doesn't matter*
- *One developer independently produces a micro-service*
- *Each micro-service is its own build, avoids trunk conflicts*
- *Deploy in a container: Tomcat, AMI or Docker, whatever...*
- *Stateless business logic. Cattle, not pets.*
- *Stateful cached data access layer using replicated ephemeral instances*

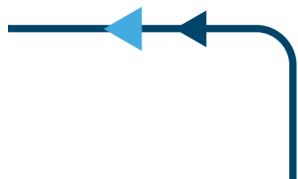
[http://en.wikipedia.org/wiki/Conway's\\_law](http://en.wikipedia.org/wiki/Conway's_law)







*What's Missing?*





# *Advanced Microservices Topics*



*Failure injection testing*

*Versioning, routing*

*Binary protocols and interfaces*

*Timeouts and retries*

*Denormalized data models*

*Monitoring, tracing*

*Simplicity through symmetry*



# *Timeouts and Retries*



*Connection timeout vs. request timeout confusion*

*Usually setup incorrectly, global defaults*

*Systems collapse with “retry storms”*

*Timeouts too long, too many retries*

*Services doing work that can never be used*



# Timeouts and Retries

*Bad config: Every service defaults to 2 second timeout, two retries*



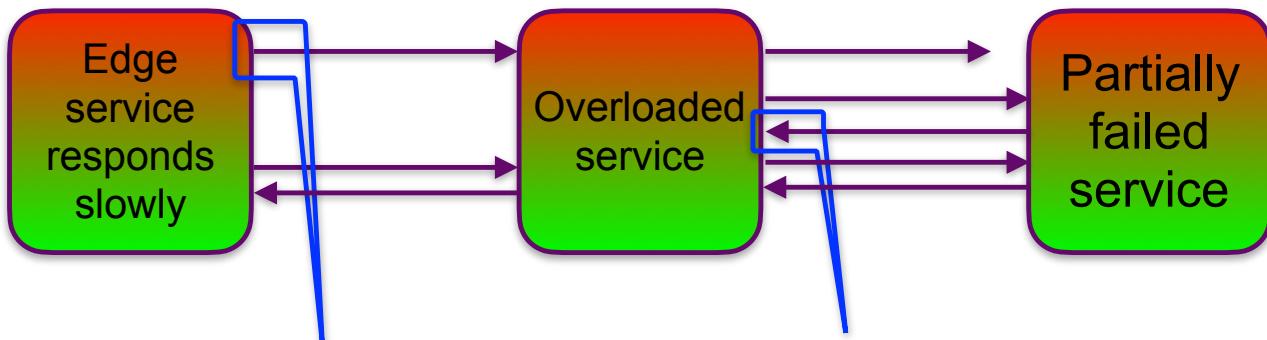
*If anything breaks, everything upstream stops responding*



*Retries add unproductive work*

# Timeouts and Retries

*Bad config: Every service defaults to 2 second timeout, two retries*



*First request from Edge timed out so it ignores the successful response and keeps retrying. Middle service load increases as it's doing work that isn't being consumed*



# *Timeout and Retry Fixes*



*Cascading timeout budget  
Static settings that decrease from the edge  
or dynamic budget passed with request*

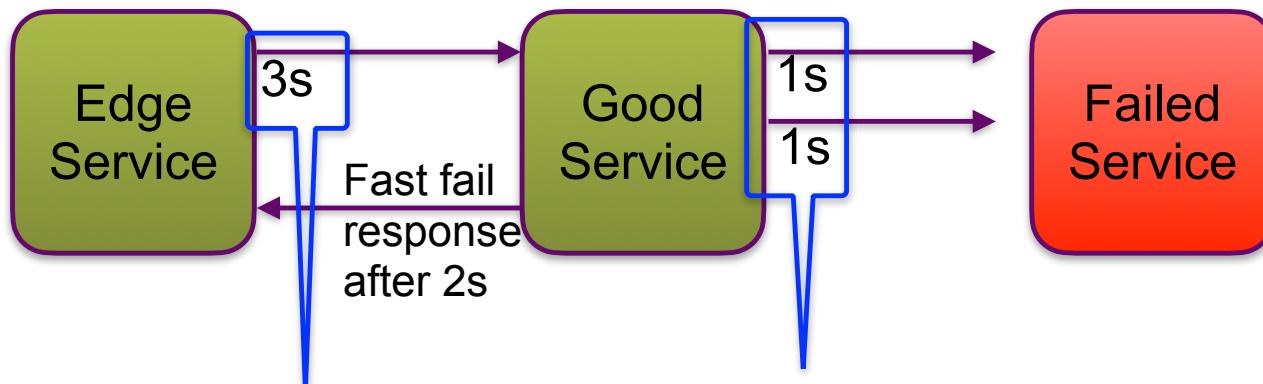
*How often do retries actually succeed?  
Don't ask the same instance the same thing  
Only retry on a different connection*



# *Timeouts and Retries*



*Budgeted timeout, one retry*



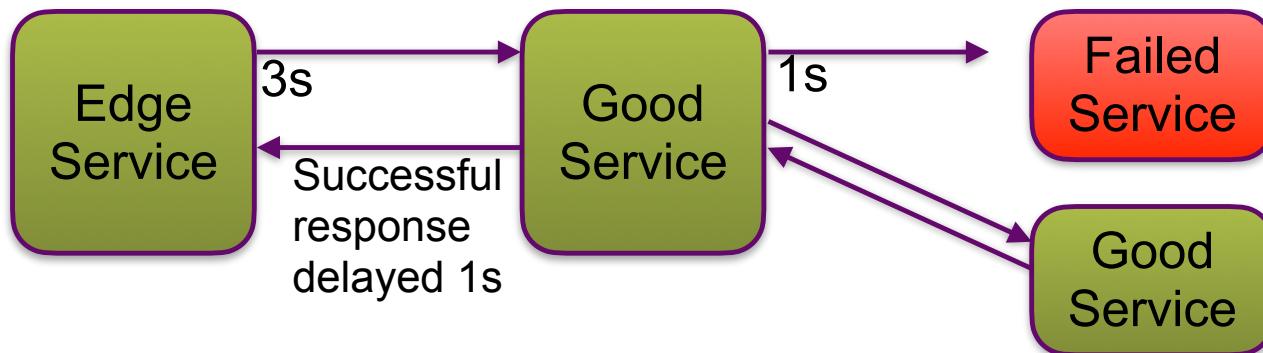
*Upstream timeout must always be longer than total downstream timeout \* retries delay*

*No unproductive work while fast failing*

# *Timeouts and Retries*

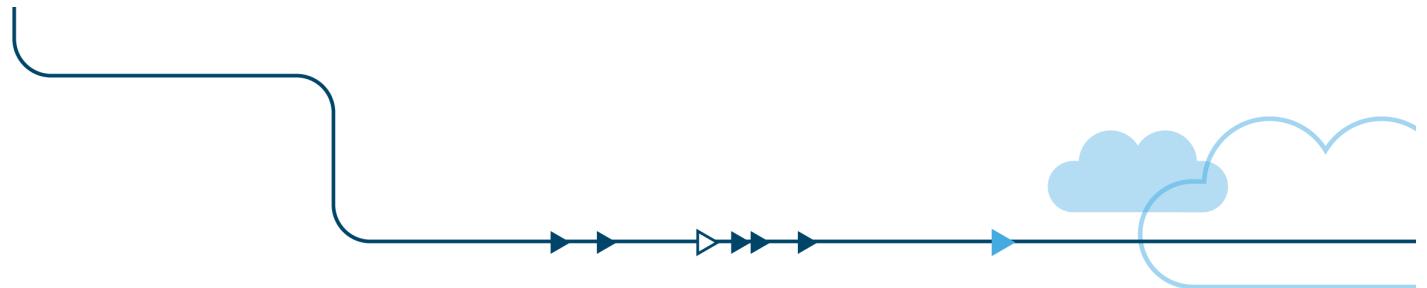


*Budgeted timeout, failover retry*



*For replicated services with multiple instances  
never retry against a failed instance*

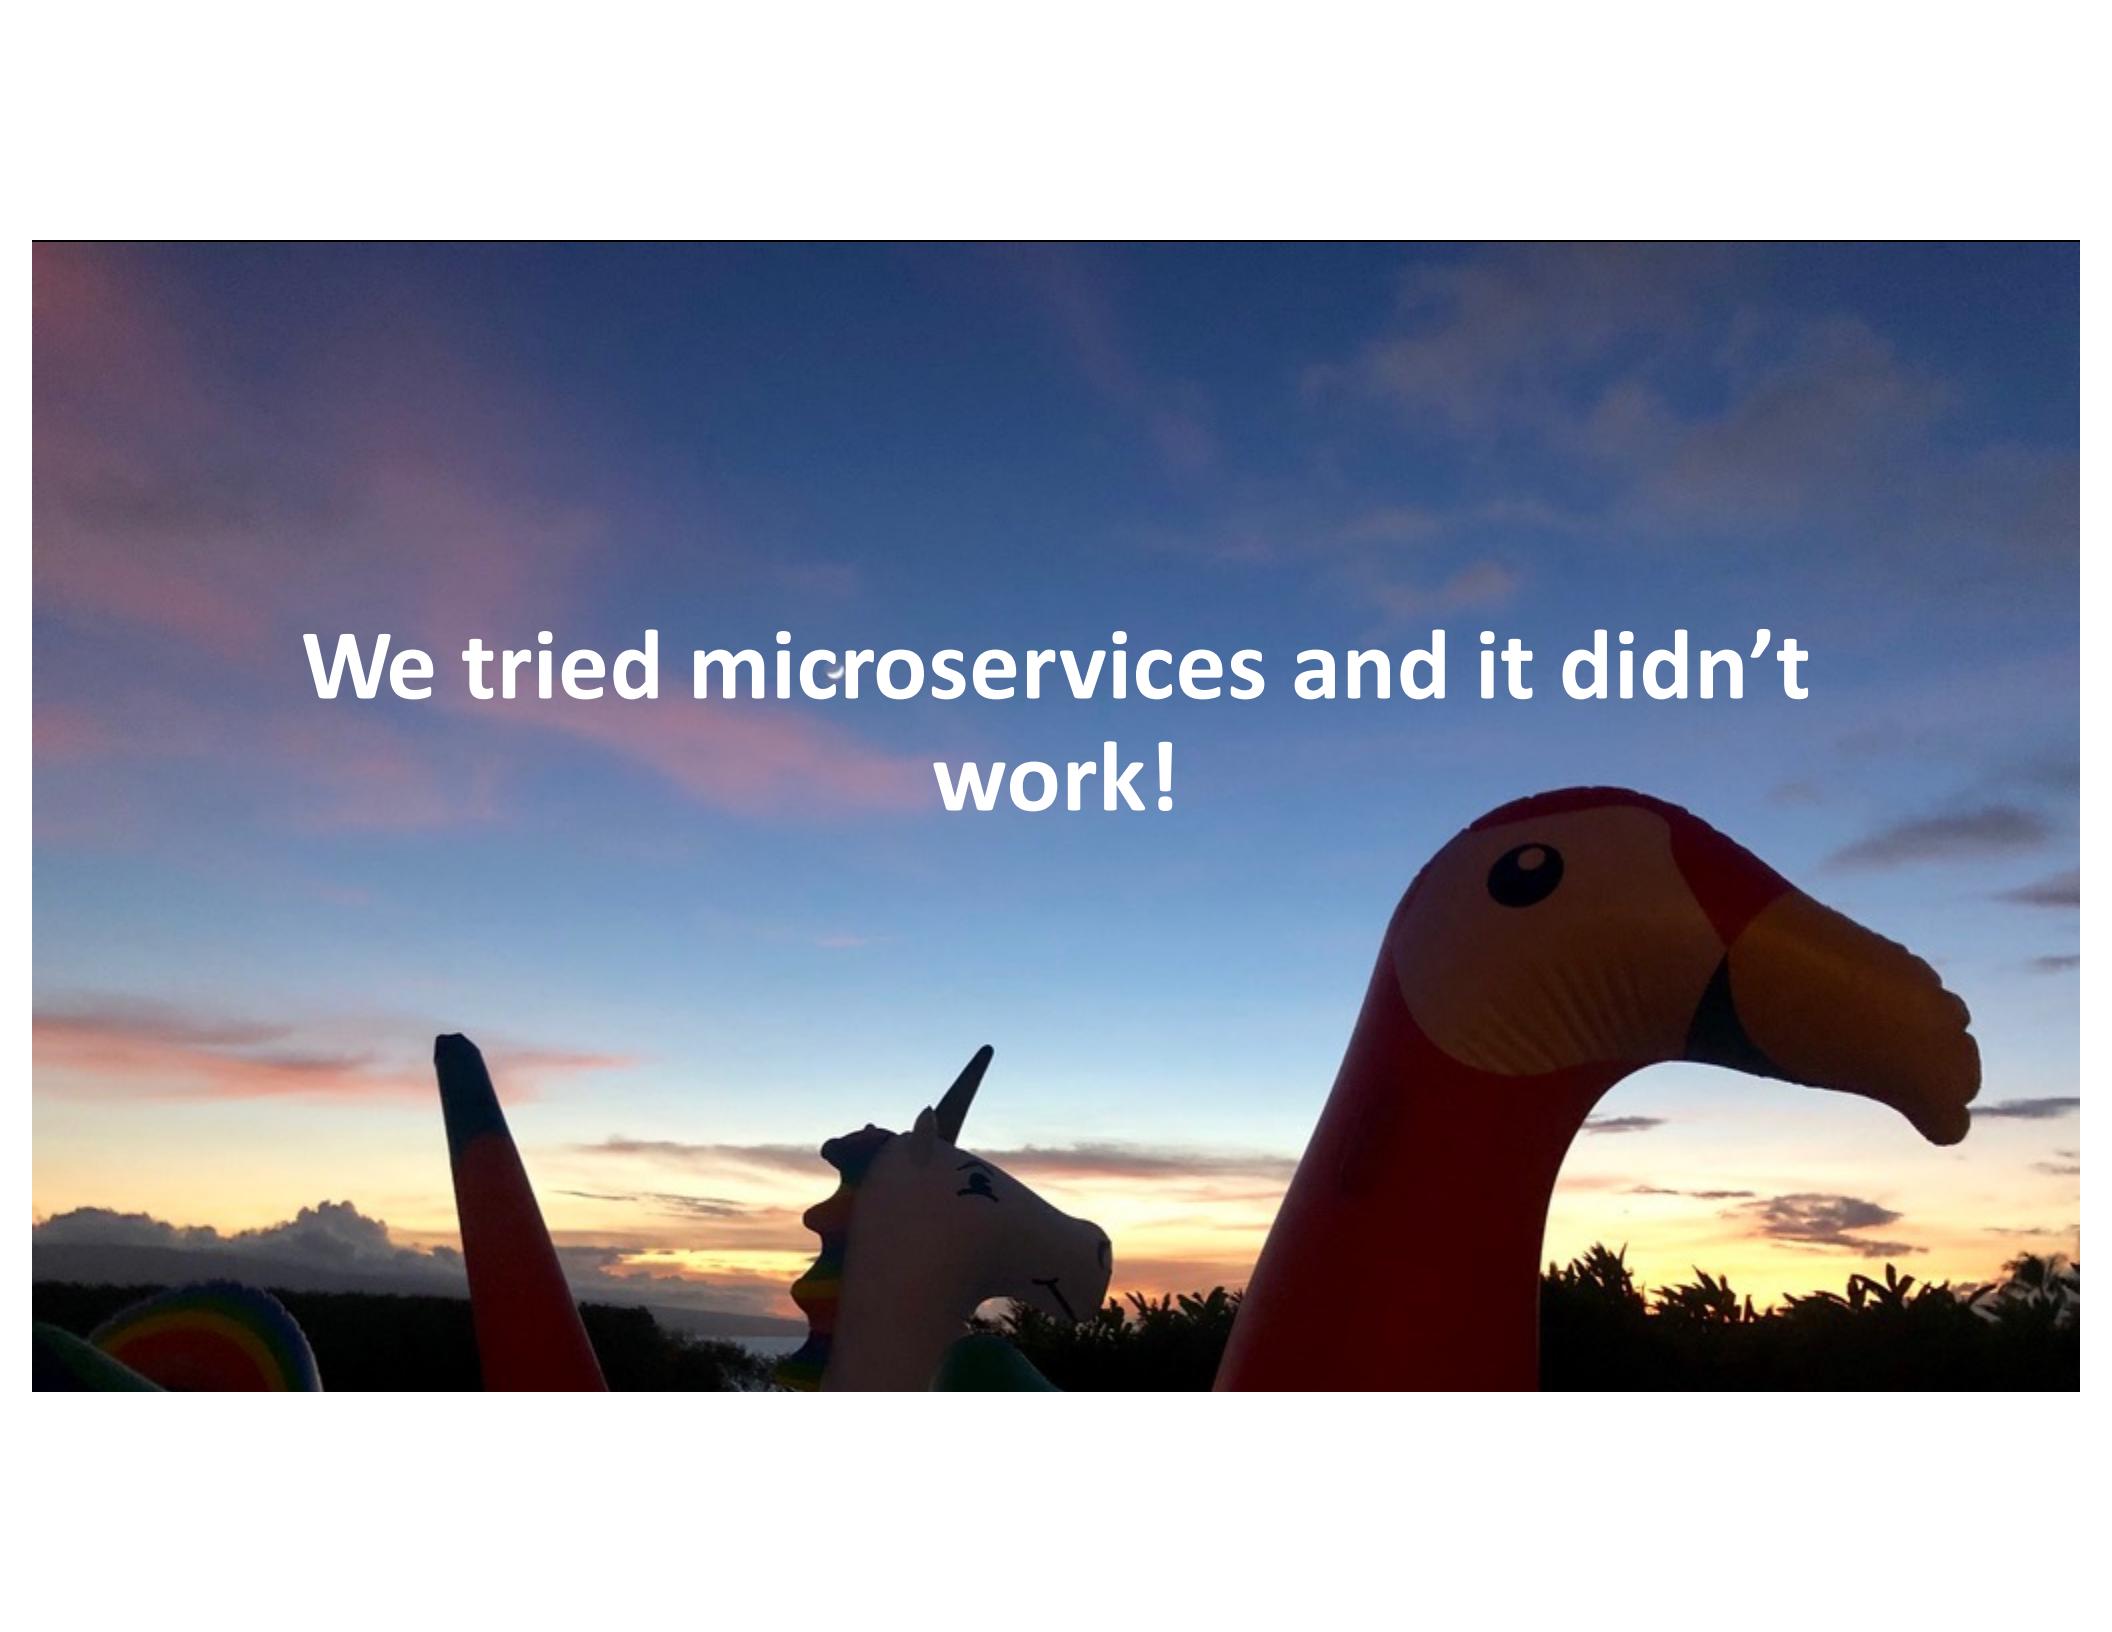
*No extra retries or unproductive work*  
These lessons weren't learned, most microservices frameworks still suffer from retry storms as far as I can tell



“We see the world as increasingly more complex and chaotic because we use inadequate concepts to explain it. When we understand something, we no longer see it as chaotic or complex.”

*Jamshid Gharajedaghi - 2011  
Systems Thinking: Managing Chaos and Complexity: A Platform for Designing Business Architecture*



A photograph of a sunset over a tropical landscape. In the foreground, there are several large, colorful inflatables shaped like a unicorn and a parrot. The sky is filled with warm orange and yellow hues from the setting sun, transitioning into a darker blue at the top.

We tried microservices and it didn't work!

A photograph of two large, colorful unicorn inflatables against a sunset sky. One unicorn is white with a rainbow horn and tail, and the other is orange with a blue horn and tail. They are positioned in front of a vibrant sunset with clouds transitioning from orange to blue.

Are you a unicorn or a donkey with a  
carrot tied to it's forehead?

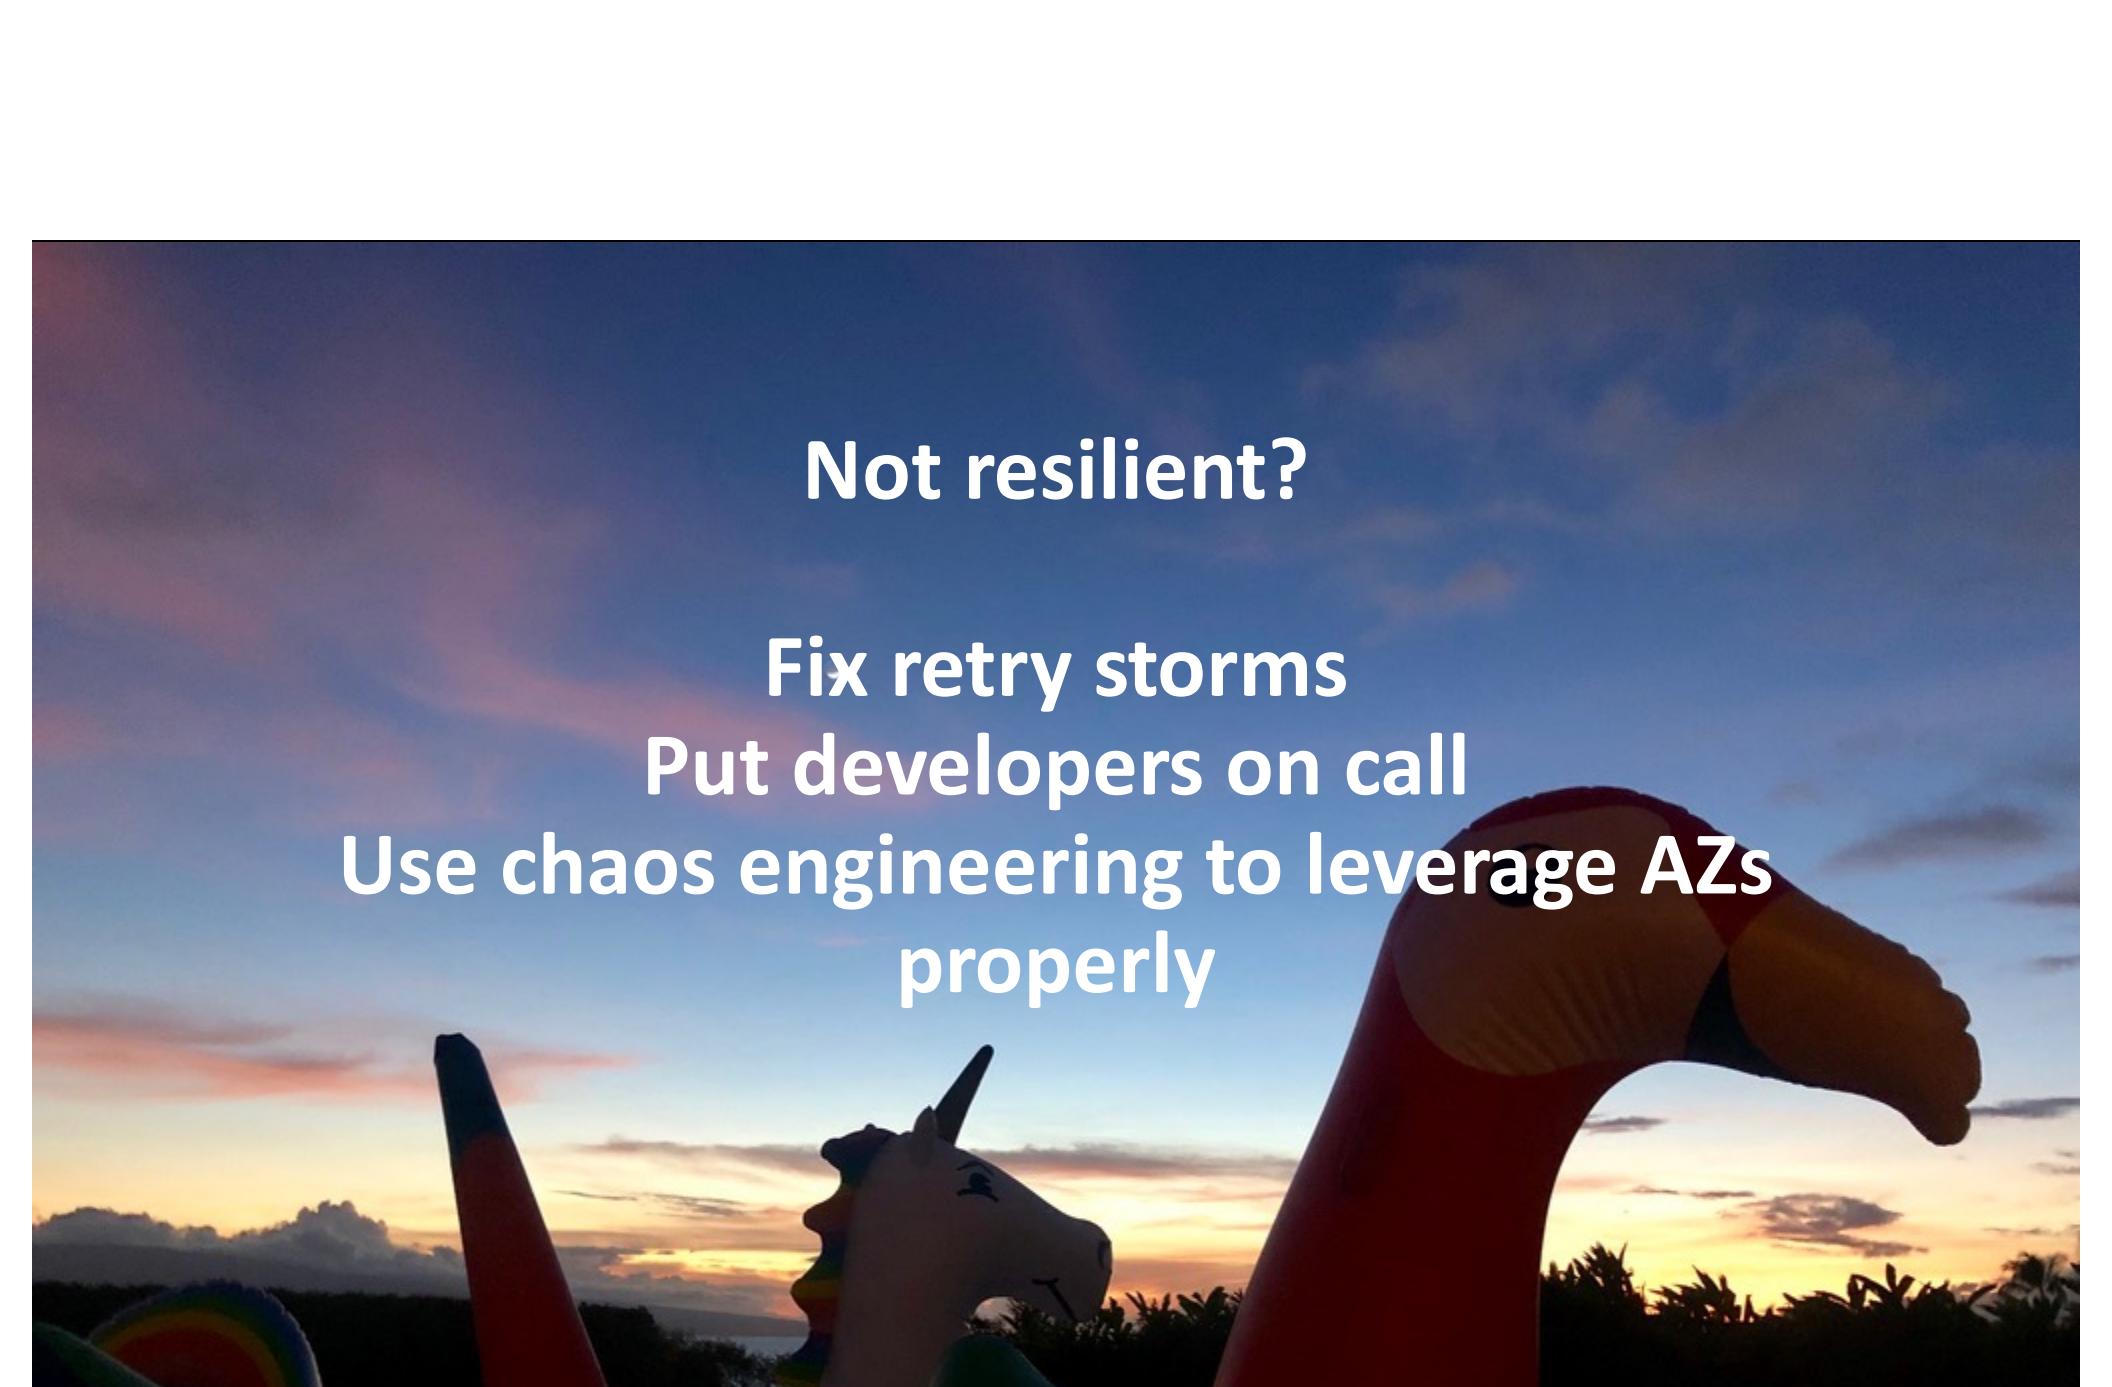
A photograph of a large, colorful unicorn statue made of what appears to be inflated plastic or fabric. The unicorn is white with a rainbow-colored mane and tail, and it has a small horn. It is positioned in front of a beautiful sunset or sunrise sky with orange, yellow, and blue hues. The silhouette of trees and hills is visible in the background.

**Microservices is one of many tactics  
that reinforce each other to form a  
system that is fast, resilient, scalable  
and efficient**

- Too slow?

**Stop using chatty/inefficient protocols  
Pick your system boundaries better**



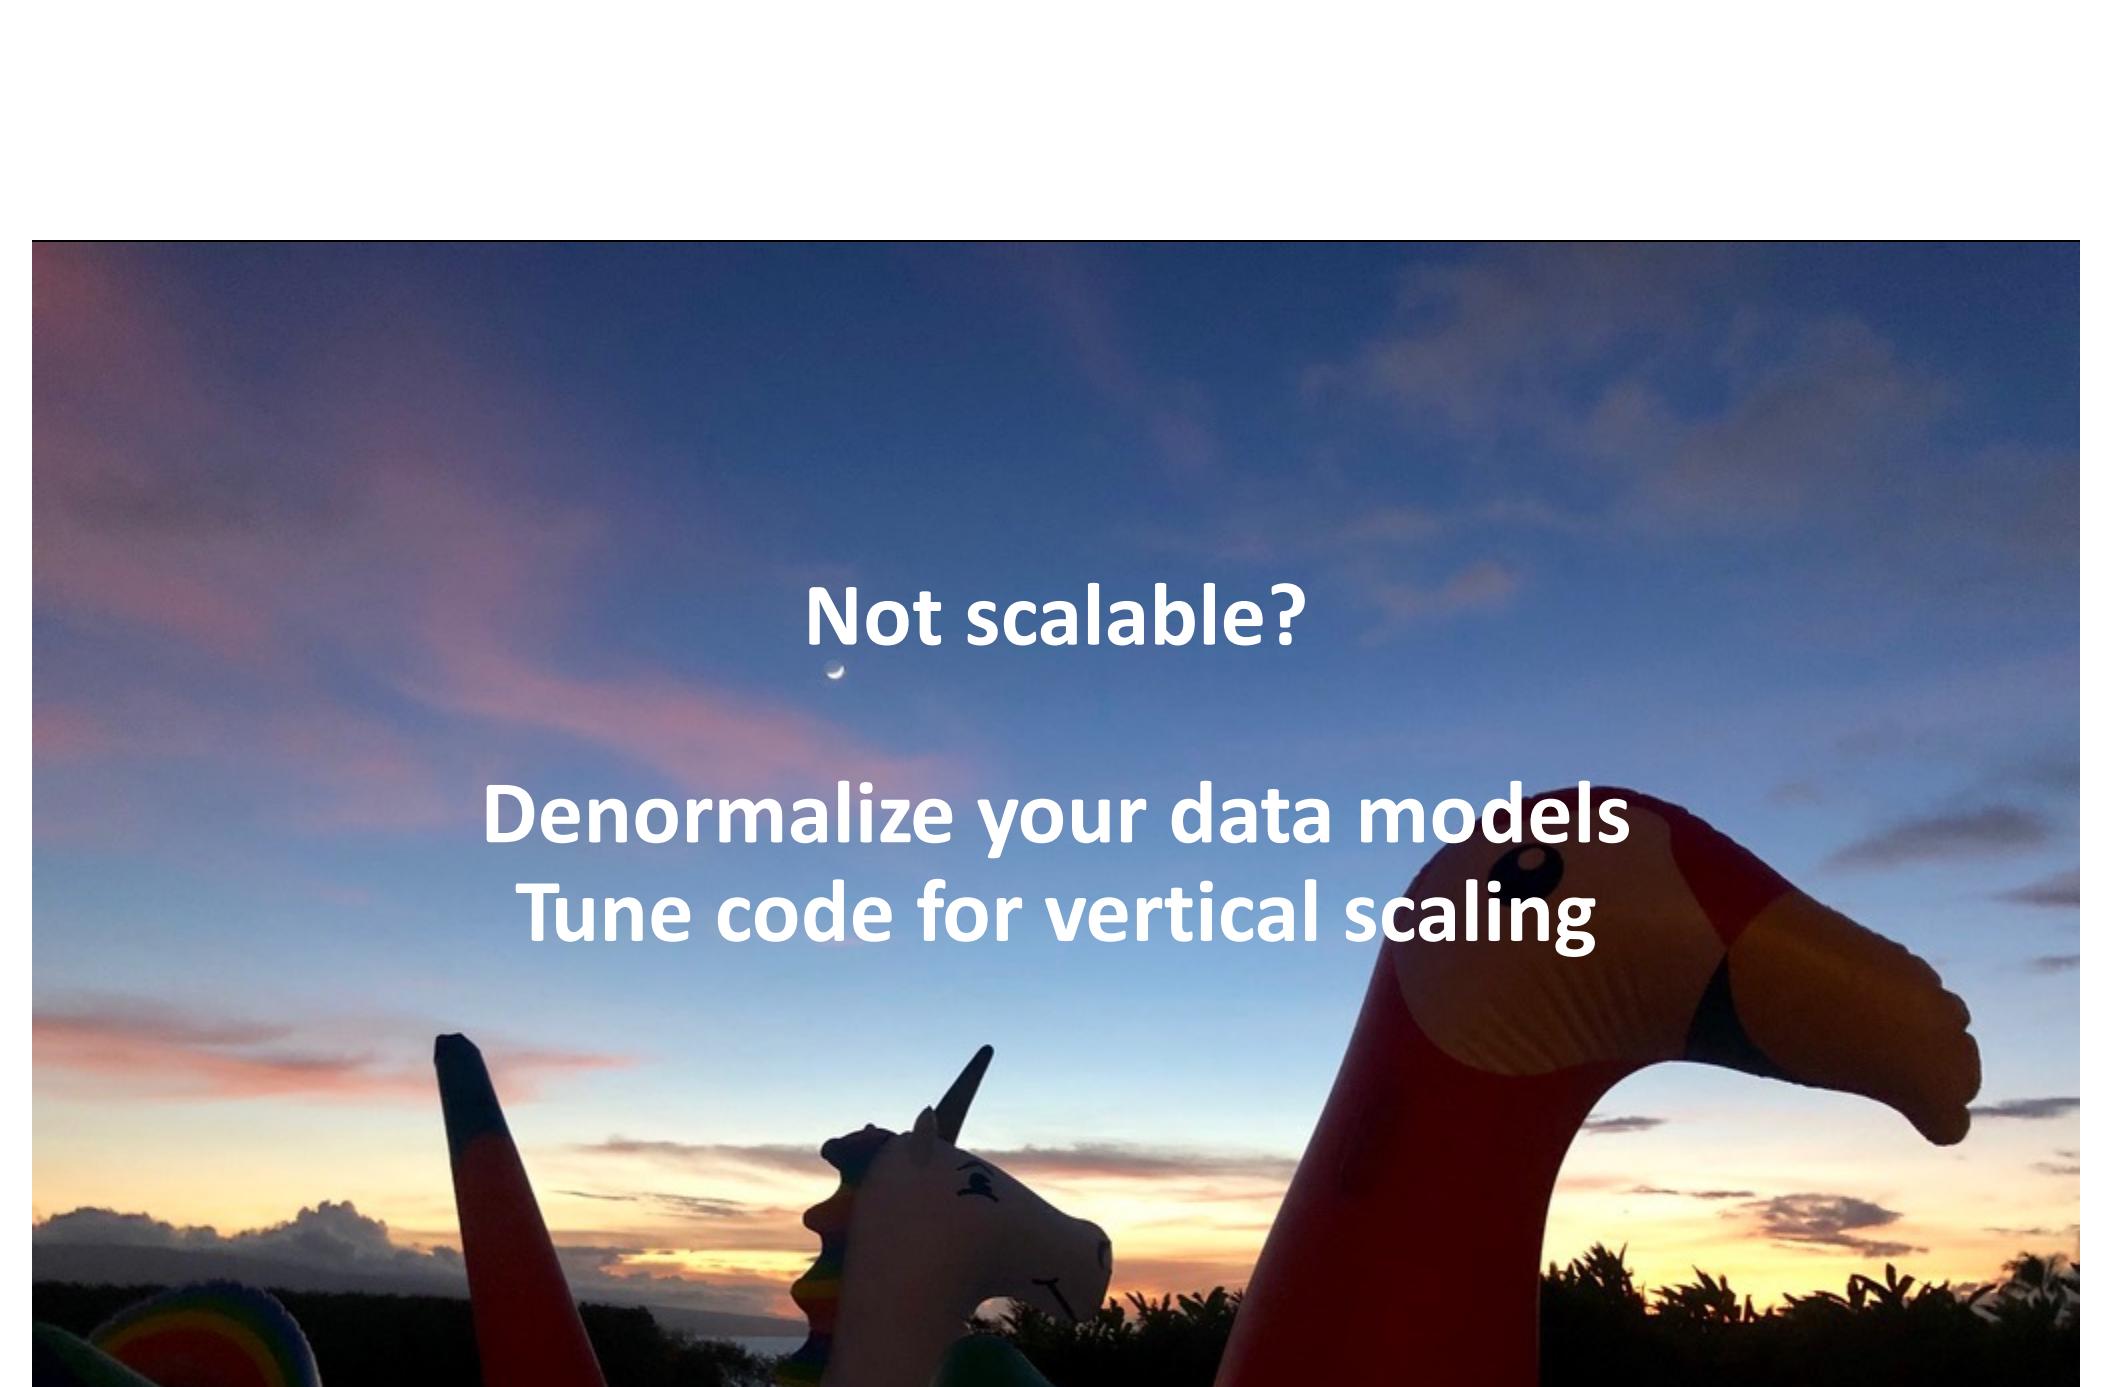
A photograph of a large, colorful unicorn statue made of what appears to be painted metal or wood. The unicorn is white with a rainbow-colored mane and tail, and it has a small horn. It is positioned in front of a vibrant sunset or sunrise sky with orange, yellow, and blue hues. The horizon shows silhouettes of trees and hills.

**Not resilient?**

**Fix retry storms**

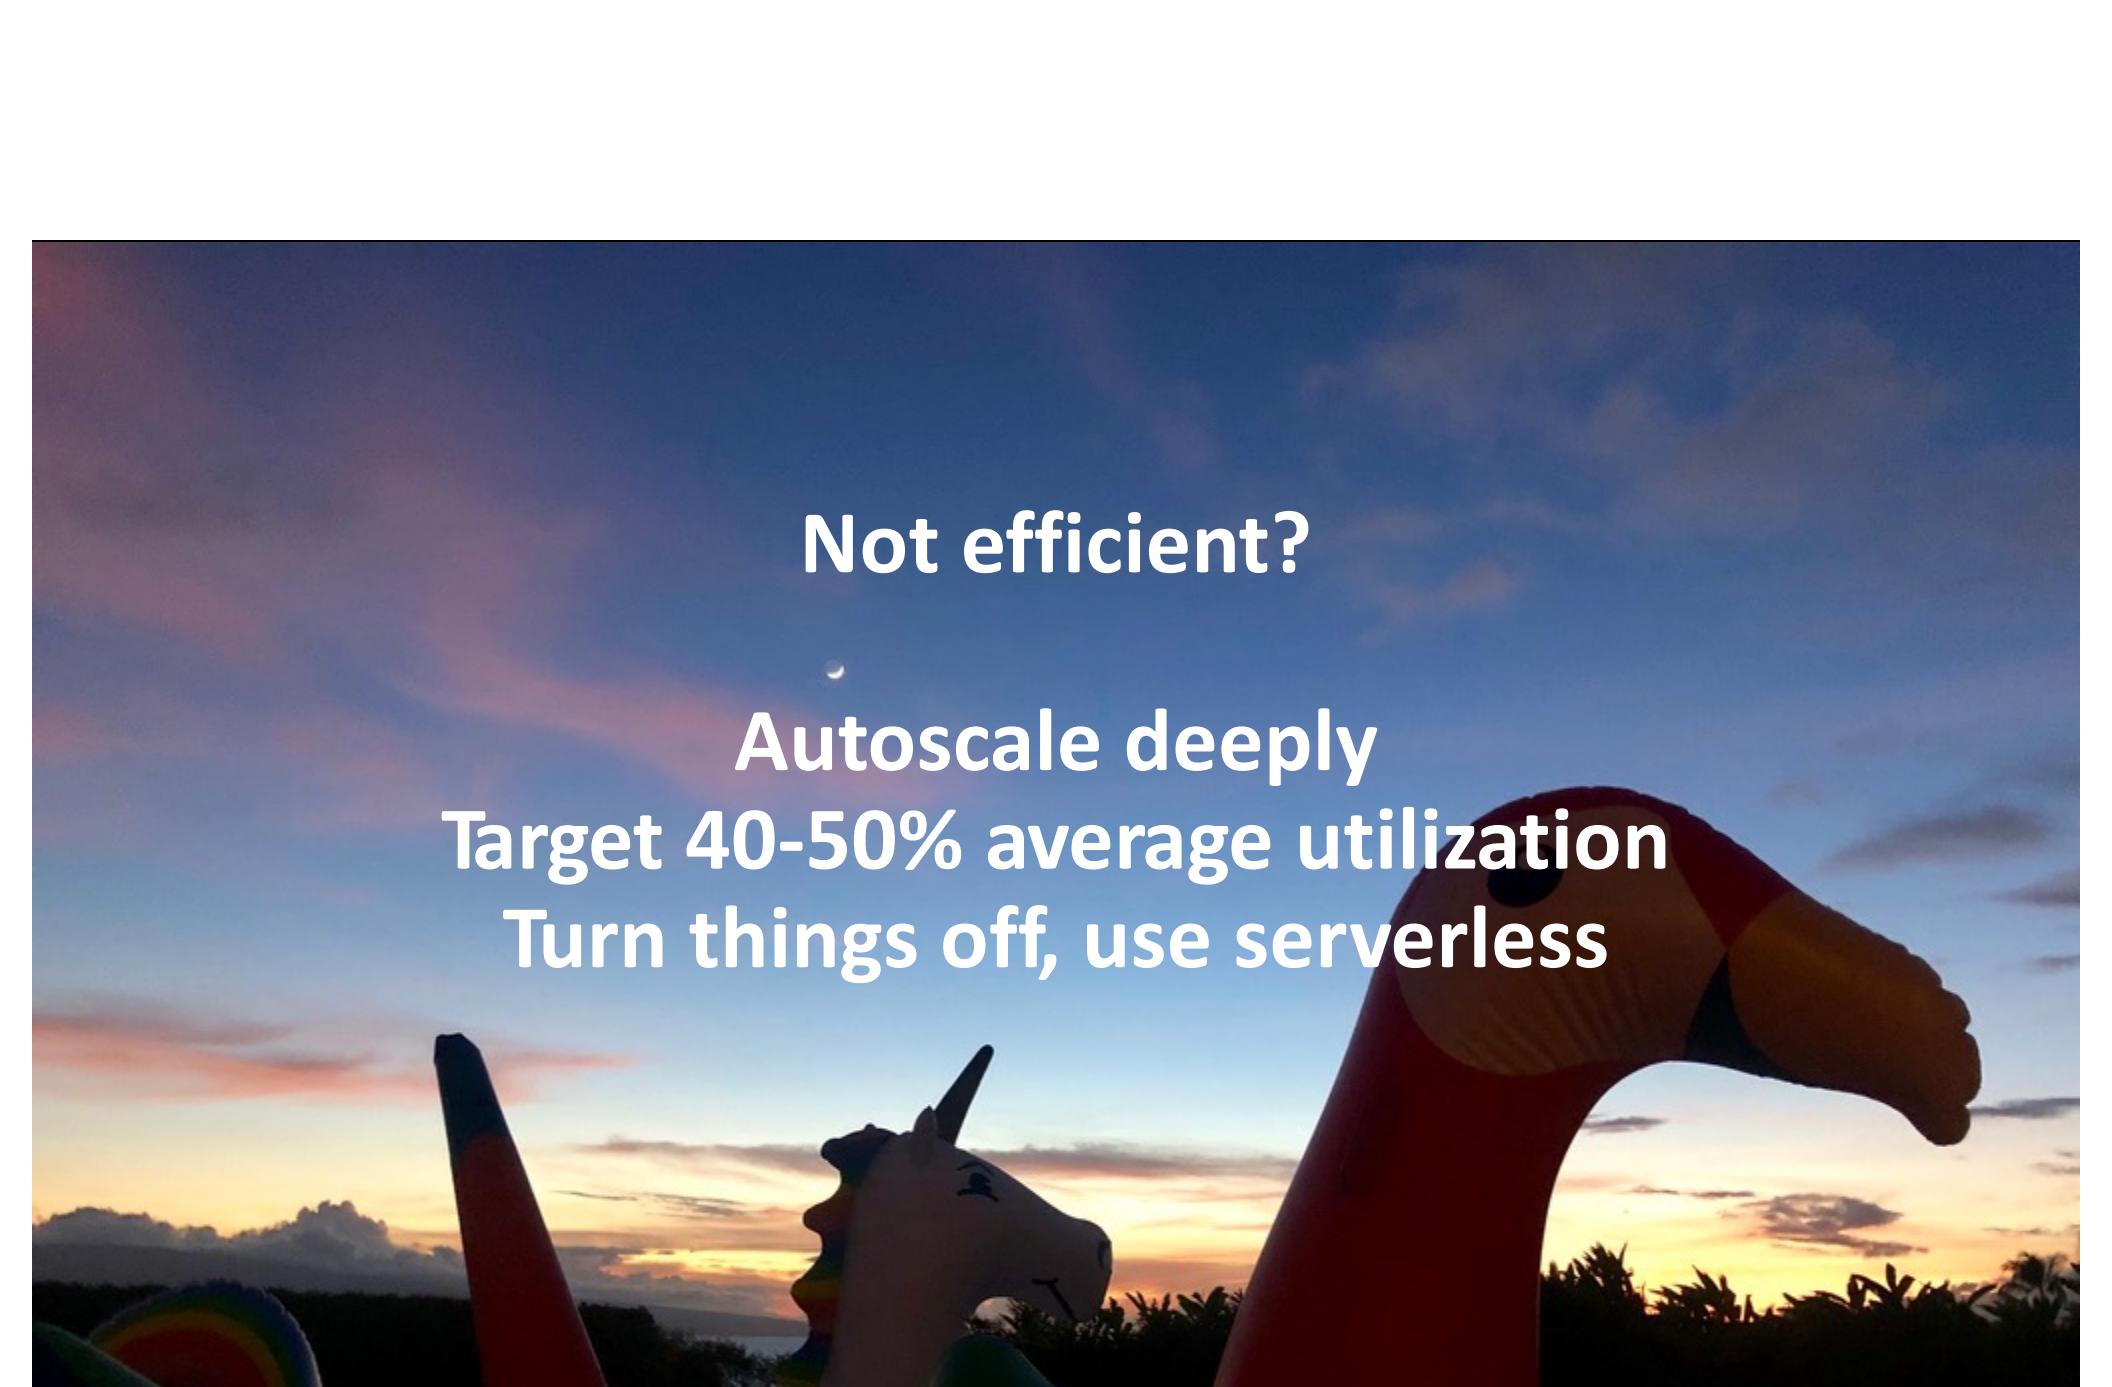
**Put developers on call**

**Use chaos engineering to leverage AZs  
properly**



Not scalable?

Denormalize your data models  
Tune code for vertical scaling

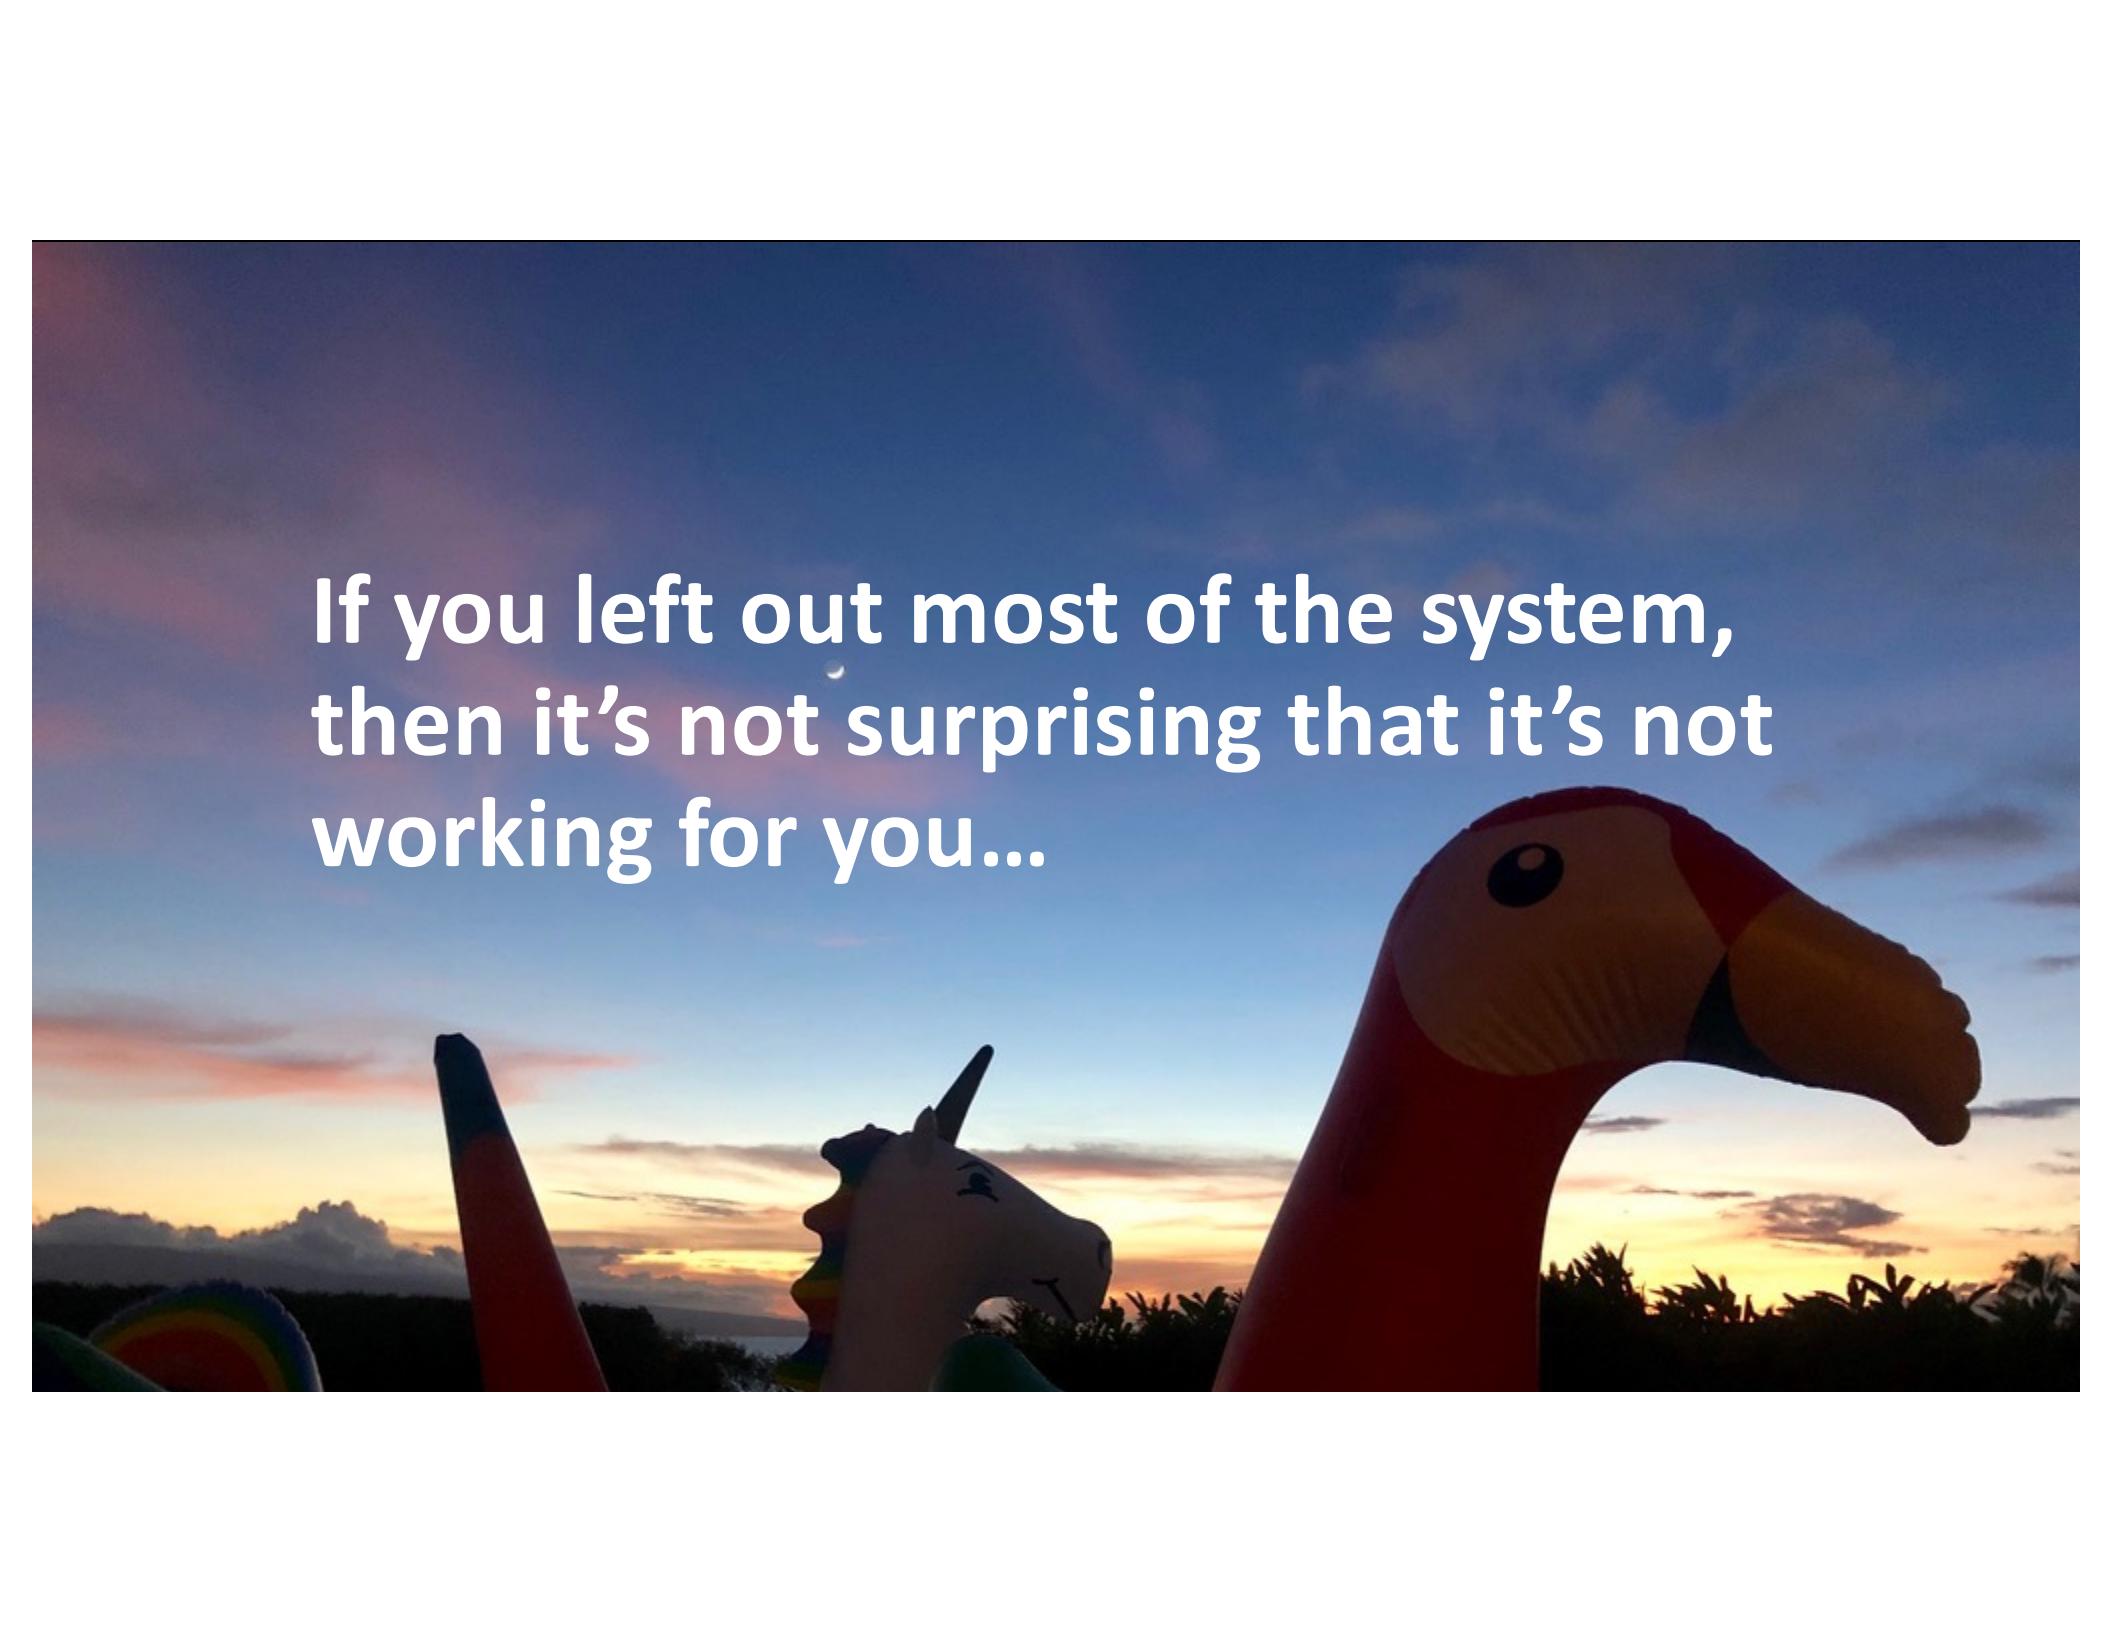
A photograph of a large, colorful unicorn statue made of what looks like painted wood or metal. The unicorn is white with a rainbow-colored mane and tail, and it has a small horn. It is positioned in front of a beautiful sunset or sunrise sky with orange, yellow, and blue hues. The background shows some silhouettes of trees and possibly a body of water.

**Not efficient?**

**Autoscale deeply**

**Target 40-50% average utilization**

**Turn things off, use serverless**

A photograph of a vibrant sunset over a tropical landscape. In the foreground, several large, colorful inflatables are visible, including a tall orange and red structure, a white unicorn-like figure, and a red parrot. The sky is filled with warm hues of orange, yellow, and blue, with scattered clouds. The horizon shows silhouettes of trees and mountains.

If you left out most of the system,  
then it's not surprising that it's not  
working for you...



Backlog

*Adrian Cockcroft - @adrianco*

# *Systems Thinking at Netflix and Beyond*

*Netflix systematically drives  
higher talent density via  
company culture optimized  
for “fully formed adults”*

Netflix is running an Olympic sports team, not a family

***“We can’t copy Netflix because it has all those superstar engineers, we don’t have the people”***

*Fortune 100 CTO after a Netflix presentation - 2013*

***“We hired them from you, and got out of their way...”***

*Adrian Cockcroft - in response*

***“Purposeful systems, representing the systems view of development, assumes plurality in all three dimensions: function, structure, and process.”***

*Jamshid Gharajedaghi - 2011  
Systems Thinking: Managing Chaos and Complexity: A Platform for Designing Business Architecture*

See also W. Ross Ashby's Law of Requisite Variety  
Simplification and efficiency drives can go too far

*If you want to build a ship,  
don't drum up the people  
to gather wood, divide the  
work, and give orders.  
Instead, teach them to yearn  
for the vast and endless sea.*

-Antoine De Saint-Exupery,  
Author of The Little Prince

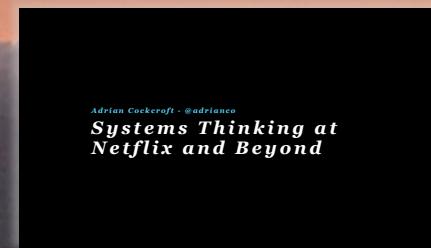
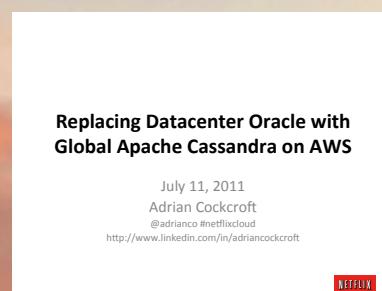
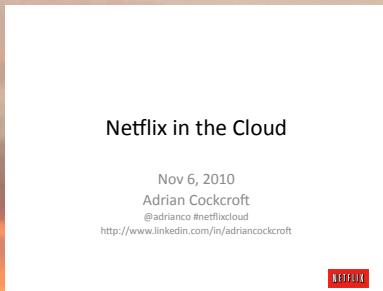
***The basic thesis of this article is that organizations which design systems (in the broad sense used here) are constrained to produce designs which are copies of the communication structures of these organizations.***

***Melvin E. Conway - Conway's Law  
How Do Committees Invent? - 1968***



Ruth and Mel Conway 2019

## This talk features excerpts from these decks

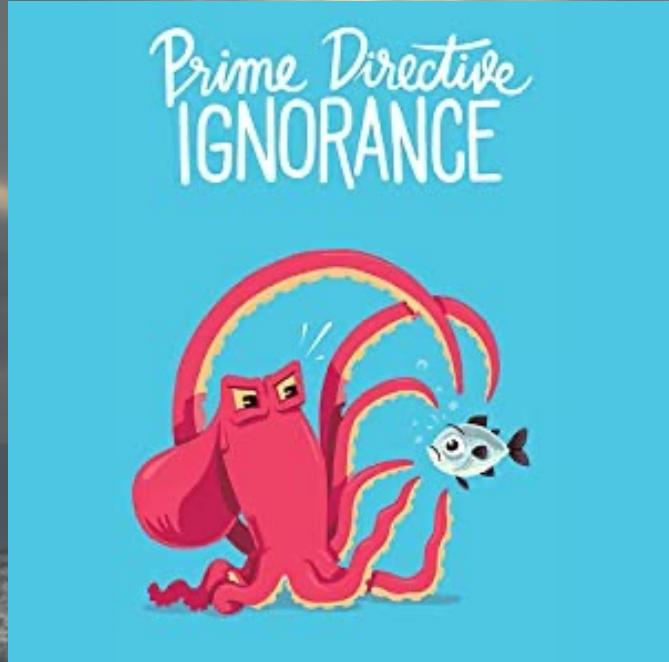


<https://www.slideshare.net/adrianco> - Netflix era decks

<https://www.slideshare.net/adriancockcroft> - Battery Ventures decks

<https://github.com/adrianco/slides> - AWS decks and later including a pdf of these slides!

<https://www.youtube.com/@adriancockcroft>



**Antipattern: You didn't protect unprepared organizations from the dangers of introducing advanced technology, knowledge and values, before they are ready...**

The background image is a wide-angle photograph of a tropical sunset. The sky is filled with large, billowing clouds that are lit from below by the setting sun, creating a warm orange and yellow glow. The horizon shows silhouettes of hills or mountains. In the foreground, dark silhouettes of palm trees and other tropical foliage are visible against the bright sky.

# Microservices Retrospective - What we learned (and didn't learn) from Netflix

Adrian Cockcroft | Qcon London | March 2023

[@adrianco@mastodon.social](mailto:@adrianco@mastodon.social)