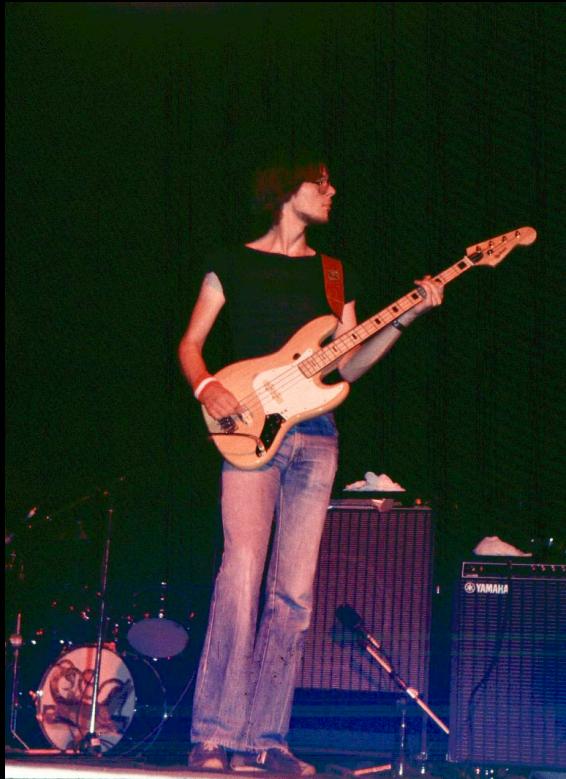


# Adrian's Greatest Hits, B-Sides and Re-issues

Adrian Cockcroft | YOW | December 2022

[@adrianco@mastodon.social](mailto:@adrianco@mastodon.social)



Adrian in 1982

<https://soundcloud.com/adrian-cockcroft/black-tiger-dont-look-back>

Like a classic rock band on their annual  
“farewell and we aren’t dead yet tour”

I know you just want to hear the hits, but I  
want to sneak in some stuff I like too...

(Headbanging optional)

## Setlist

Black Tiger - Don't Look Back - 1982

Netflix in the Cloud - QconsF 2010, Cassandra Summit 2011

Microservices - Various MicroXchg Berlin talks

Cloud Trends - GigaOM Structure - 2016 reissue

Communicating Sequential Goroutines - Gophercon - 2016

Lego spaceships and the kitchen sink - AWS - 2017-2019

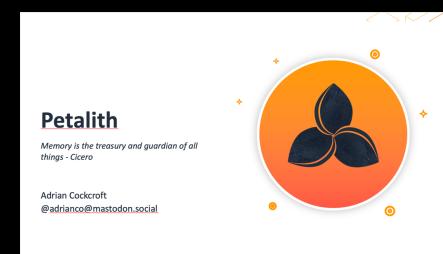
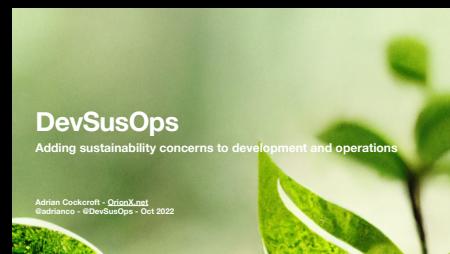
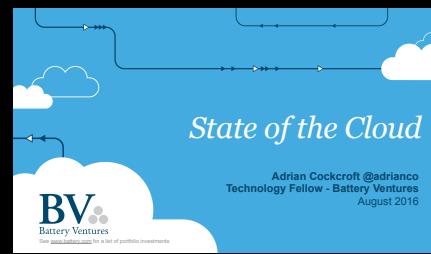
DevSusOps - a track from the new album you don't want to hear

Petalith - An experimental new work in progress

## Encore

Bottleneck Analysis - GOTO Aarhus - 2013

# This talk features "singles" from these "albums"



# Netflix in the Cloud

Nov 6, 2010

Adrian Cockcroft

@adrianco #netflixcloud

<http://www.linkedin.com/in/adriancockcroft>



# We stopped building our own datacenters

Capacity growth rate is accelerating, unpredictable  
Product launch spikes - iPhone, Wii, PS3, XBox  
Datacenter is large inflexible capital commitment



## Leverage AWS Scale “the biggest public cloud”

AWS investment in tooling and automation

AWS zones for high availability, scalability

AWS skills are common on resumes...



*“The cloud lets its users focus on delivering differentiating business value instead of wasting valuable resources on the **undifferentiated heavy lifting** that makes up most of IT infrastructure.”*



Werner Vogels  
Amazon CTO



# What, Why and How?

The details...



# Synopsis

- The Goals
  - Faster, Scalable, Available and Productive
- Anti-patterns and Cloud Architecture
  - The things we wanted to change and why
- Cloud Bring-up Strategy
  - Developer Transitions and Tools
- Roadmap and Next Steps



# Goals

- Faster
  - Lower latency than the equivalent datacenter web pages and API calls
  - Measured as mean and 99<sup>th</sup> percentile
  - For both first hit (e.g. home page) and in-session hits for the same user
- Scalable
  - Avoid needing any more datacenter capacity as subscriber count increases
  - No central vertically scaled databases
  - Leverage AWS elastic capacity effectively
- Available
  - Substantially higher robustness and availability than datacenter services
  - Leverage multiple AWS availability zones
  - No scheduled down time, no central database schema to change
- Productive
  - Optimize agility of a large development team with automation and tools
  - Leave behind complex tangled datacenter code base (~8 year old architecture)
  - Enforce clean layered interfaces and re-usable components



# Old Datacenter vs. New Cloud Arch

Central SQL Database → Distributed Key/Value NoSQL

Sticky In-Memory Session → Shared Memcached Session

Chatty Protocols → Latency Tolerant Protocols

Tangled Service Interfaces → Layered Service Interfaces

Instrumented Code → Instrumented Service Patterns

Fat Complex Objects → Lightweight Serializable Objects

Components as Jar Files → Components as Services



# Tangled Service Interfaces

- Datacenter implementation is exposed
  - Oracle SQL queries mixed into business logic
- Tangled code
  - Deep dependencies, false sharing
- Data providers with sideways dependencies
  - Everything depends on everything else

*Anti-pattern affects productivity, availability*



# Untangled Service Interfaces

- New Cloud Code With Strict Layering
  - Compile against interface jar
  - Can use spring runtime binding to enforce
- Service interface **is** the service
  - Implementation is completely hidden
  - Can be implemented locally or remotely
  - Implementation can evolve independently



# Untangled Service Interfaces

Two layers:

- SAL - Service Access Library
  - Basic serialization and error handling
  - REST or POJO's defined by data provider
- ESL - Extended Service Library
  - Caching, conveniences
  - Can combine several SALs
  - Exposes faceted type system (described later)
  - Interface defined by data consumer in many cases



# Boundary Interfaces

- Isolate teams from external dependencies
  - Fake SAL built by cloud team
  - Real SAL provided by data provider team later
  - ESL built by cloud team using faceted objects
- Fake data sources allow development to start
  - e.g. Fake Identity SAL for a test set of customers
  - Development solidifies dependencies early
  - Helps external team provide the right interface



# One Object That Does Everything

- Datacenter uses a few big complex objects
  - Movie and Customer objects are the foundation
  - Good choice for a small team and one instance
  - Problematic for large teams and many instances
- False sharing causes tangled dependencies
  - Unproductive re-integration work

*Anti-pattern impacting productivity and availability*



# An Interface For Each Component

- Cloud uses faceted Video and Visitor
  - Basic types hold only the identifier
  - Facets scope the interface you actually need
  - Each component can define its own facets
- No false-sharing and dependency chains
  - Type manager converts between facets as needed
  - `video.asA(PresentationVideo)` for www
  - `video.asA(MerchableVideo)` for middle tier



Response to 2010 talk was a mixture of incomprehension and confusion. Most people thought we were crazy and would be back in our datacenters when it failed...

# **Replacing Datacenter Oracle with Global Apache Cassandra on AWS**

July 11, 2011

Adrian Cockcroft

@adrianco #netflixcloud

<http://www.linkedin.com/in/adriancockcroft>



Get stuck with wrong config

Wait Wait File tickets

Ask permission Wait Wait

Wait Things We Don't Do Wait

Run out of space/power

Plan capacity in advance

Have meetings with IT Wait





Netflix could not  
build new  
datacenters fast  
enough

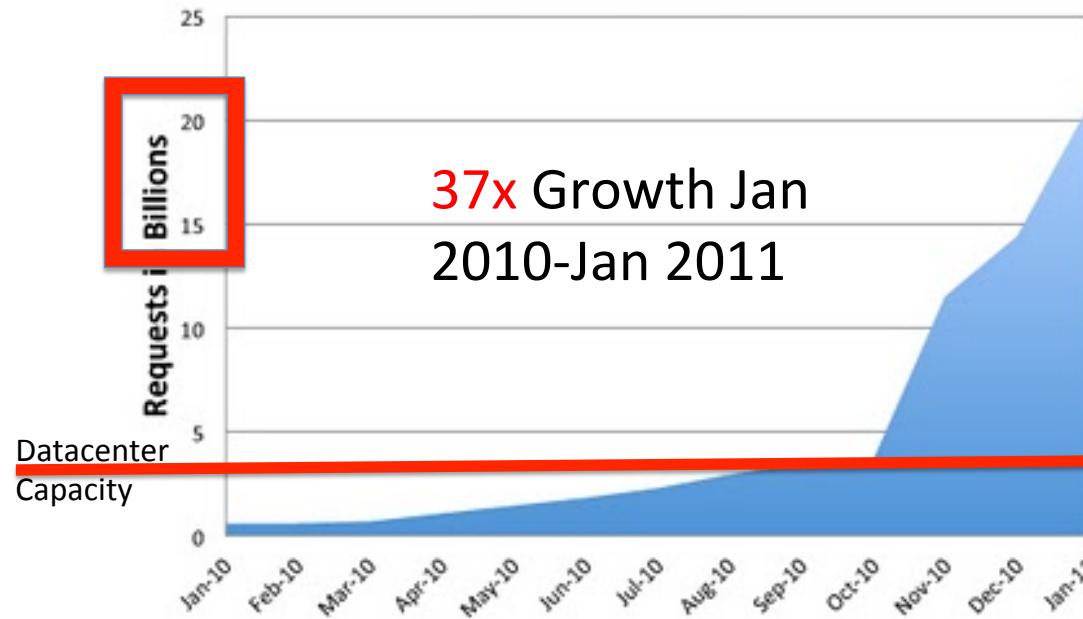
Capacity growth is accelerating, unpredictable  
Product launch spikes - iPhone, Wii, PS3, XBox



# Out-Growing Data Center

<http://techblog.netflix.com/2011/02/redesigning-netflix-api.html>

Netflix API : Growth in Requests



NETFLIX

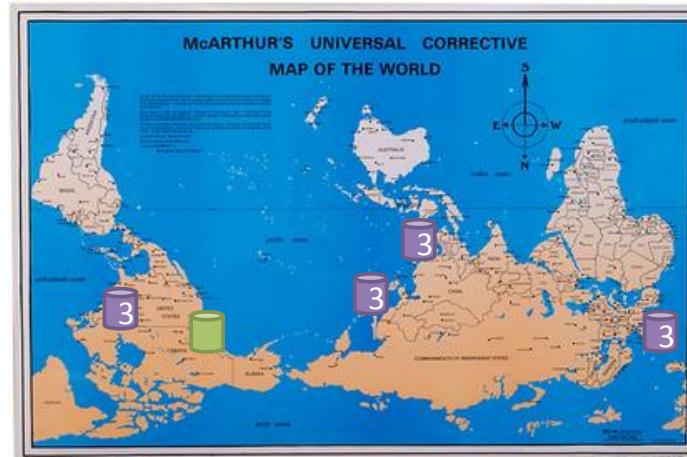
# High Availability

- Cassandra stores 3 local copies, 1 per zone
  - Synchronous access, durable, highly available
  - Read/Write One fastest, least consistent - ~1ms
  - Read/Write Quorum 2 of 3, consistent - ~3ms
- AWS Availability Zones
  - Separate buildings
  - Separate power etc.
  - Close together



# Remote Copies

- Cassandra duplicates across AWS regions
  - Asynchronous write, replicates at destination
  - Doesn't directly affect local read/write latency
- Global Coverage
  - Business agility
  - Follow AWS...
- Local Access
  - Better latency
  - Fault Isolation



NETFLIX

# Cassandra Archive

Appropriate level of paranoia needed...

- Archive could be un-readable
  - Base on restored S3 backup and BI extracted data
- Archive could be stolen
  - Encrypt archive
- AWS East Region could have a problem
  - Copy data to AWS West
- Production AWS Account could have an issue
  - Separate Archive account with no-delete S3 ACL
- AWS S3 could have a global problem
  - Create an extra copy on a different cloud vendor



# Chaos Monkey



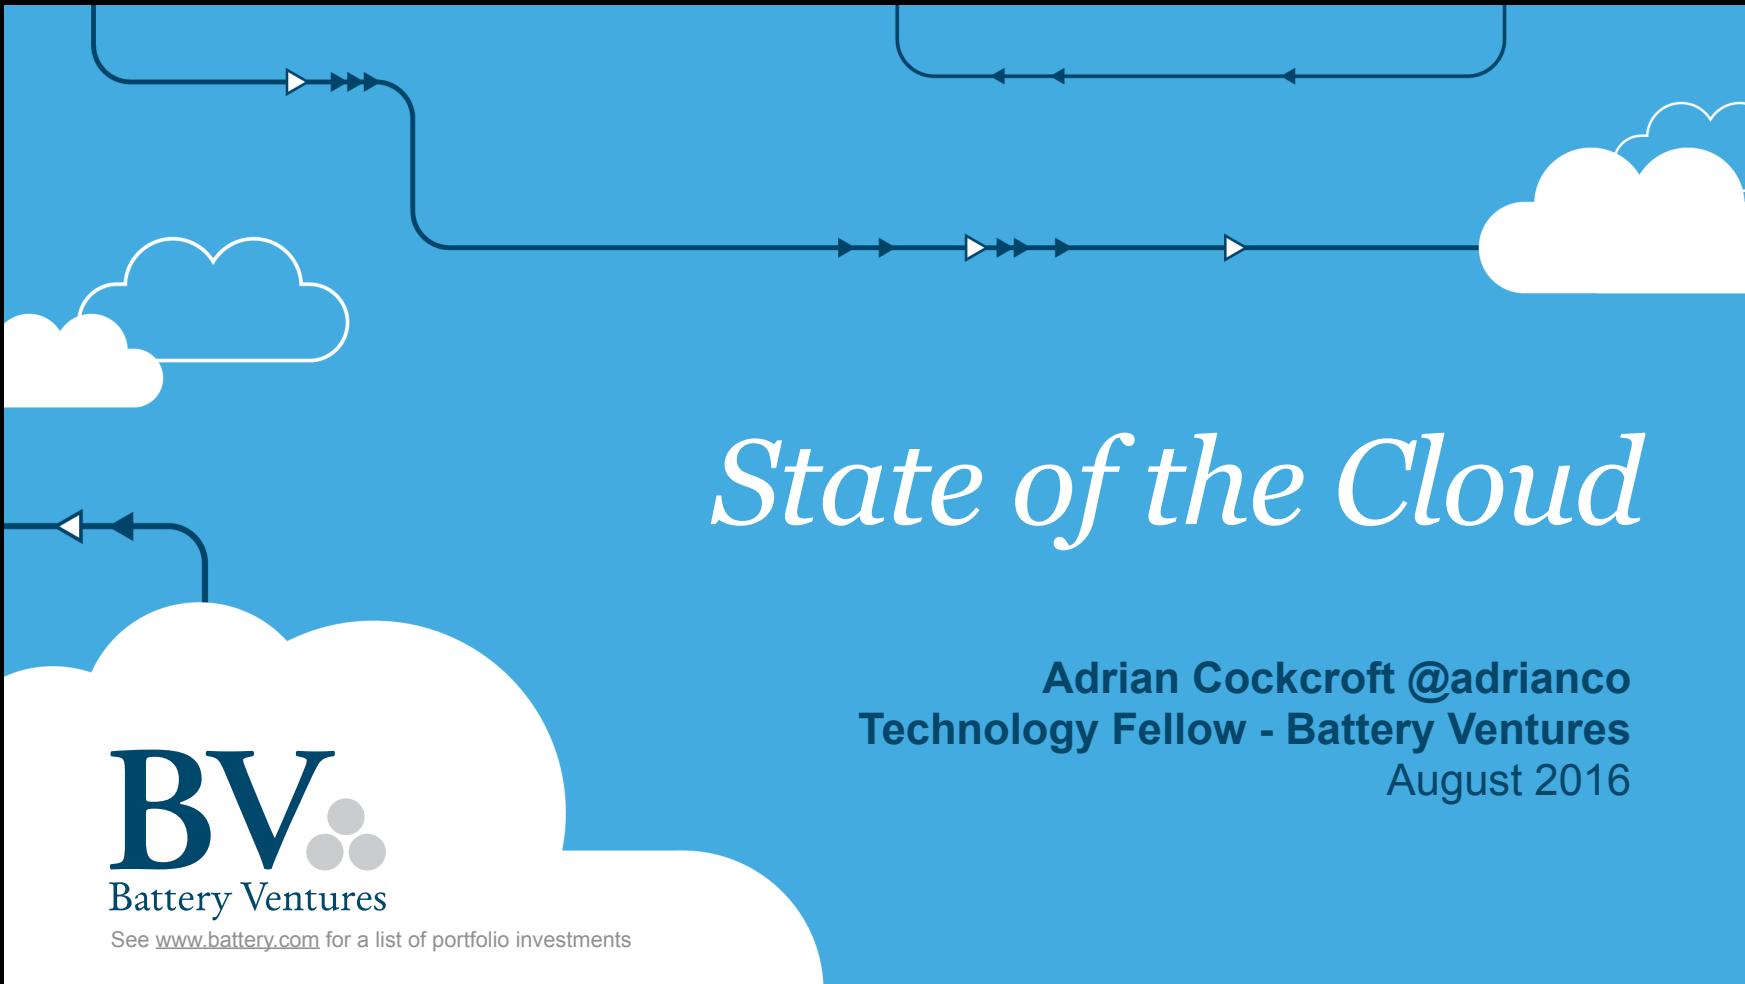
- Make sure systems are resilient
  - Allow any instance to fail without customer impact
- Chaos Monkey hours
  - Monday-Thursday 9am-3pm random instance kill
- Application configuration option
  - Apps now have to opt-out from Chaos Monkey
- Computers (Datacenter or AWS) randomly die
  - Fact of life, but too infrequent to test resiliency

Architecture design control

Be sure you can auto-scale down!

NETFLIX

Response to 2011 progress was that Netflix was a Unicorn,  
and while it might work for us, it wasn't relevant to others

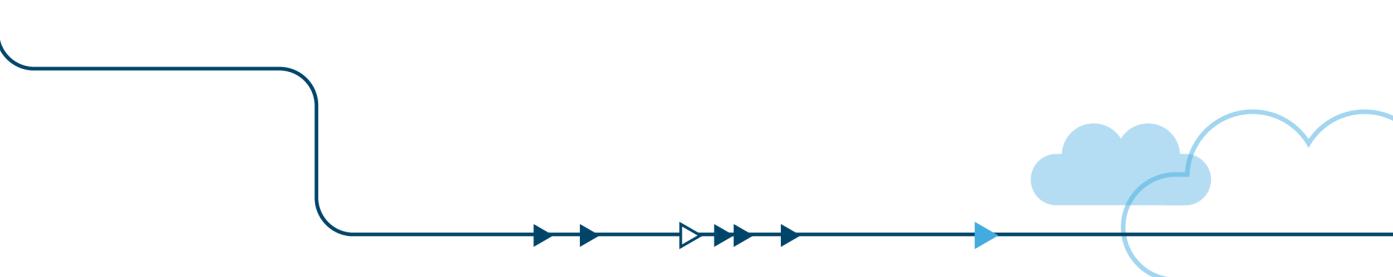


Battery Ventures

See [www.battery.com](http://www.battery.com) for a list of portfolio investments

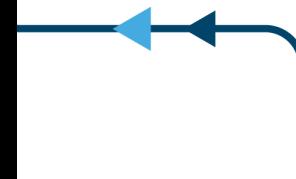
# *State of the Cloud*

Adrian Cockcroft @adrianco  
Technology Fellow - Battery Ventures  
August 2016



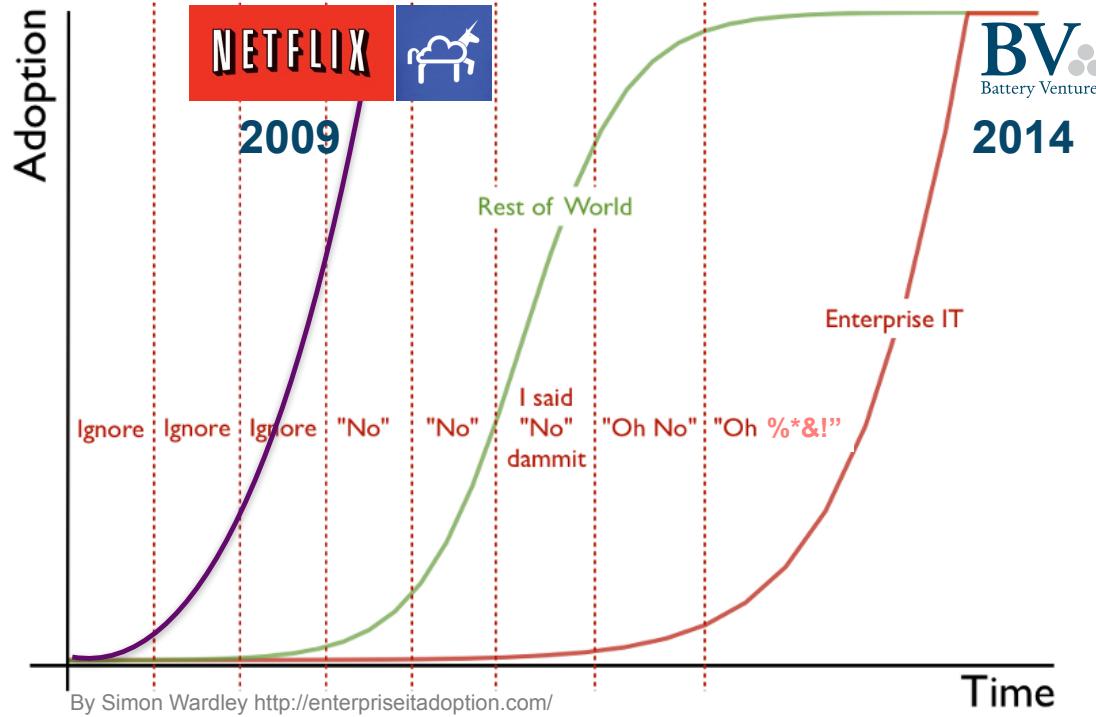
## *Previous Cloud Trend Updates*

*GigaOM Structure May 2014  
D&B Cloud Innovation July 2015  
GigaOM Structure November 2015*



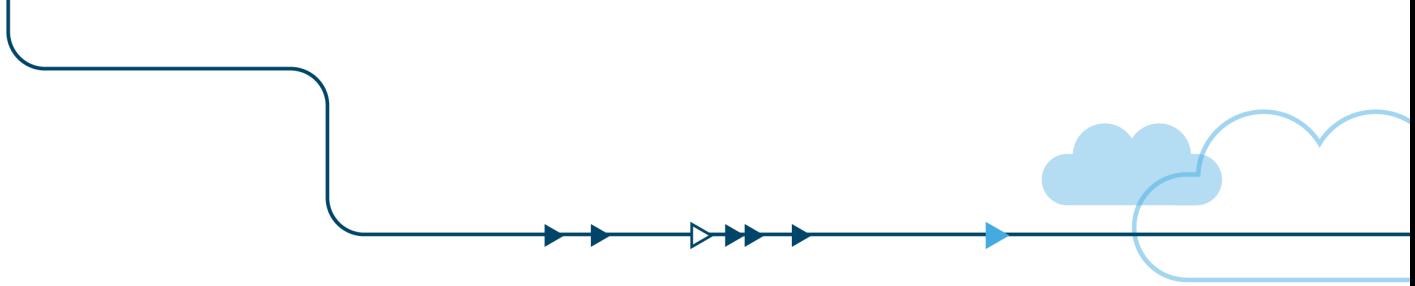
***Trends from 2014: Noted as appropriate***

# Why am I here?



@adrianco's job at the intersection of cloud and Enterprise IT, looking for disruption and opportunities.

Disruptions in 2016 coming from serverless computing and teraservices.



# *In 2014 Enterprises finally embraced public cloud and in 2015 serious deployments are under way.*



Lydia Leong  
@cloudpundit

Oct 2014



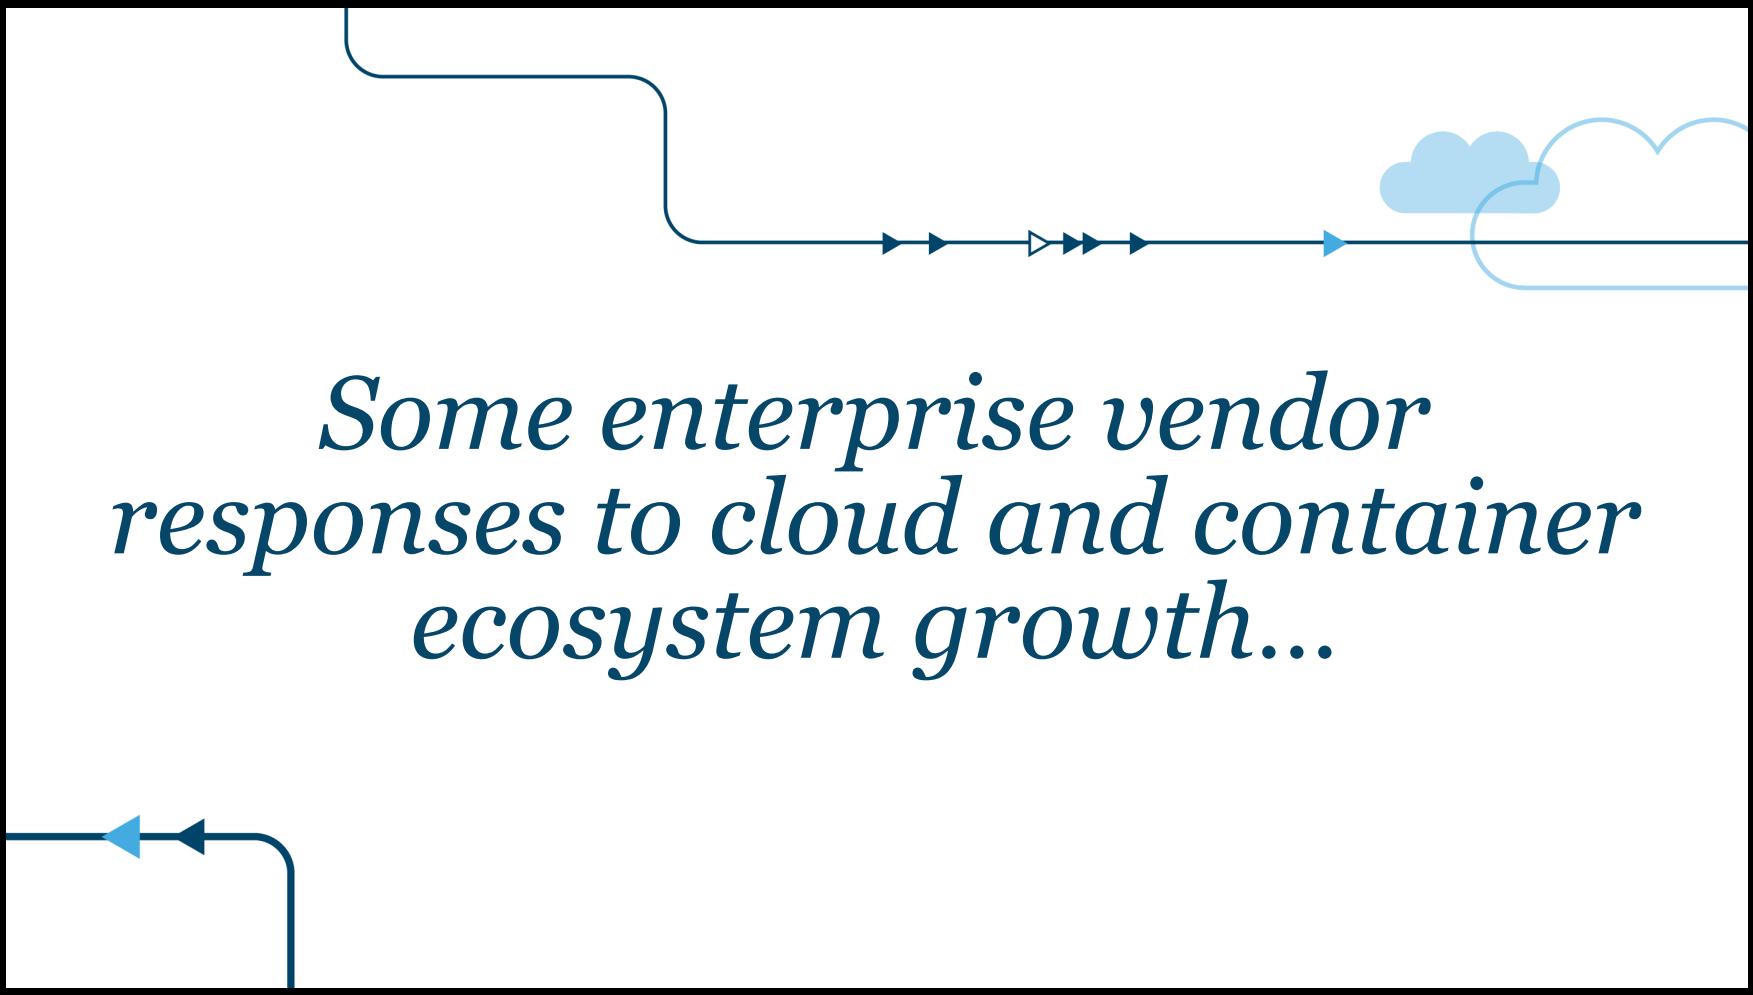
What a difference a year makes. My #GartnerSYM 1:1s this year, everyone's already comfortably using IaaS (overwhelmingly AWS, bit of Azure).

Lydia Leong @cloudpundit · Oct 7 Oct 2015

We're really seeing serious movement of the banks to the cloud at this point. Huge sea change in attitudes.

adrian cockcroft @adrianco

"We can operate more securely on AWS than we can in our own data centers" Rob Alexander of CapitalOne #reinvent



*Some enterprise vendor  
responses to cloud and container  
ecosystem growth...*



**DELL**

*The ship is sinking, let's re-brand as a submarine!*



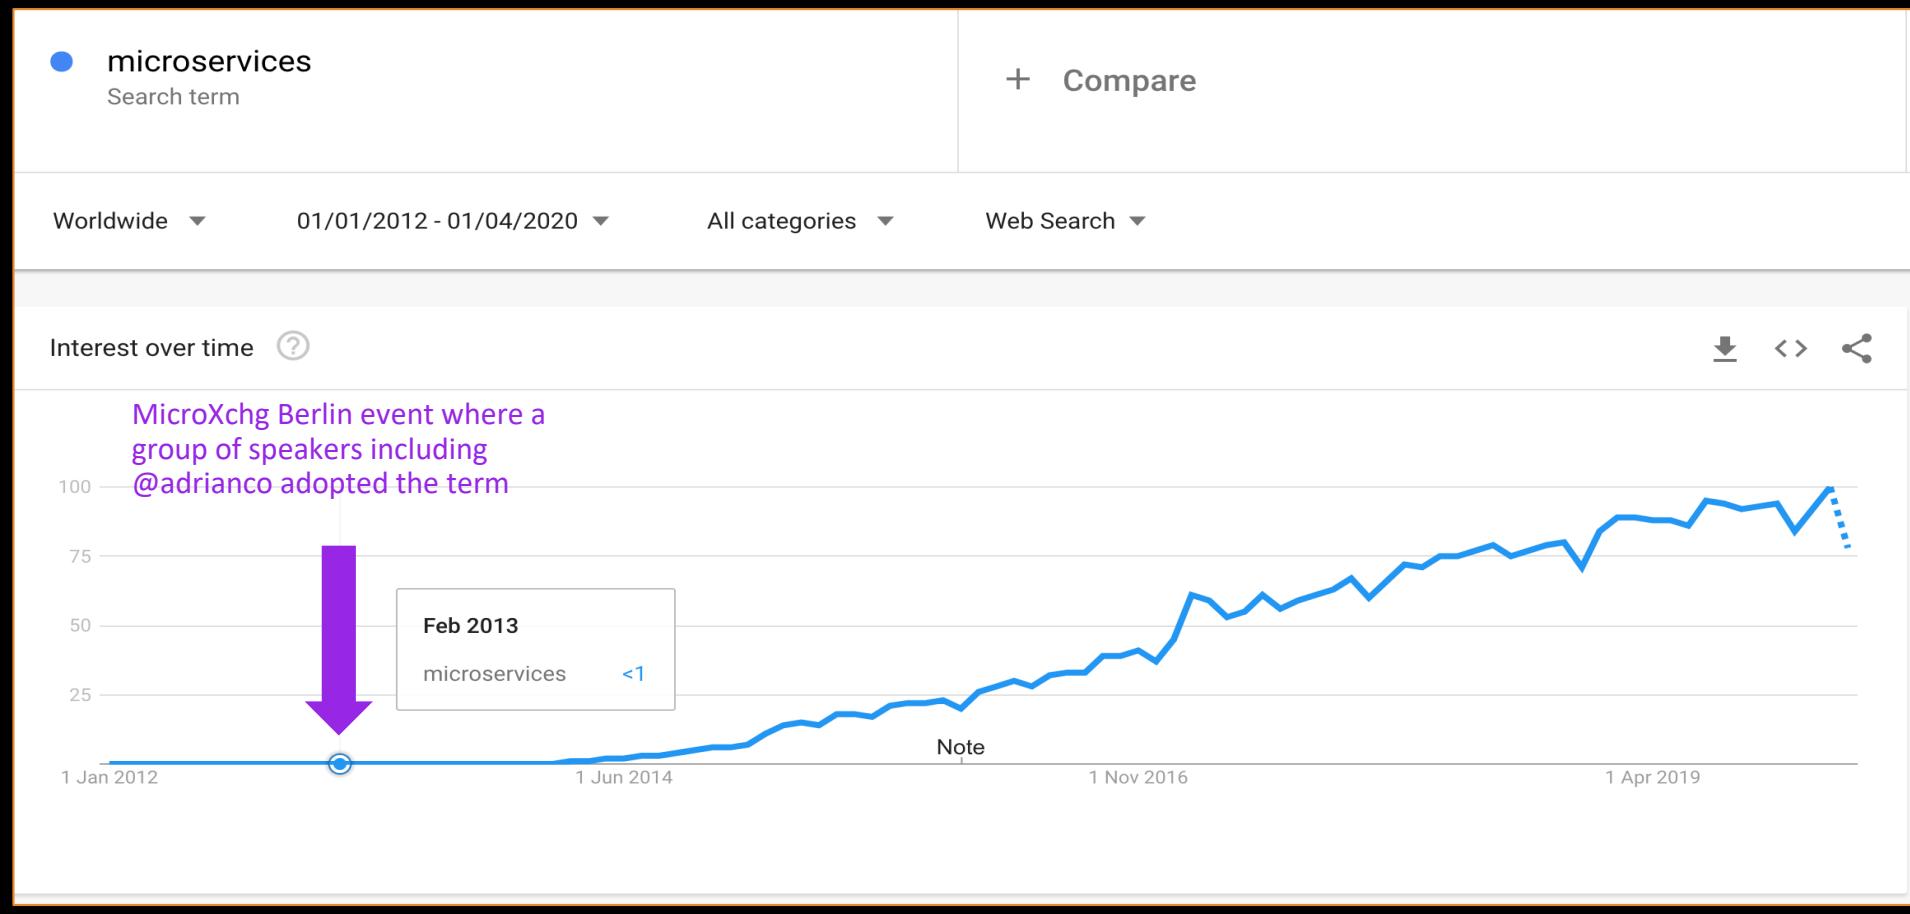
**EMC<sup>2</sup>** *The ship is sinking, let's merge with a submarine!*



*Look! we cut our ship in two really quickly!*

 Hewlett Packard  
Enterprise

# Trends: Microservices





## Typical reactions to my Netflix talks...



“You guys are crazy! Can’t believe it”  
– 2009

“What Netflix is doing won’t work”  
– 2010

It only works for ‘Unicorns’ like Netflix  
– 2011

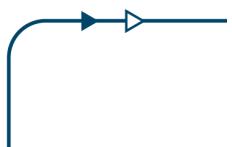
“We’d like to do that but can’t”  
– 2012

“We’re on our way using Netflix OSS code”  
– 2013



## **What I learned from my time at Netflix**



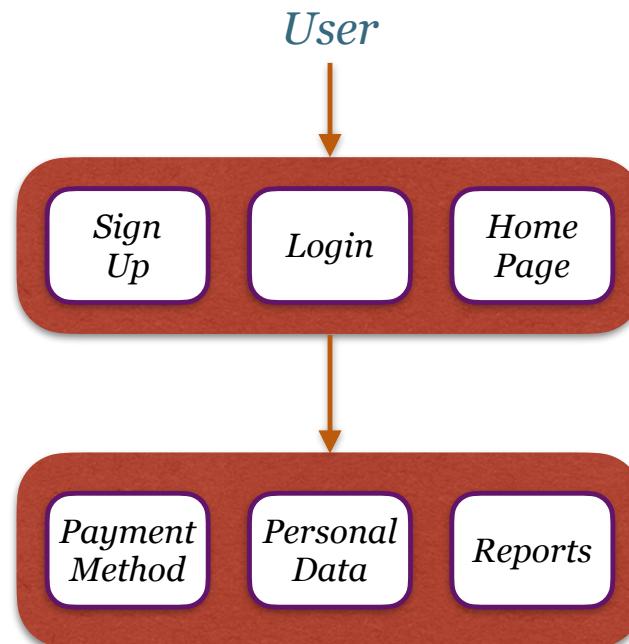
- *Speed wins in the marketplace*
  - *Remove friction from product development*
  - *High trust, low process, no hand-offs between teams*
  - *Freedom and responsibility culture*
  - *Don't do your own undifferentiated heavy lifting*
  - *Use simple patterns automated by tooling*
  - *Self service cloud makes impossible things instant*
- 

## Example Monolith:

*Monolithic application*

*Complex mix of queries*

*Monolithic “kitchen sink” database*



*Because one part of the monolithic application and database holds sensitive data all of it is subject to the most rigorous policies*

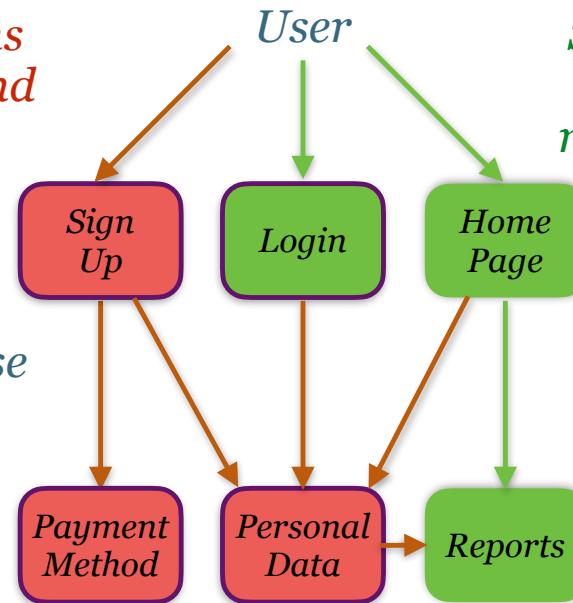
## Microservices version:

*Segregated team owns  
secure data sources and  
infrequent updates*

*Microservices  
separation of concerns*

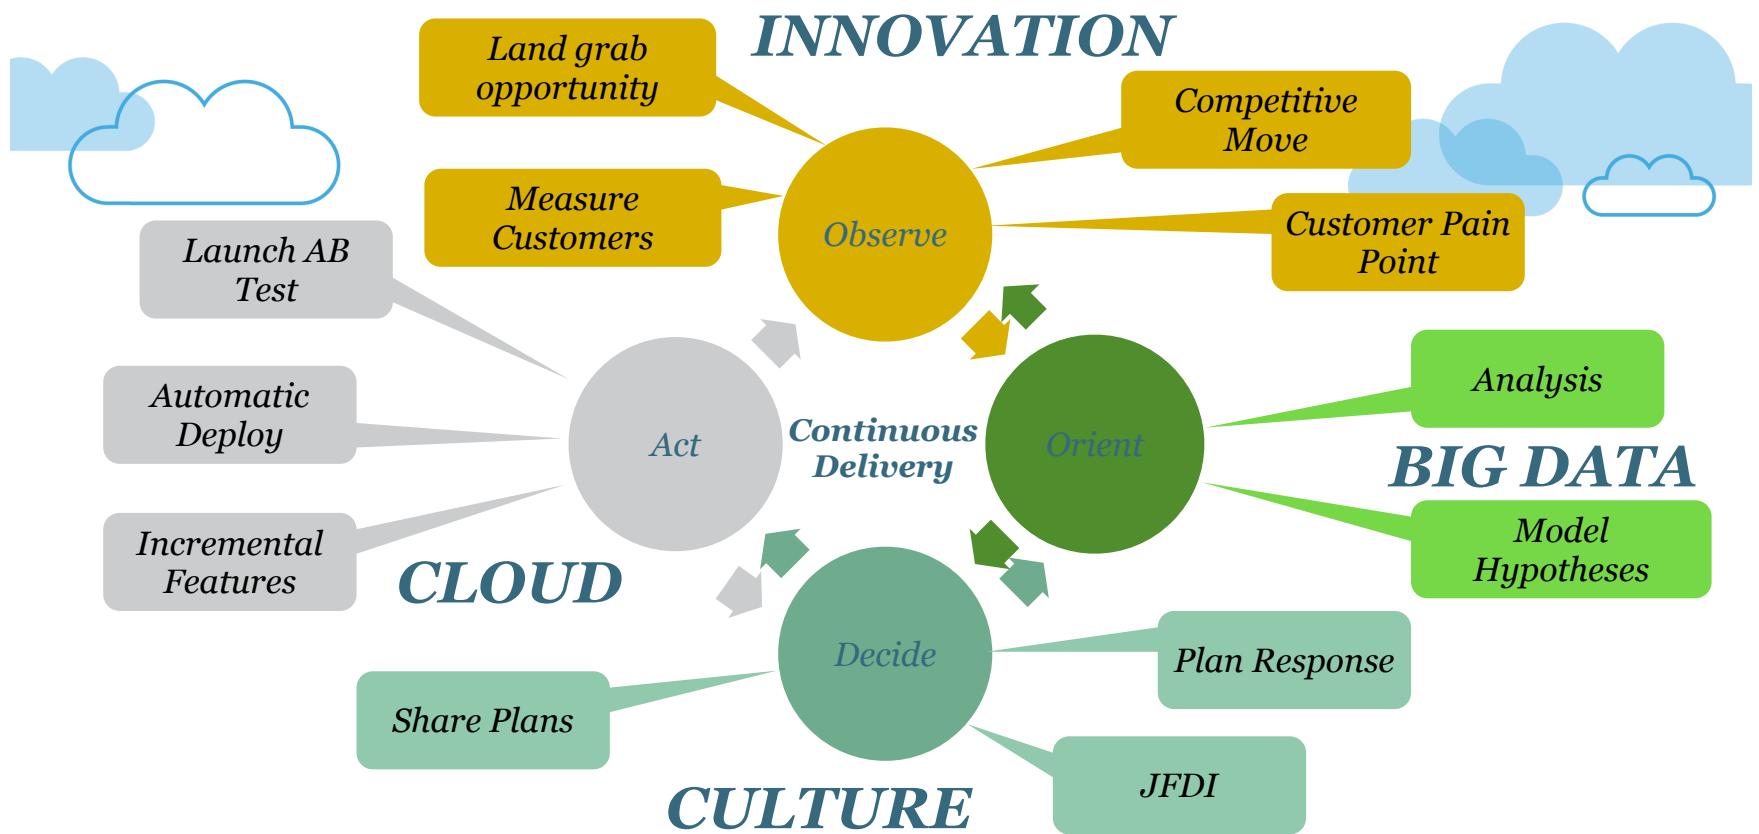
*Isolated single purpose  
connections*

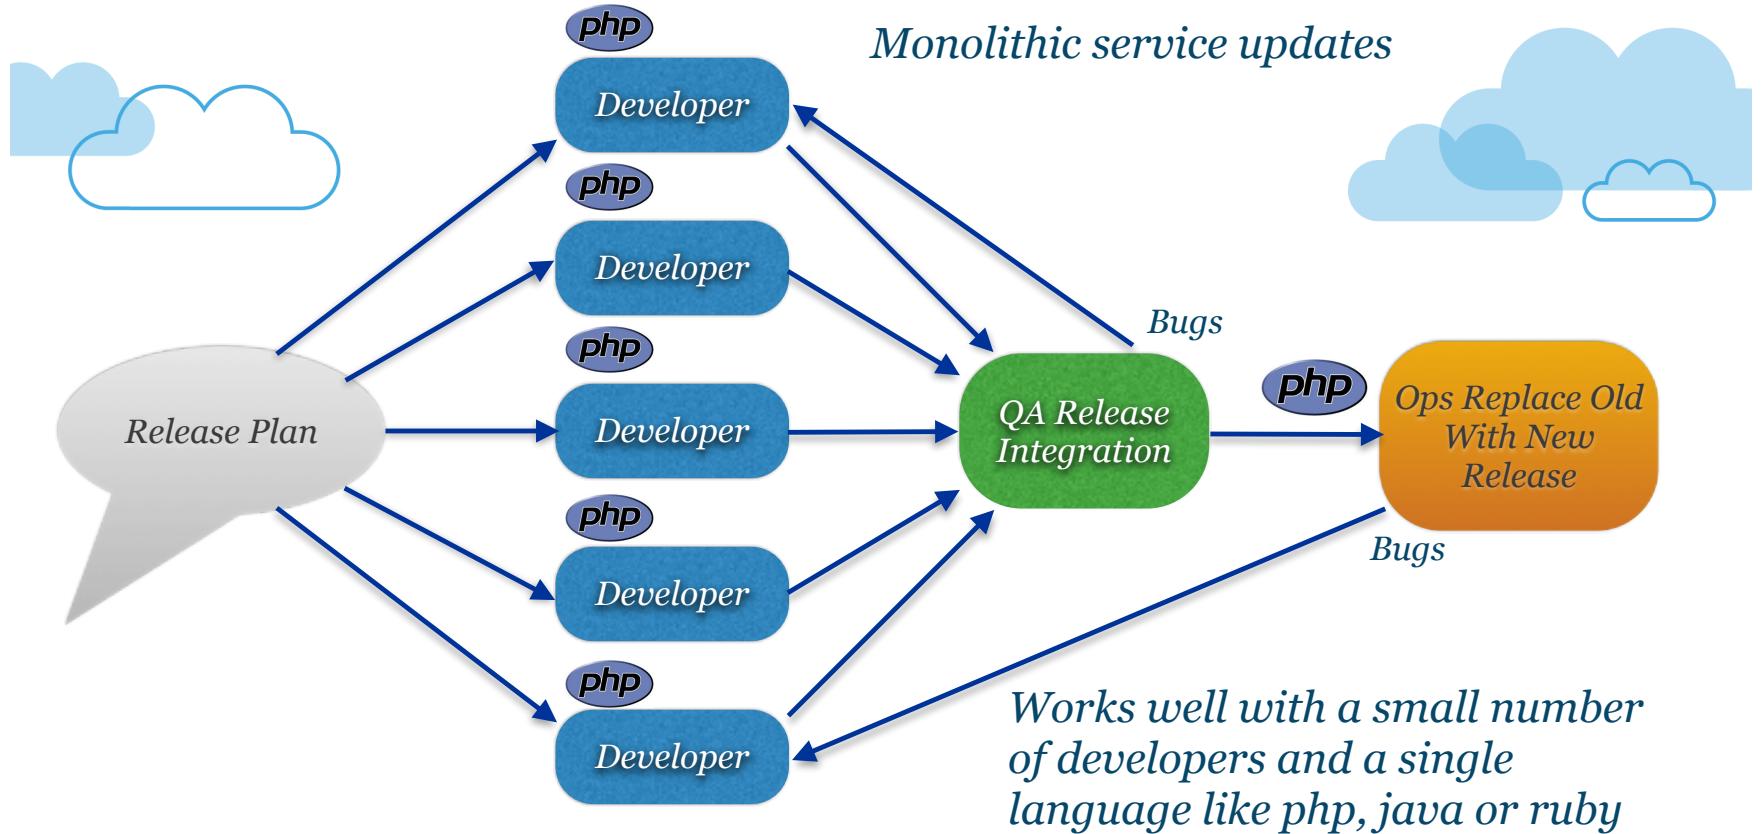
*Optimized  
datastores*

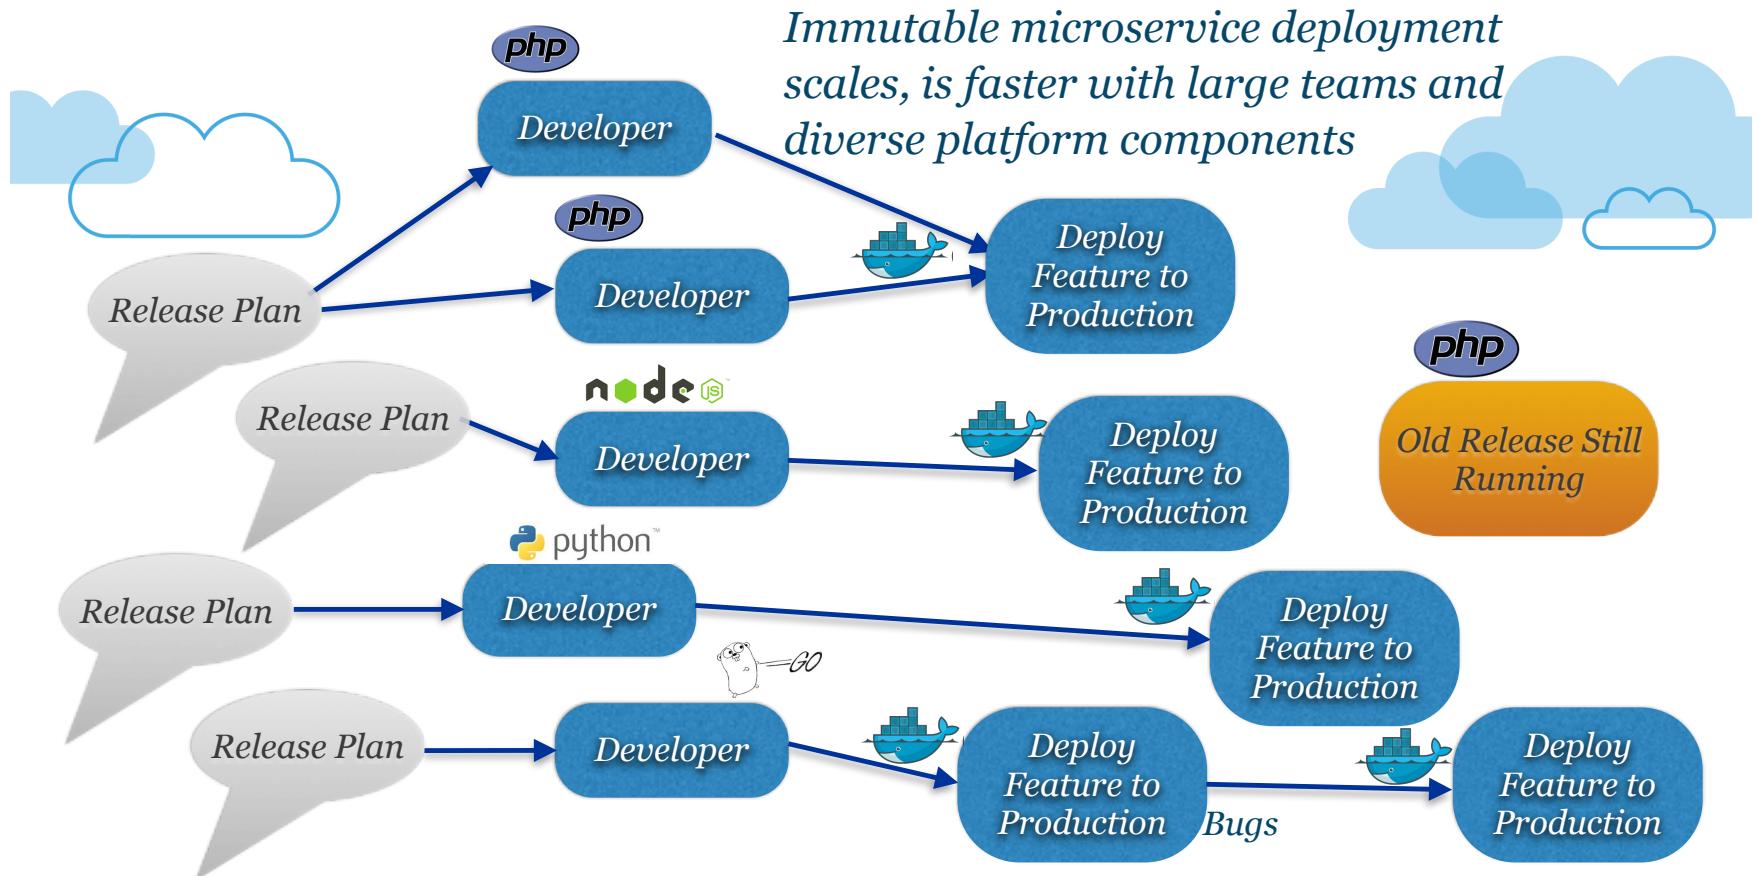


*Segregated team owns  
rapid improvement of  
most common use cases*

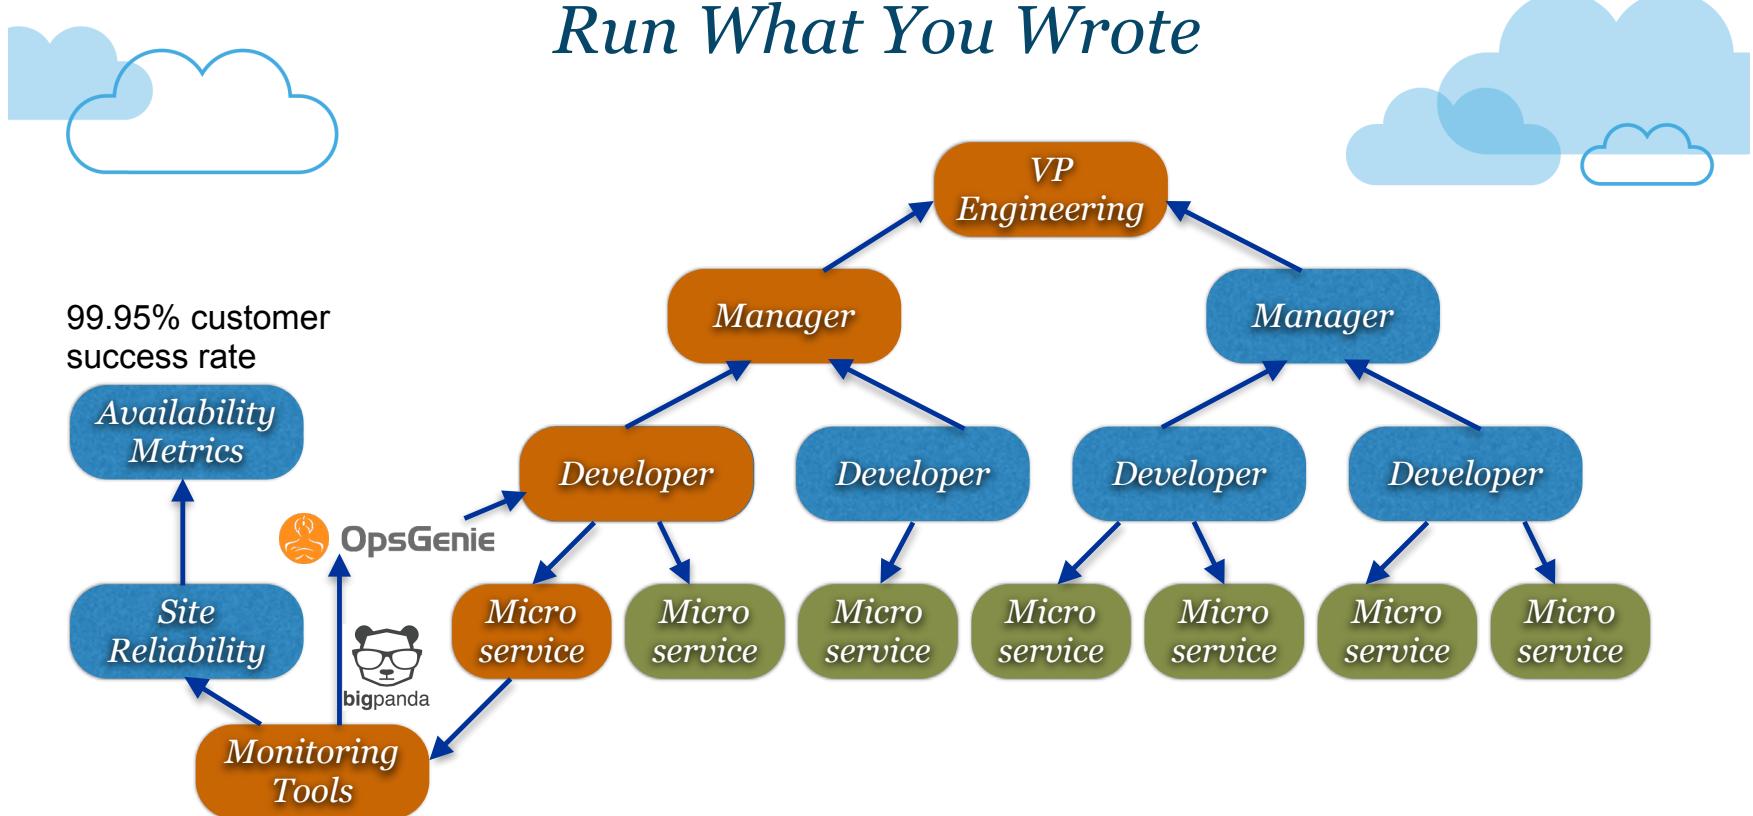
*Because each  
microservice can  
conform to the  
appropriate policy,  
demands for agility  
can be separated  
from requirements  
for security*







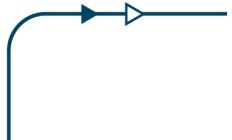
# *Run What You Wrote*



## Non-Destructive Production Updates

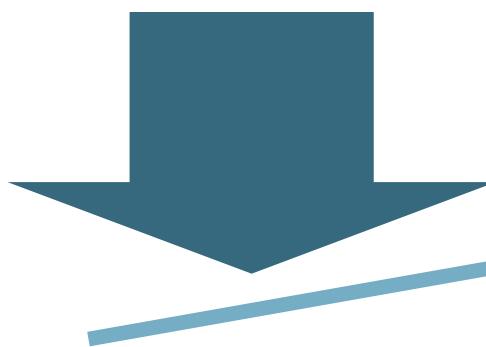


- “*Immutable Code*” Service Pattern
  - *Existing services are unchanged, old code remains in service*
  - *New code deploys as a new service group*
  - *No impact to production until traffic routing changes*
- *A|B Tests, Feature Flags and Version Routing control traffic*
  - *First users in the test cell are the developer and test engineers*
  - *A cohort of users is added looking for measurable improvement*

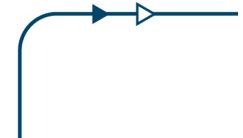
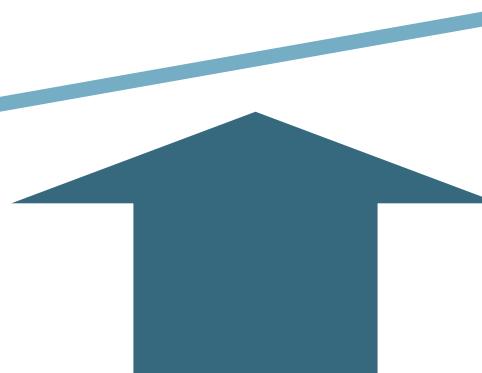


# What Happened?

*Rate of change  
increased*



*Cost and size and  
risk of change  
reduced*



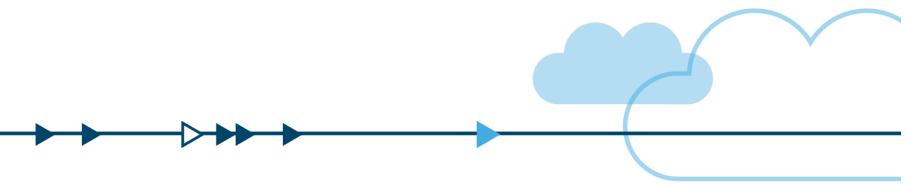
## **It's what you know that isn't so**



- *Make your assumptions explicit*
- *Extrapolate trends to the limit*
- *Listen to non-customers*
- *Follow developer adoption, not IT spend*
- *Map evolution of products to services to utilities*
- *Re-organize your teams for speed of execution*

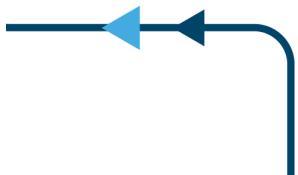


*If every service has to be updated at the same time it's not loosely coupled*



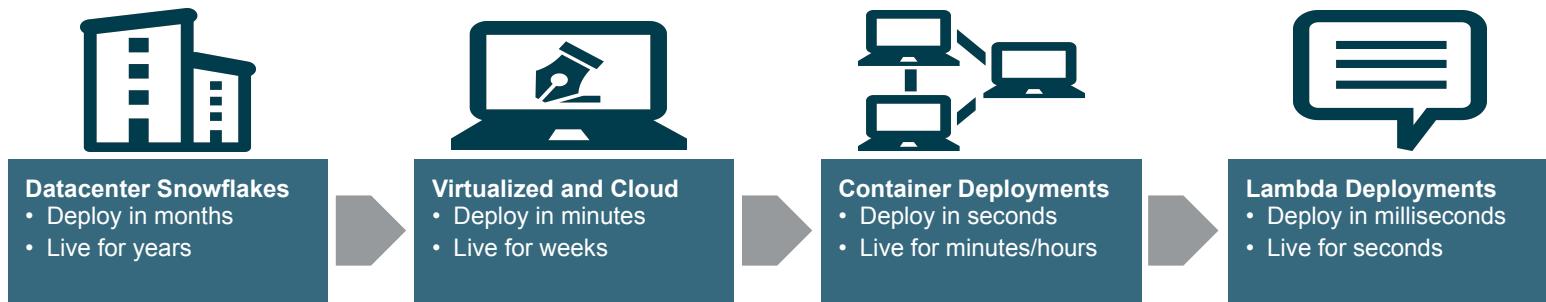
## *A Microservice Definition*

*Loosely coupled service oriented architecture with bounded contexts*

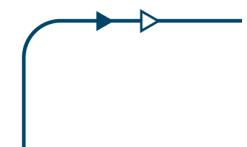


*If you have to know too much about surrounding services you don't have a bounded context. See the Domain Driven Design book by Eric Evans.*

## Speeding Up The Platform



► *AWS Lambda is leading exploration of serverless architectures in 2016*

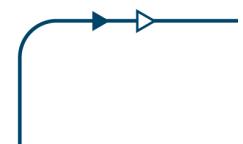


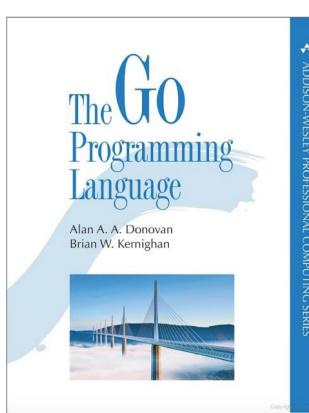
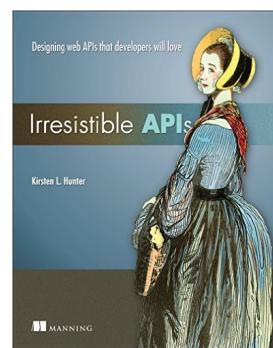
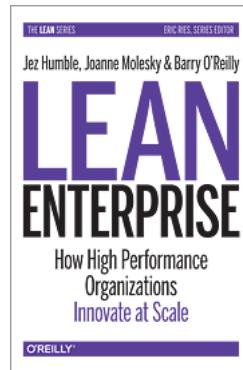
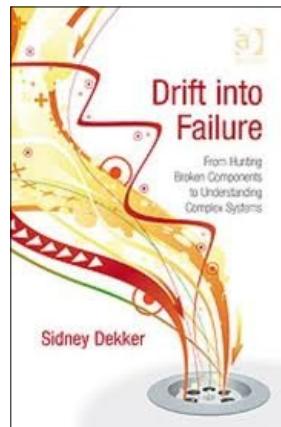
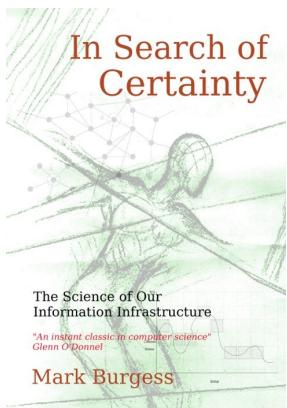
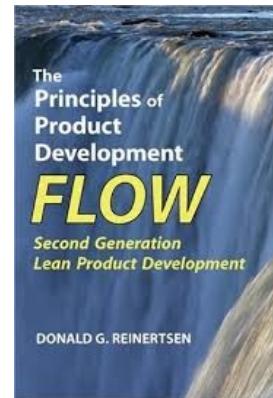
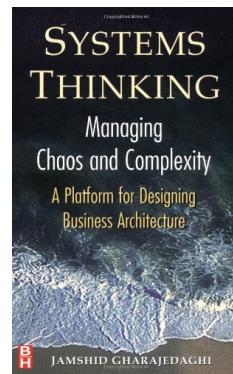
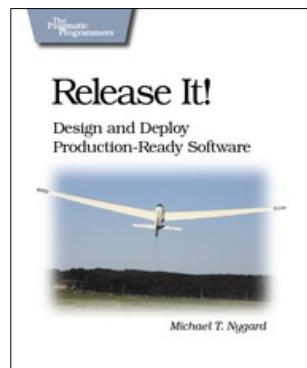
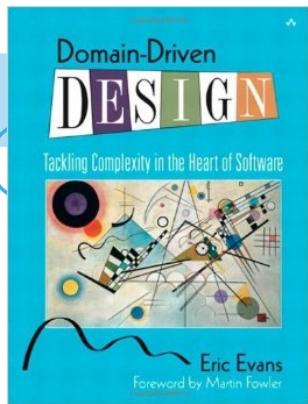
## Separate Concerns with Microservices

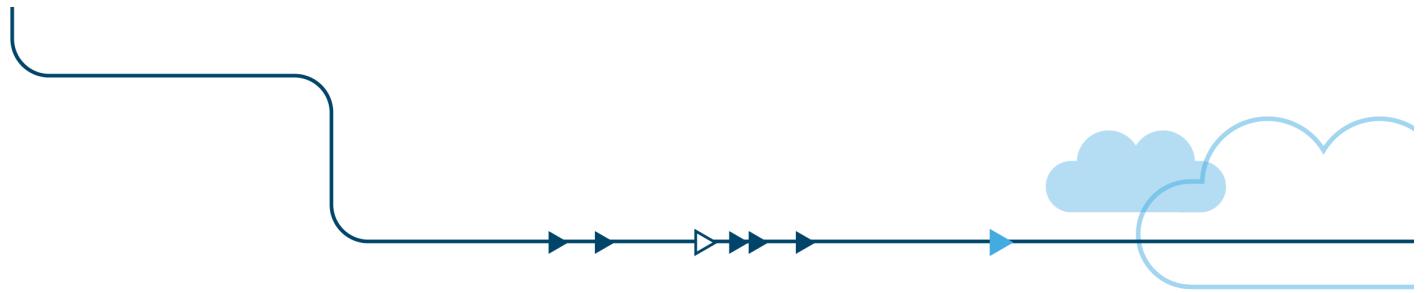


- *Invert Conway's Law – teams own service groups and backend stores*
- *One “verb” per single function micro-service, size doesn't matter*
- *One developer independently produces a micro-service*
- *Each micro-service is its own build, avoids trunk conflicts*
- *Deploy in a container: Tomcat, AMI or Docker, whatever...*
- *Stateless business logic. Cattle, not pets.*
- *Stateful cached data access layer using replicated ephemeral instances*

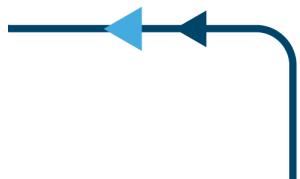
[http://en.wikipedia.org/wiki/Conway's\\_law](http://en.wikipedia.org/wiki/Conway's_law)







*What's Missing?*





# *Advanced Microservices Topics*



*Failure injection testing*

*Versioning, routing*

*Binary protocols and interfaces*

*Timeouts and retries*

*Denormalized data models*

*Monitoring, tracing*

*Simplicity through symmetry*





# *Timeouts and Retries*



*Connection timeout vs. request timeout confusion*

*Usually setup incorrectly, global defaults*

*Systems collapse with “retry storms”*

*Timeouts too long, too many retries*

*Services doing work that can never be used*



# Timeouts and Retries

*Bad config: Every service defaults to 2 second timeout, two retries*



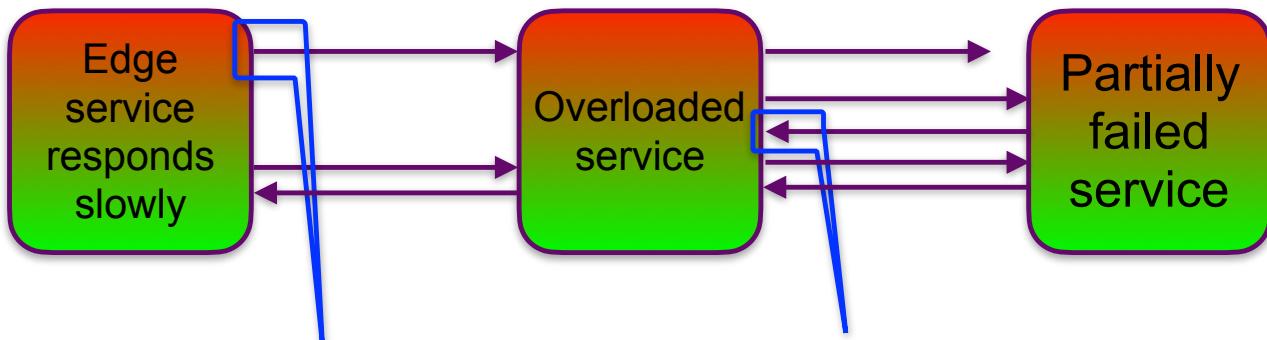
*If anything breaks, everything upstream stops responding*



*Retries add unproductive work*

# Timeouts and Retries

*Bad config: Every service defaults to 2 second timeout, two retries*



*First request from Edge timed out so it ignores the successful response and keeps retrying. Middle service load increases as it's doing work that isn't being consumed*



# *Timeout and Retry Fixes*



*Cascading timeout budget  
Static settings that decrease from the edge  
or dynamic budget passed with request*

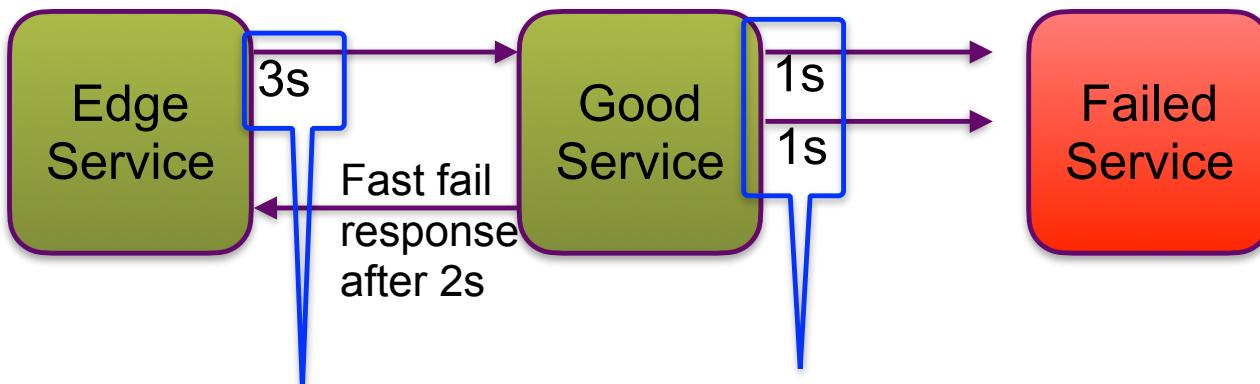
*How often do retries actually succeed?  
Don't ask the same instance the same thing  
Only retry on a different connection*



# *Timeouts and Retries*



*Budgeted timeout, one retry*

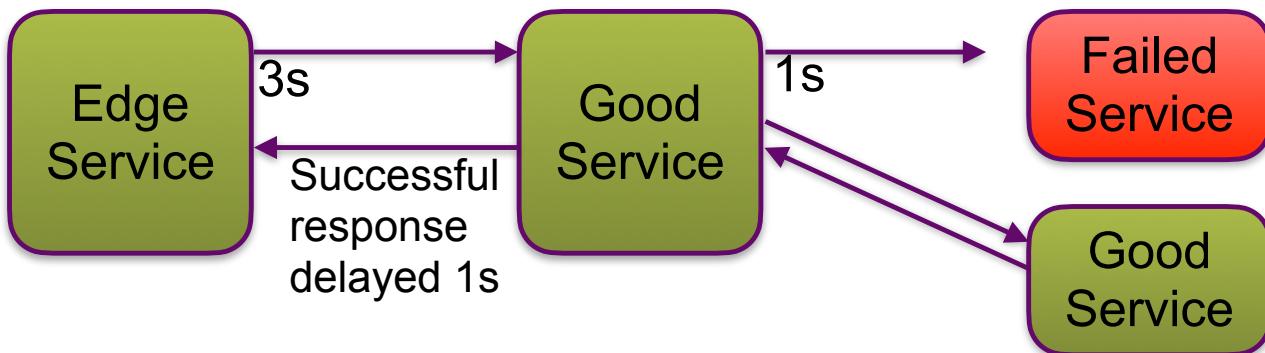


*Upstream timeout must always be longer than total downstream timeout \* retries delay*

*No unproductive work while fast failing*

# *Timeouts and Retries*

*Budgeted timeout, failover retry*



*For replicated services with multiple instances  
never retry against a failed instance*

*No extra retries or unproductive work*

# *Communicating Sequential Goroutines*

Adrian Cockcroft @adrianco  
Technology Fellow - Battery Ventures  
July 2016



# Agenda

- 1978 *Communicating Sequential Processes*
- 1983 *Occam*
  - How Channels Work*
- 1992 *Pi-Calculus*
- 2013 *The Life of Occam-Pi*
- 2006 *Occam-Pi based simulation*
  - Pi-Calculus ideas in Go*
  - Go Applications*

Programming  
Techniques

S. L. Graham, R. L. Rivest  
Editors

# Communicating Sequential Processes

C.A.R. Hoare  
The Queen's University  
Belfast, Northern Ireland

---

This paper suggests that input and output are basic primitives of programming and that parallel composition of communicating sequential processes is a fundamental program structuring method. When combined with a development of Dijkstra's guarded command, these concepts are surprisingly versatile. Their use is illustrated by sample solutions of a variety of familiar programming exercises.

**Key Words and Phrases:** programming, programming languages, programming primitives, program structures, parallel programming, concurrency, input, output, guarded commands, nondeterminacy, coroutines, procedures, multiple entries, multiple exits, classes, data representations, recursion, conditional critical regions, monitors, iterative arrays

**CR Categories:** 4.20, 4.22, 4.32

*“...the concepts and notations introduced in this paper (although described in the next section in the form of a programming language fragment) should not be regarded as suitable for use as a programming language, either for abstract or for concrete programming. They are at best only a partial solution to the problems tackled.”*

## *CSP Issues:*

*Not a full language  
Hard to read  
Process addressing*

# *David May's Occam Language*

*Extremely simple and elegant  
implementation of CSP as a language*

*Adds named channels  
Designed as the assembly language for Transputer hardware*

Occam is intended to be the smallest language which is adequate for its purpose; however, suggestions for further simplification would be welcome.

by David May

INMOS Limited  
Whitefriars  
Lewins Mead  
Bristol BS1 2NP  
UK  
(0272) 290861  
Telex 444723

INMOS Corporation  
P O Box 16000  
Colorado Springs  
Colorado 80935  
USA  
(303) 630-4000  
TWX 910 920 4904

1. Introduction

VLSI technology allows a whole microcomputer with memory, processor and communications to be constructed as a single device. Such a device can be used as a component in the construction of high performance systems. It provides a 'building block' for such systems, in the same way that logic gates provide a 'building block' for digital systems.

New programming languages are needed to allow systems consisting of many interconnected microcomputers to be designed and programmed. Occam is such a language, developed by INMOS. It is a simple language based on the concepts of concurrency and communication. It may be used both to describe the structure of a system in terms of connected microcomputers, and to program the individual microcomputers.

An initial version of occam has been produced, and this is described below. Further information, and an implementation of occam, is available from INMOS.

SIGPLAN Notices, Vi8 #4, April, 1983

# Prime Number Sieve Translated from CSP to Occam

```

PROC eratosthenes(CHAN output)
CHAN sieve[101], print[102]:
PAR
SEQ
    print[0]!2
    SEQ n = [1 FOR 4999]
        sieve[0]!(n+n)+1
        sieve[0]!-1
    PAR i = [1 FOR 100]
    var p, mp, m:
    SEQ
        sieve[i-1]? p
        print[i]! p
        mp := p
        WHILE m <> -1
            SEQ
                sieve[i-1]? m
                WHILE m > mp
                    mp := mp + p
                IF
                    m = mp
                    SKIP
                    m < mp
                    sieve[i]! m
    var n:
    SEQ
        sieve[100]? n
        print[101]! n
        sieve[100]? ANY
    var n:
    SEQ i= [0 FOR 101]
    SEQ
        sieve[i]? n
        output! n
:

```

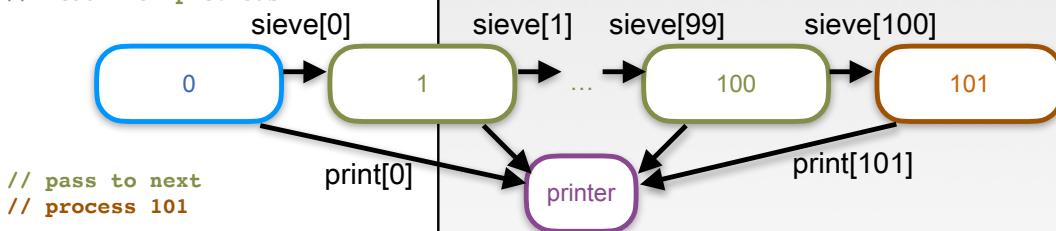
// process 0  
// start primes with 2  
// for n=3;n<10000;n+=2 {send n}  
// explicit evaluation order  
// send -1 as terminator  
// create 100 sieve processes

// first thing read is prime  
// send it to be printed  
// mp is a multiple of p  
// look for terminator

// read from previous

// pass to next  
// process 101

// last prime in line  
// discard terminator  
// printer process



*Naming shifts from processes to channels*

# *Comparing Occam and Go*

## *Parallel Channel Assignment*

```
PROC main(CHAN out)
  VAR x,y:
  SEQ
    x := 1
    CHAN c:
    PAR
      c!x
      c?y
      out!y
:
```

```
func main() {
  var x, y int
  x = 1
  c := make(chan int)
  var wg sync.WaitGroup
  wg.Add(2)
  go func() { defer wg.Done(); c <- x }()
  go func() { defer wg.Done(); y = <-c }()
  wg.Wait()
  fmt.Println(y)
}
```

*Pi-Calculus*  
*Robin Milner 1992*

# A Calculus of Mobile Processes, I

ROBIN MILNER

*University of Edinburgh, Scotland*

JOACHIM PARROW

*Swedish Institute of Computer Science, Kista, Sweden*

AND

DAVID WALKER

*University of Warwick, England*

We present the  $\pi$ -calculus, a calculus of communicating systems in which one can naturally express processes which have changing structure. Not only may the component agents of a system be arbitrarily linked, but a communication between neighbours may carry information which changes that linkage. The calculus is an extension of the process algebra CCS, following work by Engberg and Nielsen, who added mobility to CCS while preserving its algebraic properties. The  $\pi$ -calculus gains simplicity by removing all distinction between variables and constants; communication links are identified by *names*, and computation is represented purely as the communication of names across links. After an illustrated description of how the  $\pi$ -calculus generalises conventional process algebras in treating mobility, several examples exploiting mobility are given in some detail. The important examples are the encoding into the  $\pi$ -calculus of higher-order functions (the  $\lambda$ -calculus and combinatory algebra), the transmission of processes as values, and the representation of data structures as processes. The paper continues by presenting the algebraic theory of *strong bisimilarity* and *strong equivalence*, including a new notion of equivalence indexed by *distinctions*—i.e., assumptions of inequality among names. These theories are based upon a semantics in terms of a labeled transition system and a notion of *strong bisimulation*, both of which are expounded in detail in a companion paper. We also report briefly on work-in-progress based upon the corresponding notion of *weak bisimulation*, in which internal actions cannot be observed. © 1992 Academic Press, Inc.

*We present the  $\pi$ -calculus, a calculus of communicating systems in which one can naturally express processes which have changing structure. Not only may the component agents of a system be arbitrarily linked, but a communication between neighbours may carry information which changes that linkage.*

In this paper we do not present the basic semantics of the calculus; this is done in our companion paper (Milner, Parrow, and Walker, 1989), in the same style as in CCS, namely as a labeled transition system defined by structural inference rules. In that paper the notions of strong bisimulation and strong equivalence are also defined; the latter is a congruence relation, so it may be understood as (strong) semantic equality. Here, we shall rely somewhat upon analogy with the transitions of CCS agents. In particular, we assume simple transitions such as

$$(\dots + \bar{y}x.P + \dots) | (\dots + y(z).Q + \dots) \xrightarrow{\tau} P|Q\{x/z\}$$

and simple equations such as

$$(y)(\bar{y}x.P | y(z).Q) = \tau.(y)(P | Q\{x/z\})$$

*It's easy to show that...*

*This paper is incomprehensible!*

## *A triumph of notation over comprehension.*

*Simple equations such as...*

26

MILNER, PARROW, AND WALKER

It is illuminating to see how the encoding of a particular example behaves. Consider  $(\lambda xx)N$ ; first, we have

$$\llbracket \lambda xx \rrbracket v \equiv v(x)(w).\bar{x}w.$$

So, assuming  $x$  not free in  $N$ ,

$$\begin{aligned} \llbracket (\lambda xx)N \rrbracket u &\equiv (v)(\llbracket \lambda xx \rrbracket v | (x)\bar{v}xu.x(w).\llbracket N \rrbracket w) \\ &\equiv (v)(v(x)(w).\bar{x}w | (x)\bar{v}xu.x(w).\llbracket N \rrbracket w) \\ &= \tau.(v)(x)(v(w).\bar{x}w | \bar{v}u.x(w).\llbracket N \rrbracket w) \\ &= \tau.\tau.(v)(x)(\bar{x}u | x(w).\llbracket N \rrbracket w) \\ &= \tau.\tau.\tau.(v)(x)(\mathbf{0} | \llbracket N \rrbracket u) = \tau.\tau.\tau.\llbracket N \rrbracket u. \end{aligned}$$

More generally, it is easy to show that

$$\llbracket (\lambda xM)N \rrbracket u \approx \llbracket M\{N/x\} \rrbracket u, \quad (36)$$

## Life of occam-Pi

Peter H. WELCH

*School of Computing, University of Kent, UK*

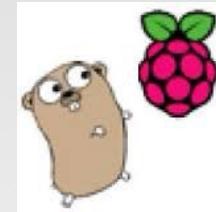
*p.h.welch@kent.ac.uk*

**Abstract.** This paper considers some questions prompted by a brief review of the history of computing. Why is programming so hard? Why is concurrency considered an “advanced” subject? What’s the matter with *Objects*? Where did all the Maths go? In searching for answers, the paper looks at some concerns over fundamental ideas within *object orientation* (as represented by modern programming languages), before focussing on the concurrency model of communicating processes and its particular expression in the *occam* family of languages. In that focus, it looks at the history of *occam*, its underlying philosophy (*Ockham’s Razor*), its semantic foundation on Hoare’s CSP, its principles of *process oriented design* and its development over almost three decades into *occam- $\pi$*  (which blends in the concurrency dynamics of Milner’s  $\pi$ -calculus). Also presented will be an urgent need for rationalisation – *occam- $\pi$*  is an experiment that has demonstrated significant results, but now needs time to be spent on careful review and implementing the conclusions of that review. Finally, the future is considered. In particular, is there a future?

**Keywords.** process, object, local reasoning, global reasoning, *occam-pi*, concurrency, compositionality, verification, multicore, efficiency, scalability, safety, simplicity

*...looks at the history of occam, its underlying philosophy (*Ockham’s Razor*), its semantic foundation on Hoare’s CSP, its principles of process oriented design and its development over almost three decades into occam- $\pi$  (which blends in the concurrency dynamics of Milner’s  $\pi$ -calculus).*

*Go-π*



*Dynamic Channel Protocol  
Actor Pattern  
Partitioned Service Registry  
Logging and Tracing*

# *Dynamic Channel Protocol*

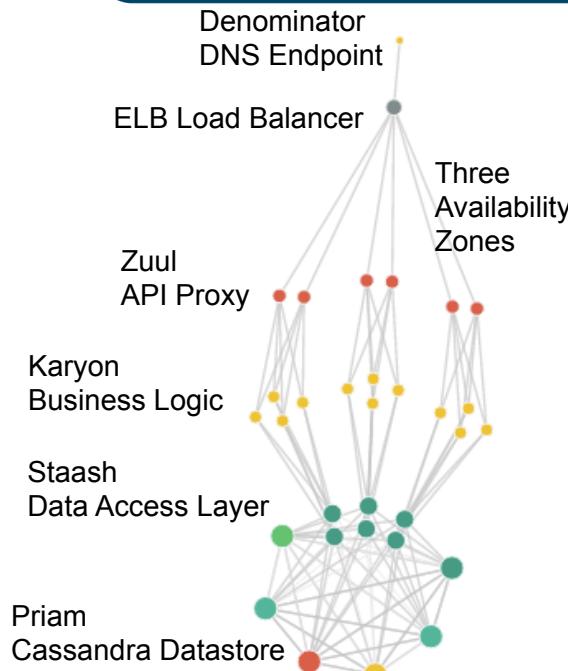
<https://github.com/adrianco/spigo/tree/master/tooling/gotocol>

Imposition/Intention? [https://en.wikipedia.org/wiki/Promise\\_theory](https://en.wikipedia.org/wiki/Promise_theory)

```
ch <-gotocol.Message{gotocol.GetRequest, listener, now, ctx, "why?"}

// Message structure used for all messages, includes a channel of itself
type Message struct {
    Imposition  Impositions // request type
    ResponseChan chan Message // place to send response messages
    Sent         time.Time   // time at which message was sent
    Ctx          Context     // message context
    Intention    string      // payload
}
```

# Simulated Microservices



*Model and visualize microservices  
Simulate interesting architectures  
Generate large scale configurations  
Eventually stress test real tools*

*Code: [github.com/adrianco/spigo](https://github.com/adrianco/spigo)  
Simulate Protocol Interactions in Go  
Visualize with D3  
See for yourself: <http://simianviz.surge.sh>  
Follow @simianviz for updates*

# Conclusions

*CSP is too limited*

*$\pi$ -Calculus syntax is incomprehensible*

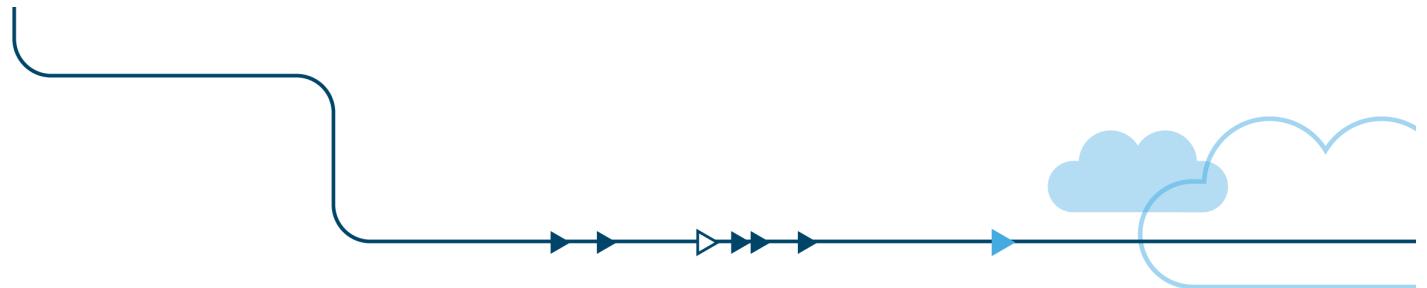
*Occam-Pi makes CSP and  $\pi$ -Calculus readable*

*Go concurrency syntax is clumsy in places but works*

*Showed some useful channel based Go- $\pi$  idioms*

*Pass channels over channels for dynamic routing*

*Go works well for actor like simulation*



“We see the world as increasingly more complex and chaotic because we use inadequate concepts to explain it. When we understand something, we no longer see it as chaotic or complex.”

*Jamshid Gharajedaghi - 2011  
Systems Thinking: Managing Chaos and Complexity: A Platform for Designing Business Architecture*

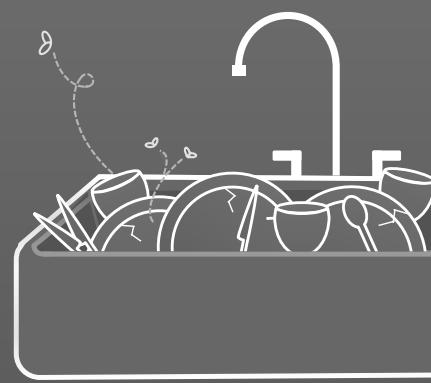


**Adrian signed to a new label (AWS) at the end of 2016 and now had a much bigger production budget for making slides look cool, and a PR department to keep him from being too controversial!**

# The New De-Normal



Monolithic  
Databases

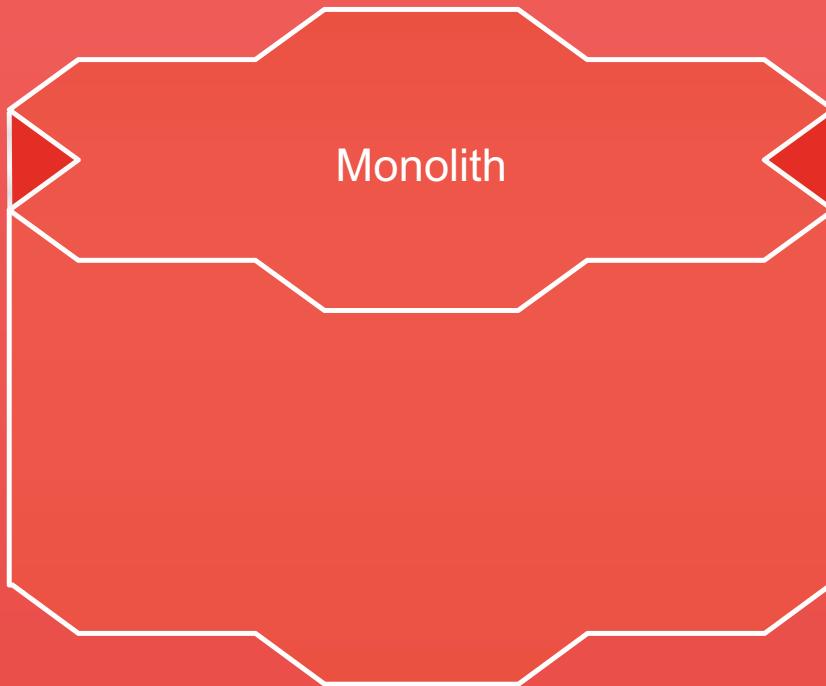


Kitchen Sink  
Analogy

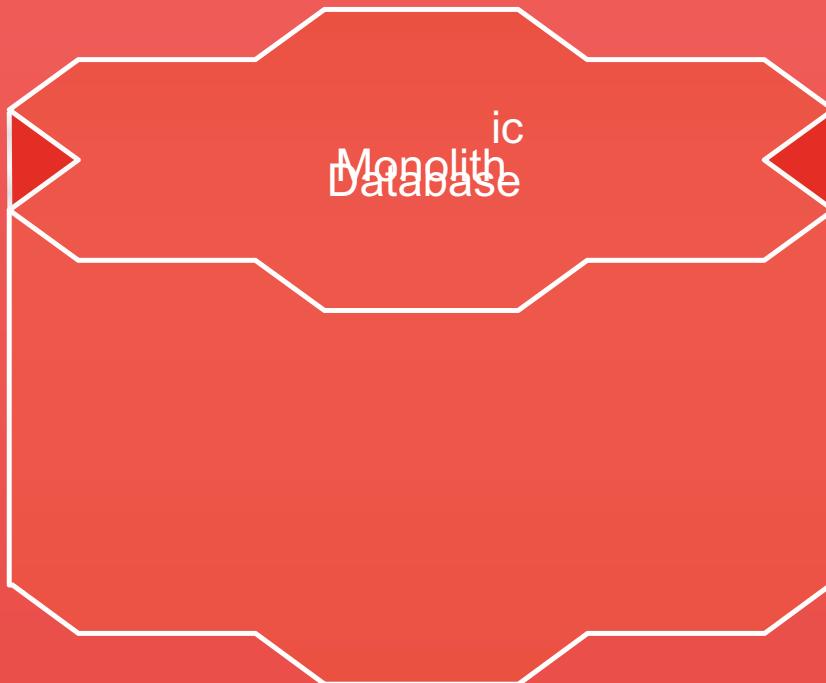


De-normalized

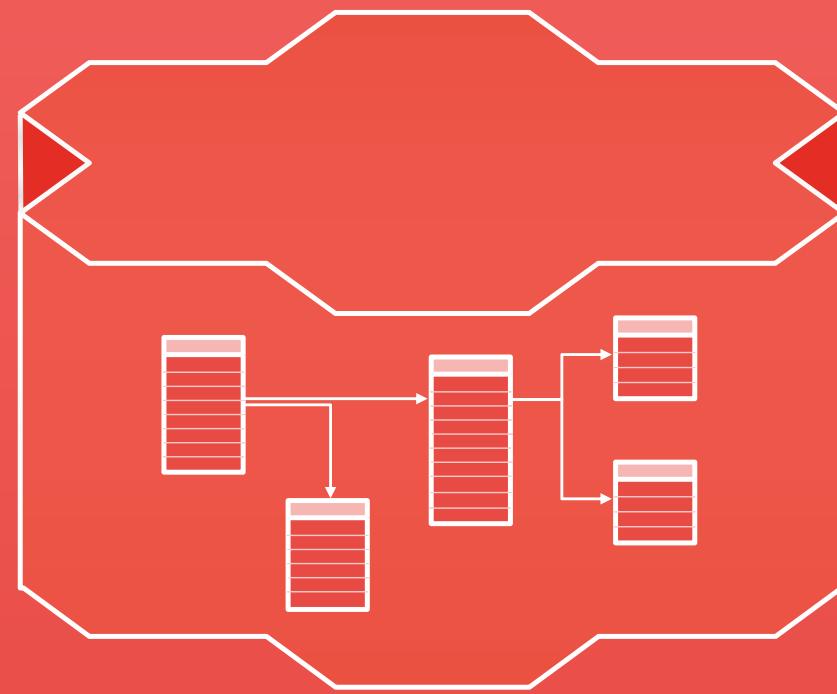
**Expensive,  
Hard to Create  
and Run**



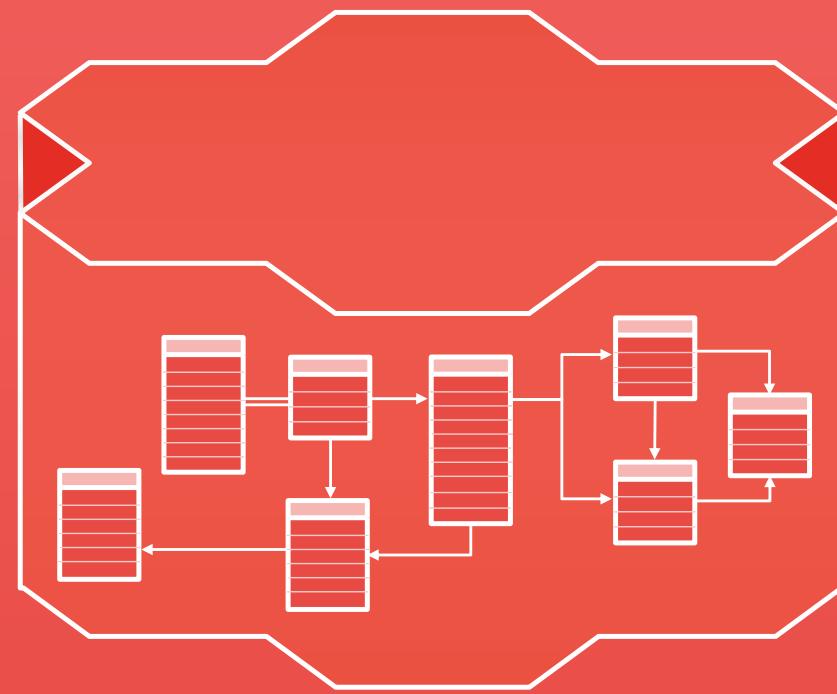
**Expensive,  
Hard to Create  
and Run**



# Database Schema Entity Relationship



# Database Schema Entity Relationship



# Database Schema Entity Relationship



# Kitchen Sink Analogy



# Kitchen Sink Cleanup



# Kitchen Sink Cleanup



# Kitchen Sink Cleanup



# Kitchen Sink Cleanup



# Kitchen Sink Cleanup



# Kitchen Sink Cleanup



# Kitchen Sink Cleanup



# Consistency Problem

How Many Complete  
Sets Are There?



# Consistency Problem

How Many Complete  
Sets Are There?



# Consistency Problem

How Many Complete  
Sets Are There?





# Adding a New Use Case





# Adding a New Use Case



# Cloud Makes it Easy to Add New Databases



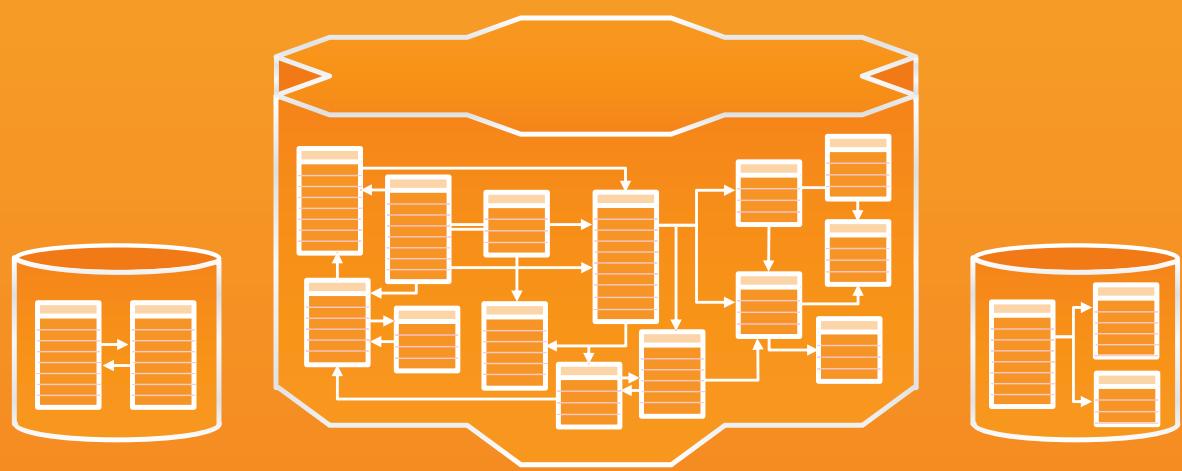
Amazon  
DMS



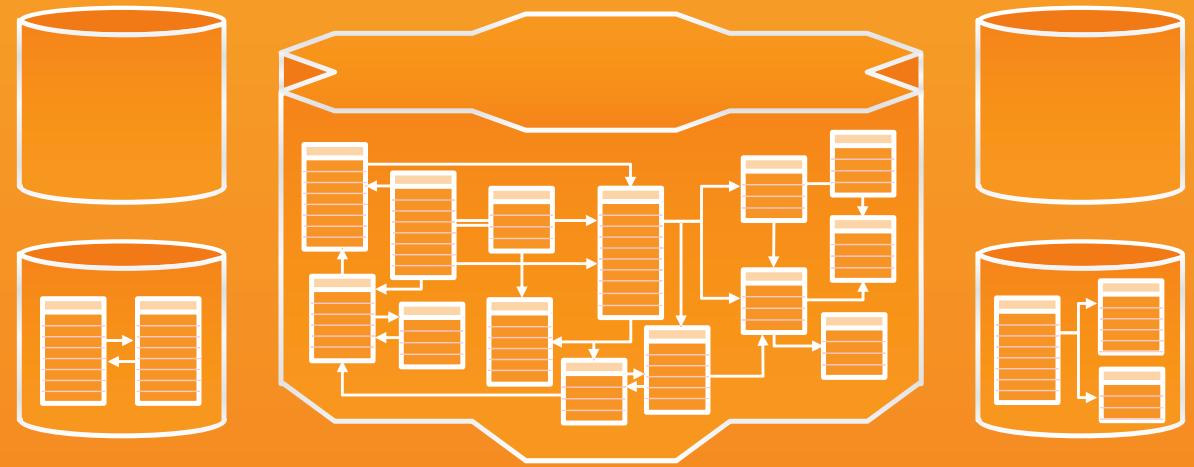
Amazon  
DynamoDB



Amazon  
RDS



# Untangle and Migrate Existing “Kitchen Sink” Schemas

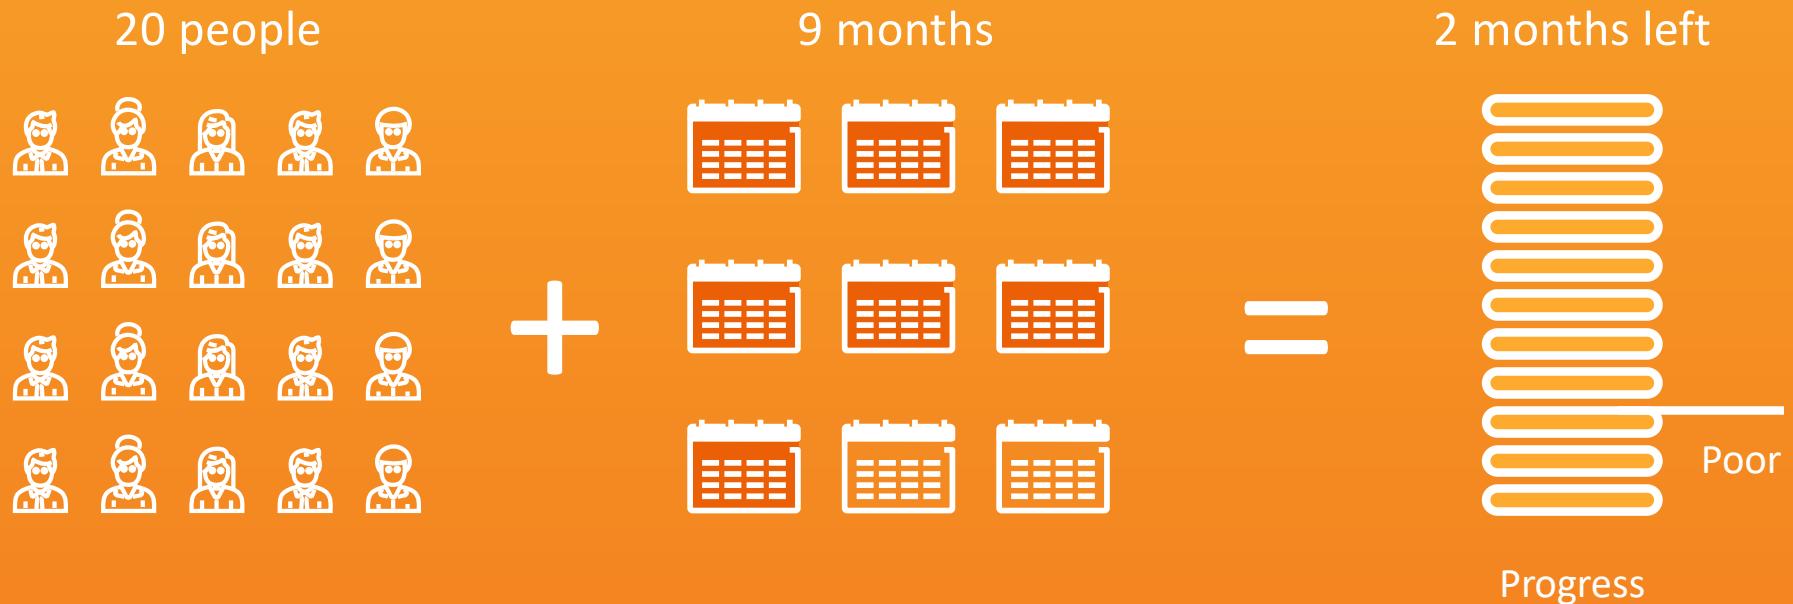


# Untangle and Migrate Existing “Kitchen Sink” Schemas

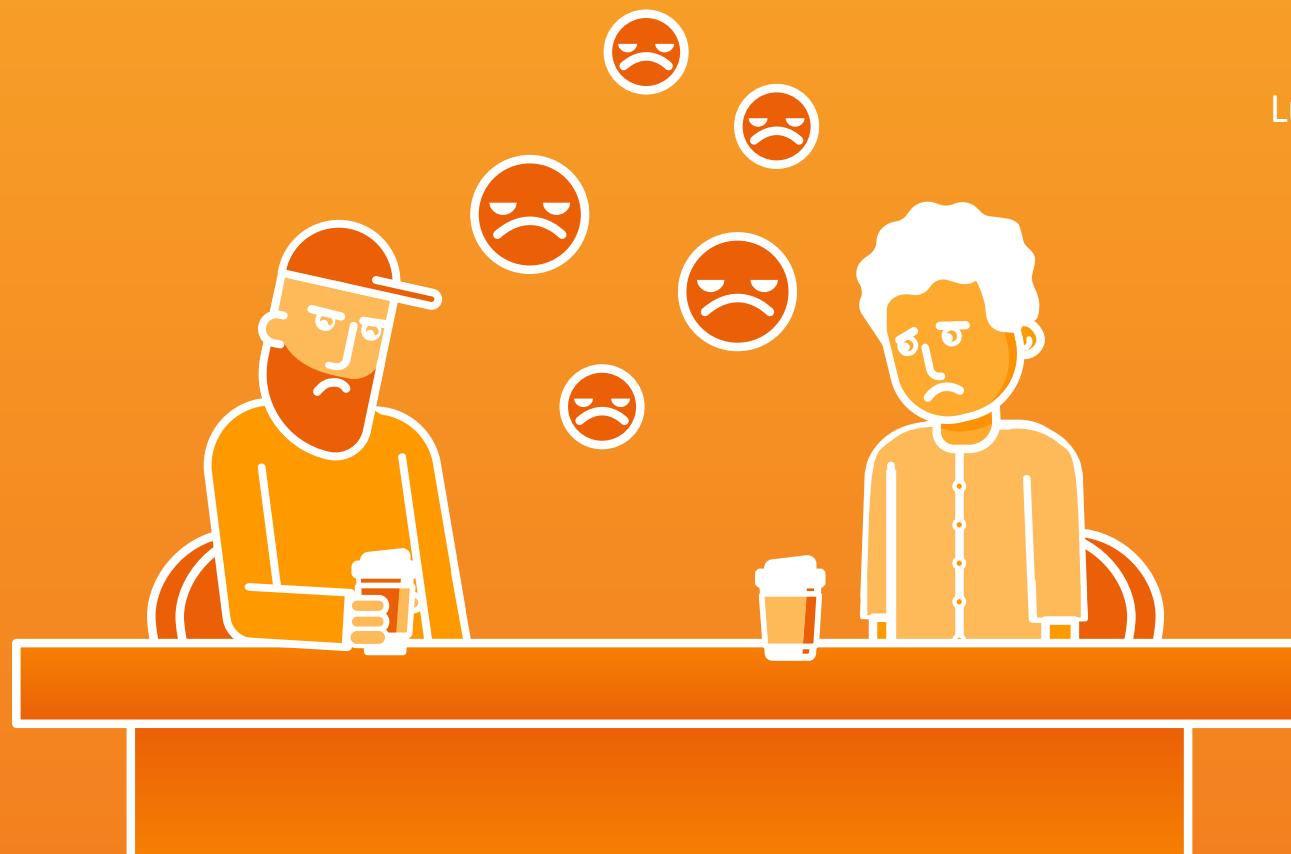


I started to collect stories, scribbled them as rough ideas on my iPad with Apple pencil, sent them to the graphic designer, and got some cool decks back...

# Conventional Development



# Conventional Development



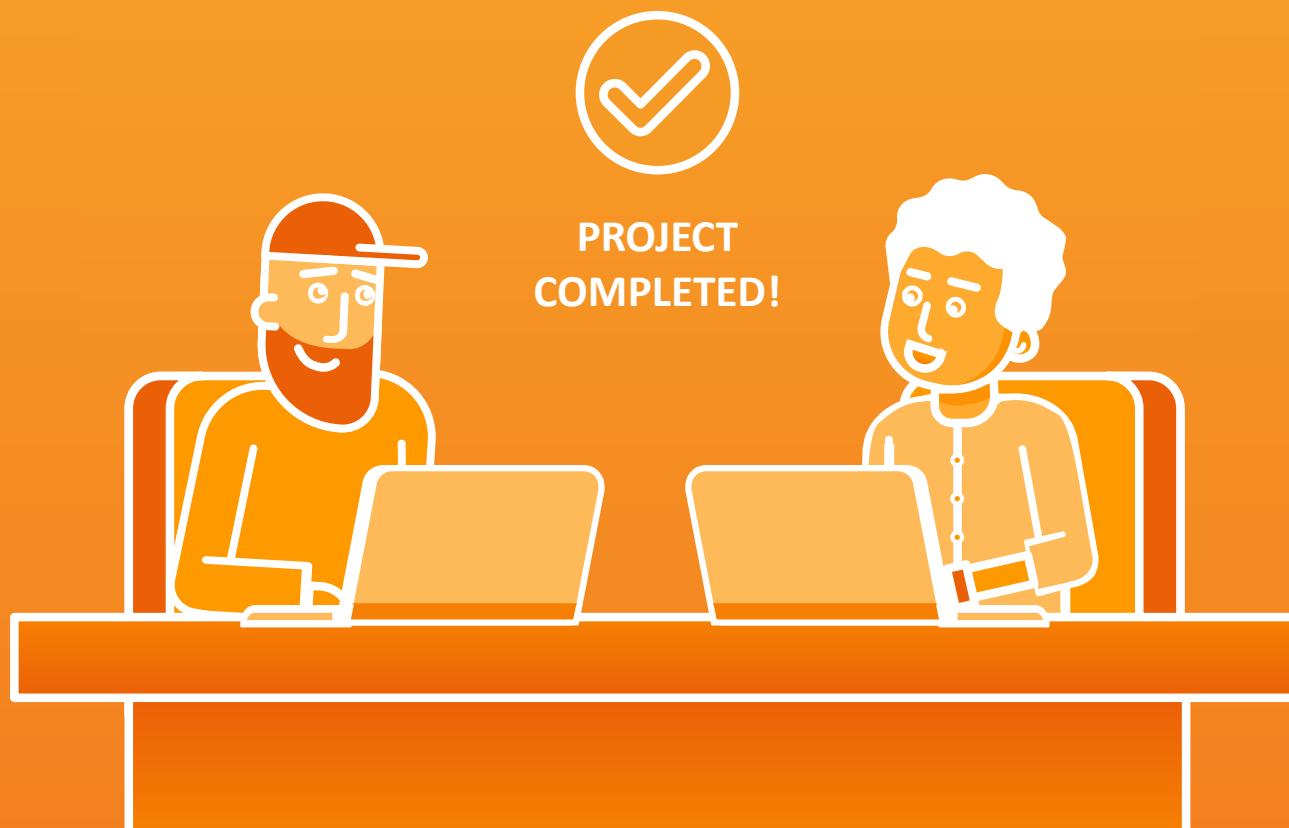
# Conventional Development



# Serverless Development



# Serverless Development



# Serverless Development

Team finally agrees  
It works and is secure



2 months left  
A whole month later



AWS Lambda



Serverless



Progress

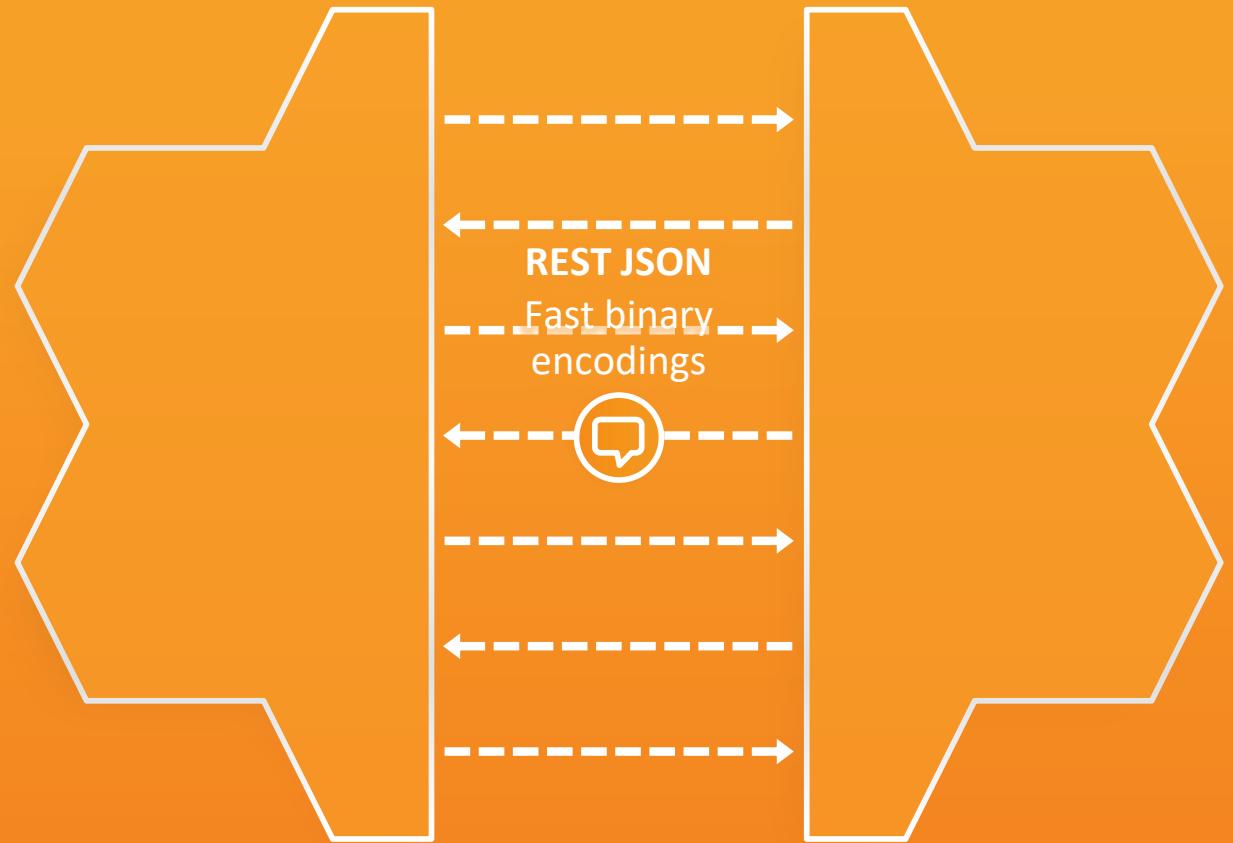
# Serverless Development



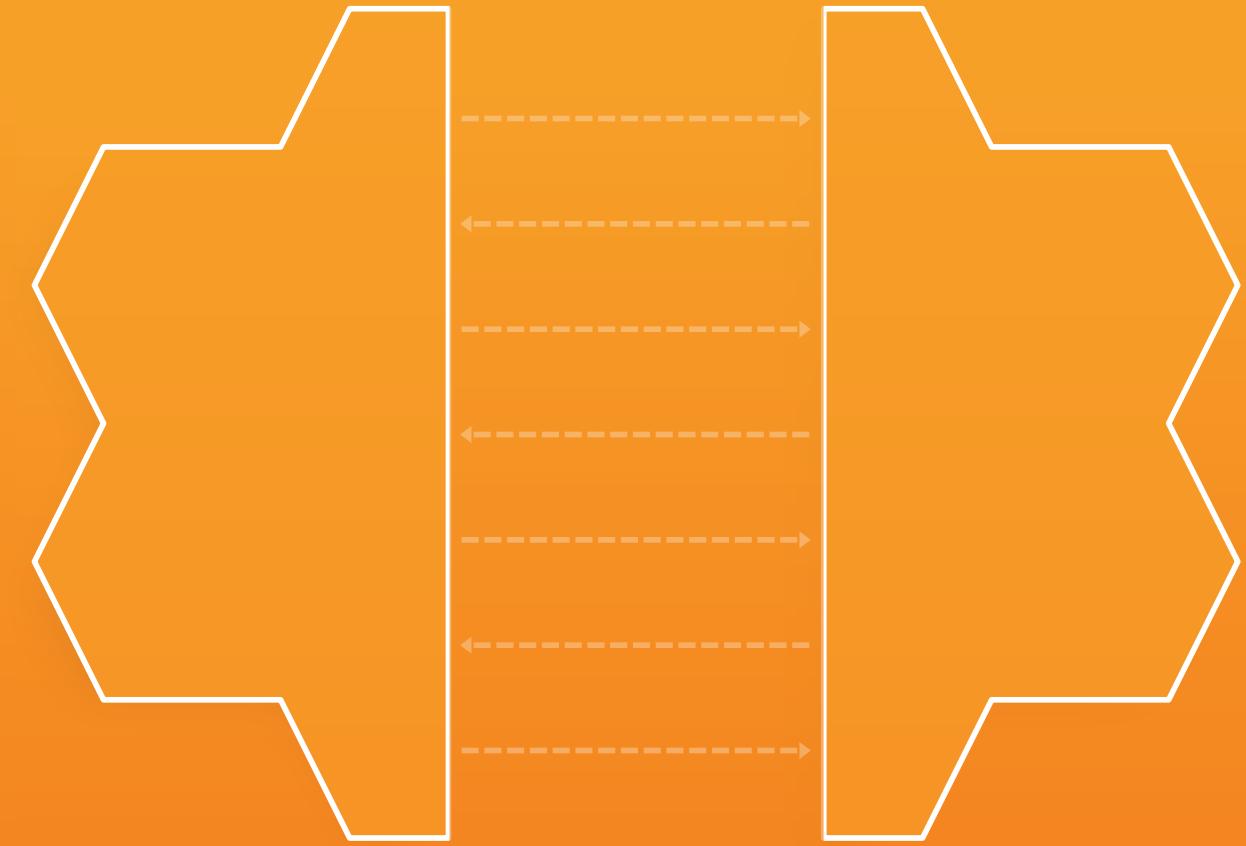
**Shipped  
application  
1 month early**

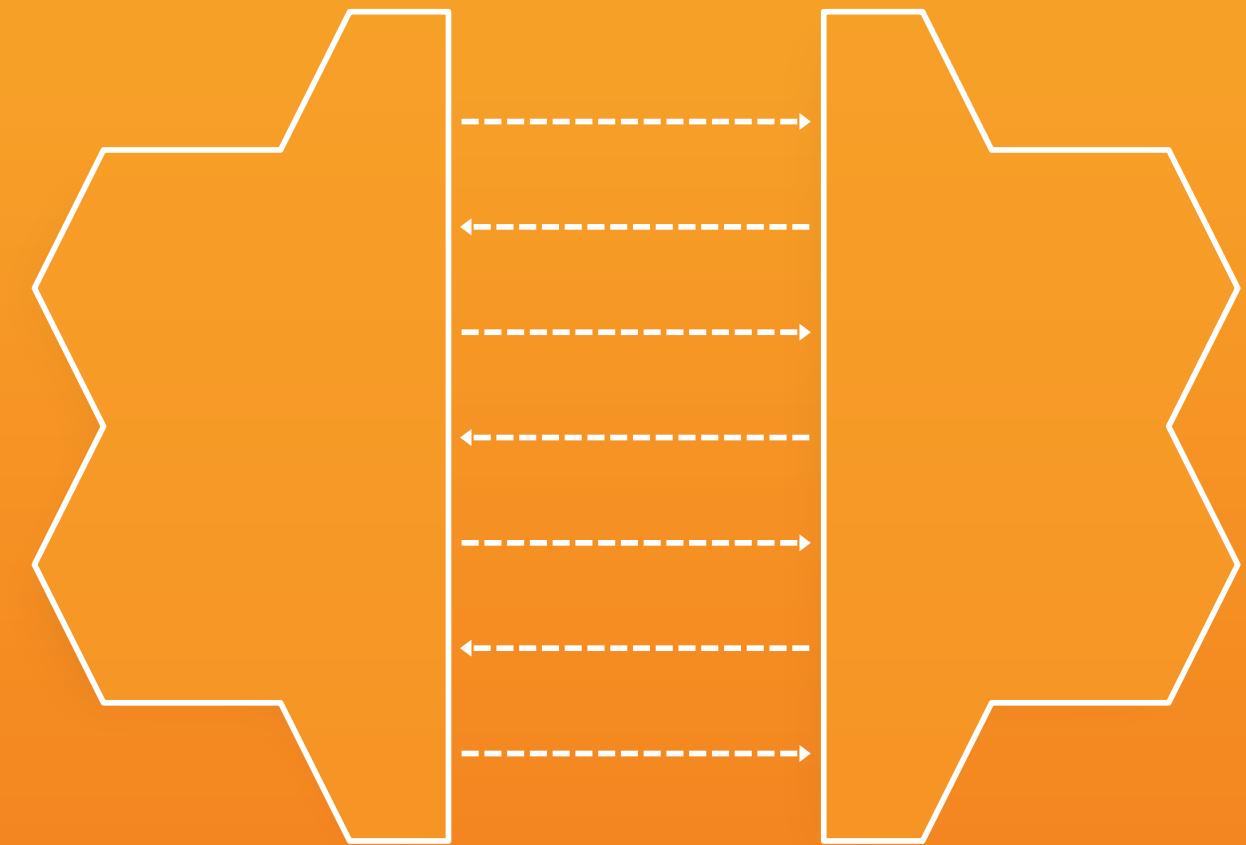
# **What is different about serverless?**

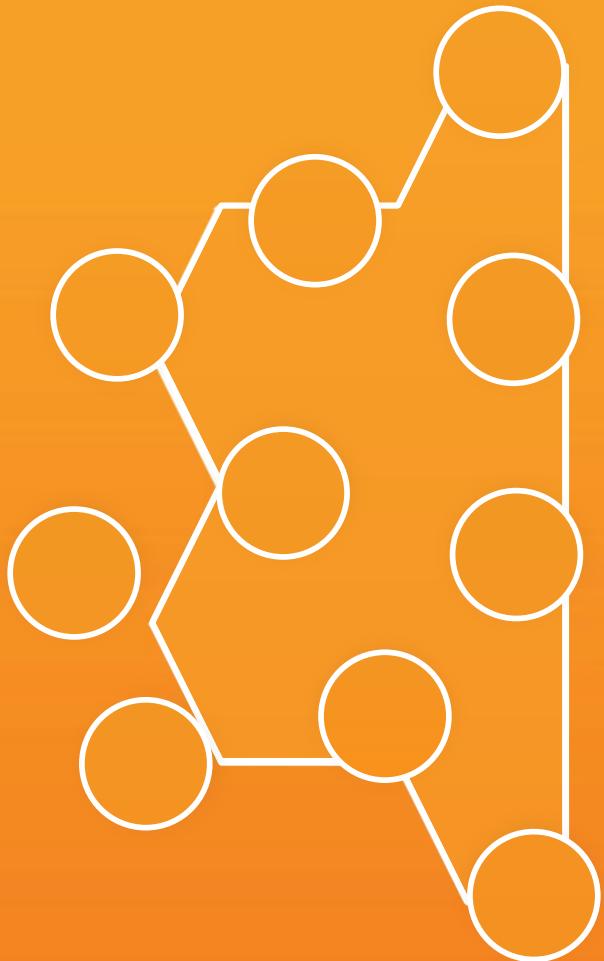
# Splitting Monoliths

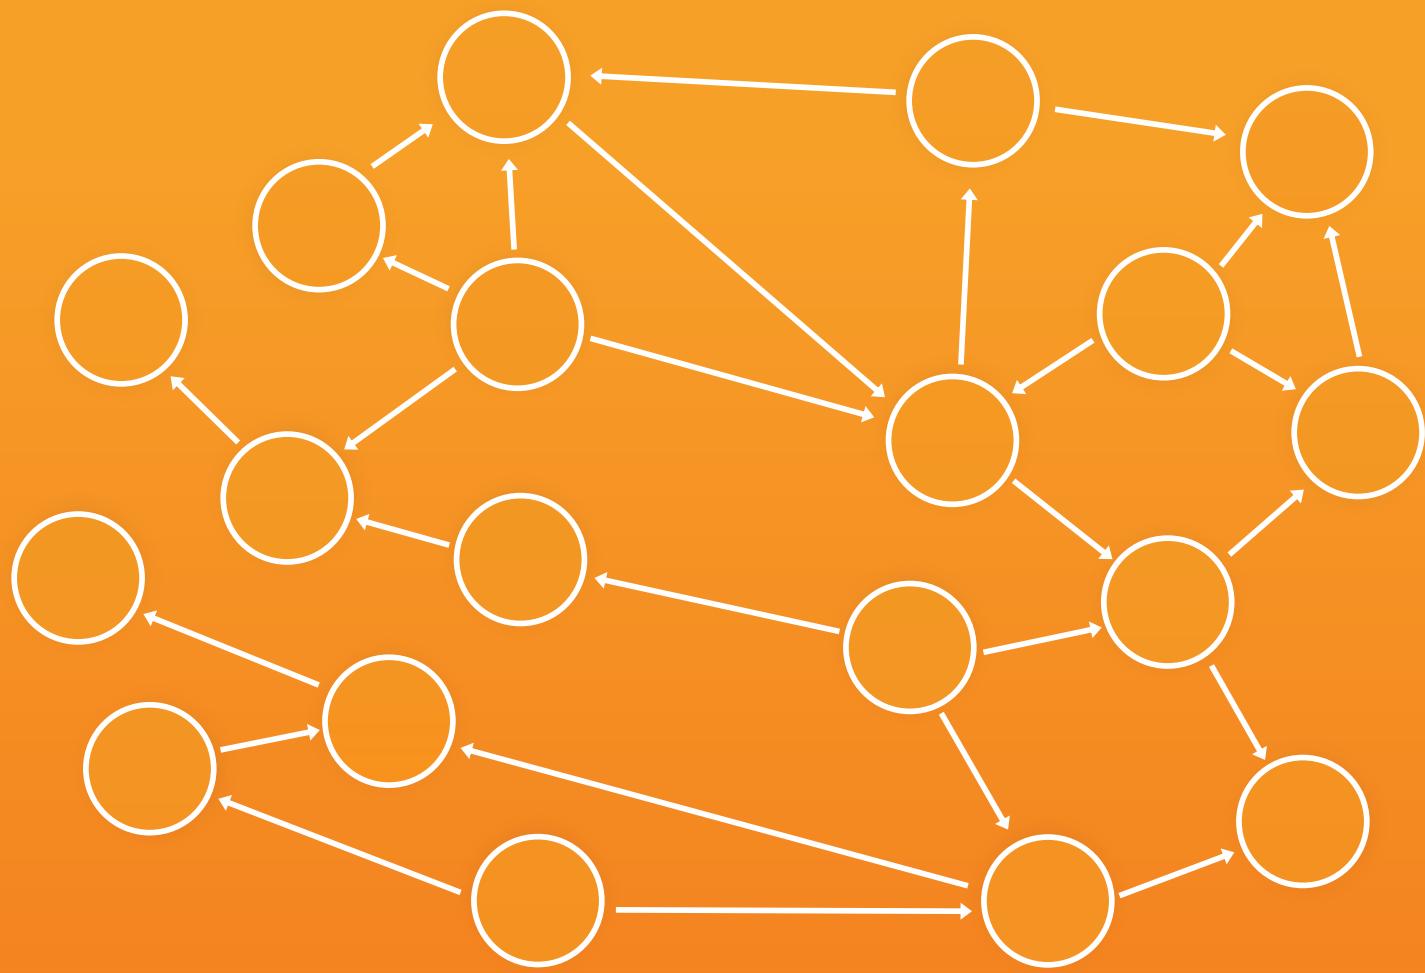


# Splitting Monoliths

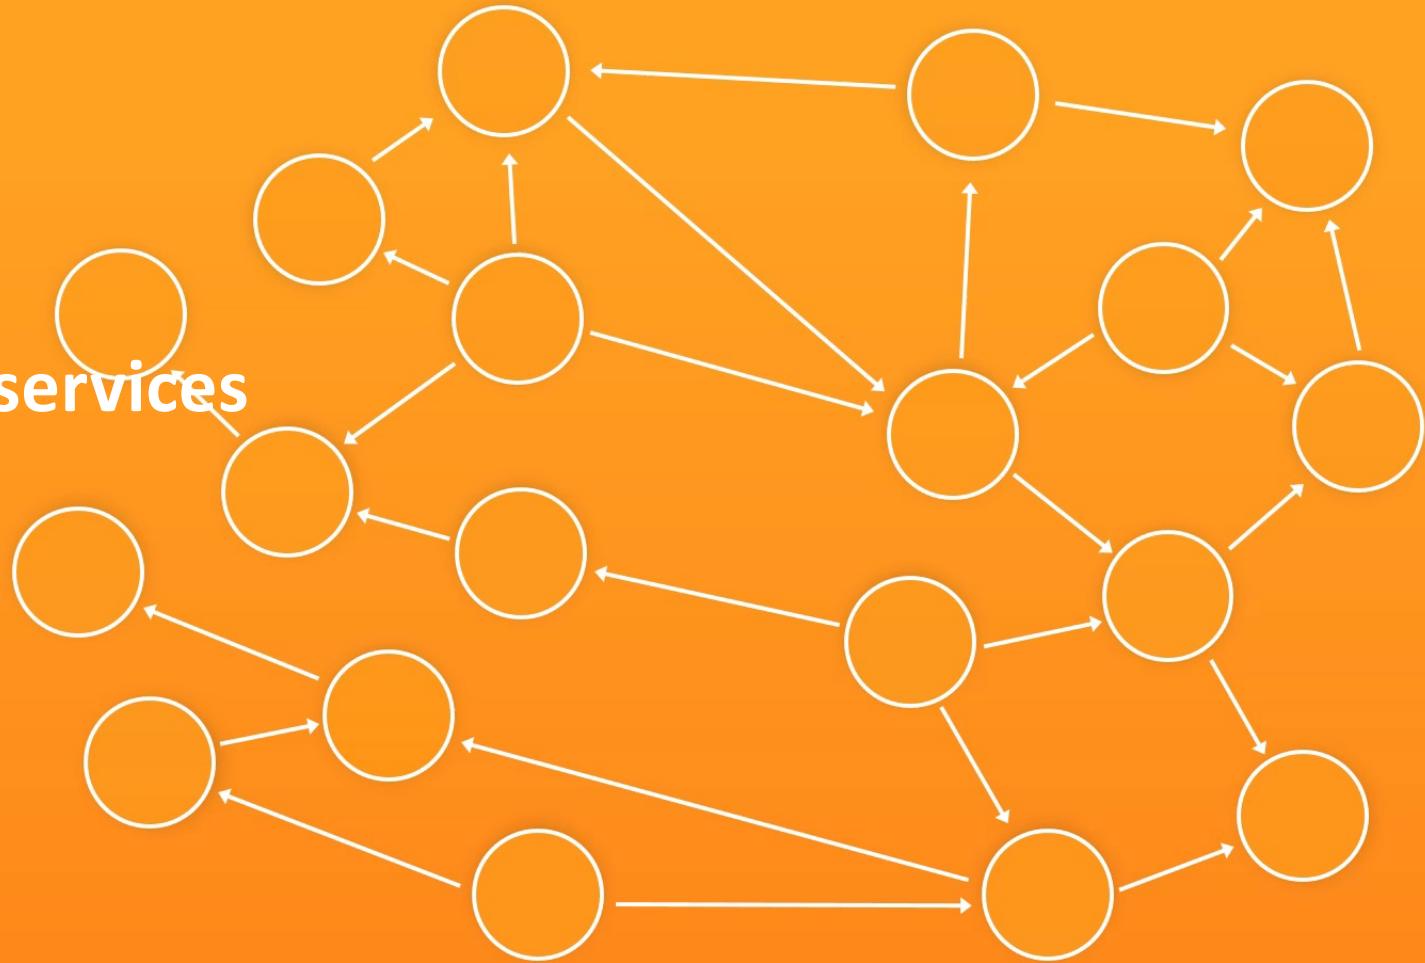








# Microservices



# Microservices to Functions

Standard building brick services provide standardized platform capabilities



Amazon API  
Gateway



Amazon S3



Amazon  
SQS



Amazon  
DynamoDB



Amazon  
Kinesis



Amazon SNS

# Microservices to Functions



# Microservices to Functions



# Microservices to Functions



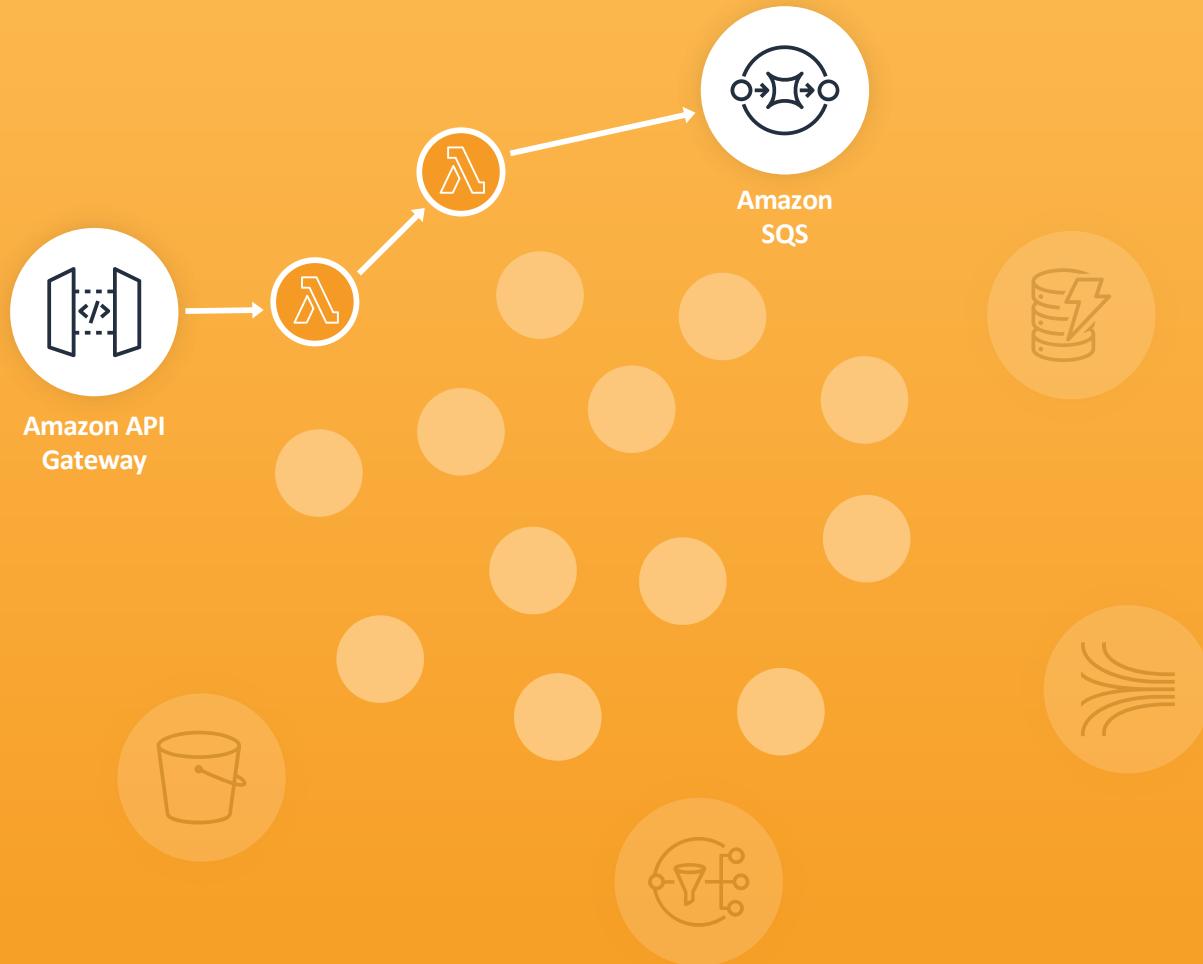
# Microservices to Ephemerical



# Microservices to Ephemeral Functions



# Microservices to Ephemeral Functions



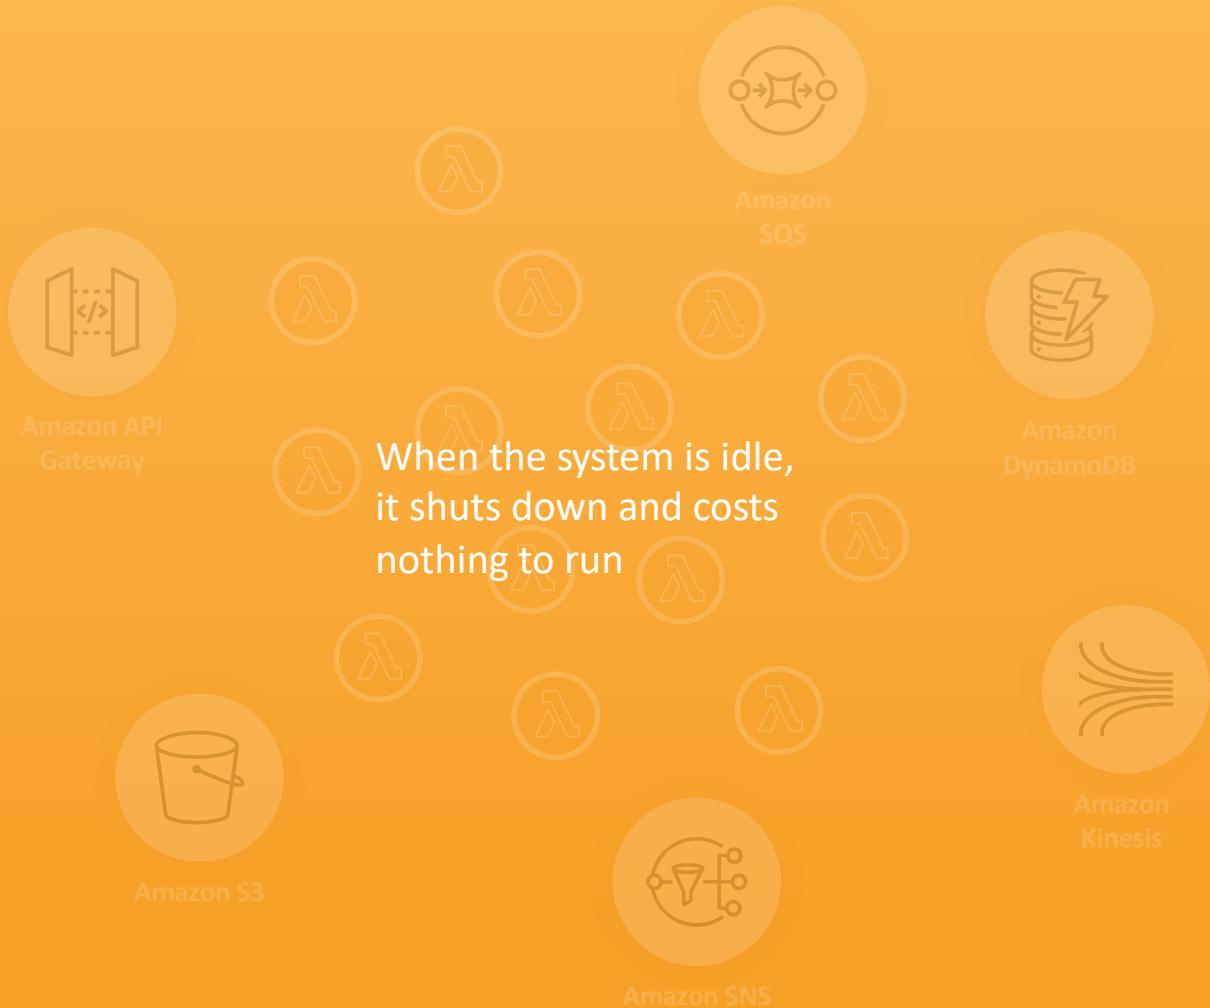
# Microservices to Ephemeral Functions



# Microservices to Ephemeral Functions



# Microservices to Ephemeral Functions

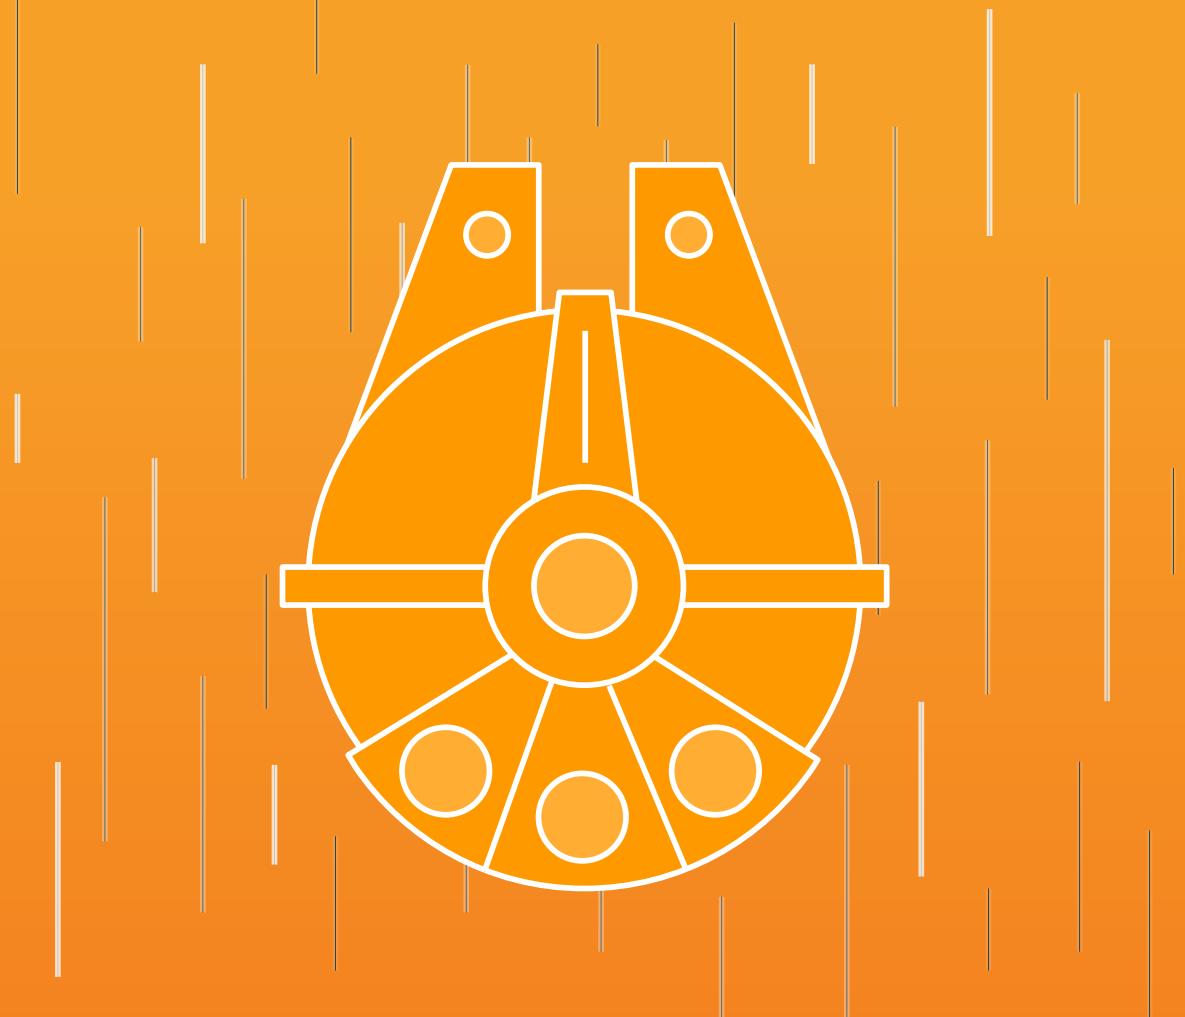


**So WHY is it so fast to  
write a serverless app?**

An analogy...



**What is the  
user need?**



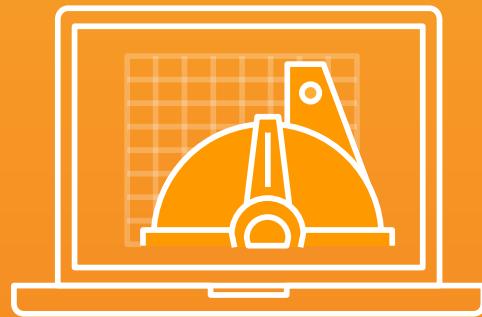
**What is the  
problem you are  
trying to solve?**



**Make a model spaceship  
quickly and **cheaply****



# Traditional Development



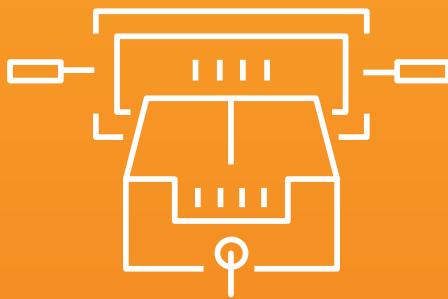
Design a prototype

# Traditional Development



Carve from  
modelling clay

# Traditional Development



Make molds

# Traditional Development



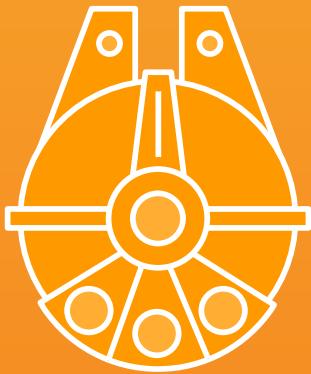
Produce injection molded parts

# Traditional Development



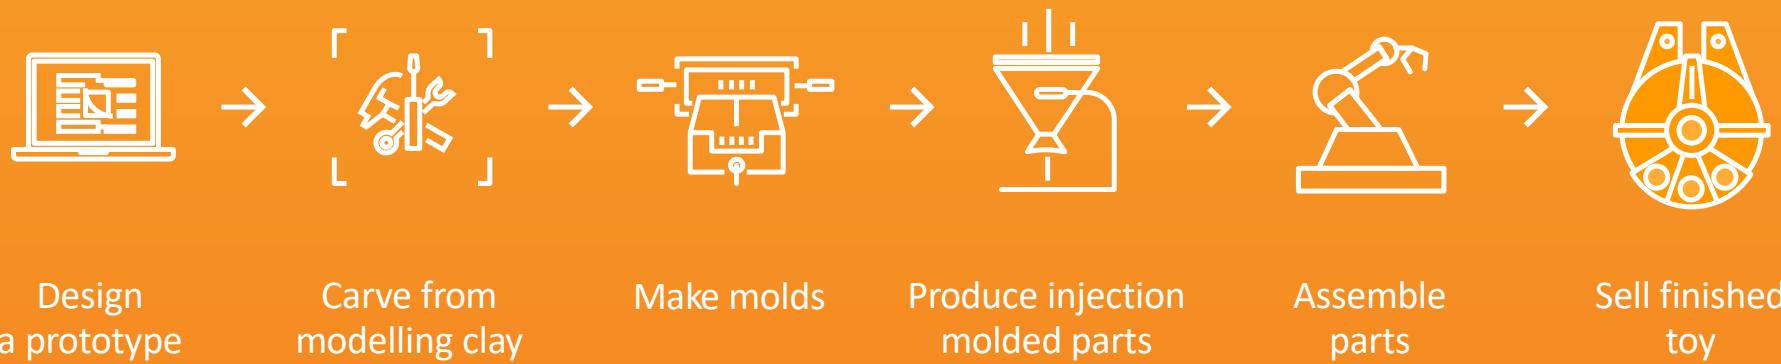
Assemble parts

# Traditional Development



Sell finished toy

# Traditional Development



# Rapid Development



Big bag of blocks

+



Instructions

+



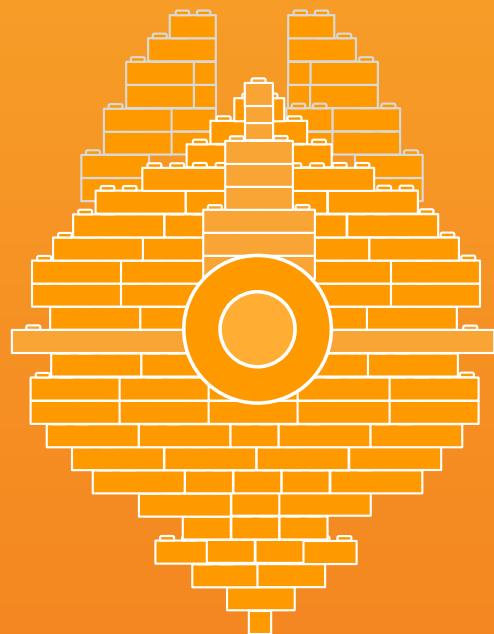
A few hours

# Rapid Development



A finished toy

# Rapid Development

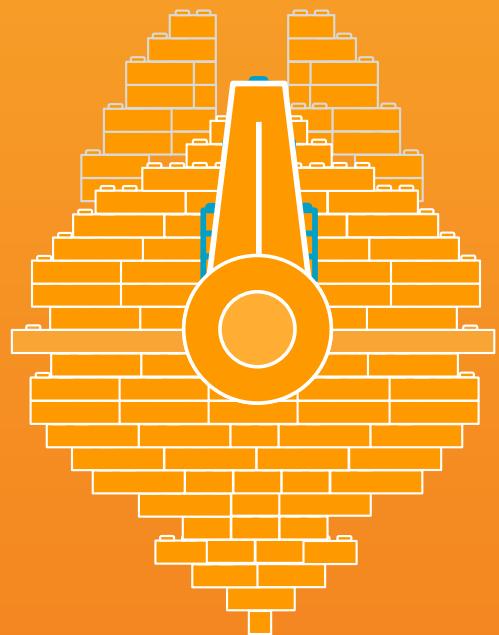


Lacks fine detail

Recognizable, but not exactly what  
was asked for

Easy to modify and extend

# Optimization



Take a group of Lego bricks...  
...and form a new custom brick

A more specialized common component

## Traditional

## Rapid Development

Full custom design

Building blocks assembly

Months of work

Hours of work

Custom components may be fragile and need to be debugged and integrated

Standard reliable components scale and are well understood and interoperable

Too many detailed choices

Need to adjust requirements to fit the patterns available

Long decision cycles

Constraints tend to reduce debate and speed up decisions

## Containers

Custom code and services

Lots of choices of frameworks  
and API mechanisms

Where needed, optimize serverless  
applications by also building services using  
containers to solve for anything serverless  
doesn't do well... yet.

## Serverless

Serverless events and functions

Standardized choices

Combine building blocks including:

- λ AWS Lambda
- ⌚ API Gateway, EventBridge
- ✉️ Amazon SNS, SQS
- 🔥 Amazon DynamoDB
- ⤓ AWS Step Functions

**So... why doesn't everyone use serverless first?**

Objections and limitations



Note: See Re:Invent 2019 SVS343

## Objections Summary

**There are answers to all of these...**

Patterns, Portability

Too hard to get started

Language support

Security

Scalability, Resilience

State handling, event processing

Startup and network latency

Limited run duration

Databases/storage interfacing

Complex configs

**Adrian retired from Amazon in the middle of  
2022 and is now working as an advisor, analyst  
and consultant via OrionX.net**

# DevSusOps

**Adding sustainability concerns to development and operations**

Adrian Cockcroft - [OrionX.net](https://OrionX.net)  
@adrianco - @DevSusOps - Oct 2022

Leave the world habitable for future generations

Market transition risks

Regulatory compliance

Physical risks to business assets

## Why does sustainability matter?

“Green” market positioning

Employee enthusiasm

Reduced costs now or in the future

Social license to operate

**What can we do about it?**

## Development

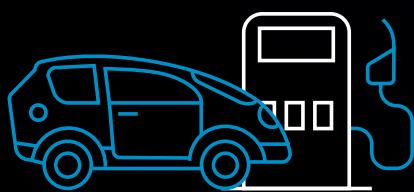
- Optimize code
- Choose faster languages and runtimes
- Efficient algorithms
- Faster implementations
- Reduce logging
- Reduce retries and work amplification

## Operations

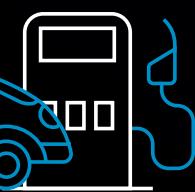
- Higher utilization
- Automation
- Relax over-specified requirements
- Archive and delete data sooner
- Deduplicate data
- Choose times and locations carefully

# Scope 1—fuel consumed

COUNTED BY WHOEVER OWNS THE FUEL WHEN IT BURNS



Car



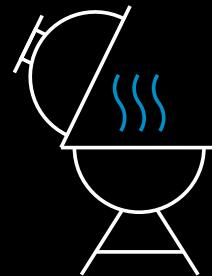
Gas pump



Fireplace

Furnace

Gas meter



Cooking

Electrify everything to take Scope 1 to zero

# Scope 2—energy used

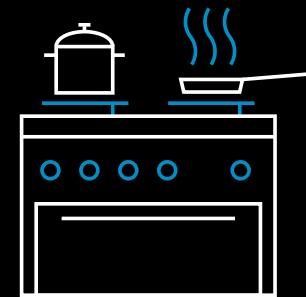
ELECTRICITY USE IS COUNTED ONLY ONCE WHERE IT IS CONSUMED



Windmill  
Power station  
Grid mix



Heat pump Solar panels Batteries Electric car



Induction range  
(gas ranges are also a big cause of indoor pollution that leads to asthma)

Change grid mix to renewable power and store renewable energy in batteries to reduce Scope 2

# Scope 3—everything else

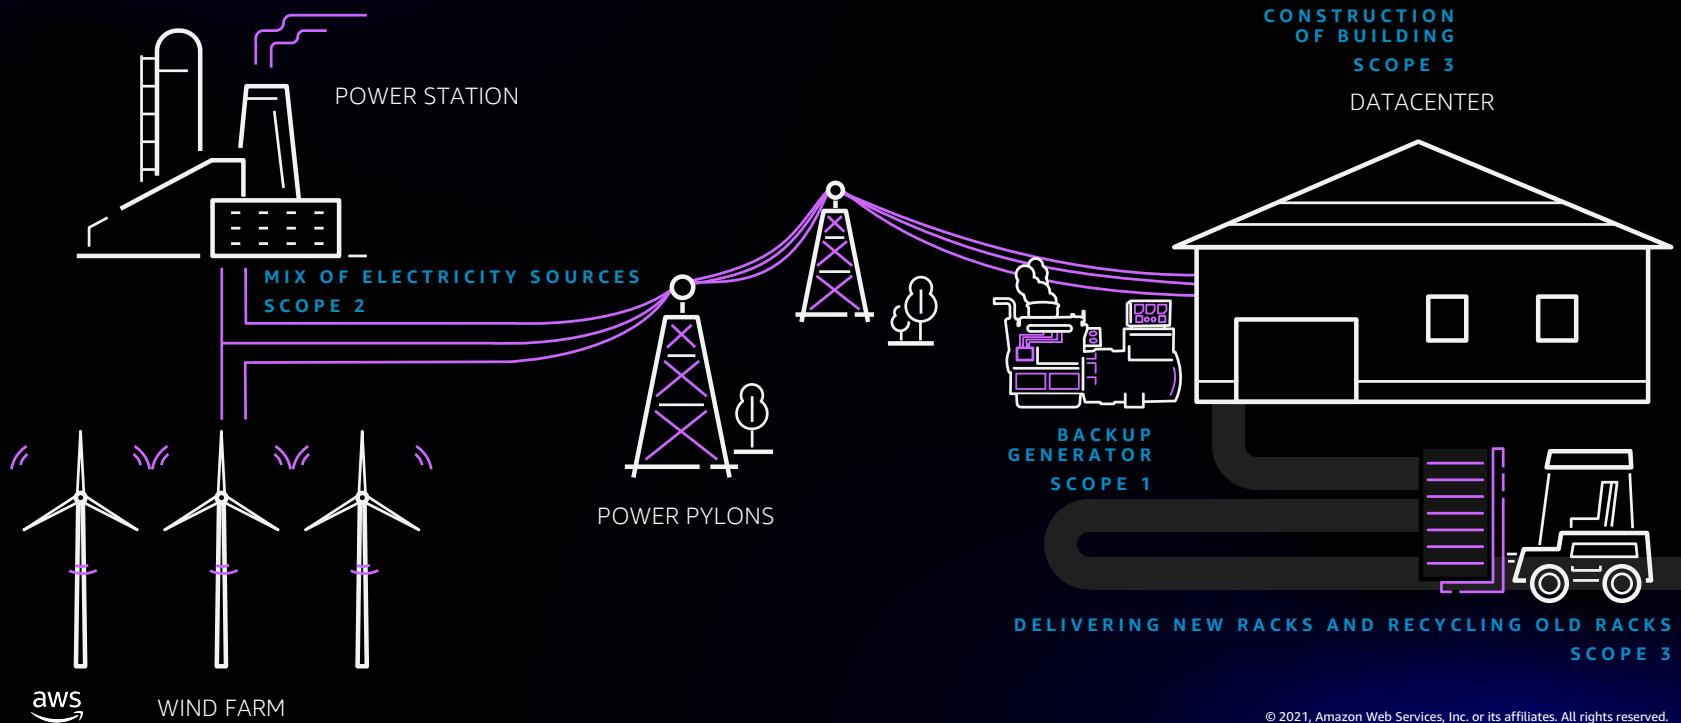
SUPPLY CHAIN AND INVESTMENTS



Scope 3 depends a lot on what kind of business you are in

# The carbon footprint of datacenters

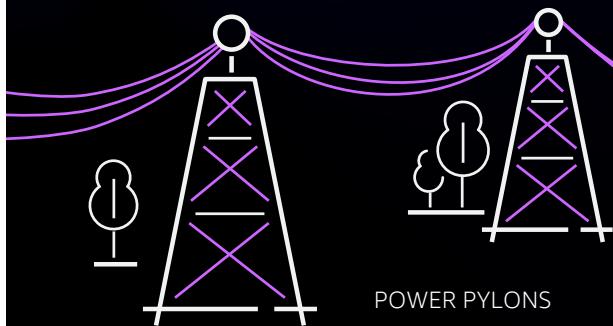
TOP-DOWN CONTEXT



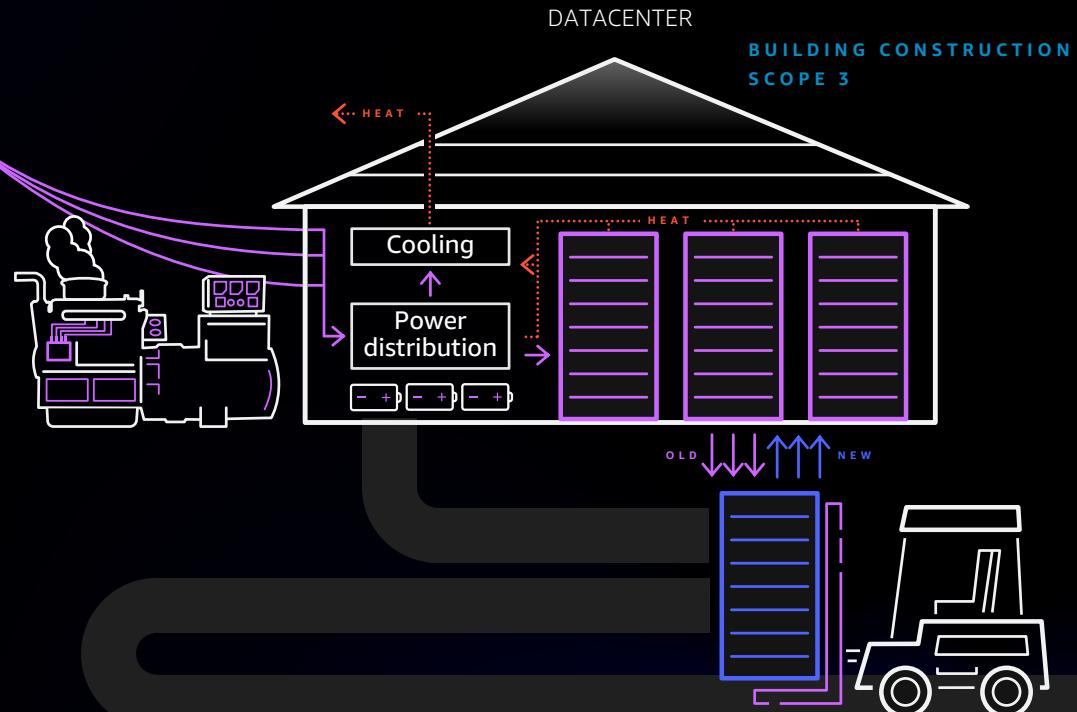
© 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# The carbon footprint of datacenters

TOP-DOWN CONTEXT



POWER PYLONS



DELIVERING NEW RACKS AND RECYCLING OLD RACKS

© 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.



**Your suppliers need to report scope 1,2,3 to you**

**You need to report scope 1,2,3 to your customers**

Greenhouse Gas Protocol - <https://ghgprotocol.org>

# **For workload optimization we need directional and proportional guidance:**

Cloud Carbon Footprint tool - Open source, uses billing data as input  
Maintains a set of reasonable estimates/guesses for carbon factors  
<https://www.cloudcarbonfootprint.org>

Green Software Foundation Software Carbon Intensity - SCI  
A model for reporting software impact per business operation  
<https://greensoftware.foundation/projects/>

AWS Well Architected Pillar for Sustainability  
Guidance on how to optimize development and operations for carbon  
<https://docs.aws.amazon.com/wellarchitected/latest/sustainability-pillar/sustainability-pillar.html>

## Where is all this going to be in a few years?

Monitoring tools will report carbon

Cloud providers will (all eventually) have detailed metrics

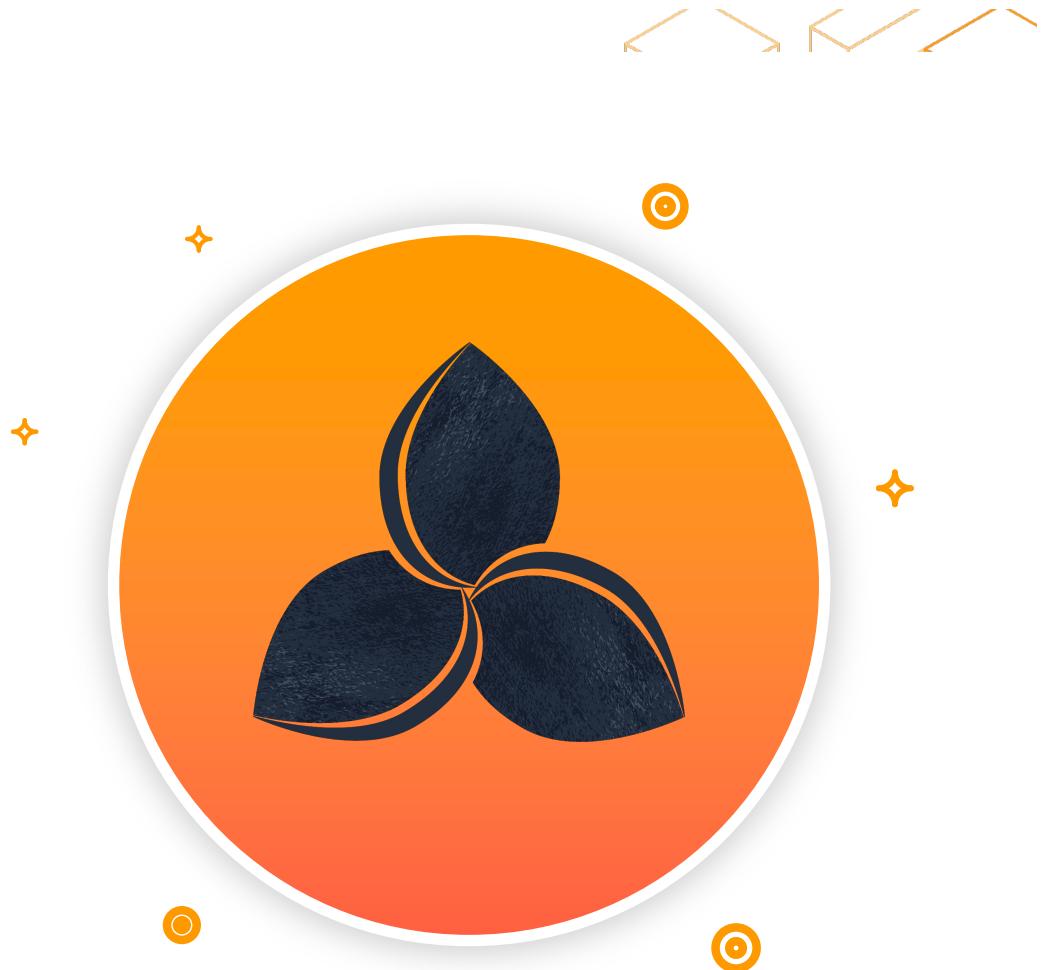
EU and US cloud regions are close to zero carbon now  
Asian regions will move to zero carbon by 2025  
(All providers have the same problem with regional policies)

# Petalith

*Memory is the treasury and guardian of all things - Cicero*

Adrian Cockcroft

@adrianco@mastodon.social



See <https://insidehpc.com/2022/12/sc22-cxl3-0-the-future-of-hpc-interconnects-and-frontier-vs-fugaku/>

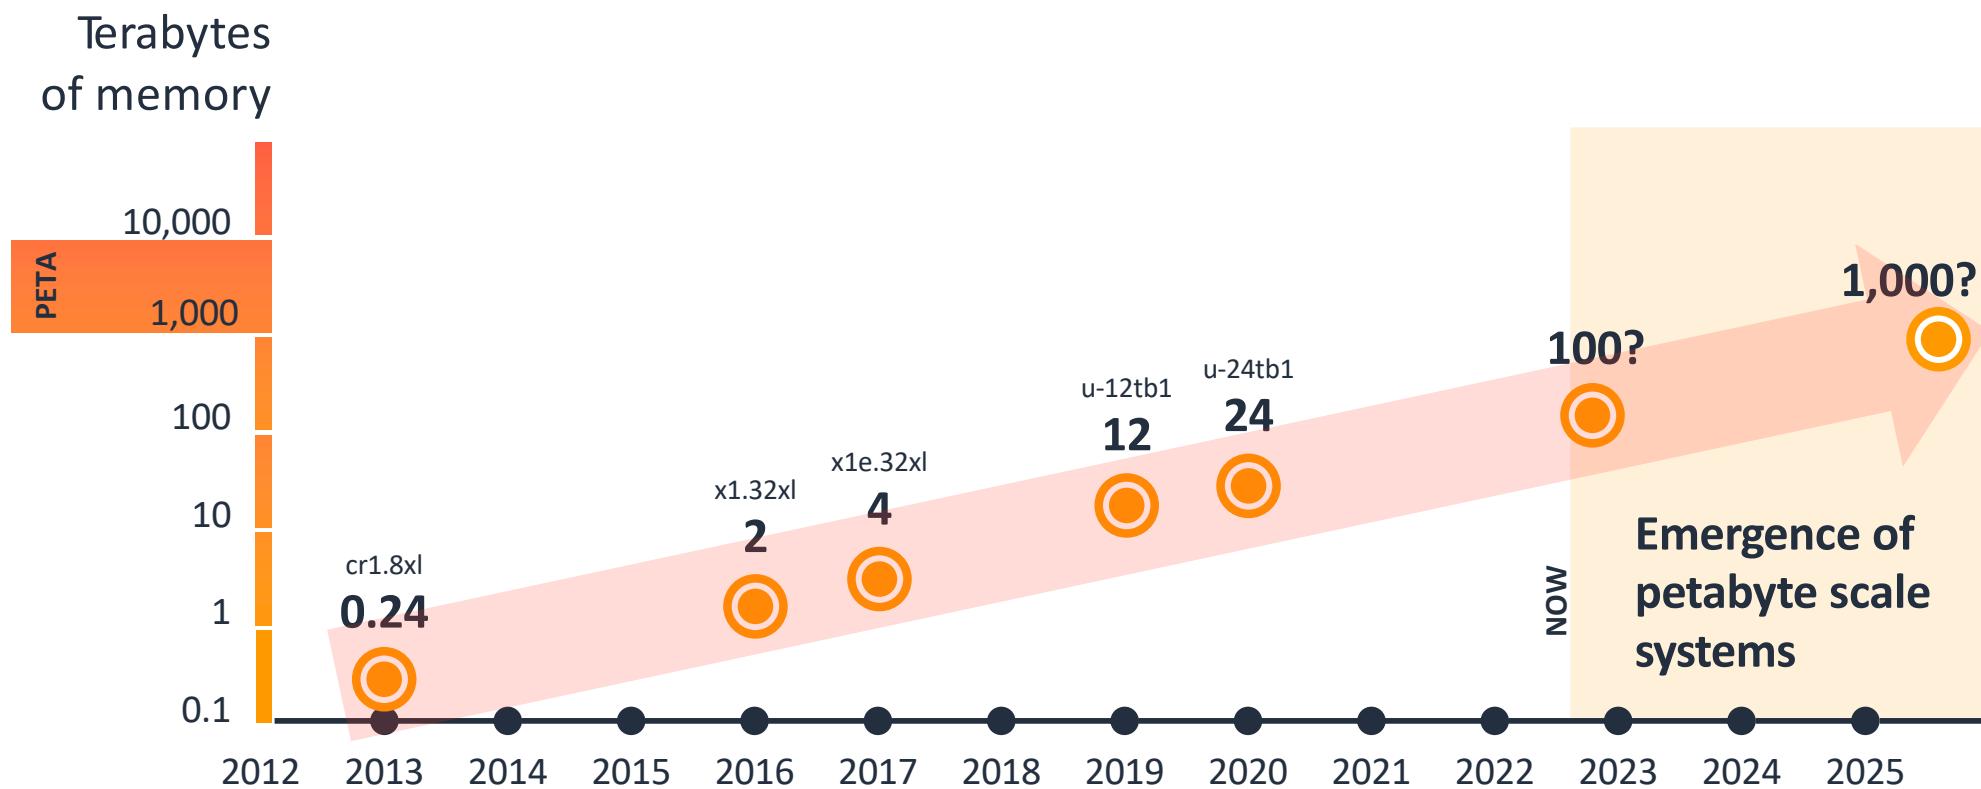


# What is big data?

Data that doesn't fit in memory on one machine

Currently about **24 TB**

## Biggest Memory Sizes Trend





# What limits speedup?

Amdahl's Law – the serial portion of a workload

In a distributed system **The Communication**

"the overall performance improvement gained by optimizing a single part of a system is limited by the fraction of time that the improved part is actually used" – Gene Amdahl - 1967

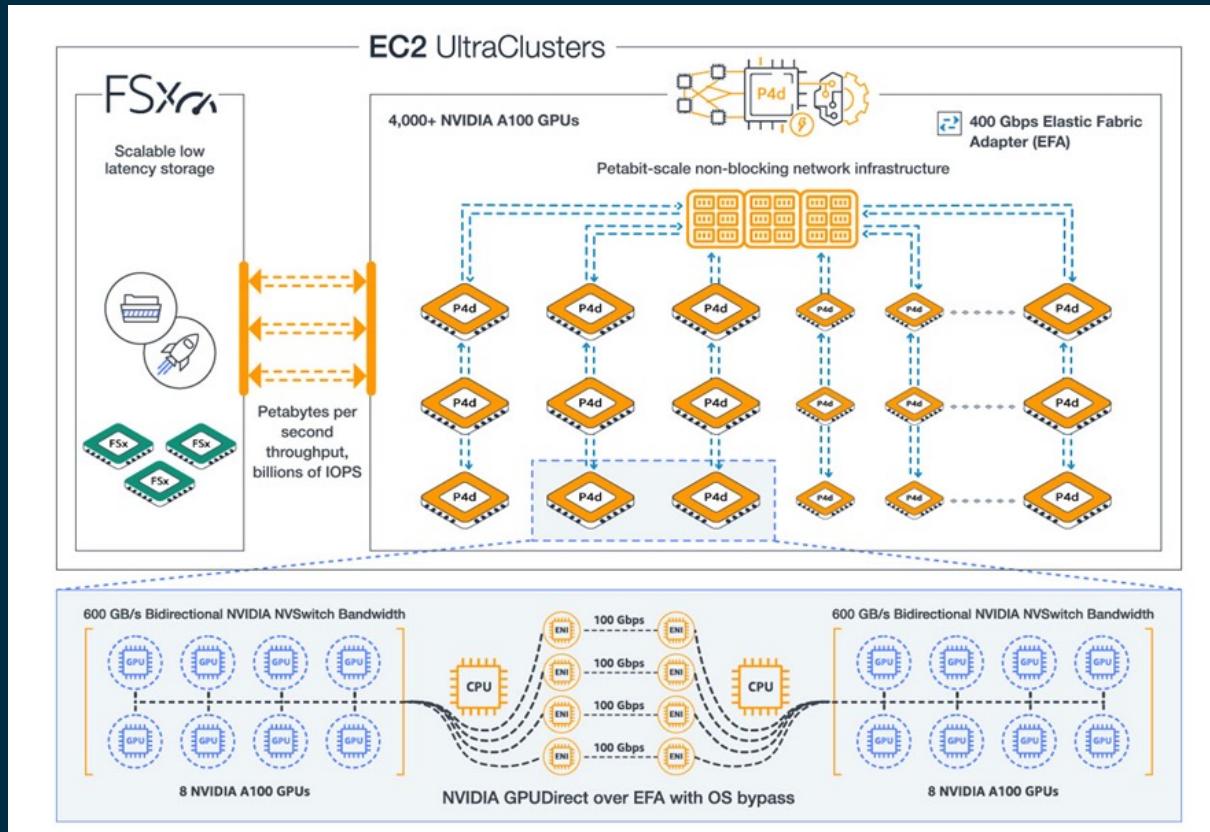
## NETWORKING

# Ethernet capacity in a single instance

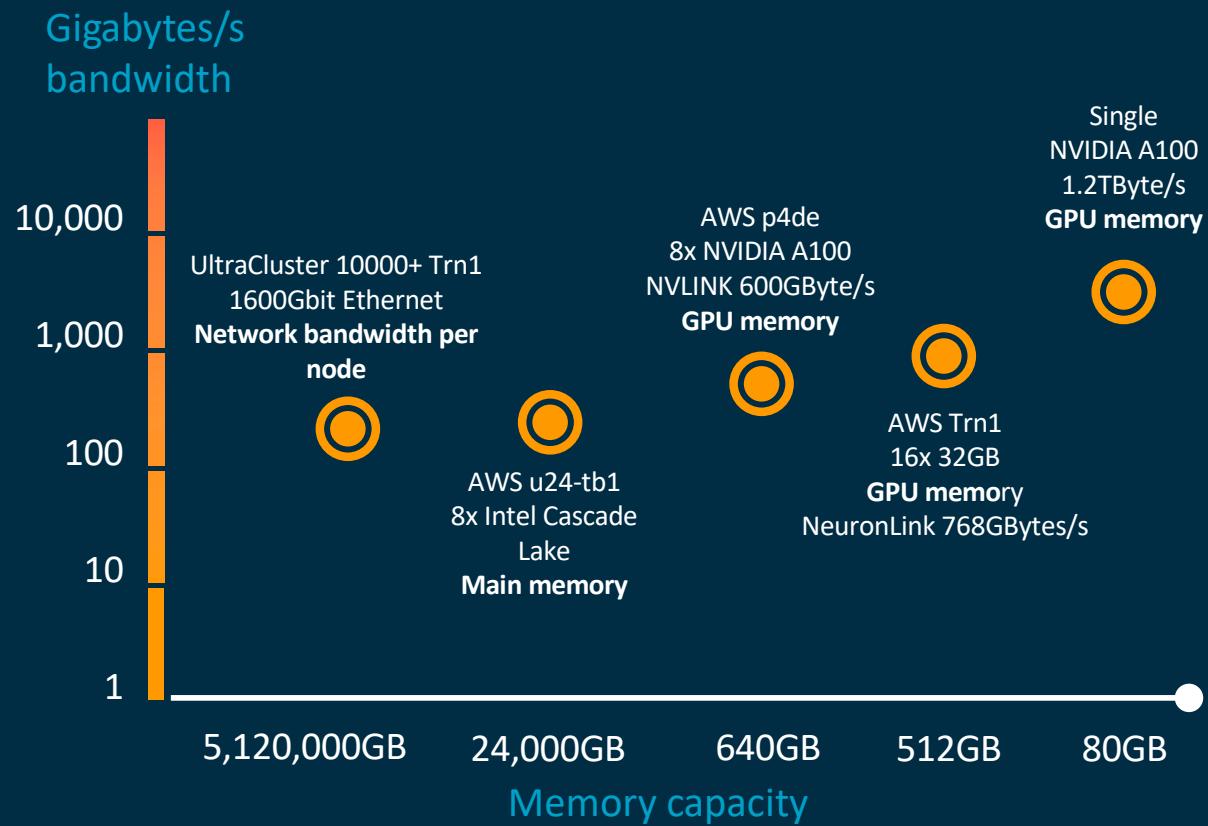


# EC2 UltraClusters P4d with NVIDIA GPUs (2022)

2023 - 10000+ Trainium GPUs on 1600Gbit links into petabit scale fabric



## Memory capacity vs. bandwidth



**Engulf your data in  
memory to reduce  
overhead, if it fits**



# How do we communicate?

By encoding transmitting receiving and decoding

How do we do it? **Really inefficiently!**



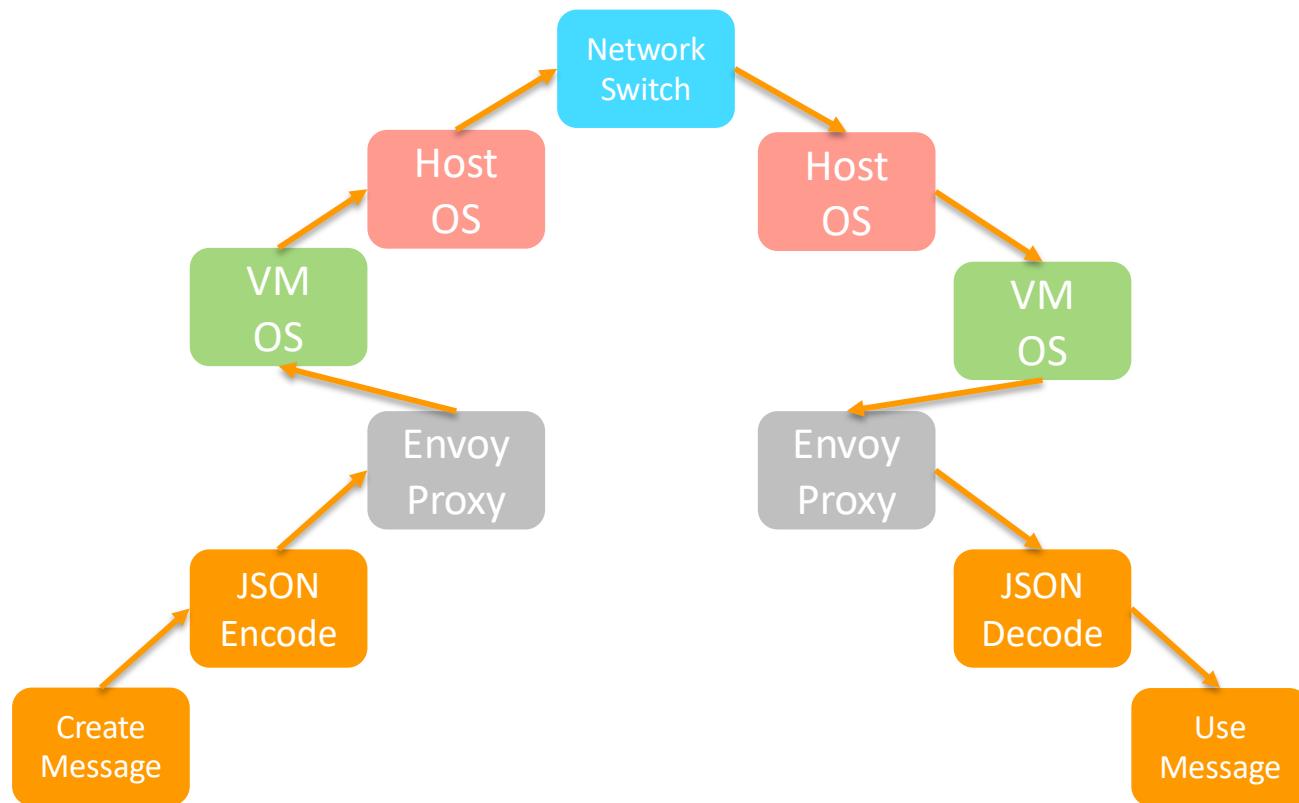
# How do systems communicate?

By encoding transmitting receiving and decoding

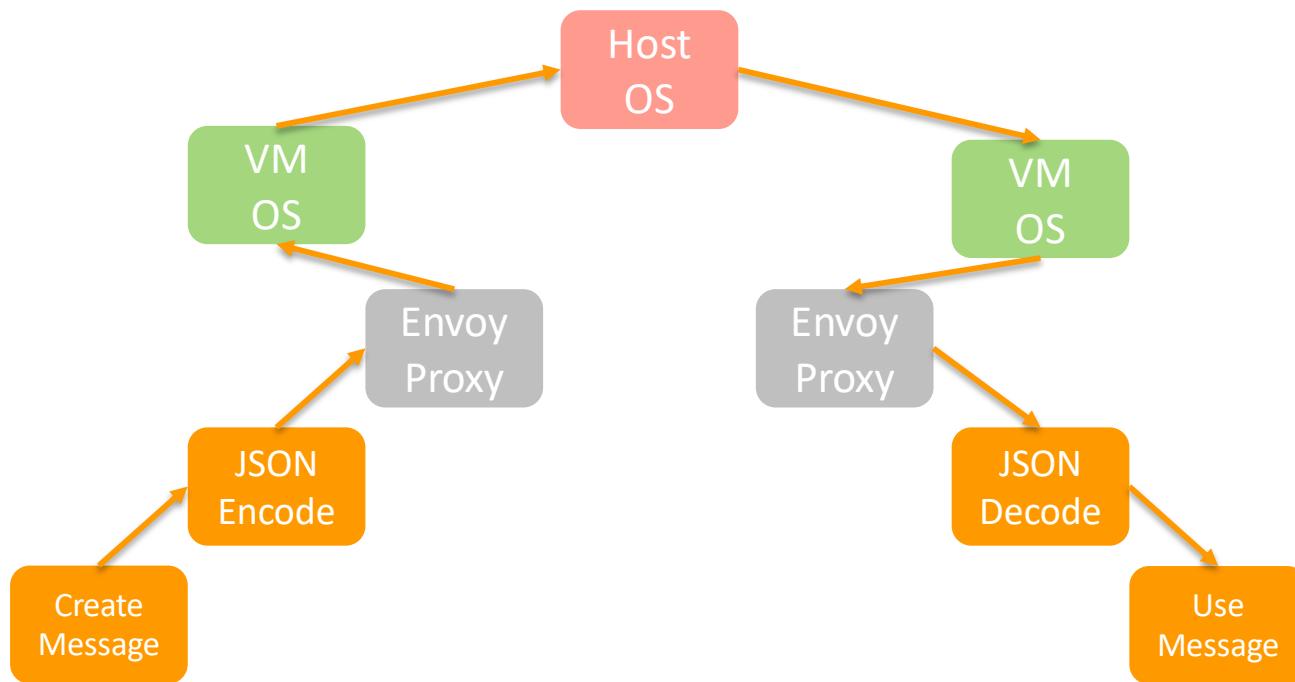
How do systems do it? **Really inefficiently!**

## Typical containerized microservice call pattern

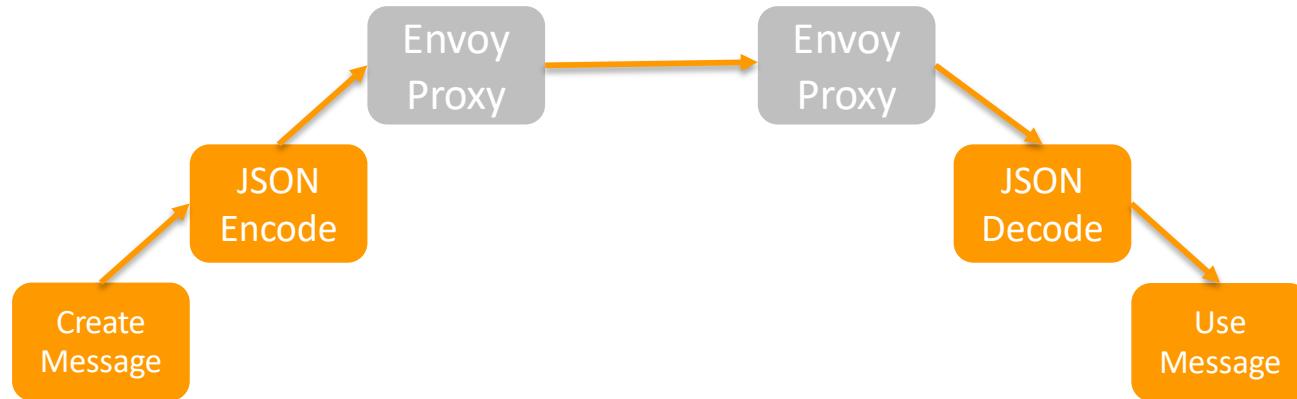
**Every step makes a new copy of the message**



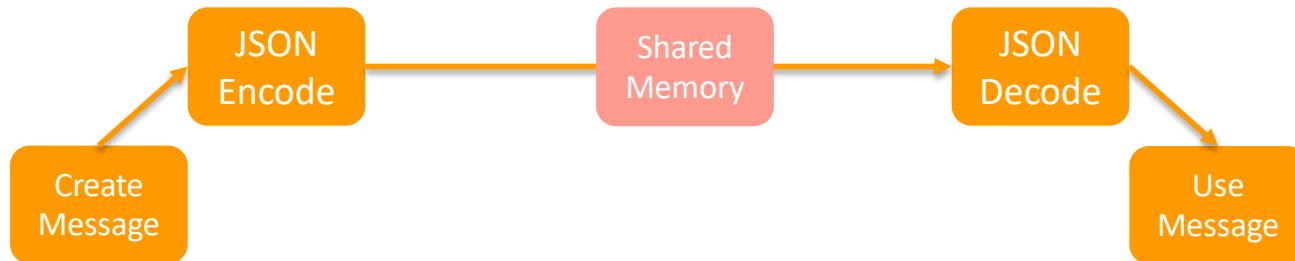
## Shortcut the network



## **Use Memory as the Network (no app code changes)**



## **Use Memory as the Network (Repackaged container sidecar)**



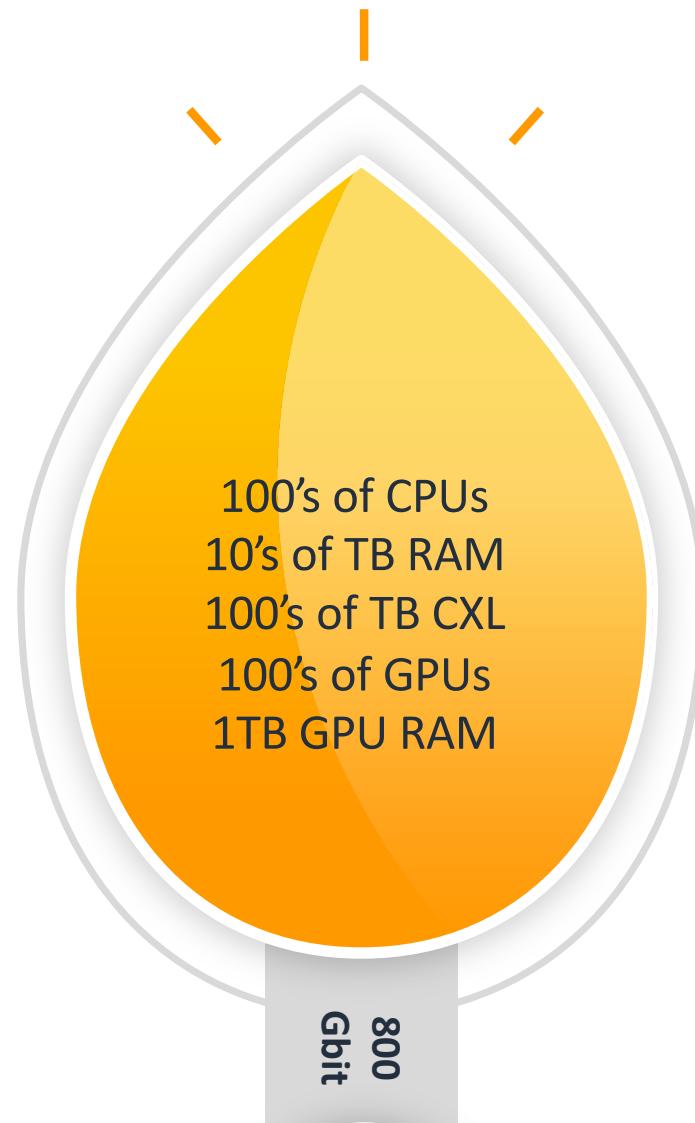
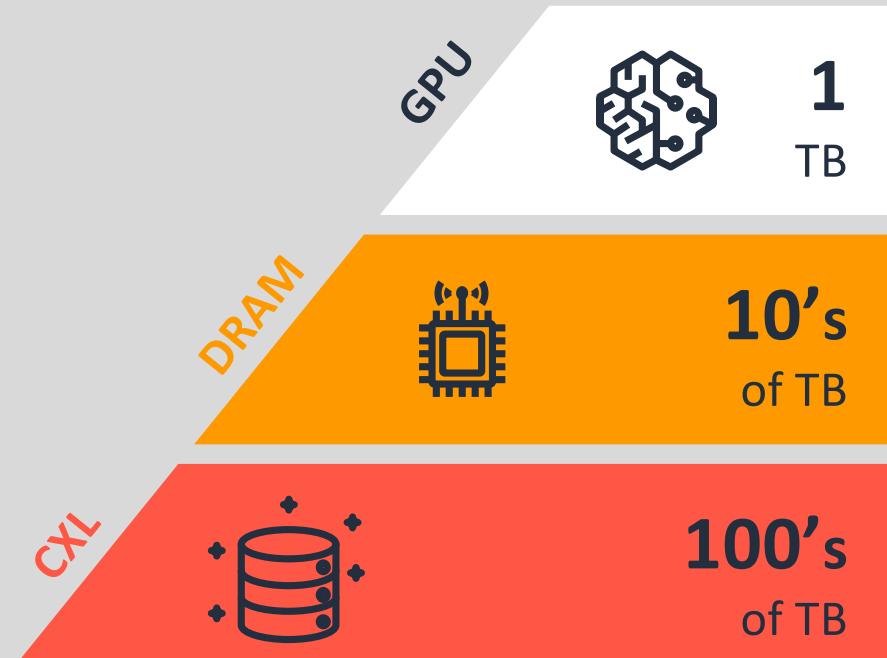
## **Use Memory as the Network (No need to encode/decode)**



IN 2023 OR SO...

## Large scale system

New memory hierarchy to manage

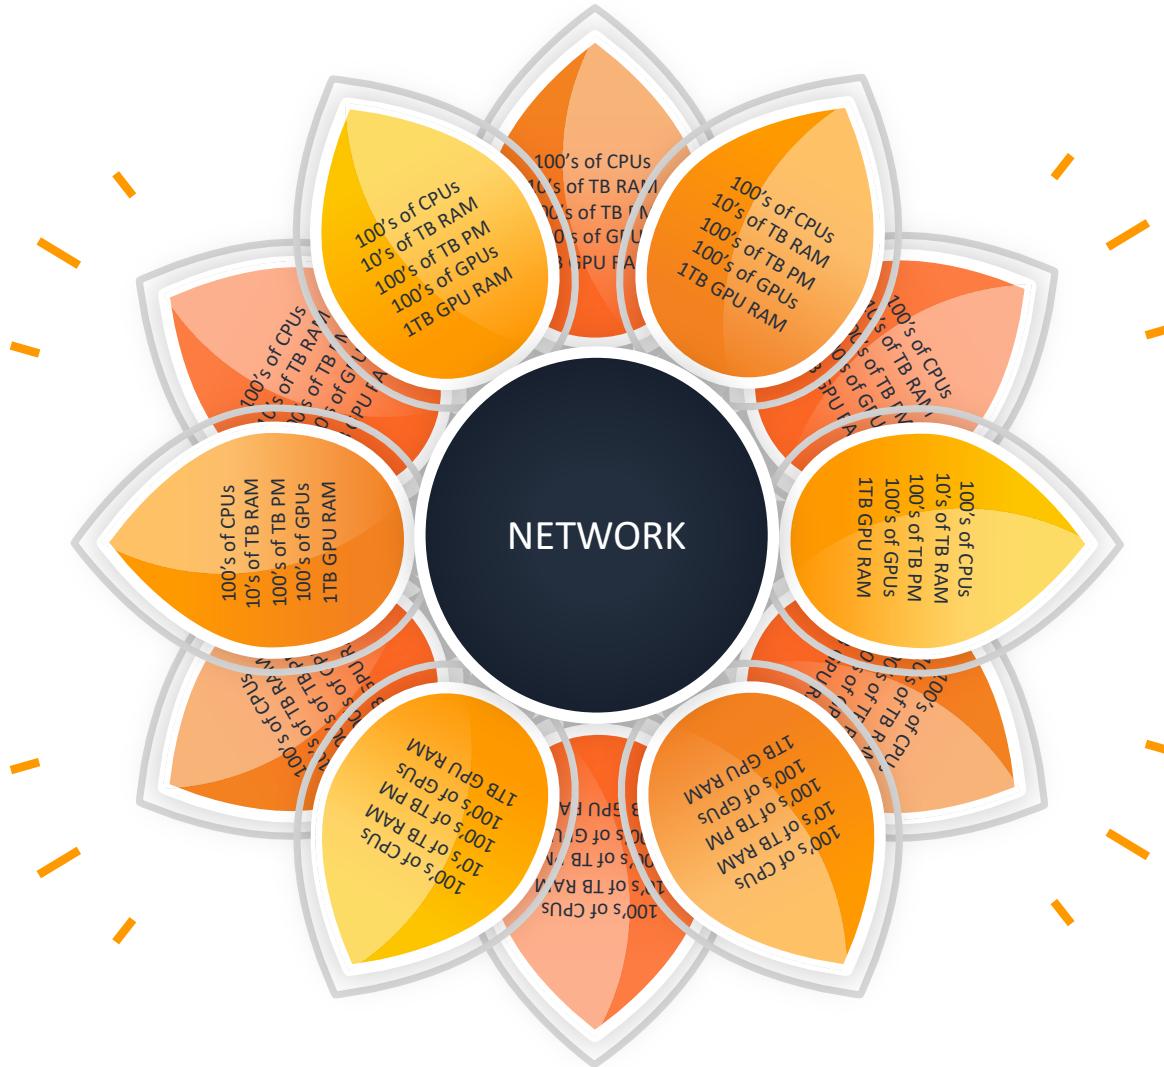


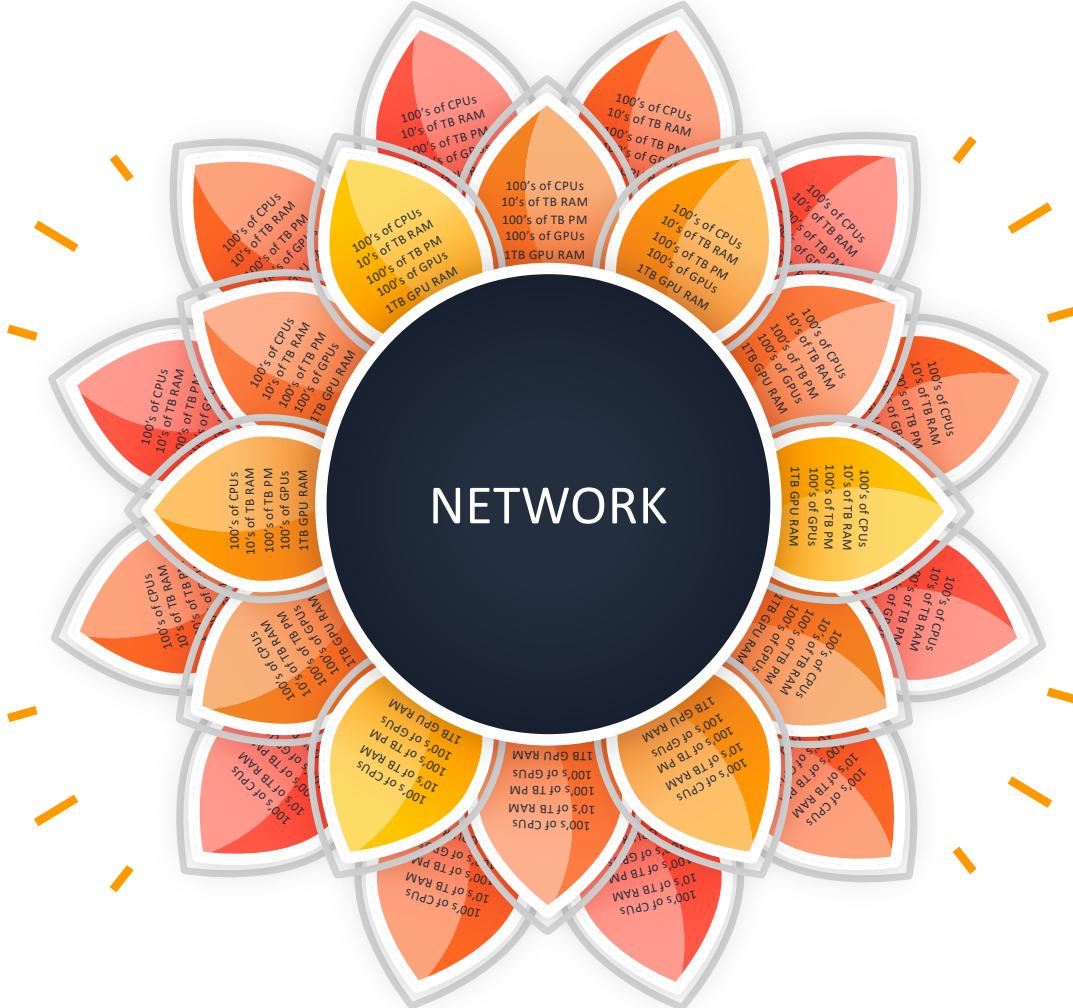


**Petalith  
architecture  
replicates data  
THREE WAYS  
for durability**

---

Via a local switch  
for lowest latency





**Dynamic Petalith  
fabric architecture**  
Shared memory in  
each petal is sized  
to suit the  
workload

Network topology  
is optimized for  
the workload



That's all folks!

Unless there's time for an encore  
and some drinks?



# Bottleneck Analysis

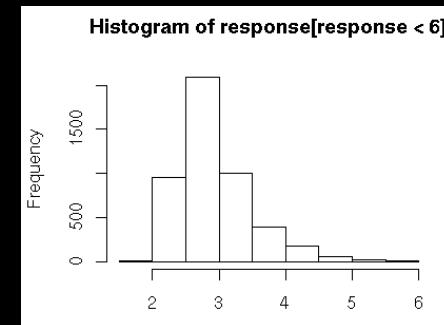
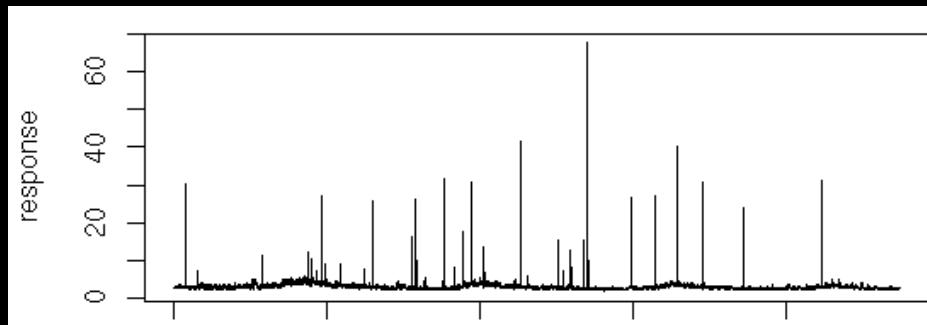
July 2013  
Adrian Cockcroft  
@adrianco

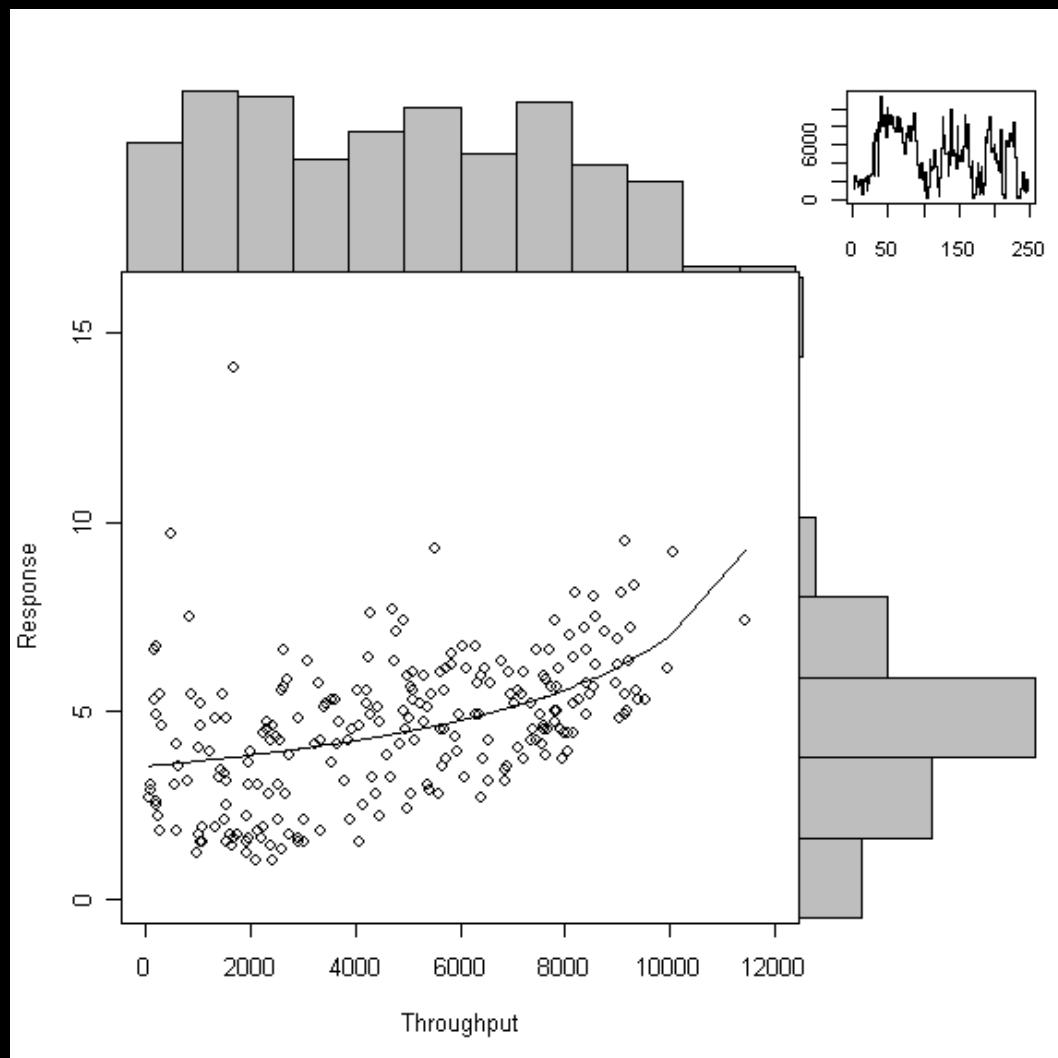




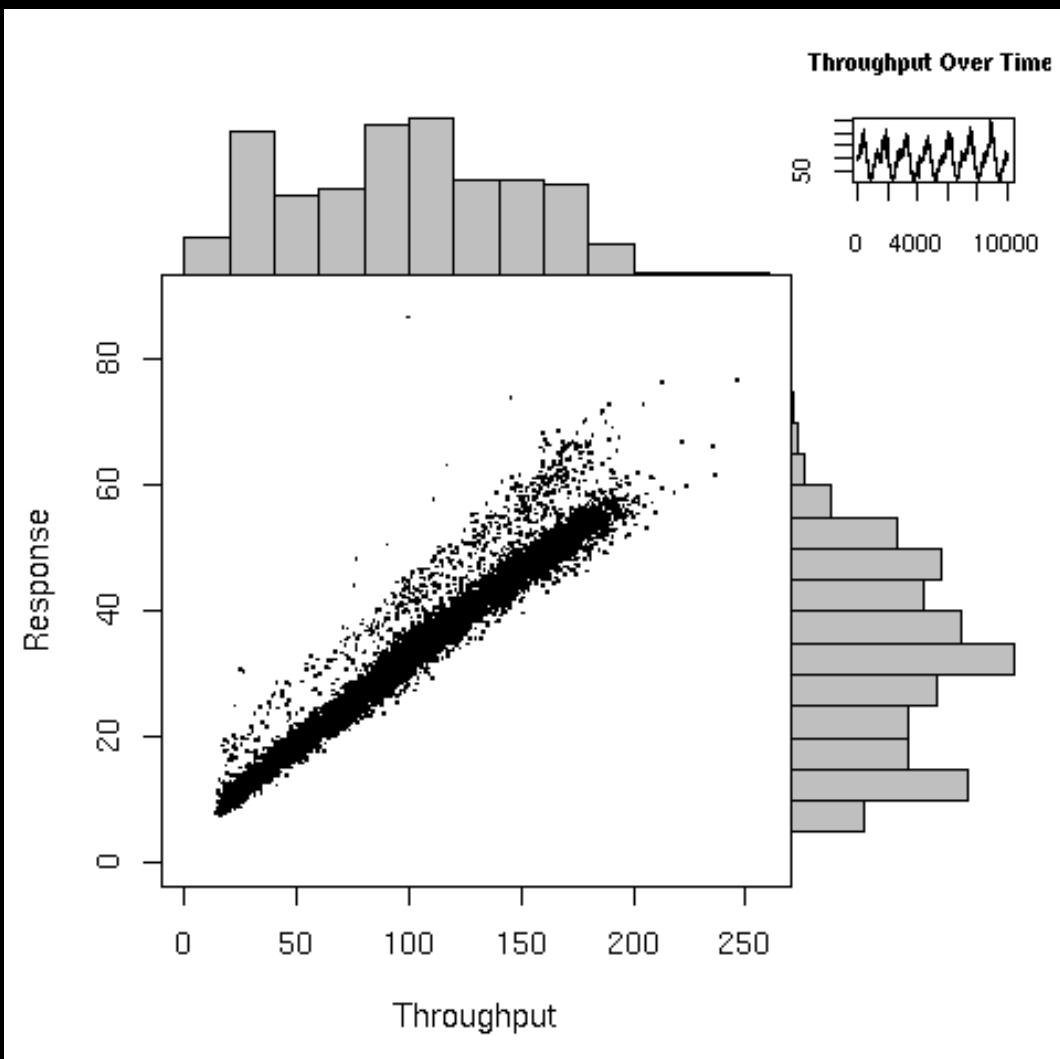
# Analysis

```
> summary(response)
   Min. 1st Qu. Median     Mean 3rd Qu.    Max.
1.909   2.550   2.820   3.086   3.214   67.680
> quantile(response,c(0.95,0.99))
   95%      99%
4.149556 6.922115
> sd(response)
1.941328
> mean(response) + 2 * sd(response)
6.968416
```





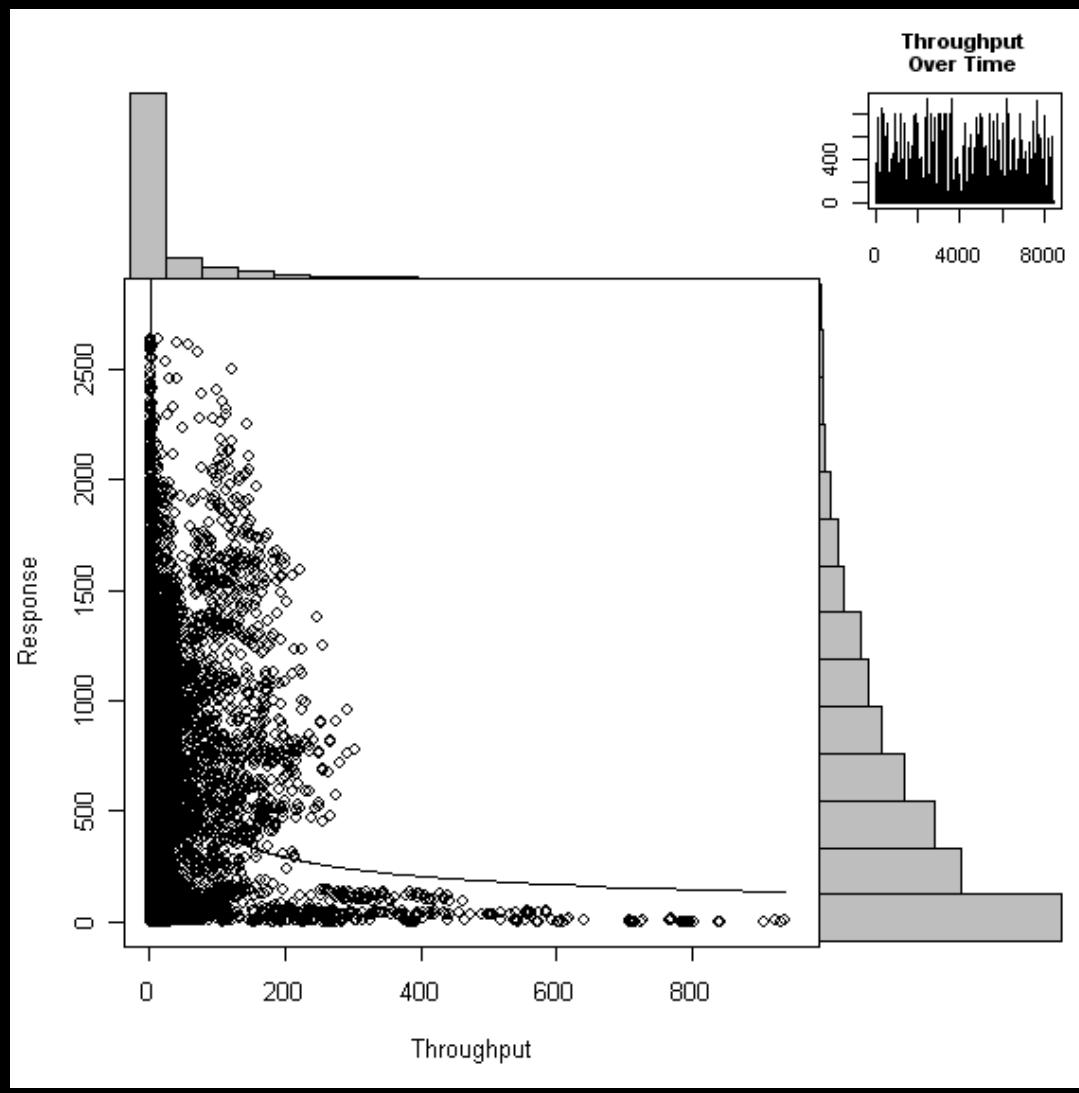


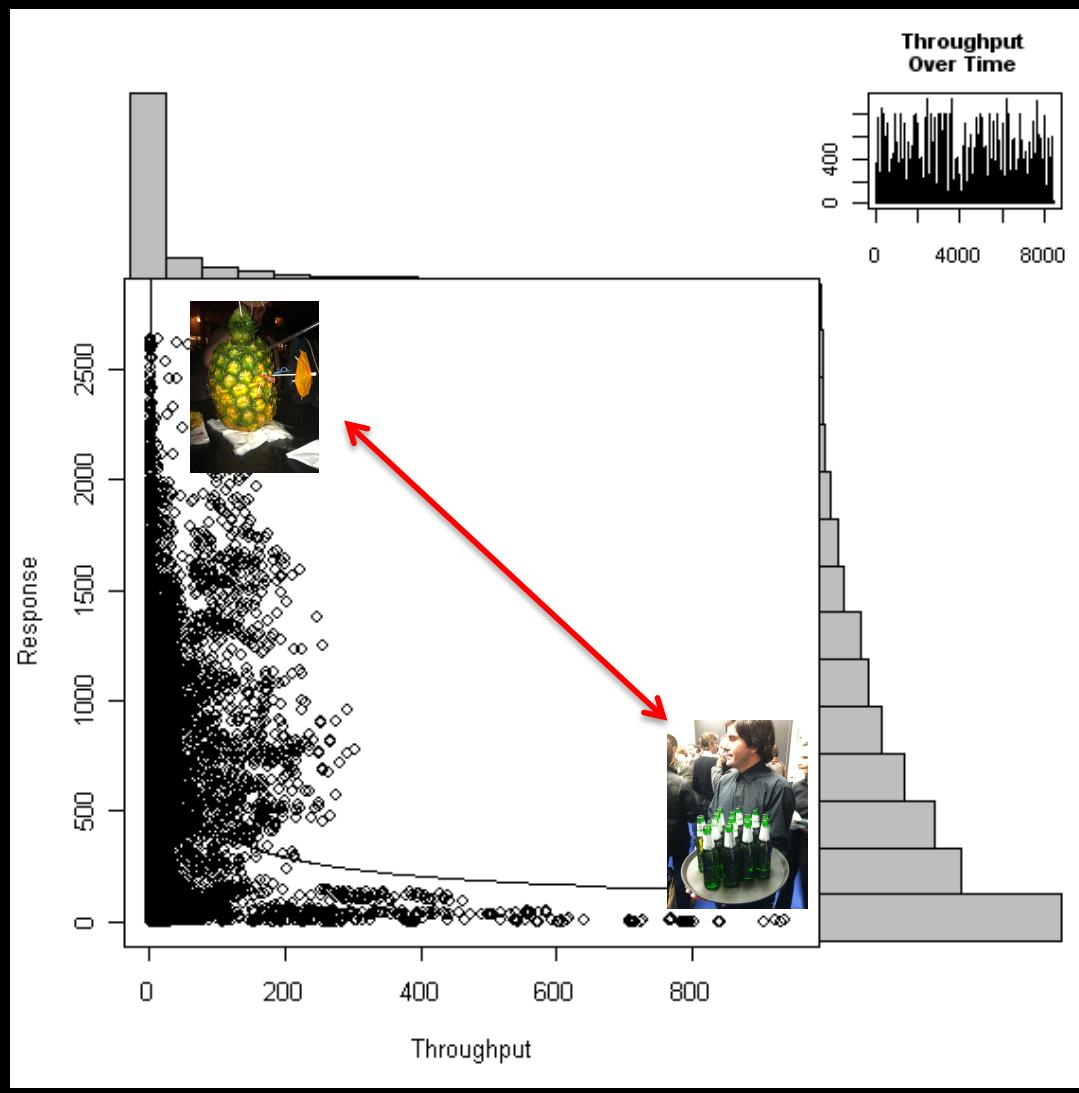




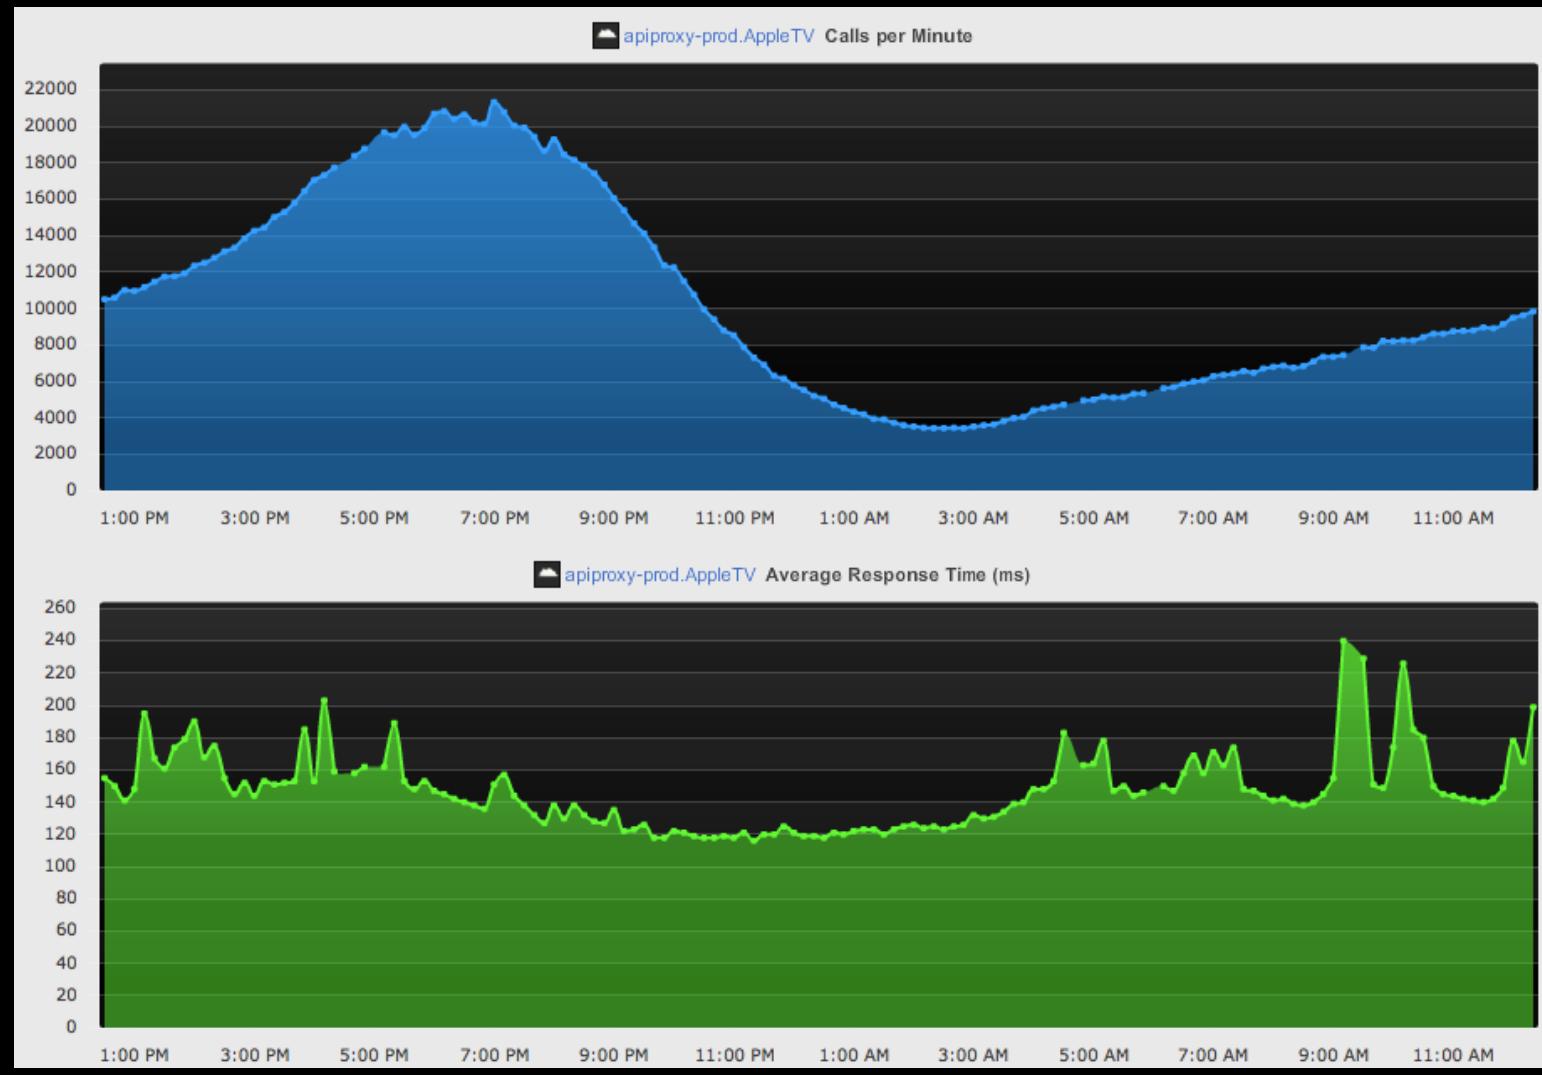


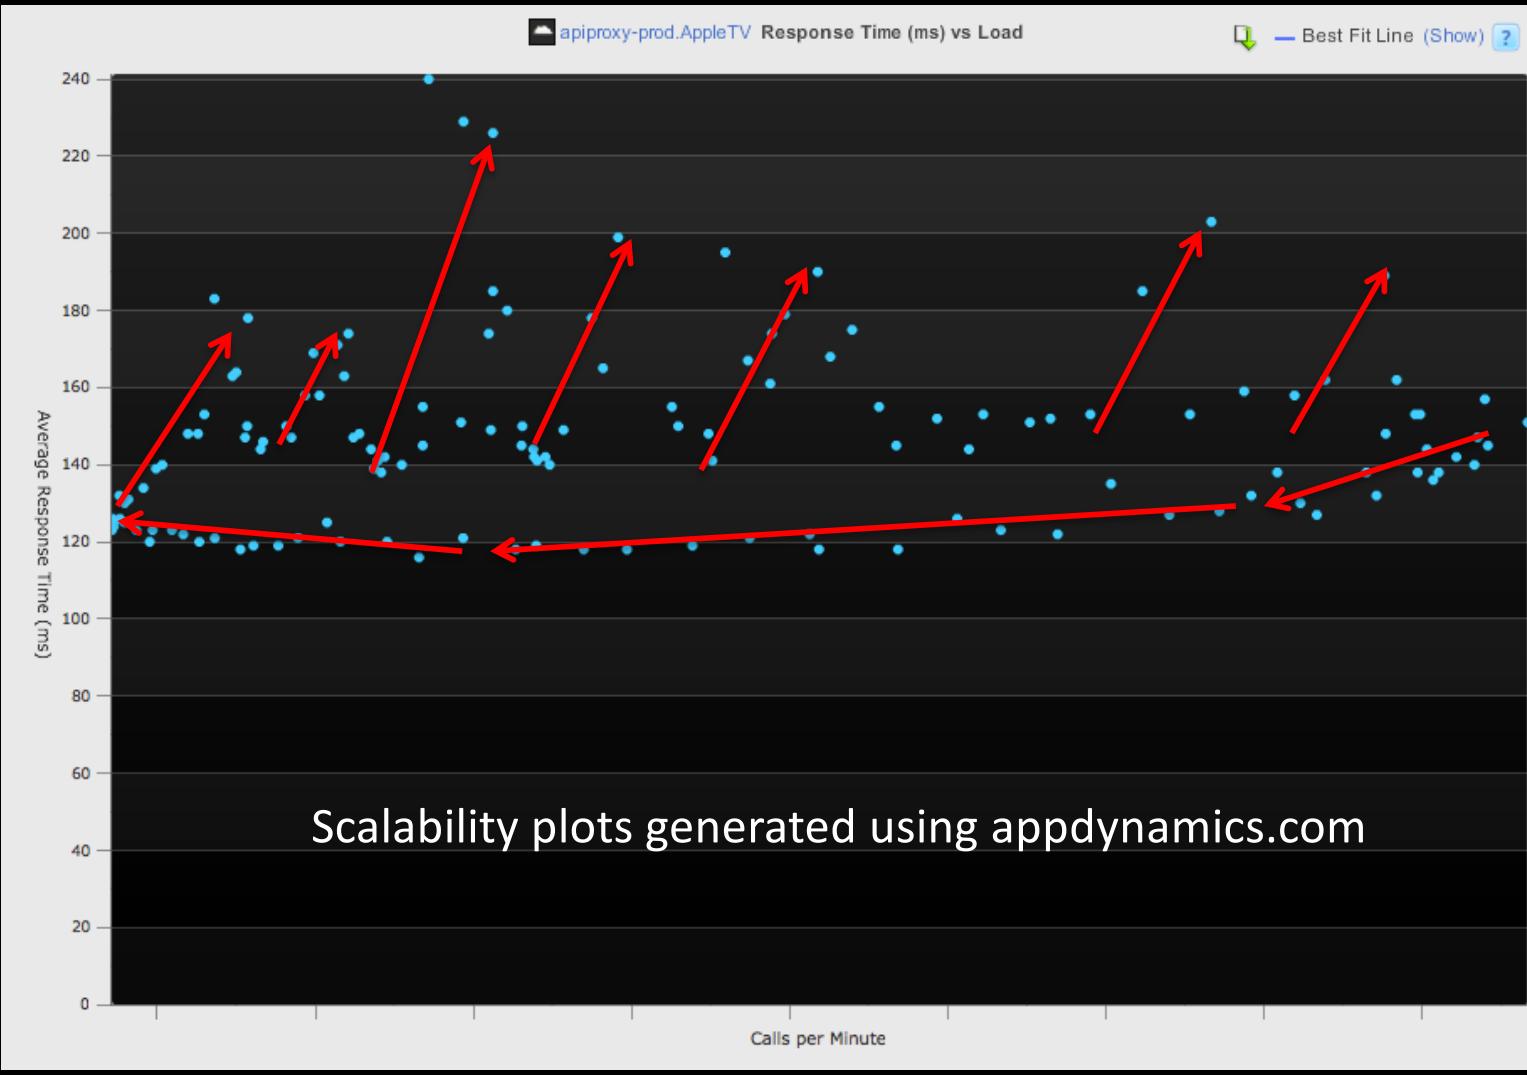
Pineapple Dalek says  
**INEBRIATE!**  
**INEBRIATE!**

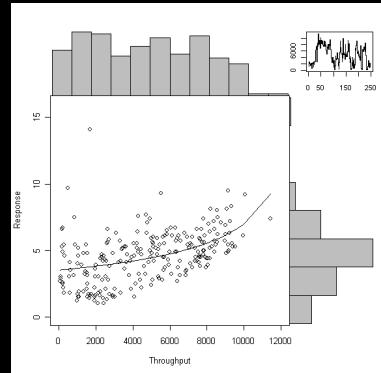




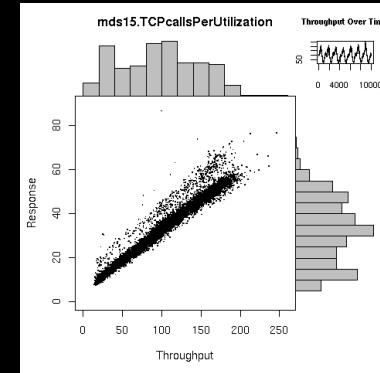




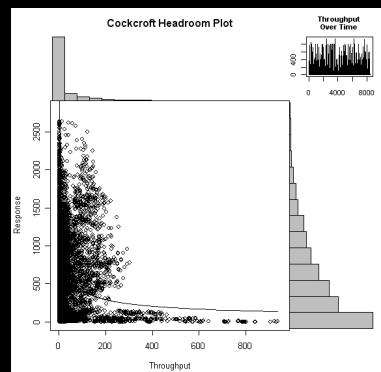




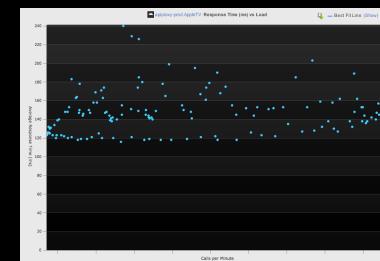
Well behaved



Lock Contention



Oscillating, thread shortage



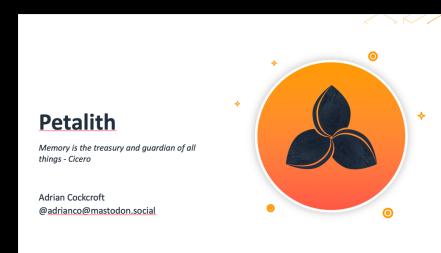
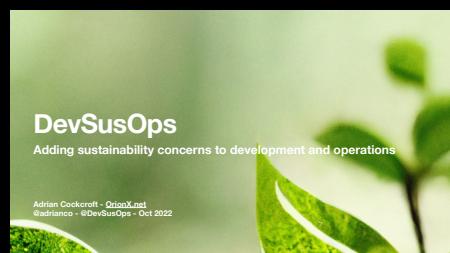
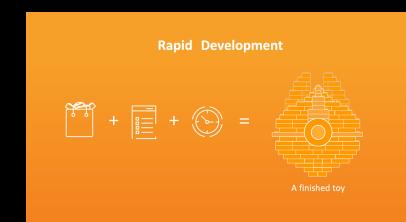
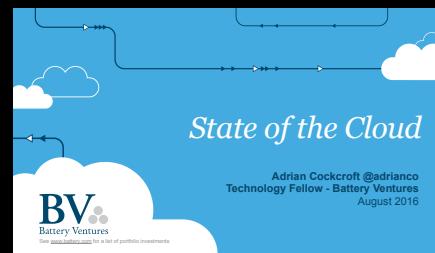
Looping autoscaled

<http://perfcap.blogspot.com/search?q=chp>

@adrianco



<https://soundcloud.com/adrian-cockcroft/black-tiger-dont-look-back?in=adrian-cockcroft/sets/black-tiger-bedford-civic-hall>  
<https://www.slideshare.net/adrianco>  
<https://www.slideshare.net/adriancockcroft>  
<https://github.com/adrianco/slides>  
<https://www.youtube.com/@adriancockcroft>



# Adrian's Greatest Hits, B-Sides and Re-issues

Thanks! Any questions?

Adrian Cockcroft | YOW | December 2022

[@adrianco@mastodon.social](mailto:@adrianco@mastodon.social)